

NoteBook

zhenqiufu

# Contents

<b>MyNote</b>	<b>3</b>
<b>I</b>	<b>4</b>
<b>1</b>	<b>5</b>
1.1 . . . . .	5
1.1.1 . . . . .	5
1.1.2 SMMU . . . . .	11
1.1.3 SVA . . . . .	22
1.2 NoC . . . . .	25
1.2.1 CHI . . . . .	25
1.2.2 NoC . . . . .	36
1.2.3 CMN700 . . . . .	38
1.2.4 CMN700 . . . . .	39
1.2.5 . . . . .	39
1.2.6 ARM GIC . . . . .	43
1.2.7 RISCV AIA . . . . .	49
1.2.8 Cortex M3 MCU . . . . .	49
1.3 ARM N2 . . . . .	53
1.4 . . . . .	55
1.4.1 ARM . . . . .	55
1.4.2 RISC-V Firmware . . . . .	59
1.4.3 Cache as RAM and On Chip Memory . . . . .	60
1.4.4 System Table . . . . .	61
1.4.5 EDK2 EVENT . . . . .	63
1.4.6 EDK2 MEMORY MAP . . . . .	74
1.5 . . . . .	74
1.5.1 . . . . .	74
1.5.2 RAID . . . . .	75
1.5.3 SSD . . . . .	77
1.5.4 . . . . .	78
1.6 . . . . .	79
1.6.1 Hypervisor . . . . .	79
1.6.2 MIT 6.824 Distributed Systems . . . . .	82
1.7 Memory consistency and cache coherence . . . . .	82
1.7.1 Memory Consistency . . . . .	82
<b>2</b>	<b>88</b>
2.1 ARM . . . . .	88
2.2 RISC-V Firmware . . . . .	93
2.3 Cache as RAM and On Chip Memory . . . . .	94
2.4 System Table . . . . .	95
2.5 EDK2 EVENT . . . . .	97
2.6 EDK2 MEMORY MAP . . . . .	109
<b>3 OS</b>	<b>110</b>
3.1 FreeRTOS . . . . .	110

<b>II</b>	<b>112</b>
<b>4 ACPI</b>	<b>113</b>
4.1 ACPI Introduction and Overview . . . . .	113
4.2 CEDT . . . . .	124
<b>5 AMBA</b>	<b>128</b>
5.1 PCIe . . . . .	128
5.2 PCIe . . . . .	128
5.2.1 . . . . .	128
5.2.2 Data Link Layer . . . . .	140
5.2.3 Transaction Layer . . . . .	146
5.2.4 PCIe . . . . .	151
5.3 ARM N2 . . . . .	157
5.3.1 PCIe . . . . .	159
5.4 PCIe . . . . .	160
5.4.1 Resource Degradation . . . . .	160
5.4.2 EDK2 PCIe OpRom . . . . .	170
5.5 PCIe . . . . .	174
5.5.1 PCIe AER . . . . .	174
5.5.2 PCIe Interrupt . . . . .	174
5.5.3 PCIe Hot-Plug . . . . .	174
5.5.4 PCIe Power Management . . . . .	174
5.5.5 <i>DPC</i> . . . . .	174
<b>6 CXL</b>	<b>195</b>
6.1 CXL . . . . .	195
6.2 CXL . . . . .	200
6.3 CXL in CMN . . . . .	205
6.4 CXL in SCP and EDK2 . . . . .	207
6.5 CXL . . . . .	214
6.6 DDR Introduction . . . . .	216
<b>III</b>	<b>223</b>
<b>7 LeetCode</b>	<b>224</b>
7.1 C basic . . . . .	224
7.2 Bit . . . . .	228
7.3 LeetCode . . . . .	229
7.4 . . . . .	232
7.5 . . . . .	241
7.6 . . . . .	242
7.7 . . . . .	244
7.8 Stack . . . . .	245
<b>8 System Design</b>	<b>249</b>
8.1 Git . . . . .	249
<b>IV</b>	<b>250</b>
<b>9</b>	<b>251</b>
<b>10 2024</b>	<b>252</b>

# MyNote

- 
- 
- 
- 

by Zhenqiu

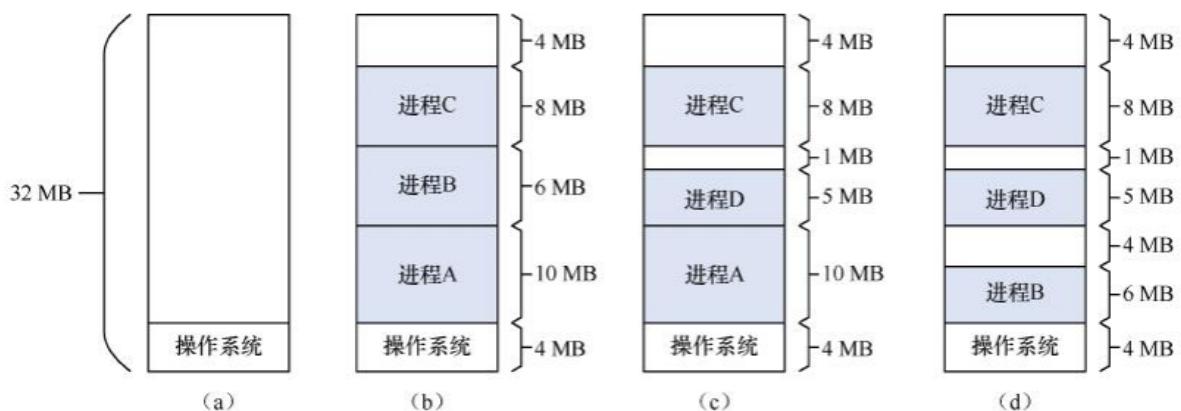
# **Part I**

# Chapter 1

## 1.1

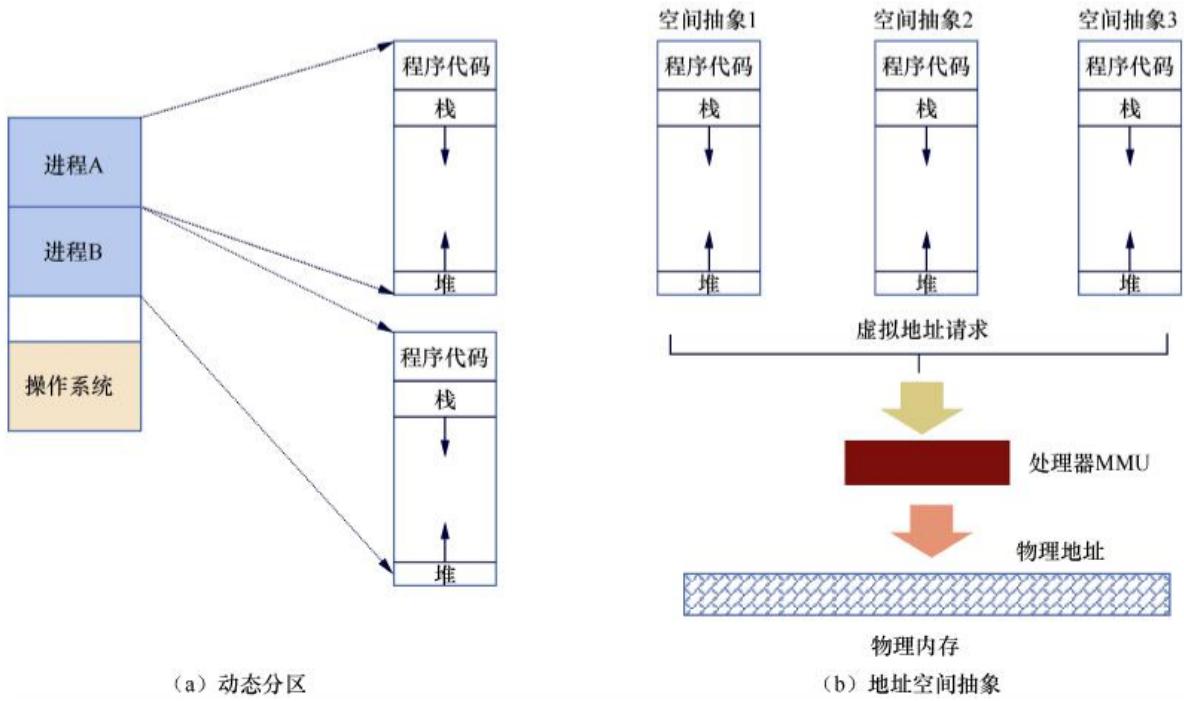
### 1.1.1

[toc]

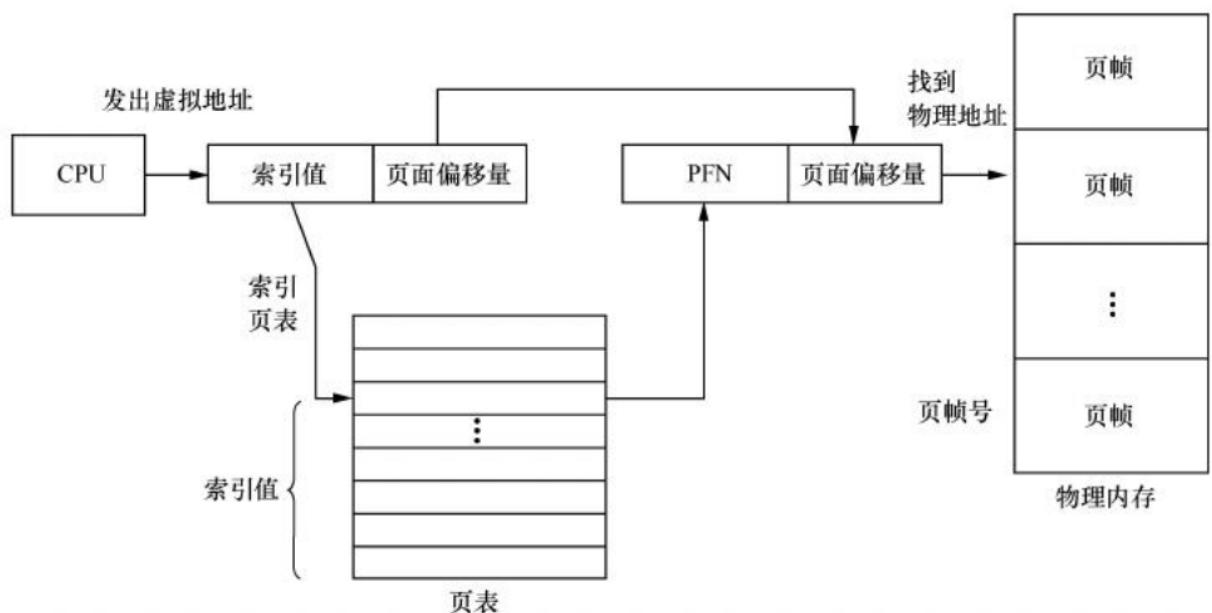


▲图 14.1 动态分区示意图

- A B
- 
-



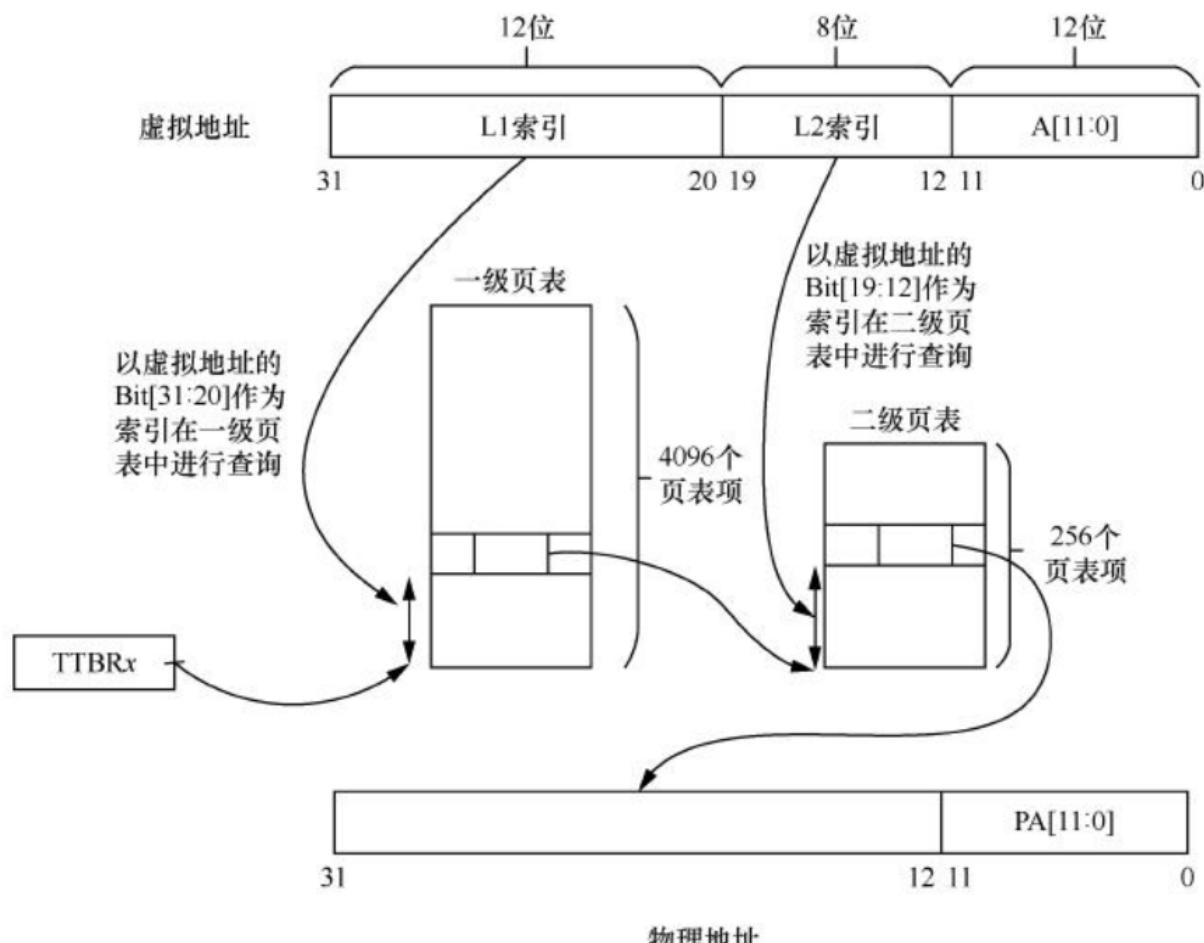
▲图 14.2 动态分区和地址空间抽象



▲图 14.3 页表查询过程

4G 4K  $2^{20}$  4B 4M  
4M table table

32 4 GB 512 MB



▲图 14.5 二级页表查询过程

14.5	VA[31:20]	12	4096	VA[19:12]	8	256	16 KB
4096		4	16 KB				

PFN

entry	5	KB
-------	---	----

ARM64 3 4

X86 5

#### ARM64

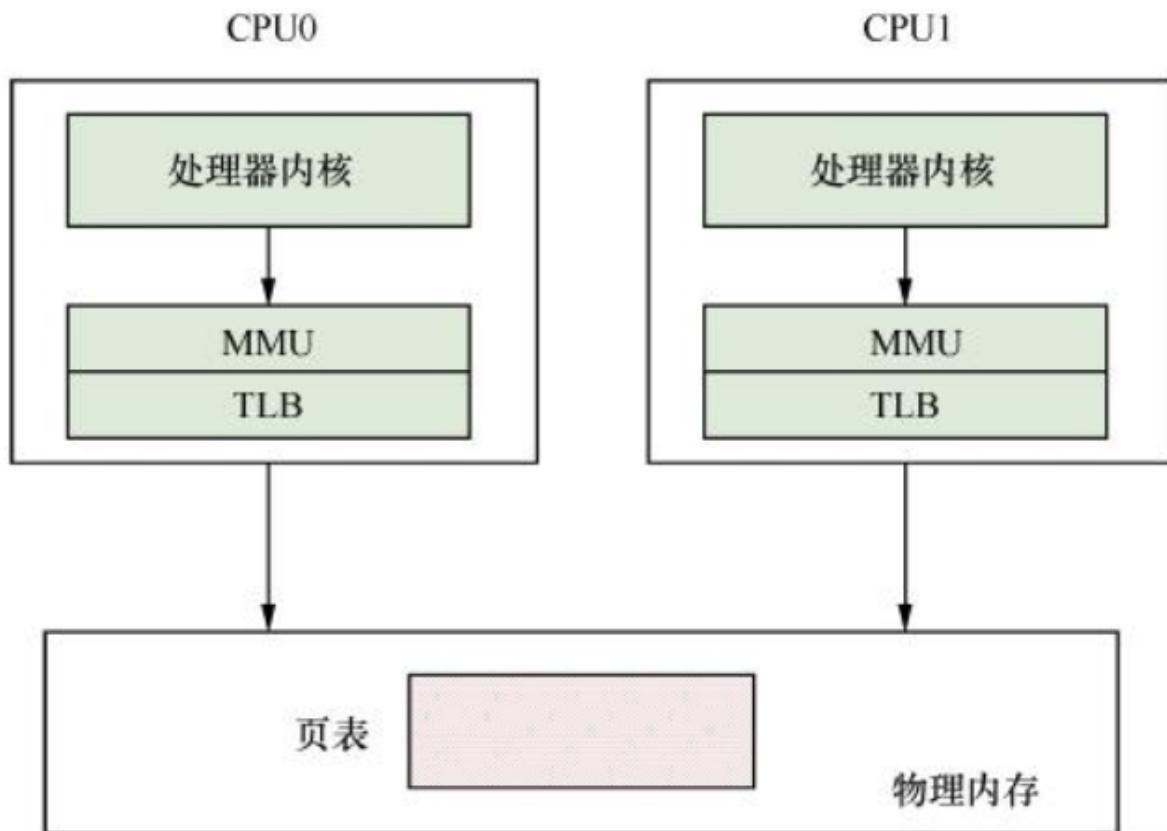
ARM64 MMU TLB Table Walk Unit TWU



▲图 14.8 ARM 处理器的内存管理体系结构

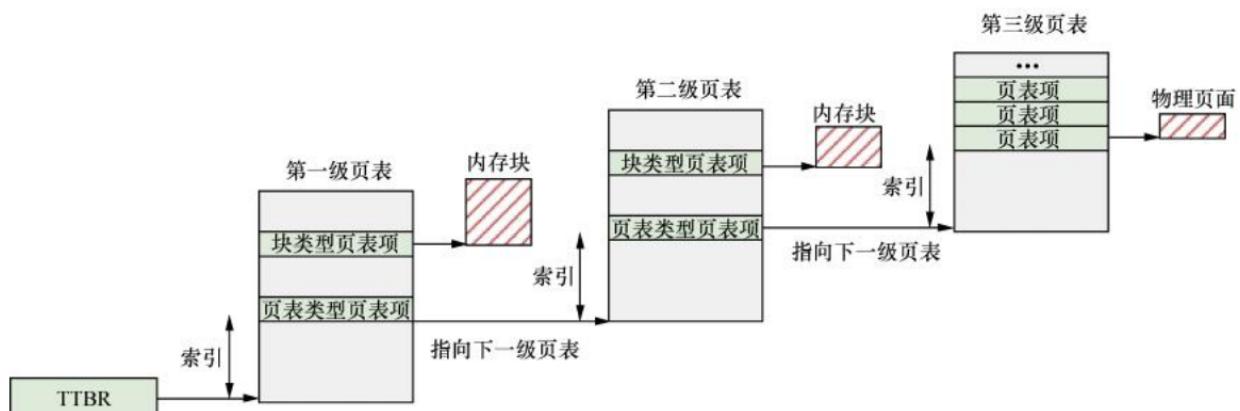
SMP Symmetric Multi-Processor

MMU TLB

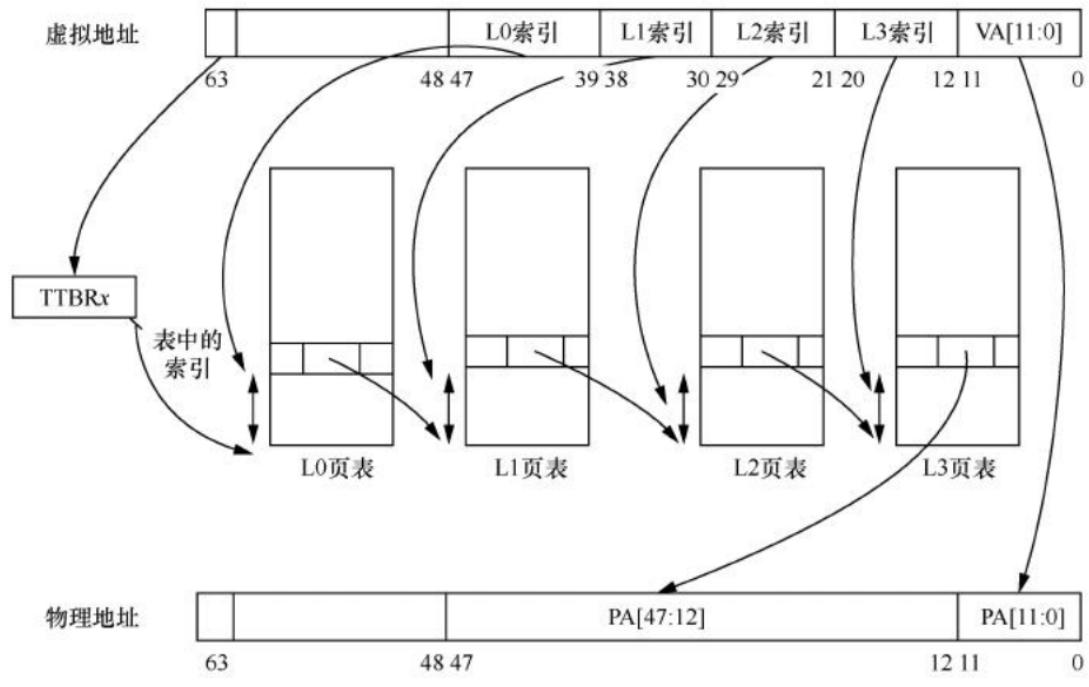


▲图 14.10 SMP 系统与 MMU

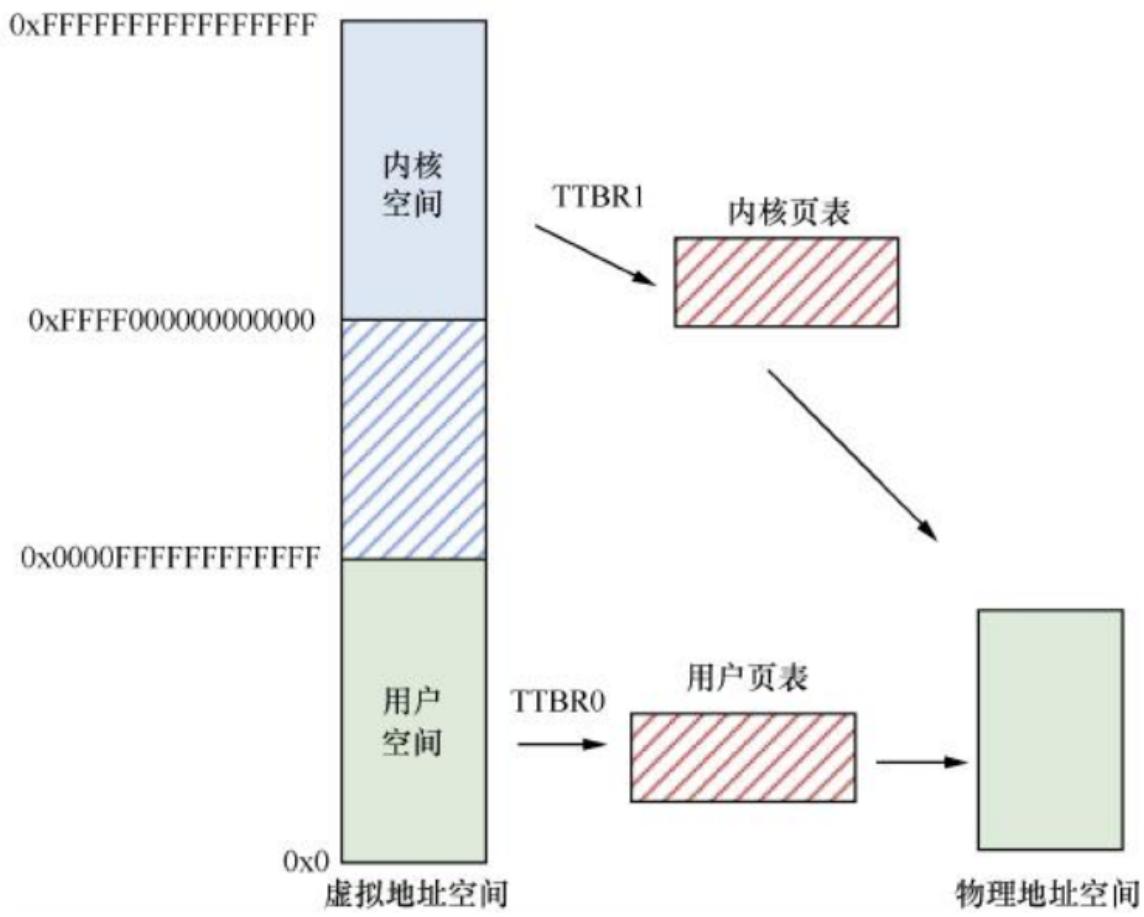
AArch64	MMU	
VA	PA	
1	Intermediate Physical Address IPA	2 IPA PA
ARMv8	4 KB 16 KB 64 KB 3	
AArch64	48 VA	256 TB
•	0x0000000000000000	0x0000FFFFFFFFFFFF
0	TTBR0_ELx	
•	0xFFFF000000000000	0xFFFFFFFFFFFFFF
1		TTBR1_ELx



▲图 14.11 页表三级映射的示意图



x86_64	AArch64	14.16	3	CPU	MMU
TTBR0	CPU	MMU	TTBR1		



▲图 14.16 两套页表

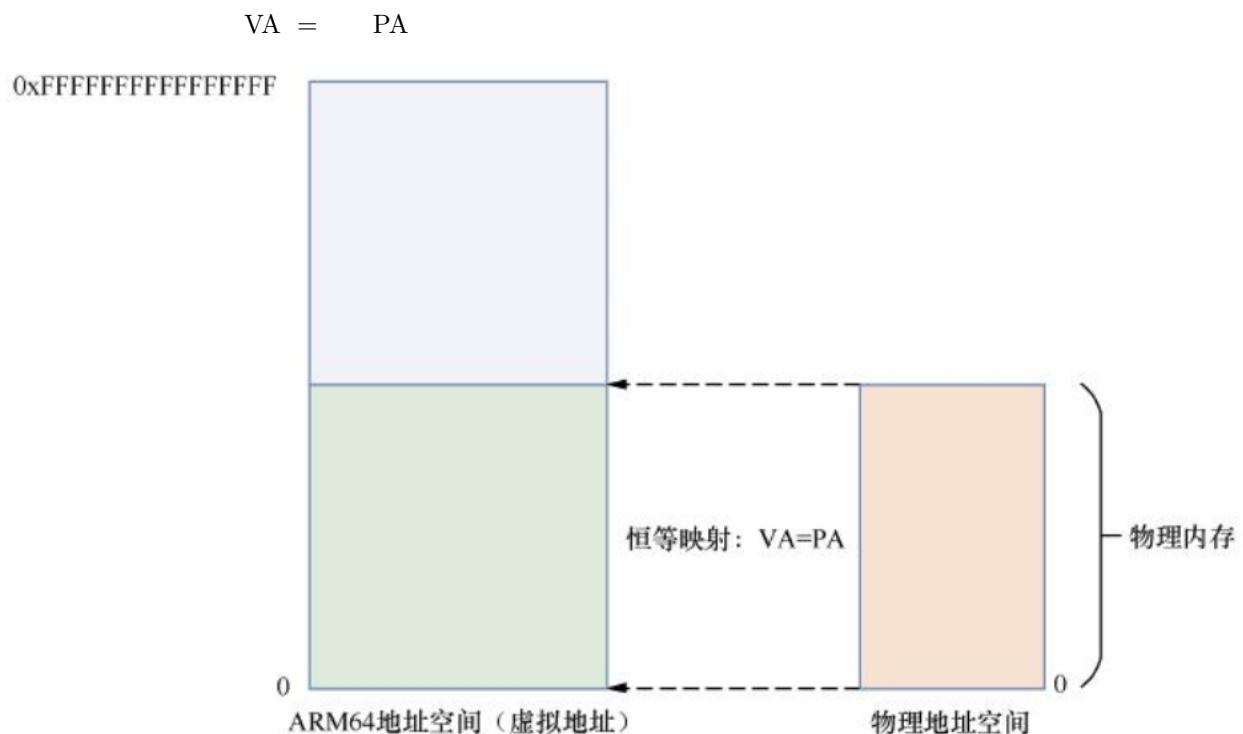


▲图 14.17 L0~L2 页表项描述符

无效的页表项	63	保留										1	0
保留的页表项	63	保留										0	1
4 KB粒度的页表项	63	51	50	48	47	输出地址[47:12]				12	11	1	0
16 KB粒度的页表项	63	51	50	48	47	输出地址[47:14]			14	13	12	11	1
64 KB粒度的页表项	63	51	50	48	47	输出地址[47:16]		16	15	14	13	12	11
L3 Entry													

▲图 14.18 L3 页表项描述符

Translation Control Register TCR                      System Control Register SCTRLR  
 Translation Table Base Register TTBR



▲图 14.24 恒等映射

### 1.1.2 SMMU

[toc]

## SMMU

(SMMU, System Memory Management Unit)

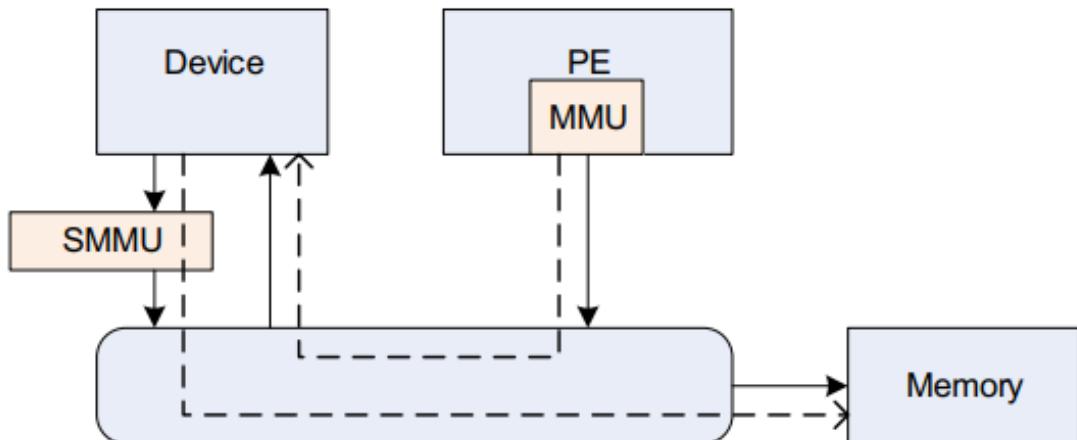
SMMU

1.
  - SMMU
  - I/O SMMU I/O
2.
  - SMMU
  - SMMU
3.
  - SMMU I/O
4.
  - SMMU
  - SMMU

## MMU-700

ARM SMMU IO IOMMU MMU-700 IP  
SMMU Spec MMU-700 TRM

## MMU SMMU



**Figure 2.1: System MMU in DMA traffic**

## SMMU

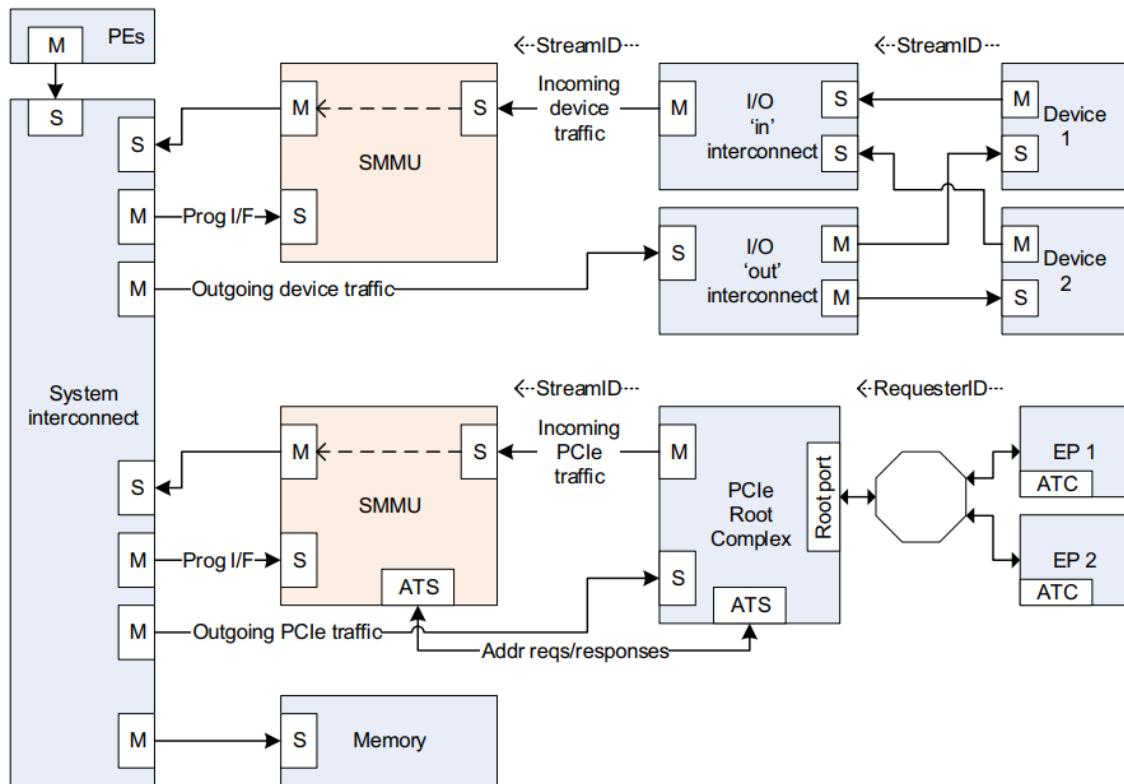


Figure 2.2: SMMU placement in an example system

- 1 SMMU DMA device      device1 VA      smmu VA PA
- 2 SMMU PCIe Root Complex (which itself hosts a network of endpoints)      PCIE      ATS

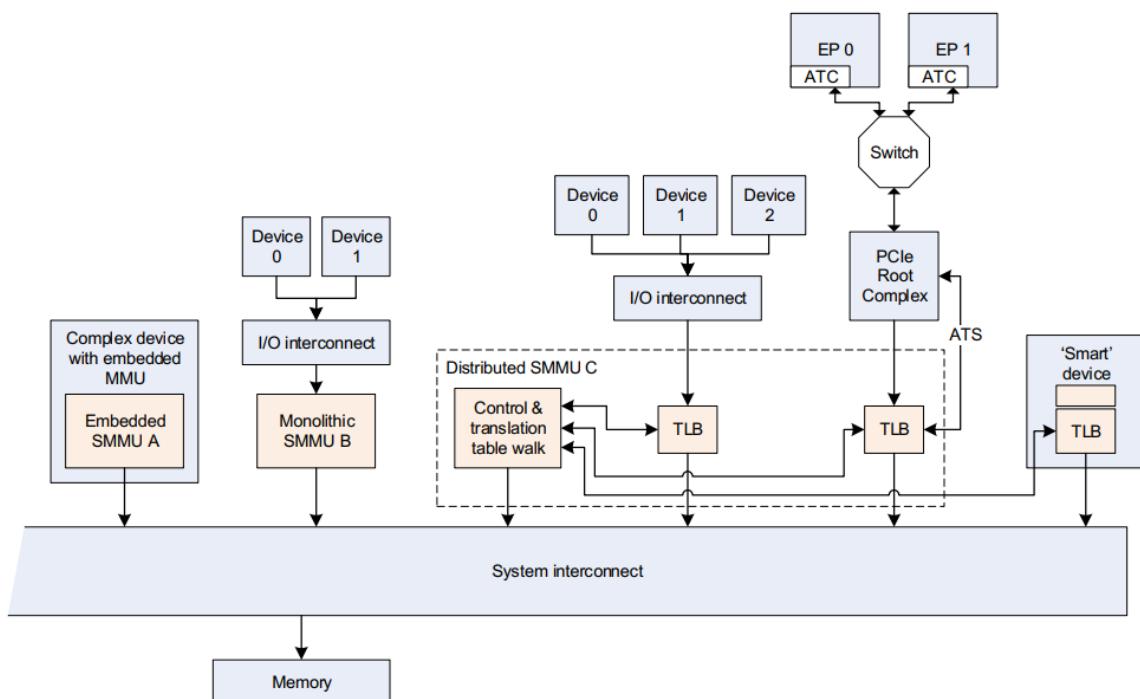


Figure 2.3: Example SMMU implementations

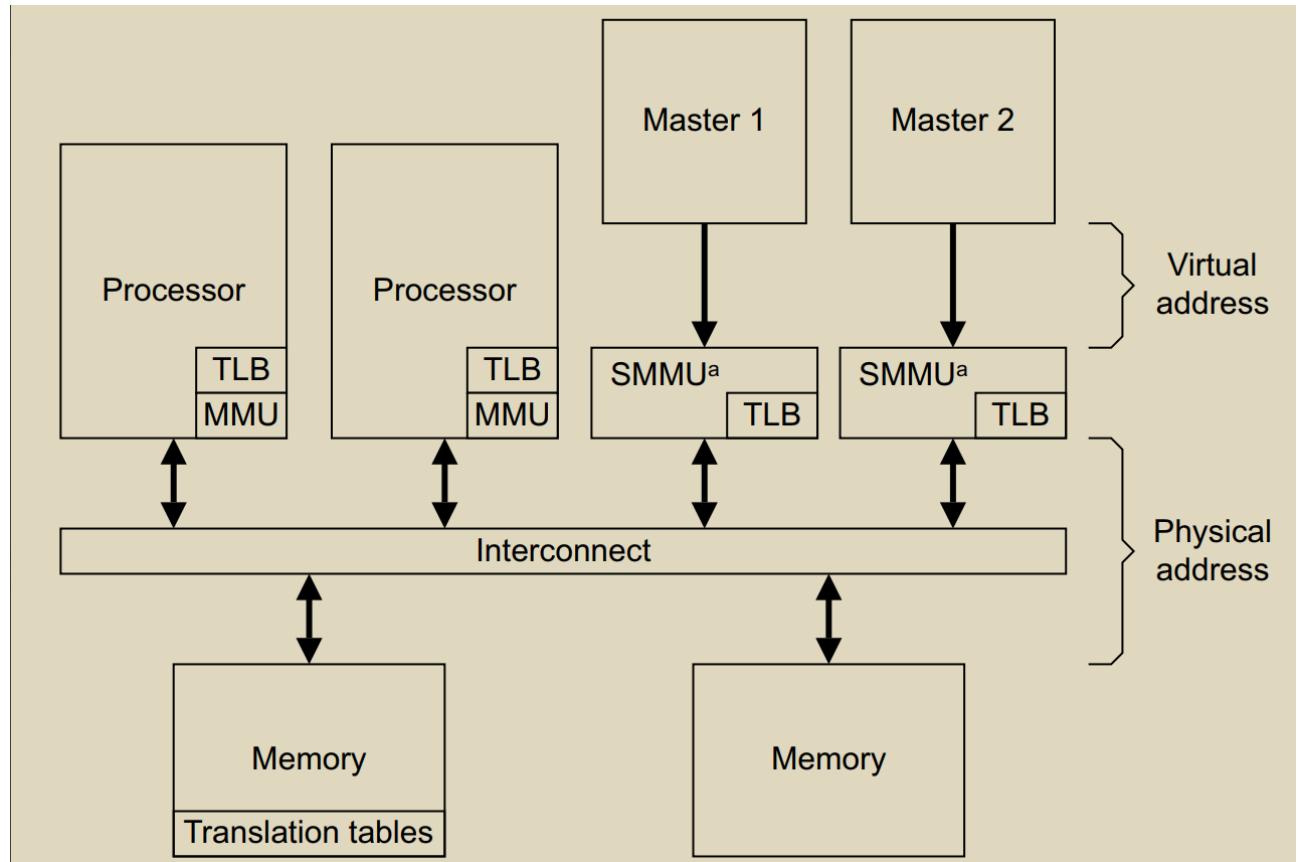
1 SMMU A complex device

2 SMMU B device I/O DMA

3 SMMU C

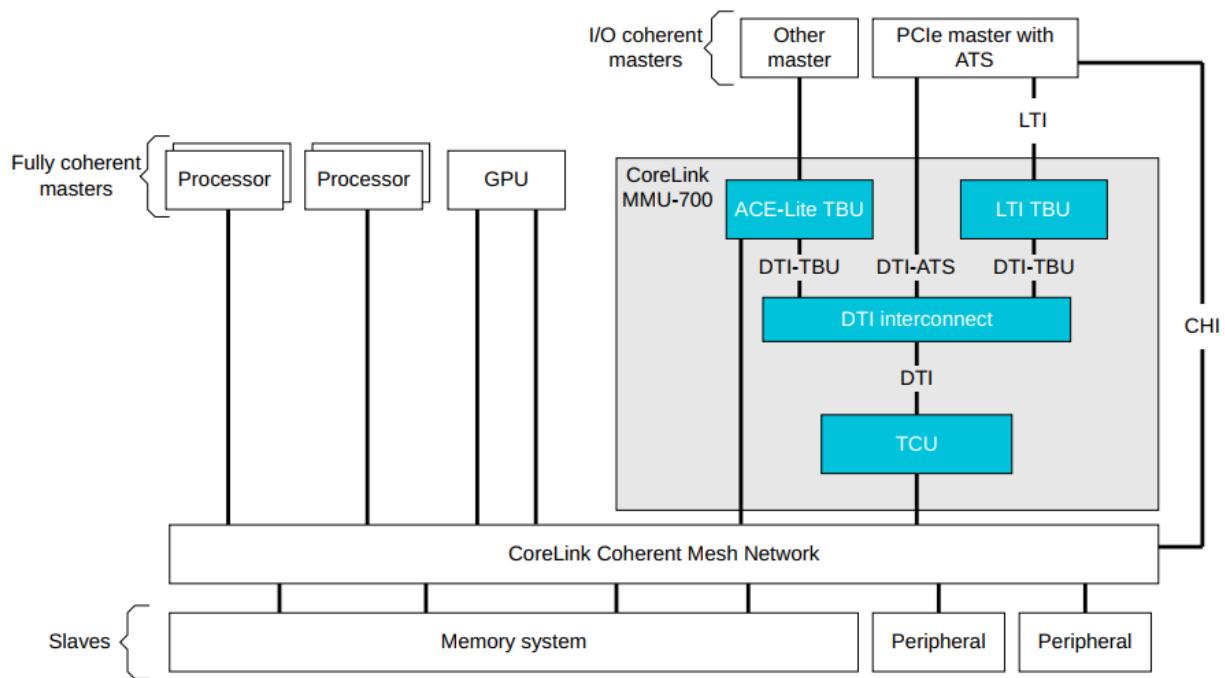
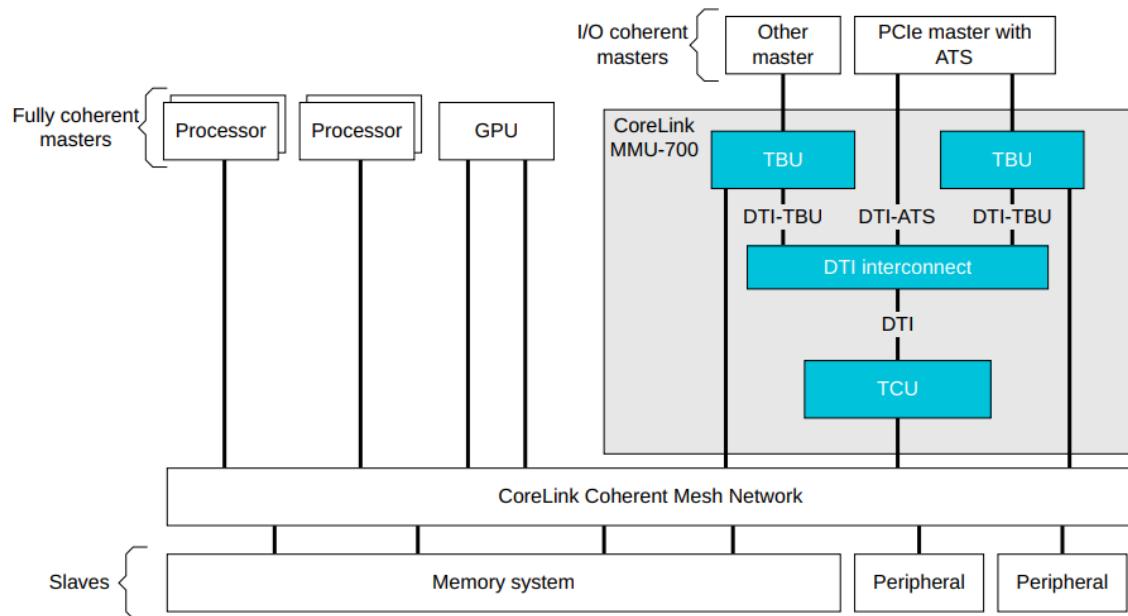
a. central translation table walker

b. device PCIe



## MMU-700

**Figure 3-1: Example system with the MMU-700**



The MMU-700 contains the following key components:

- Translation Buffer Unit (TBU)**: The TBU contains Translation Lookaside Buffers (TLBs) that cache translation tables. The MMU-700 implements a TBU that can be connected to single master or multiple masters. It is also possible to connect multiple TBUs to a single master to improve performance. These TBUs are local to the corresponding master and can be one of the following:

- ACE-Lite TBU
- LTI TBU

**Translation Control Unit (TCU)** The TCU controls and manages the address translations. The MMU-700 implements a single TCU. In MMU-700-based systems, the AMBA® DTI protocol defines the standard for communicating with the TCU. See the AMBA® DTI Protocol Specification.

**DTI interconnect** The DTI interconnect connects multiple TBUs to the TCU

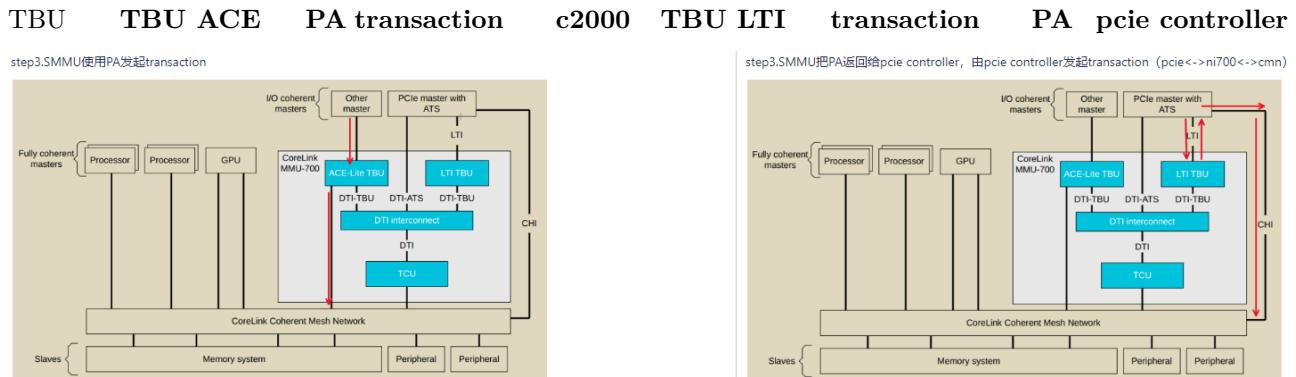
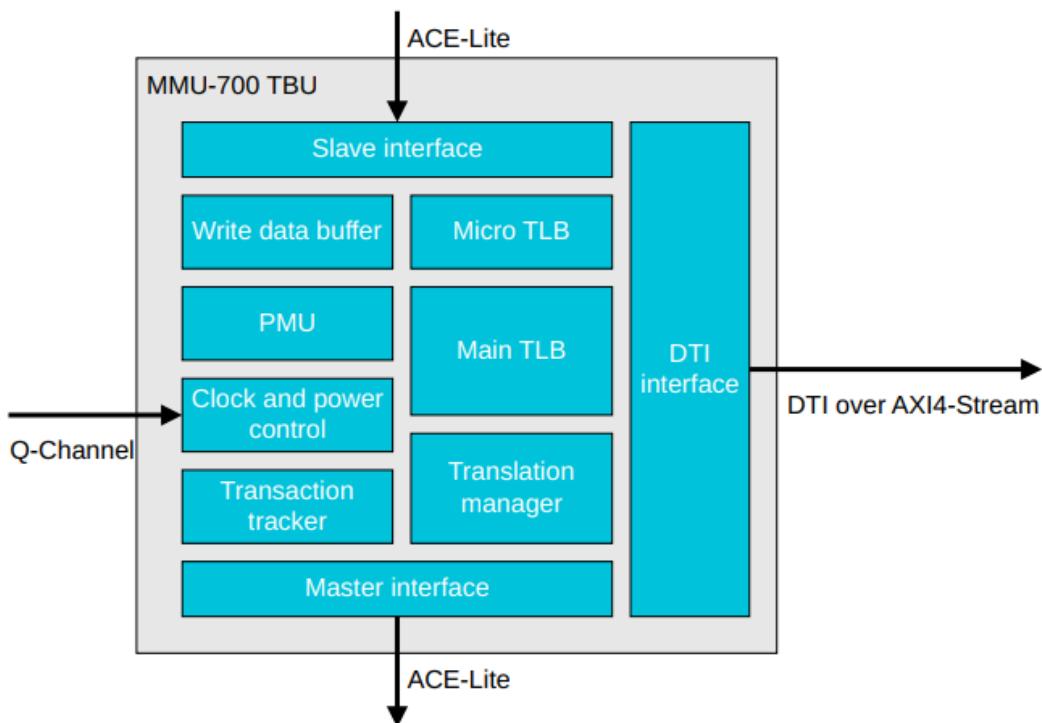
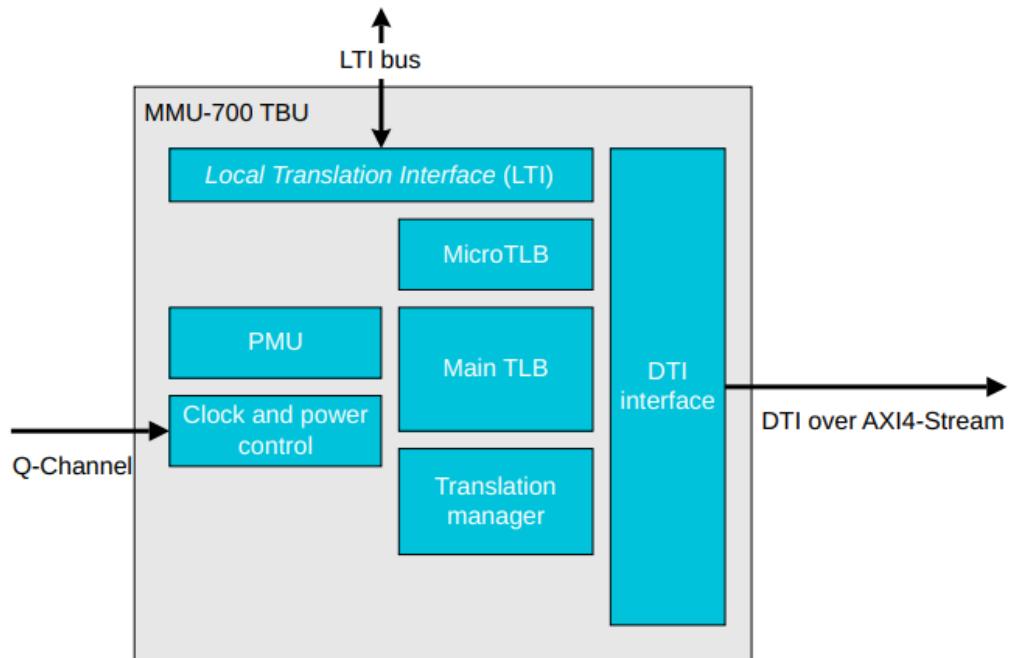


Figure 3-3: MMU-700 ACE-Lite TBU



**Figure 3-4: MMU-700 LTI TBU**



## TLB

The TBU contains Translation Lookaside Buffers (TLBs) that cache translation tables.

TBU cache micro TLB mainTLB

PA = F(streamID, substreamID, Address)

step1.master pcie controller VA transaction

step2.SMMU VA SMMU\_CRO.SMMUEN SMMU, bypass, VA PA DMA , StreamID STE;

STE.Config Stage1 , VA IPA Stage2 , STE.S1ContextPtr SubStreamID CD, Stage 1 translation table ASID, Stage 1 translation table walk, VA IPA Stage 2 ;

STE.Config Stage2 , IPA PA SMMU , STE.S2TTB Stage 2 translation table VMID, Stage 2 translation table walk, IPA PA SMMU ,

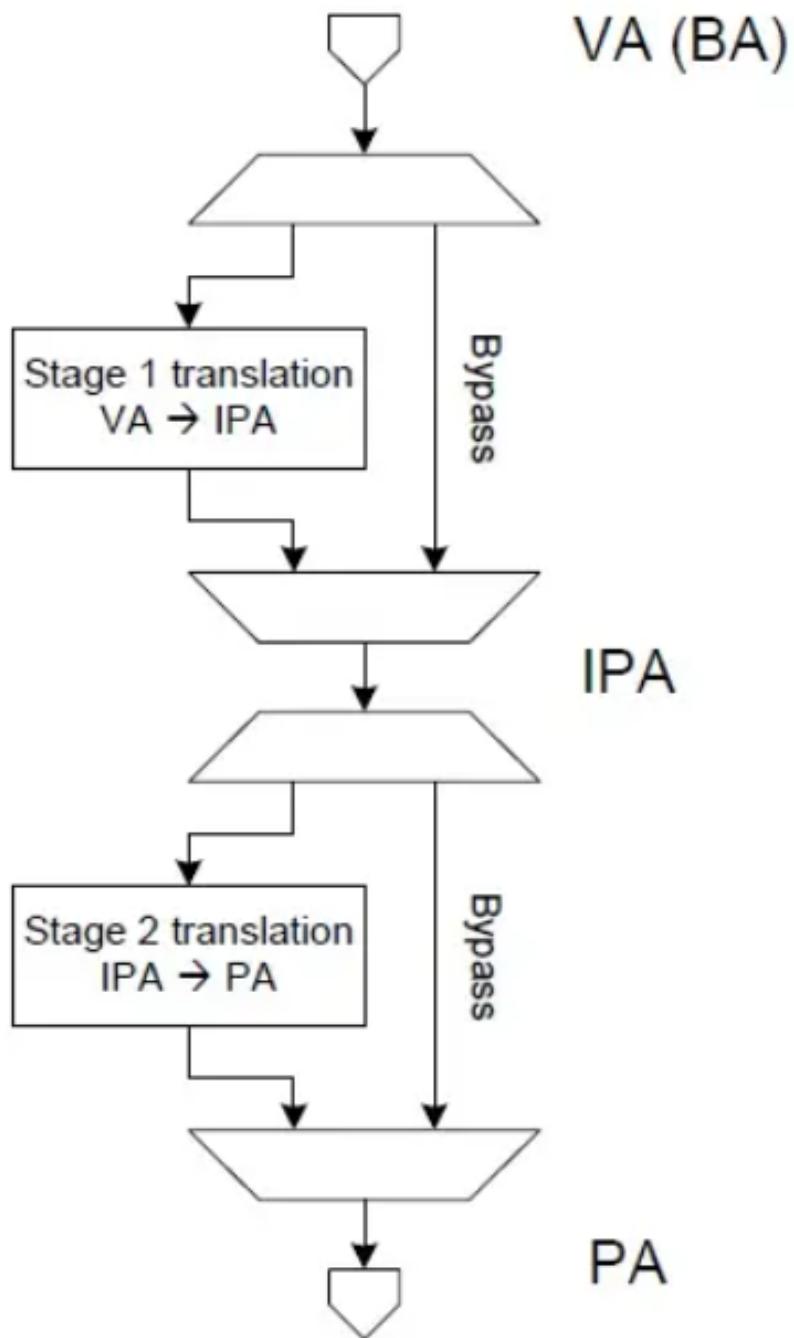
Note: Stage 1/2 translation table walk , SMMU TLB , , , translation table walk translation table walk TLB

untranslated transaction, PCIe ATS translation request/PRI

PA

- TBU TLB PA
- TBU TLB TBU TCU transaction table walk memory translation table VA PA

step3.SMMU PA transaction



SMMU enable

SMMU

StreamID	SMMU	SMMU	Endpoint	SMMU StreamID	(
	SubstreamID	PCIe PASID	)		
STE	Stream Table Entry, STE		stage2	CD Context Descriptor	.
CD	Context Descriptor,		stage1		

STE CD

STE stream table entry stage1 CD CD stage1 stage2 stage2

### 3.3.1.1 Linear Stream table

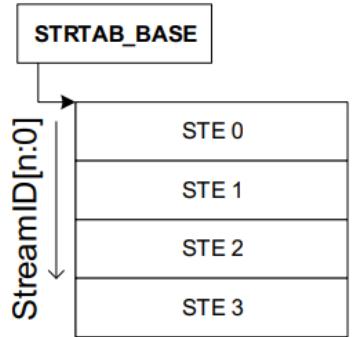


Figure 3.1: Linear Stream table

STE

### 3.3.1.2 2-level Stream table

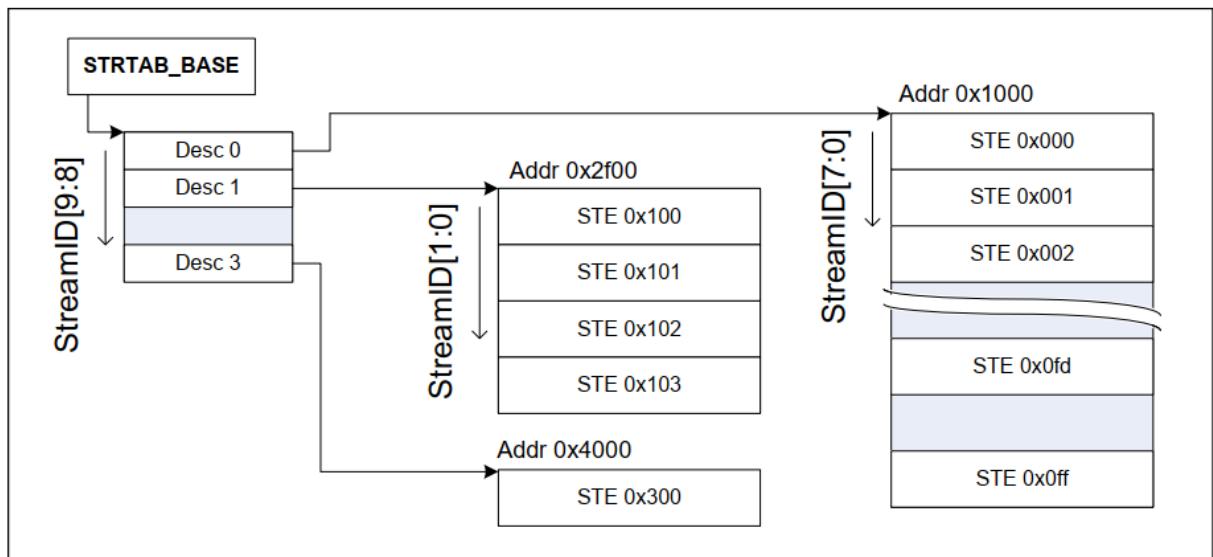
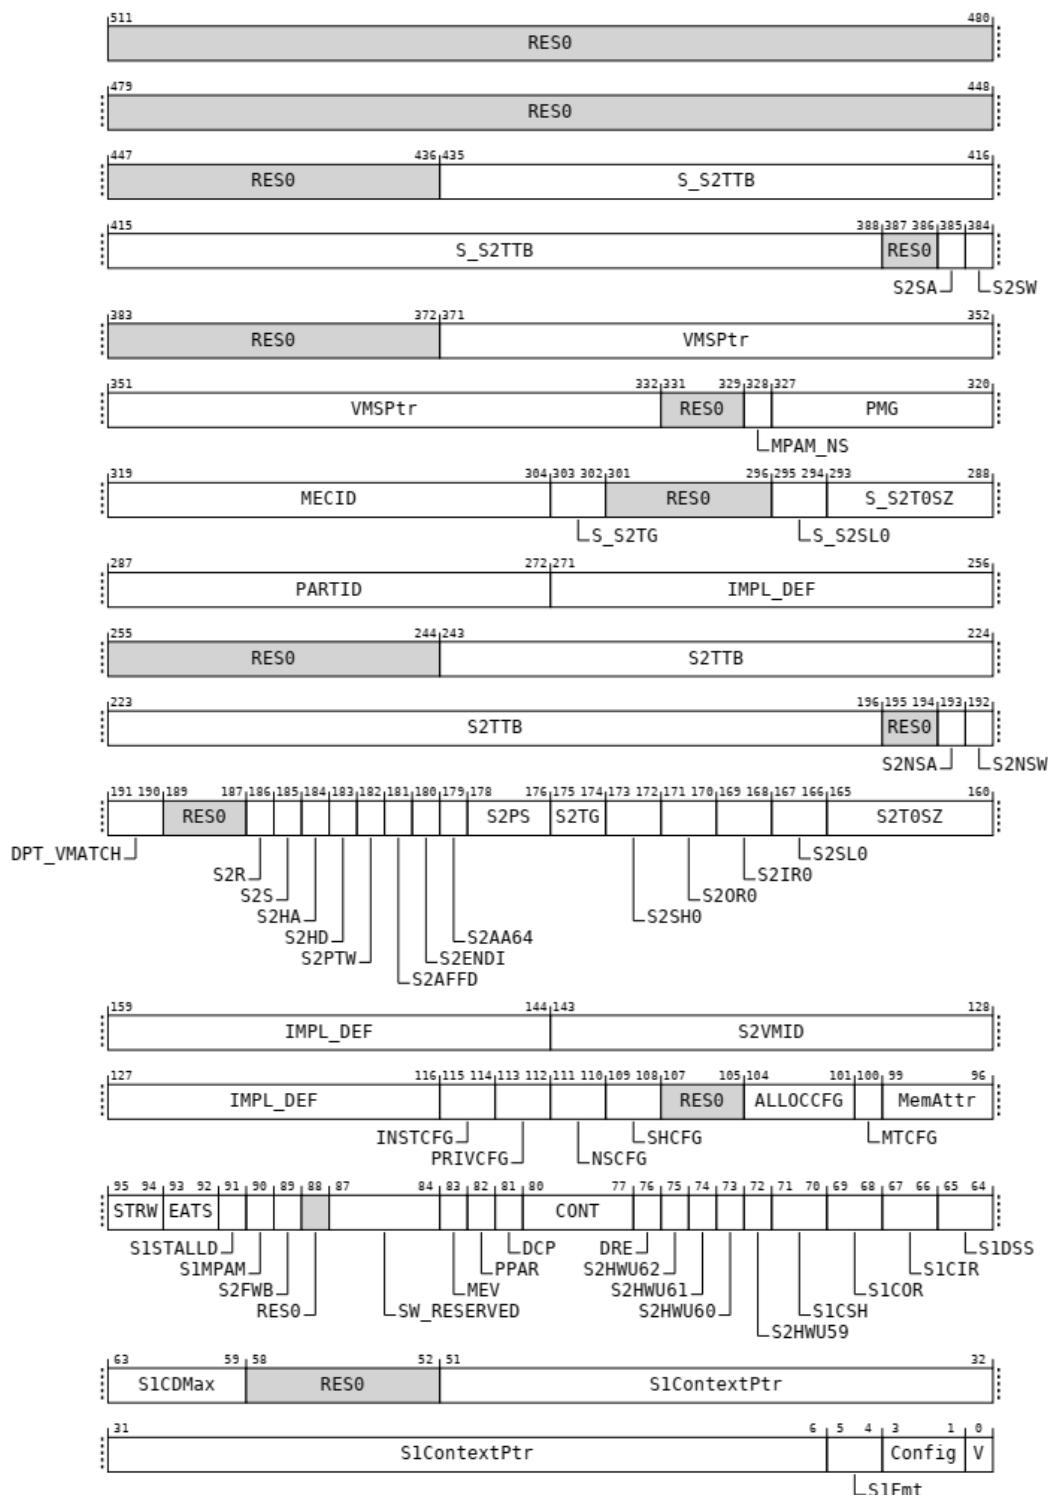


Figure 3.2: Example Two level Stream table with SPLIT == 8

STE 64B SMMU      STRTAB\_BASE      STE      STRTAB\_BASE      STE      Stream id      STE      linear  
 STRTAB\_BASE + sid \* 64      STE 64B      STE 2-level      sid      L1\_STD      STRTAB\_BASE + sid[9:8] \* 8,  
 L1\_STD 8B , L1\_STD      sid      STE l2ptr + sid[7:0] \* 64 .

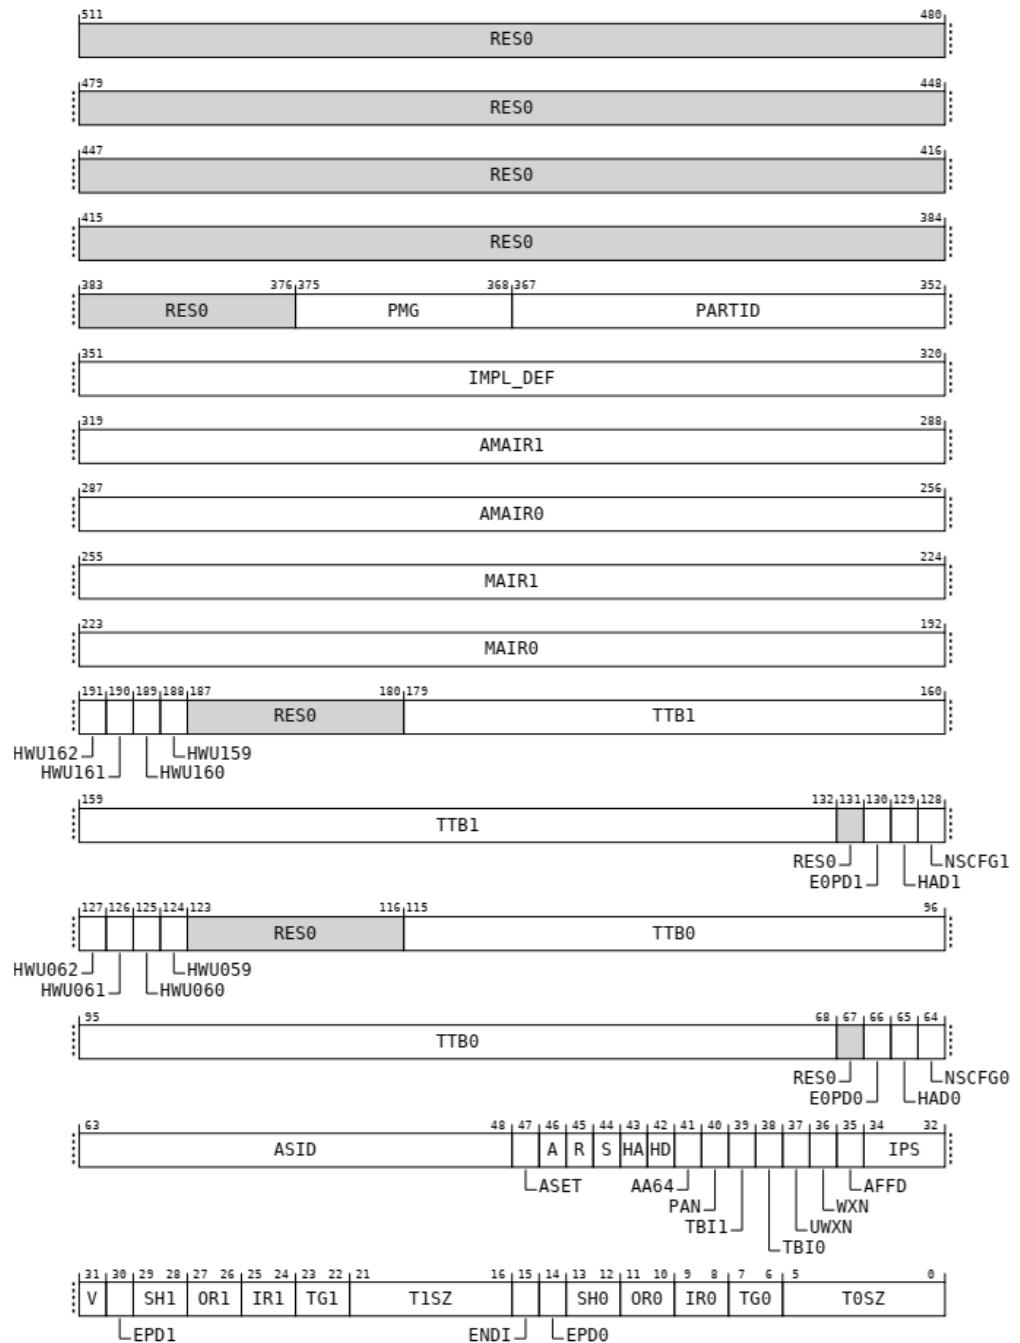
STE      stream bypass,      stage1, stage2      stage1 + stage2

## 5.2 Stream Table Entry

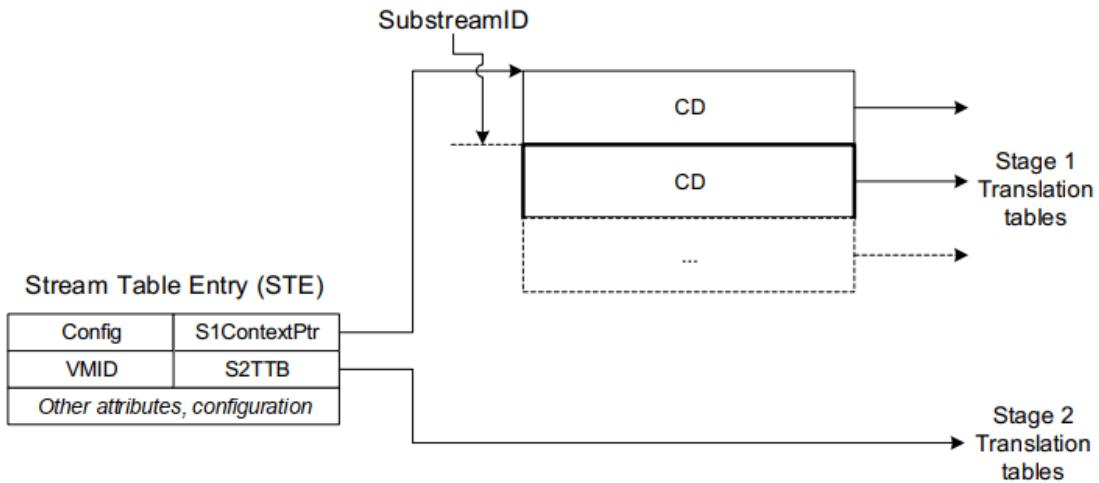


Chapter 5. Data structure formats

5.4. Context Descriptor



CD



**Figure 3.4: Multiple Context Descriptors for Substreams**

SubStreamID    CD    SubStreamID

SubStreamID

#### SMMU command queue event queue

Command Queue Event Queue SMMU 2 Queue      Command queue SMMU      command queue, SMMU command queue      Event Queue SMMU      SMMU      Event queue      Event queue

#### streamID substreamID

##### StreamID

StreamID is generated from the PCI RequesterID so that **StreamID[15:0] == RequesterID[15:0]**

##### SubstreamID

The SubstreamID is equivalent to a **PCIe PASID**. Because the concept can be applied to non-PCIe systems, it has been given a more generic name in the SMMU. The maximum size of SubstreamID, 20 bits, matches the maximum size of a PCIe PASID.

c2k StreamID N 24    StreamID[23:0]

StreamID[15:0] pcie RequesterID[15:0].

StreamID[23:16] Root Complex c2000 root complex id :

SubstreamID pcie PASID

#### 1.1.3      SVA

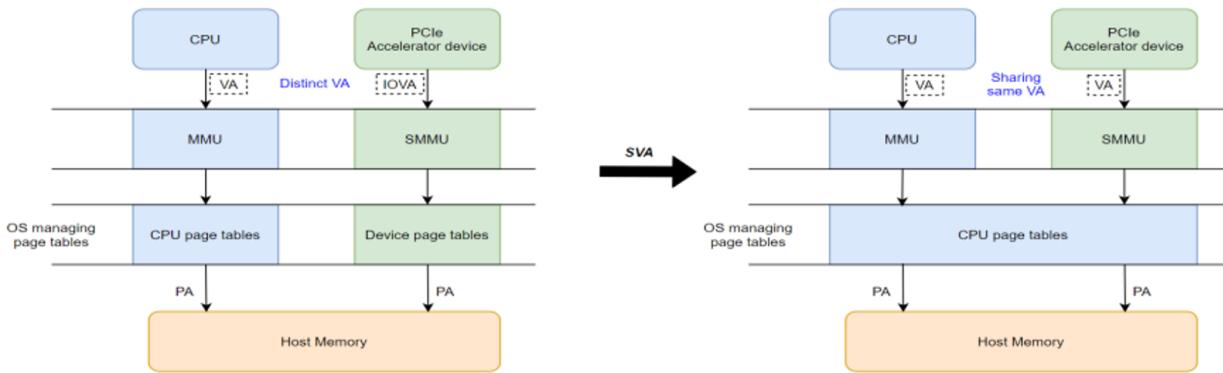
[toc]

##### SVA

SVA (Shared Virtual Addressing)  
Virtual Machine

SVA      SVM (Shared Virtual Memory)

SVA      Secure



## SVA

- CPU GPU/NPU/SmartNIC/FPGA/ASIC/Accelerator
- device cpu
- DMA cpu pin DMA page

## SVA

1. PASID
2. IO page fault PCIe PRI
3. MMU IOMMU

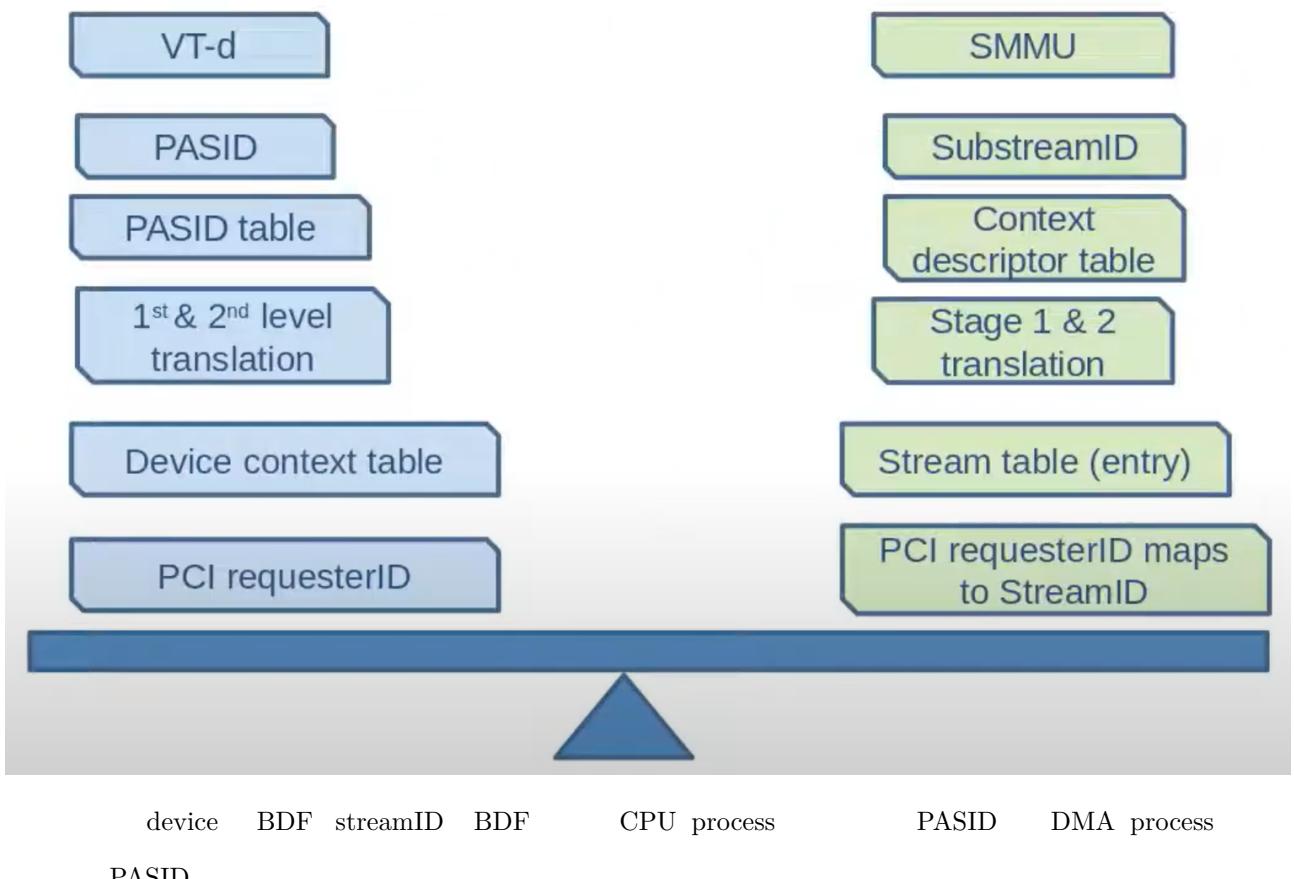
C2K PPC MMU ARM SMMU

ARM SMMU

PPC MMU

SVA MMU SMMU page table SMMU 64bit EA MMU 64bit EA SLB 78bit VA

SVA Intel ARM



## PASID

PASID (BDF)	PCIe spec Requester ID EP	PASID TLP prefix PASID ATS PRI	20 bit Requester ID PASID TLP prefix	PASID value Requester ID PASID	process PASID	Requester ID PASID prefix
----------------	---------------------------------	--------------------------------------	--	--------------------------------------	------------------	------------------------------

## 6.20.2 PASID TLP Layout

A TLP may contain at most one PASID TLP Prefix.

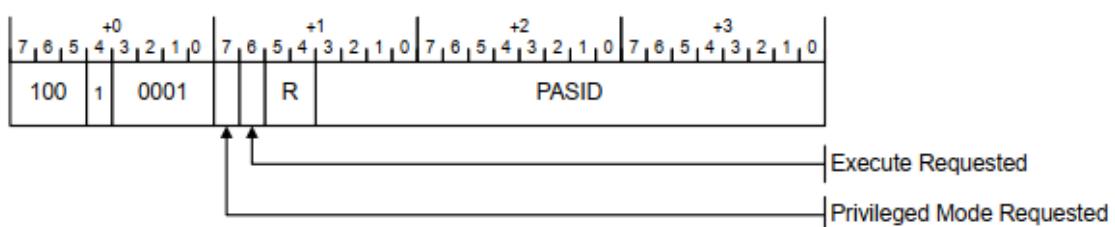
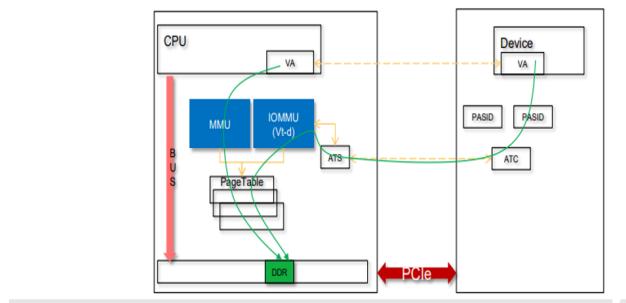


Figure 6-20 PASID TLP Prefix

## Intel IOMMU VT-d

- Intel: (Share Virtual Memory)
  - CPU access DDR through MMU
  - Device access DDR through IOMMU
  - MMU share the same page table with IOMMU
  - (used by AMD's Secure Virtual Machine in Linux)



DMA Remapping—Intel® Virtualization Technology for Directed I/O



Figure 3-5 illustrates the paging structure for translating a 48-bit address to a 4-KByte page.

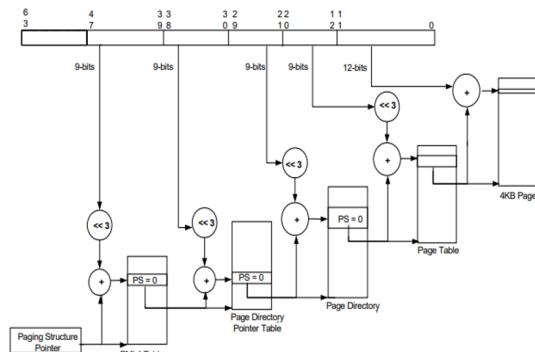
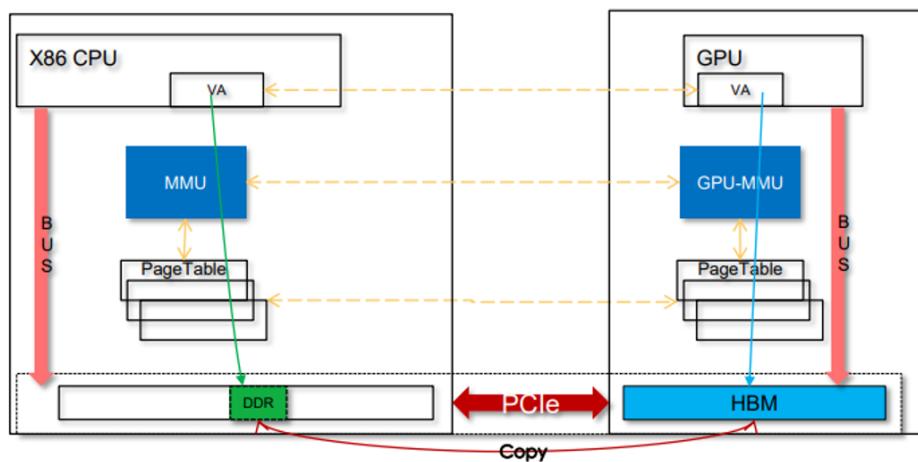


Figure 3-5. Address Translation to a 4-KByte Page

## Nvidia UVA

- Nvidia: (UVA: Unified Virtual addressing)
  - GPU has its own MMU
  - CPU can only access DDR while GPU can only access HBM
  - GPU MMU mirror the Page Table of CPU
  - When GPU access a VA not populated in HBM it will copy the data from DDR , and vice versa.



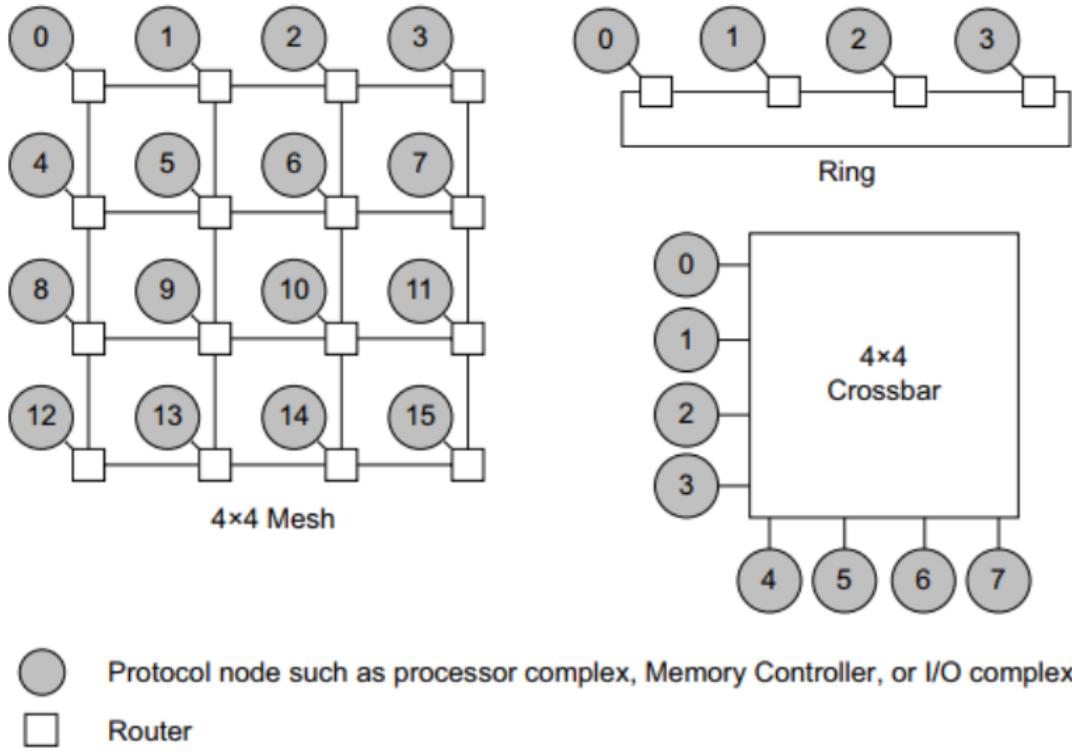
## 1.2 NoC

### 1.2.1 CHI

[toc]

#### CHI

SoC



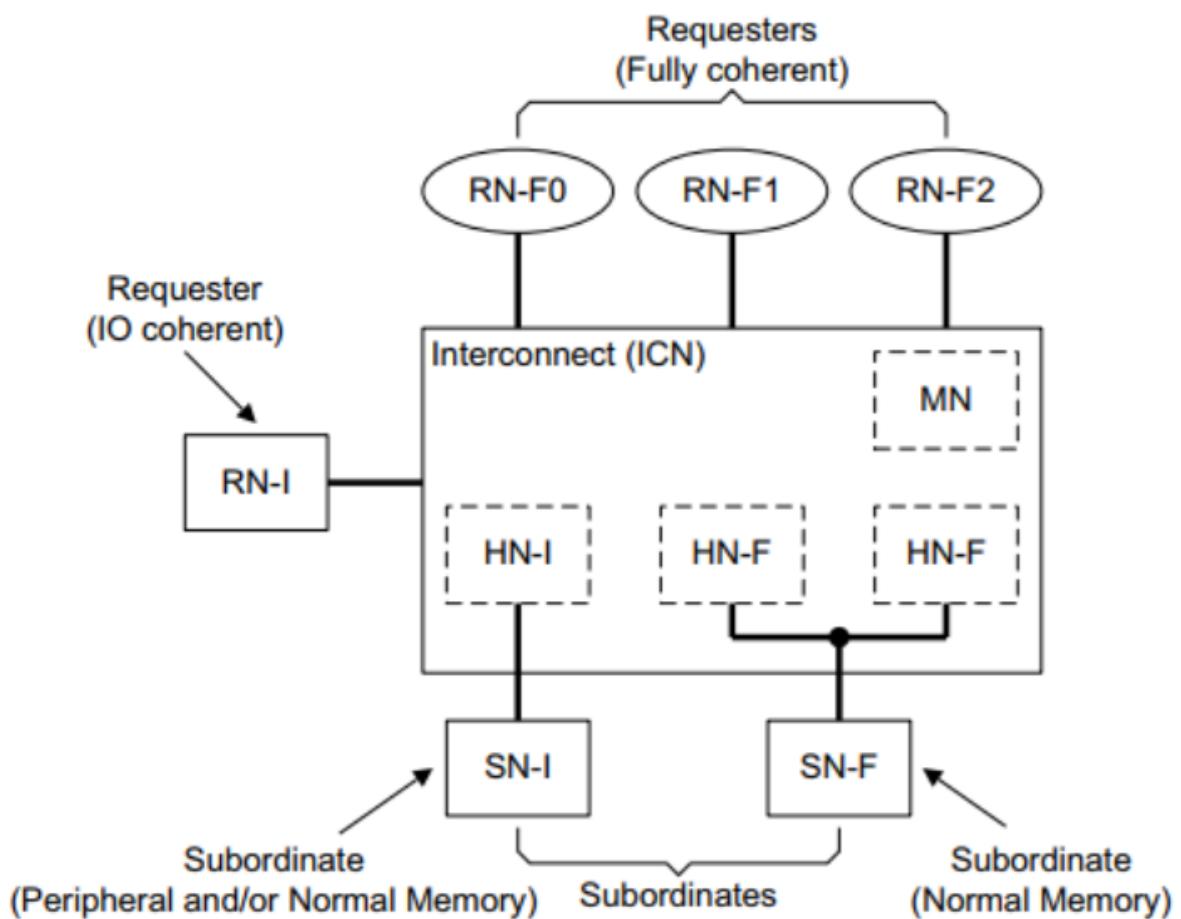
**Figure 1-1 Example interconnect topologies**

### Component

\*\*RN \*\*

\*\*HN \*\* cache CMN700 HNF SLC Combined PoS/PoC(ordering) Snoop Filter

\*\*SN \*\* HN



**Figure 1-4 Protocol node examples**

PCIe

**Table 1-1 Layers of the CHI architecture**

<b>Layer</b>	<b>Communication granularity</b>	<b>Primary function</b>
Protocol	Transaction	<p>The Protocol layer is the topmost layer in the CHI architecture. The function of the Protocol layer is to:</p> <ul style="list-style-type: none"> <li>• Generate and process requests and responses at the protocol nodes.</li> <li>• Define the permitted cache state transitions at the protocol nodes that include caches.</li> <li>• Define the transaction flows for each request type.</li> <li>• Manage the protocol level flow control.</li> </ul>
Network	Packet	<p>The function of the Network layer is to:</p> <ul style="list-style-type: none"> <li>• Packetize the protocol message.</li> <li>• Determine the source and target Node IDs required to route the packet over the interconnect to the required destination and add to the packet.</li> </ul>
Link	Flit	<p>The function of the Link layer is to:</p> <ul style="list-style-type: none"> <li>• Provide flow control between network devices.</li> <li>• Manage link channels to provide deadlock-free switching across the network.</li> </ul>

\*\*Protocol      \*\* transaction

\*\*Network    \*\* transaction    packet    SAM    flit

\*\*Link      \*\* Flit    4 6 channel

## Protocol

**Transaction**    CHI transaction    read write snoop atomic snoop

Type	Transaction Type	Description
Read	<b>ReadNoSnoop</b>	
Read	<b>ReadOnce</b>	
Read	<b>ReadShared</b>	
Read	<b>ReadClean</b>	
Read	<b>ReadNotSharedDirty</b>	
Read	<b>ReadUnique</b>	
Snoop	<b>CleanUnique</b>	
Snoop	<b>CleanShared</b>	
Snoop	<b>MakeUnique</b>	
Snoop	<b>Evict</b>	
Snoop	<b>StashOnce</b>	
Snoop	<b>StashUnique</b>	
Snoop	<b>DVM Operations</b>	DVM    TLB
Write	<b>WriteNoSnoop</b>	
Write	<b>WriteUnique</b>	
Write	<b>WriteLineUnique</b>	
Write	<b>WriteBack</b>	
Write	<b>WriteClean</b>	

Type	Transaction Type	Description
Write	<b>WriteEvictFull</b>	
Atomic	<b>Atomic Operations</b>	- -

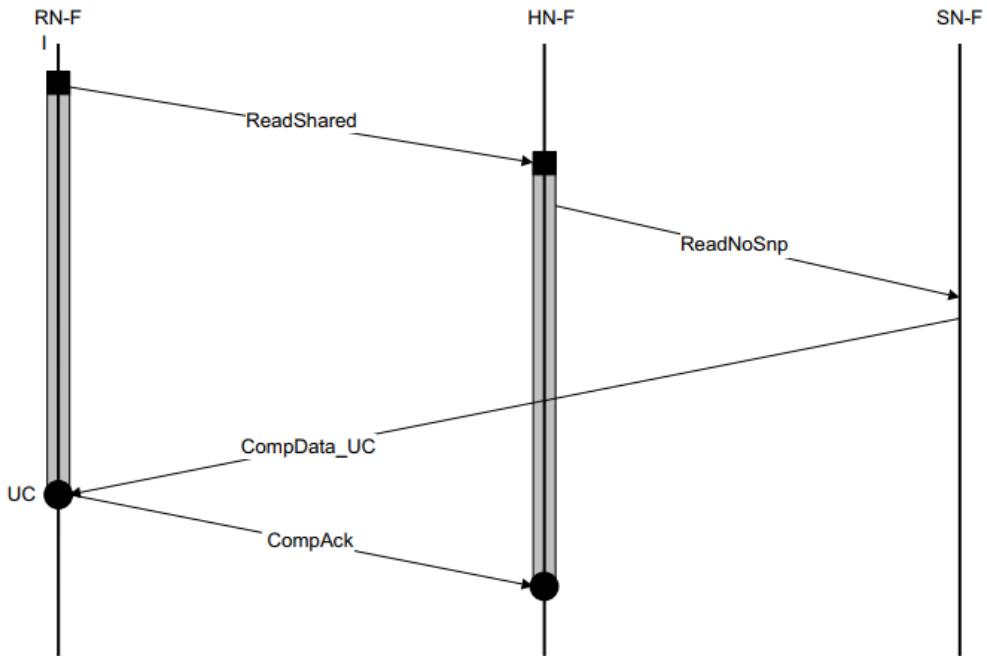
transaction flow

### 5.1.1 Read transactions with DMT and without snoops

For Read transactions without snoops, this specification recommends the use of *Direct Memory Transfer* (DMT).

[Figure 5-1](#) shows an example DMT transaction flow using the ReadShared transaction.

In [Figure 5-1](#), a response from SN-F to HN-F is not required because CompAck from the Requester is used to deallocate the request at Home.



**Figure 5-1 DMT Read transaction example without snoops**

1. RN-F sends a Read request to HN-F.
2. HN-F sends a Read request to SN-F
- The ID field values in the Read request are based on where the data response is to be sent. Data can be sent to the Requester or to the HN-F.
3. SN-F sends a data response directly to RN-F.
4. RN-F sends CompAck to HN-F as the request is ReadShared and requires CompAck to complete the transaction.

## Network

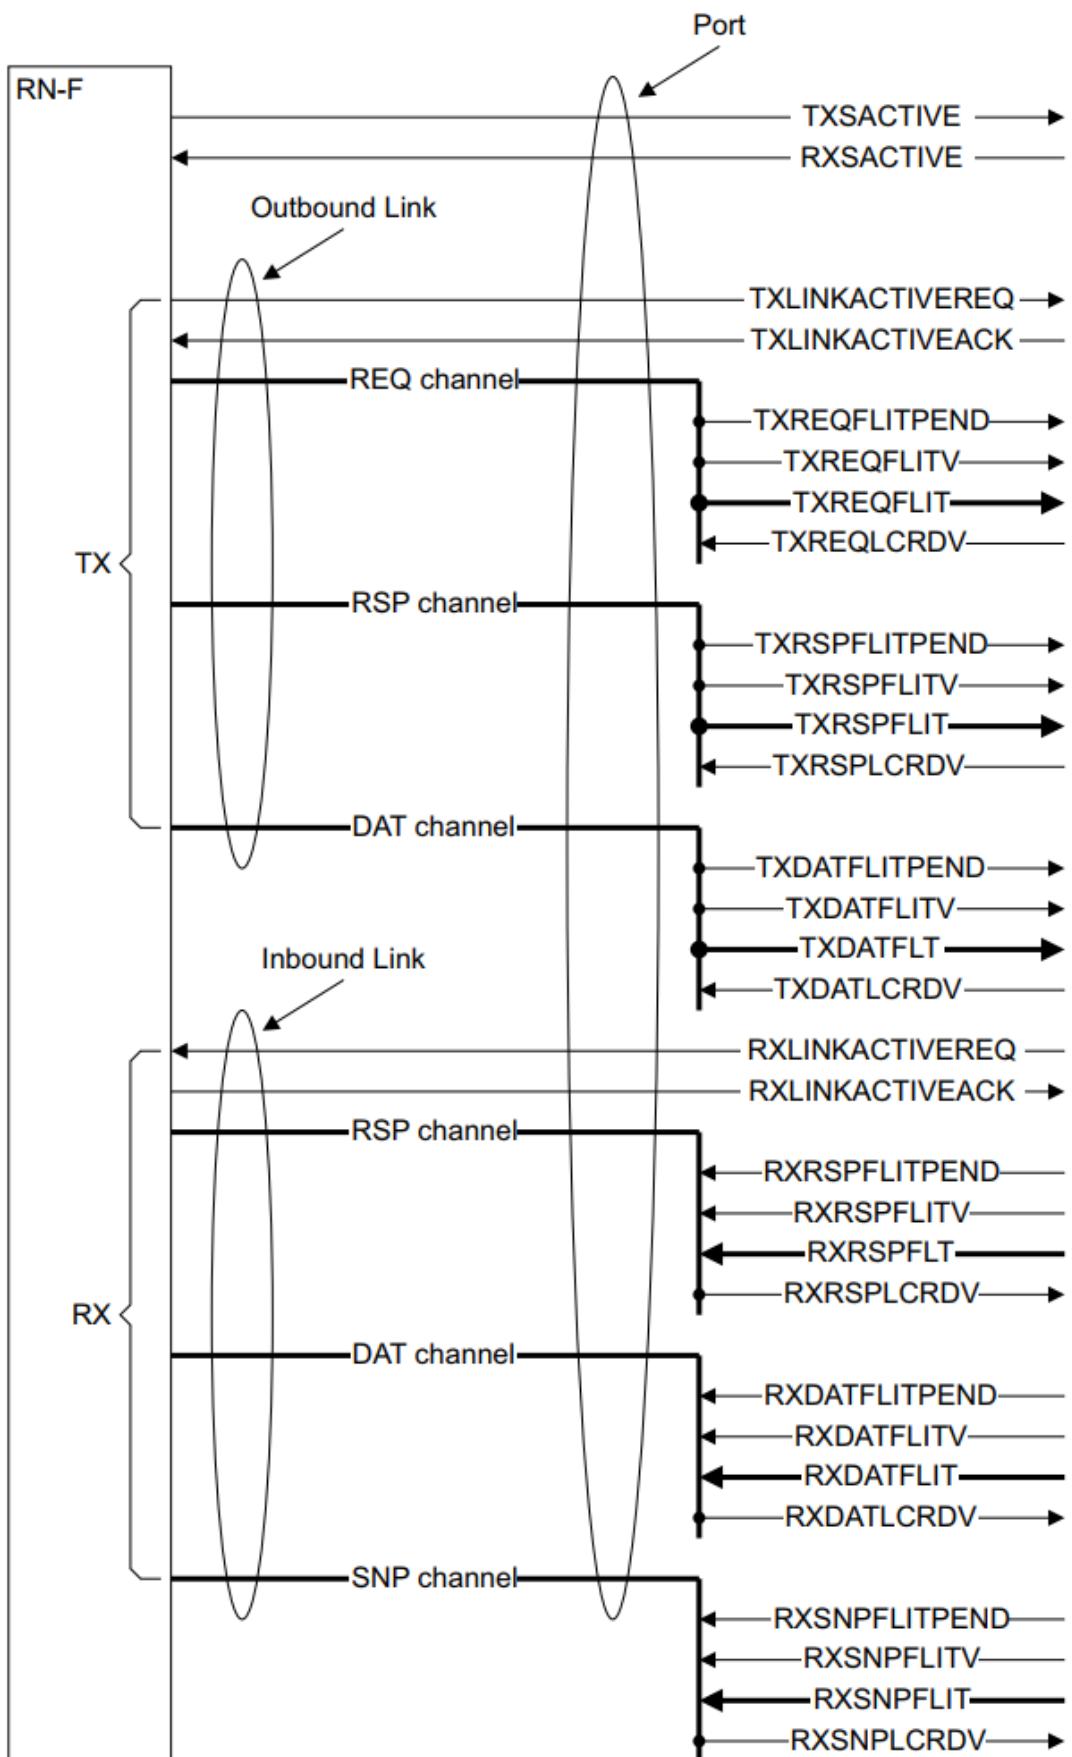
transaction SAM packet srcid tgtid  
SAM RNSAM, HNSAM  
power of 2 hashing hierachical hashing

**Link**

4 6 channel

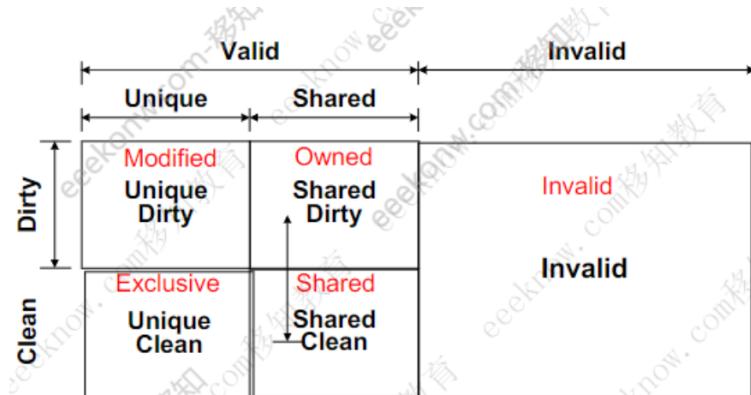
**TX: RSP DATA REQ**

**RX: RSP DATA SNP**



**Figure 13-4 Relationship between links, channels, and port**

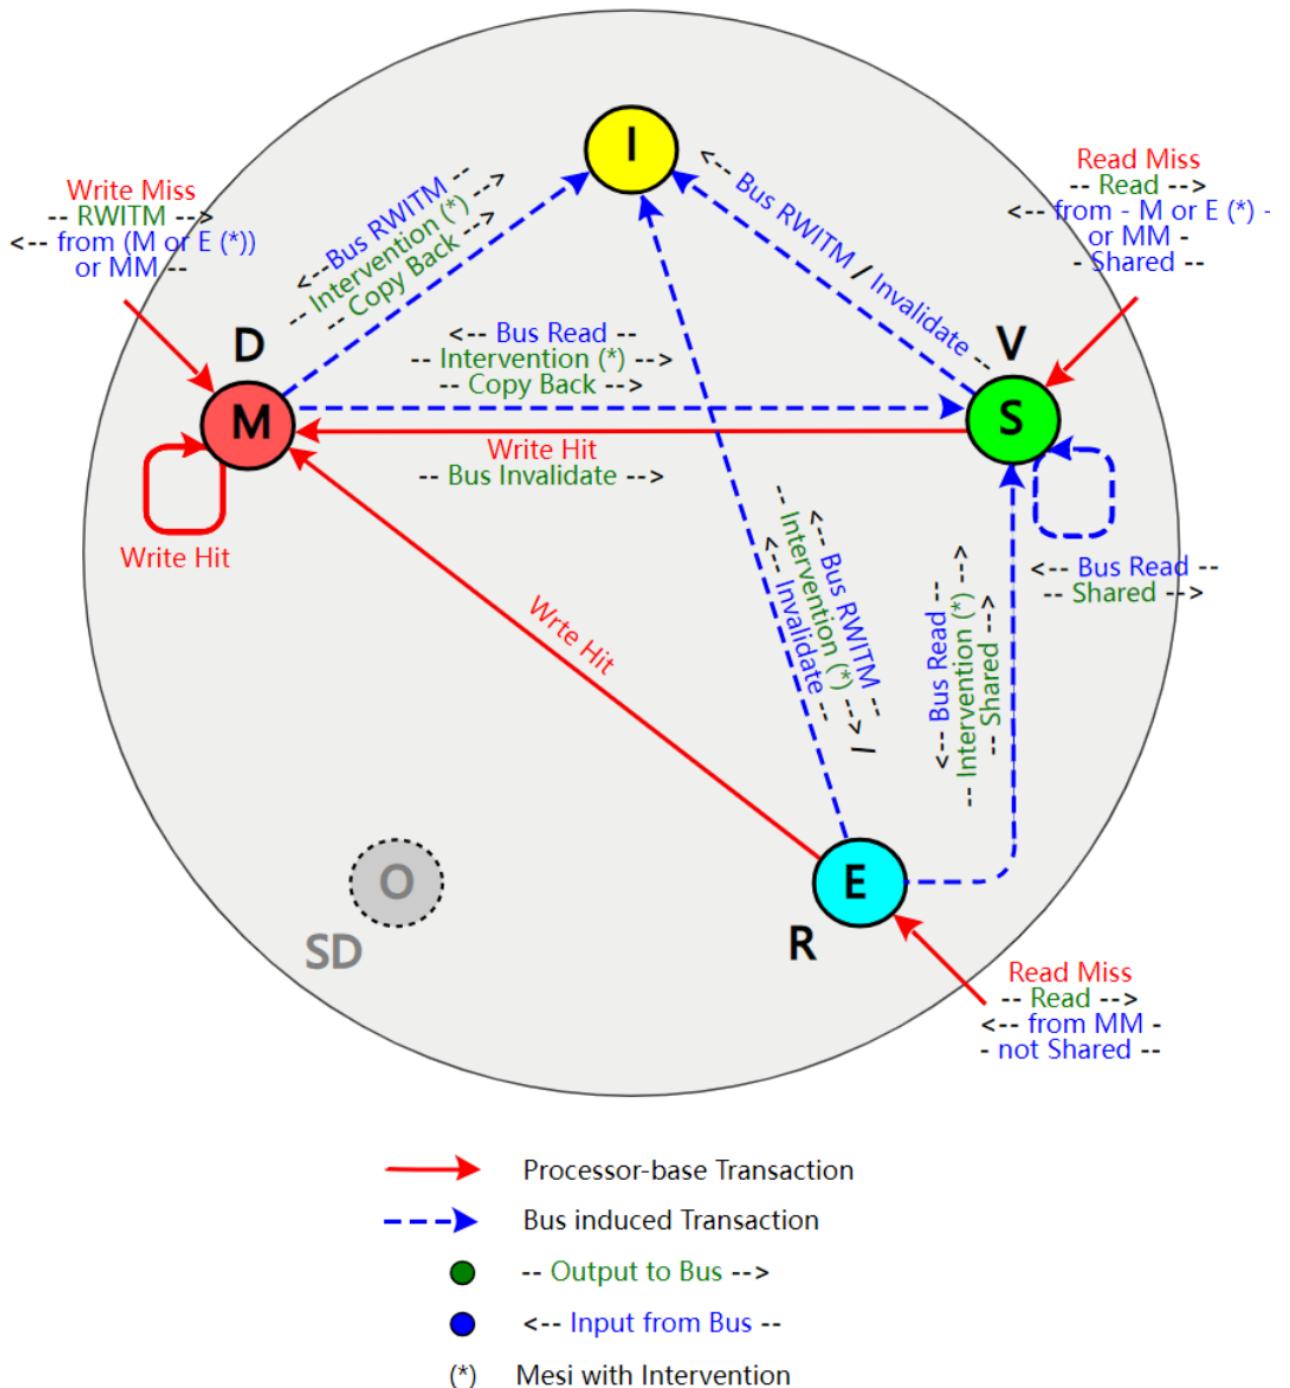
cache cache



Cache line states

- Unique** – The cache line is only in this cache
- Shared** – The cache line may be in multiple cache
- Dirty** – The cache line has been modified respect to main memory and this cache must be written back to main memory eventually updated
- Clean** – The cache line has the same value as main memory or the value of shared dirty data written to main memory
- Invalid** – The cache line is not valid

## MESI Protocol State Transaction Diagram

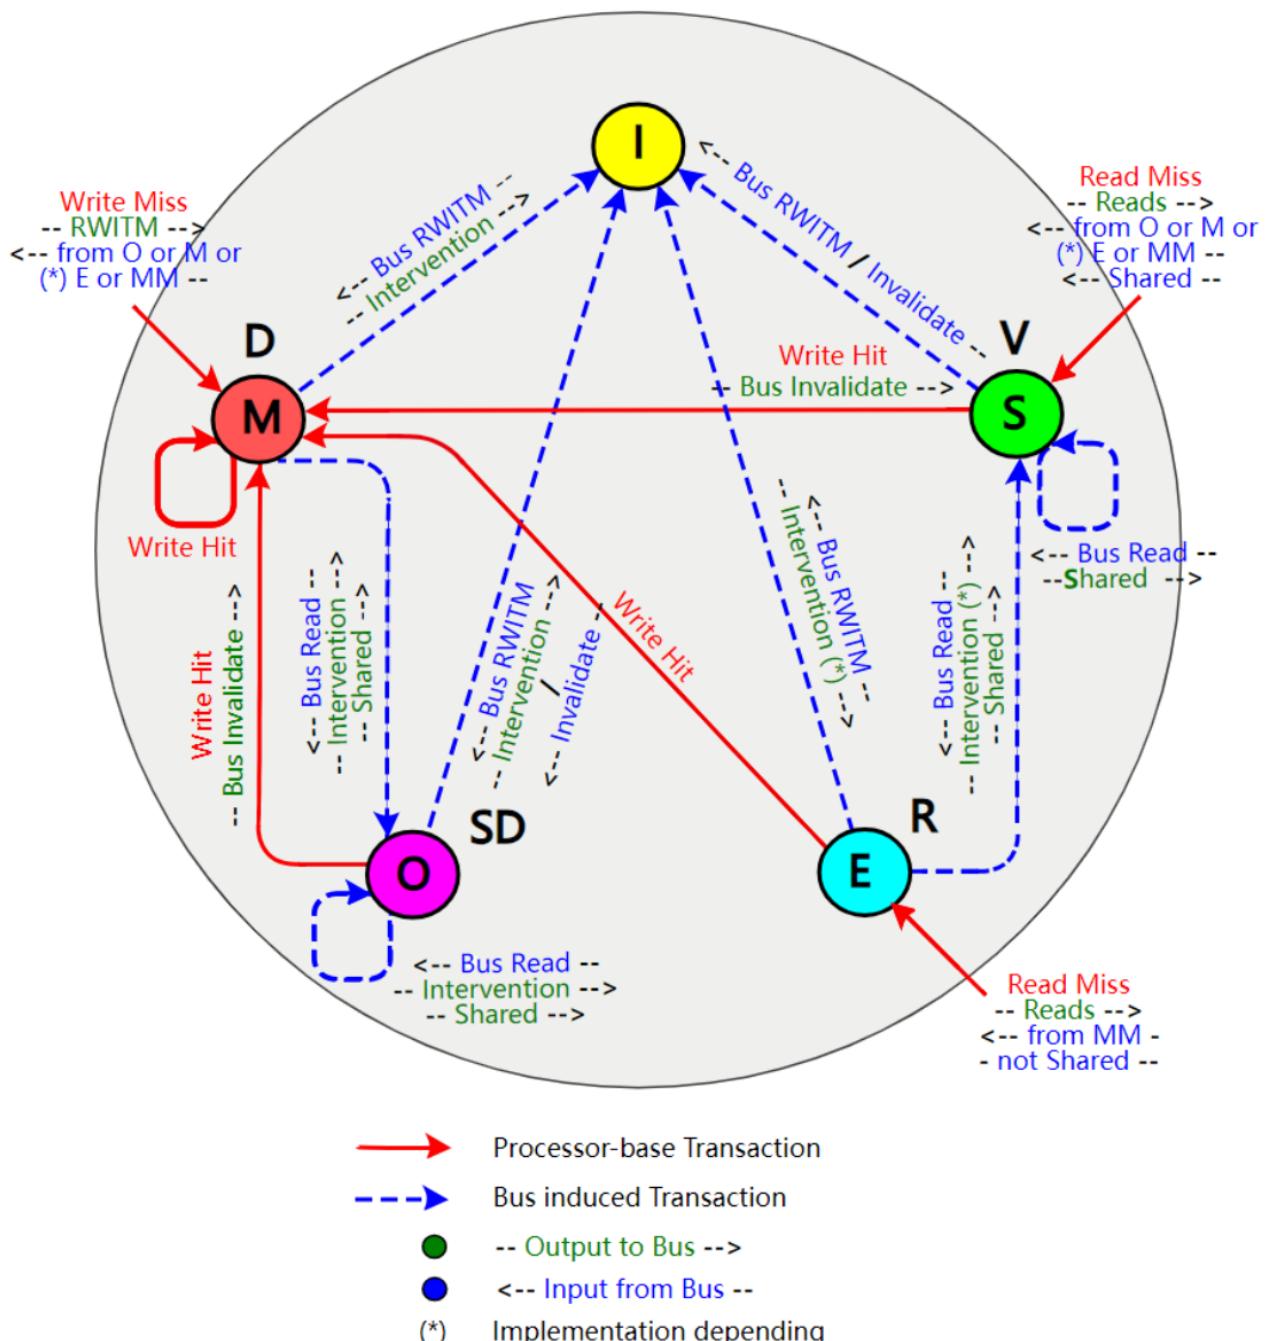


MESI

By Ferruccio Zulian - Milan, Italy

MOESI	MESI	MESI	M	CPU	CPU	CPU	MOESI
Owned O	M		M	O	MOESI	O	S
MESI	S	MOESI	O	O		O	MOESI
S		CPU	O	CPU		O	S

## MOESI Protocol State Transaction Diagram



By Ferruccio Zulian - Milan, Italy

	address	data
DDR	0x100	0xa
Cache 1	0x100	0xa
Cache 2	0x100	0xa

core 2 cache

	address	data
DDR	0x100	0xa
Cache 1	0x100	0xa
Cache 2	0x100	<b>0xa --&gt; 0xb</b>

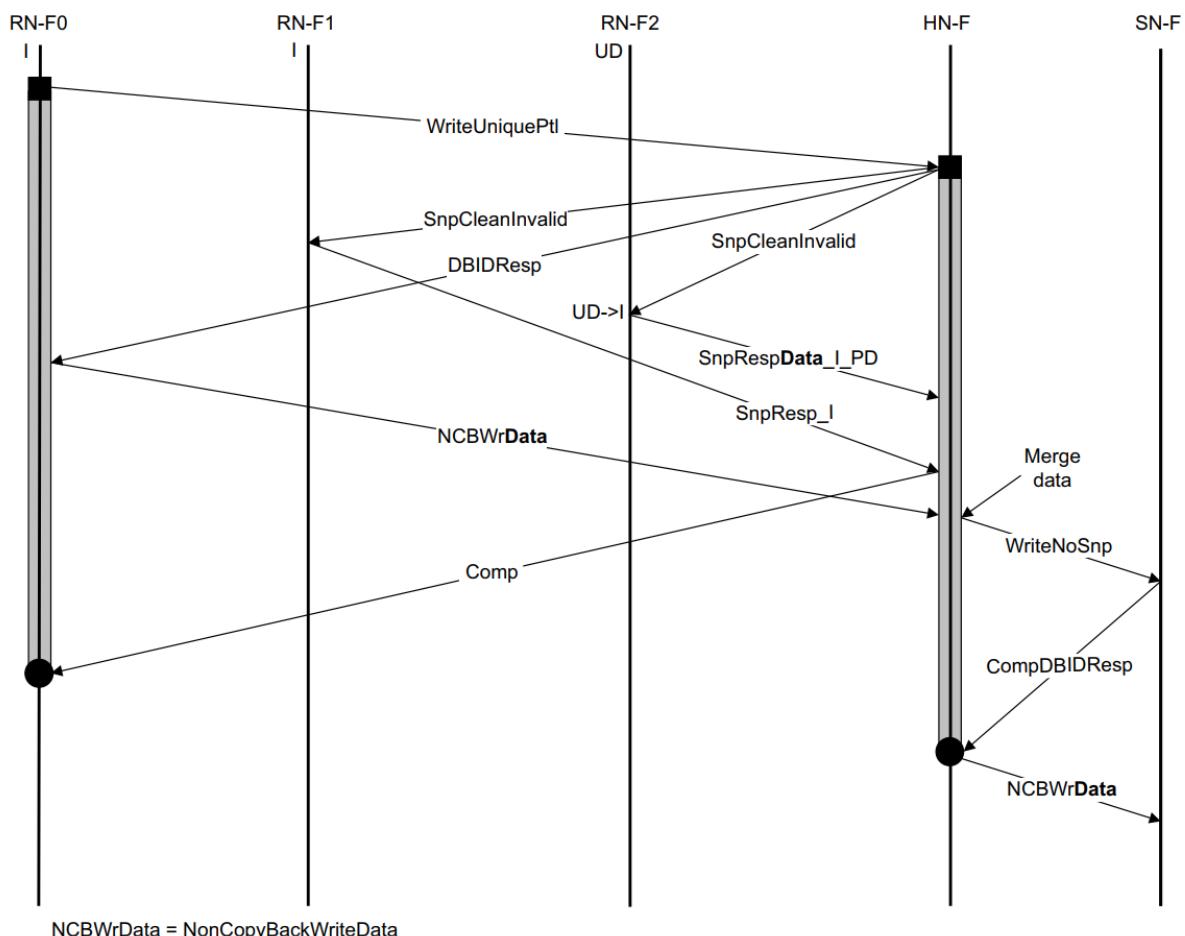
HN directory snoop cache ddr

### 5.3.2 Write transaction with snoop and separate responses

An example of this type of flow is a WriteUniquePtl transaction.

The Comp\_I response from HN-F must be sent when the coherency activity is complete at HN-F.

Figure 5-16 shows the transaction flow. The copy of data being transferred is marked in bold.



CHI -  
based

snooping directory-based AMBA CHI Coherent Hub Interface

directory-

- 
- 
- /

- 
- 
- CHI
- 
- 
- /
- 

- 
- 
- 
- 

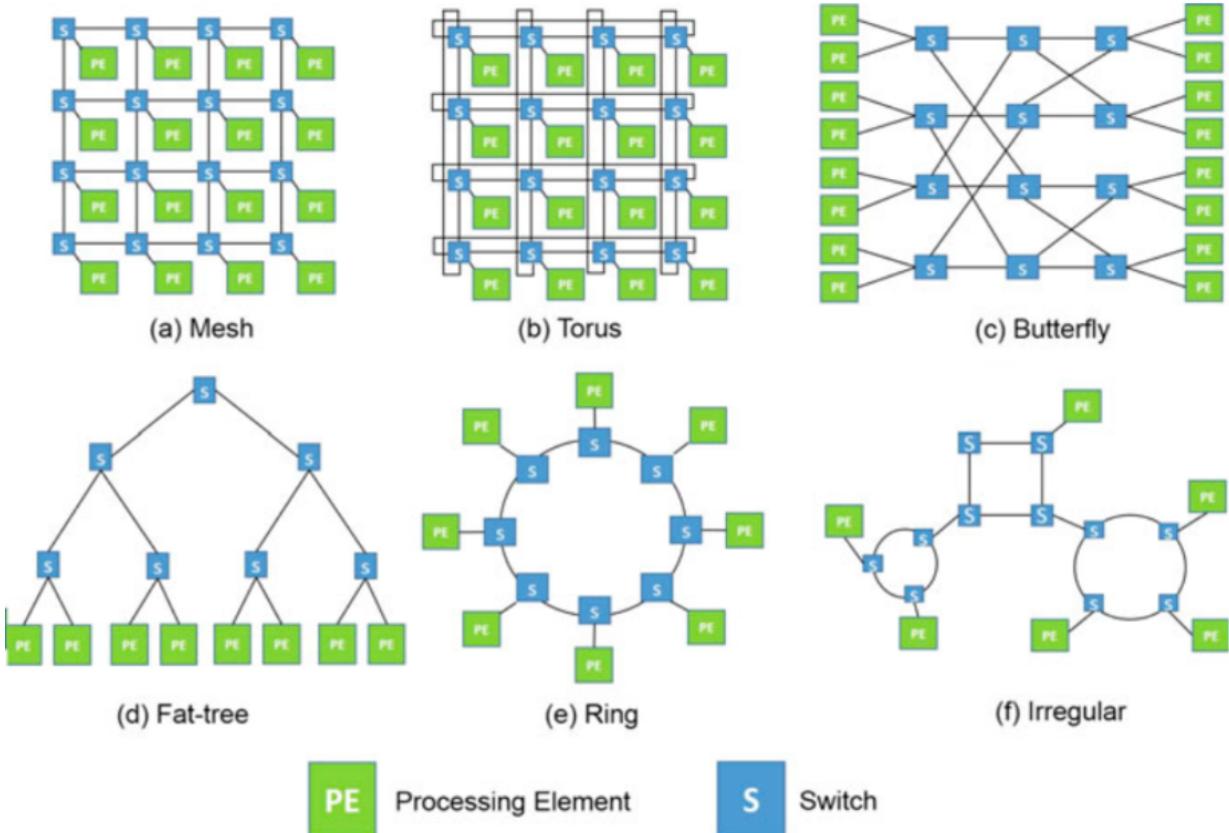
- DVM

### 1.2.2 NoC

[toc]

CPU      On-Chip Interconnect

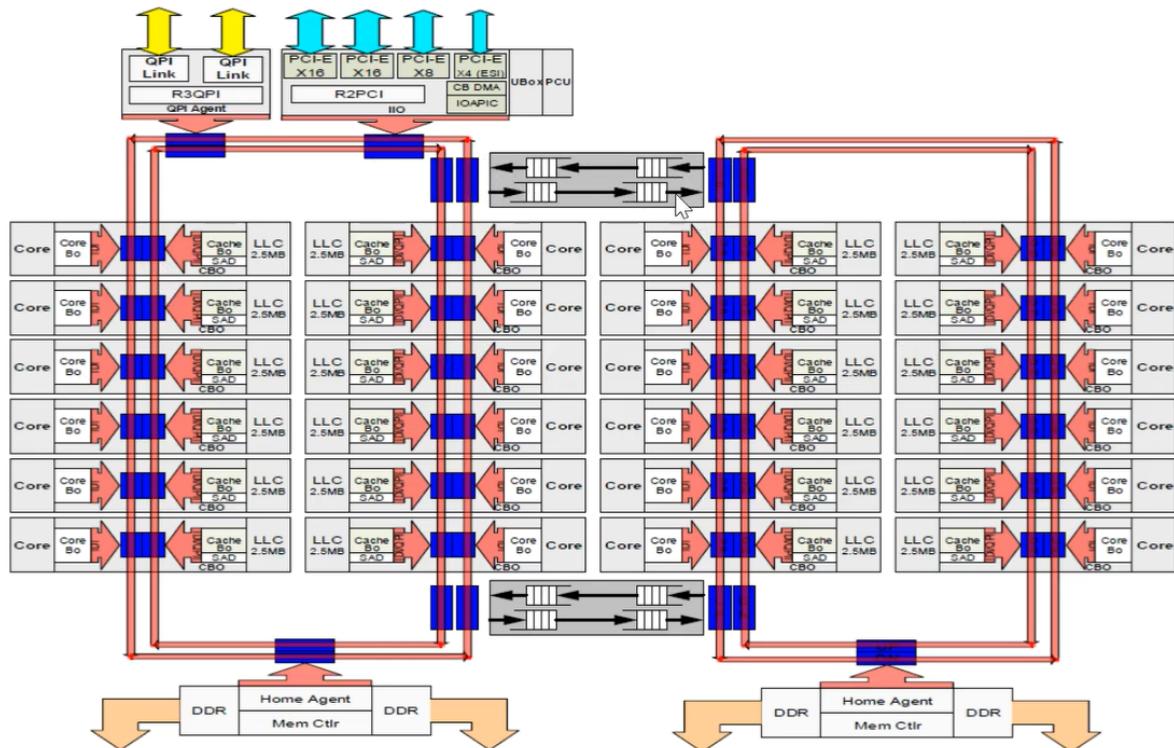
1. Shared Bus a. b. c.
2. Crossbar Switch a. b. c.
3. Ring Interconnect a. b. c.
4. Mesh    Torus    Mesh and Torus Interconnects a. b. c.
5. Network-on-Chip, NoC a. b. c.



1.

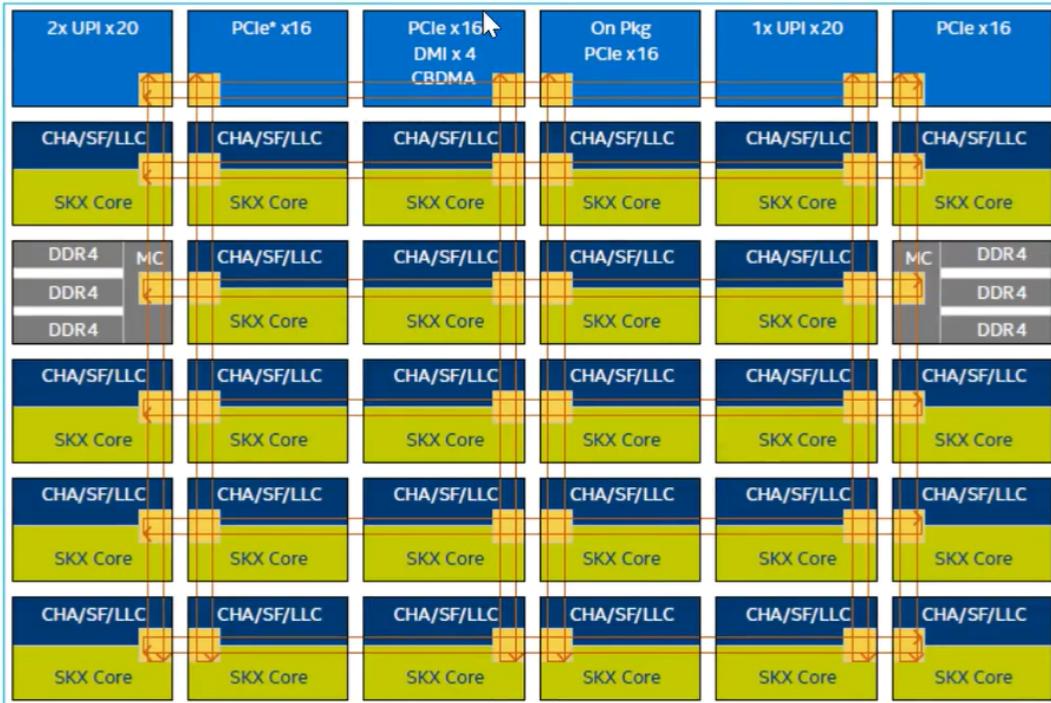
Intel Broadwell      ring      core      core      FIFO

## Broadwell EX 24-core die



Intel Skylake      Mesh

# Skylake-SP 28-core die



CHA – Caching and Home Agent ; SF – Snoop Filter; LLC – Last Level Cache ;  
SKX Core – Skylake Server Core ; UPI – Intel® UltraPath Interconnect

1. —
2. Efficient Barrier Implementation on the POWER8 Processor
- 3.

### 1.2.3 CMN700

[toc]

#### CMN-700 HNF/HNI/HNP

**Home Node**      HN Home Node      HN      SAM      HNF    HN

**HNF**      DRAM      HNF

HNF

1. **System Level Cache** The System Level Cache (SLC) is a last-level cache. The SLC allocation policy is exclusive for data lines, except where sharing patterns are detected and pseudo-inclusive for code lines, as indicated by the RN-Fs. All code lines can be allocated into the SLC on the initial request. When MTE is enabled, SLC stores data and tags.
2. **Combined PoS/PoC** The combined Point-of-Serialization/Point-of-Coherency (PoS/PoC) is responsible for the ordering of all memory requests sent to the HN-F. Ordering includes serialization of multiple outstanding requests and actions to the same line, and request ordering as required by the RN-F.
3. **Snoop Filter** The Snoop Filter (SF) tracks cachelines that are present in the RN-Fs. It reduces snoop traffic in the system by favoring directed snoops over snoop broadcasts when possible. This approach substantially reduces the snoop response traffic that might otherwise be required.

snoop		Snoop Filter
Ans	HN	invalid

**HNI** The I/O coherent Home Node (HN-I) is a Home Node for all CHI transactions targeting AMBA subordinate devices

The HN-I acts as a proxy for all the RNs of CMN-700, converting CHI transactions to ACE5-Lite transactions. The HN-I includes support for the correct ordering of Arm device types.

The HN-I does not support caching of any data read from or written to the downstream ACE5-Lite I/O subordinate subsystem. Any cacheable request that is sent to the HN-I does not result in any snoops being sent to RN-Fs in the system. Instead, the request is converted to the appropriate ACE5-Lite read or write command and sent to the downstream ACE5-Lite subsystem.

If an RN-F caches data that is read from or written to the downstream ACE5-Lite I/O subordinate subsystem, coherency is not maintained. Any subsequent access to that data reads from or writes to the ACE5-Lite I/O subordinate subsystem directly, ignoring the cached data

There are HN-I types with extra functionality. These types include:

- HN-T HN-I that has a debug trace controller and DVM Node. CMN-700 can have zero or more HN-I and HN-T instances.
- HN-D HN-I that has a debug trace controller, DVM Node, and configuration subordinate. CMN-700 must have exactly one HN-D instance.
- HN-P HN-I that is optimized for peer-to-peer PCIe traffic.
- HN-V A device that includes an HN-I and DVM Node (DN).

**HNP** The I/O coherent Home Node with PCIe optimization (HN-P) is a device that includes the HN-I functionality and dedicated trackers for PCIe peer-to-peer traffic. HN-P can only be used to connect to PCIe subordinates.

HNF	HNI	HNP
System Level Cache Combined PoS/PoC Snoop Filter	CHI    ACE5-Lite	HNI    PCIe P2P
Yes		PCIe
	ACE5-Lite	PCIe
NA	AMBA	PCIe
		PCIe

## 1.2.4 CMN700

### 1.2.5

[toc]

ARMv7	ARMv7	0x00000000	4	ARMv8	EL3	EL2	EL1
ARMv8	64	32	64		128	32	

**ARM M7 Arm v7 32**

地址	异常	进入模式
0x00000000	复位	管理模式
0x00000004	未定义指令	未定义模式
0x00000008	软件中断	管理模式
0x0000000C	中止(预取)	中止模式
0x00000010	中止(数据)	中止模式
0x00000014	保留	保留
0x00000018	中断 IRQ	中断模式
0x0000001C	快中断 FIQ	快中断模式

Address      Exception Type

0x00000000	Reset
0x00000004	Undefined Instruction
0x00000008	Software Interrupt (SWI)
0x0000000C	Prefetch Abort
0x00000010	Data Abort
0x00000014	Reserved
0x00000018	IRQ
0x0000001C	FIQ

Interrupt      I/O

- 1.
- 2.                  IRQ
- 3.
- 4.                  NVIC GIC
- 5.

- 
- 
- 

**Exception**

- 1.
- 2.
- 3.
- 4.

5.

- •  
•

---

Interrupt	Exception
-----------	-----------

---

ARM

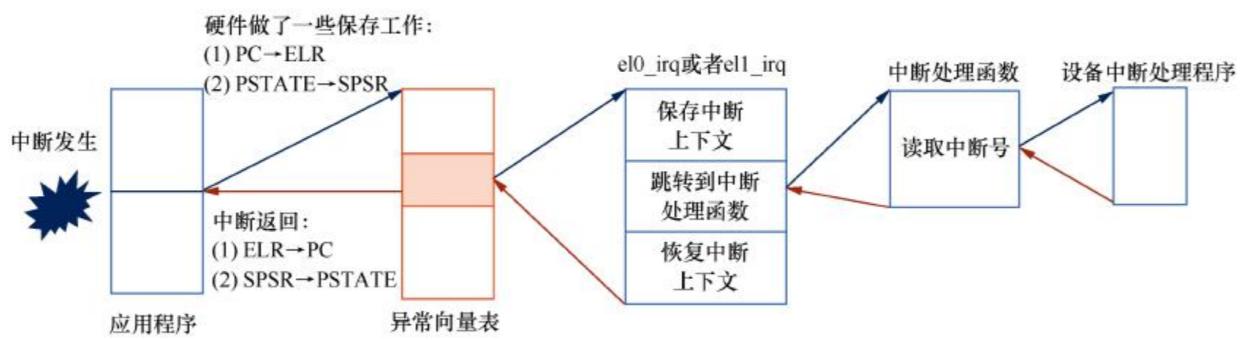
ARM64 —nIRQ nFIQ

ARM64      ARM      IRQ Interrupt Request   FIQ Fast Interrupt Request

PSTATE CPU I IRQ F FIQ

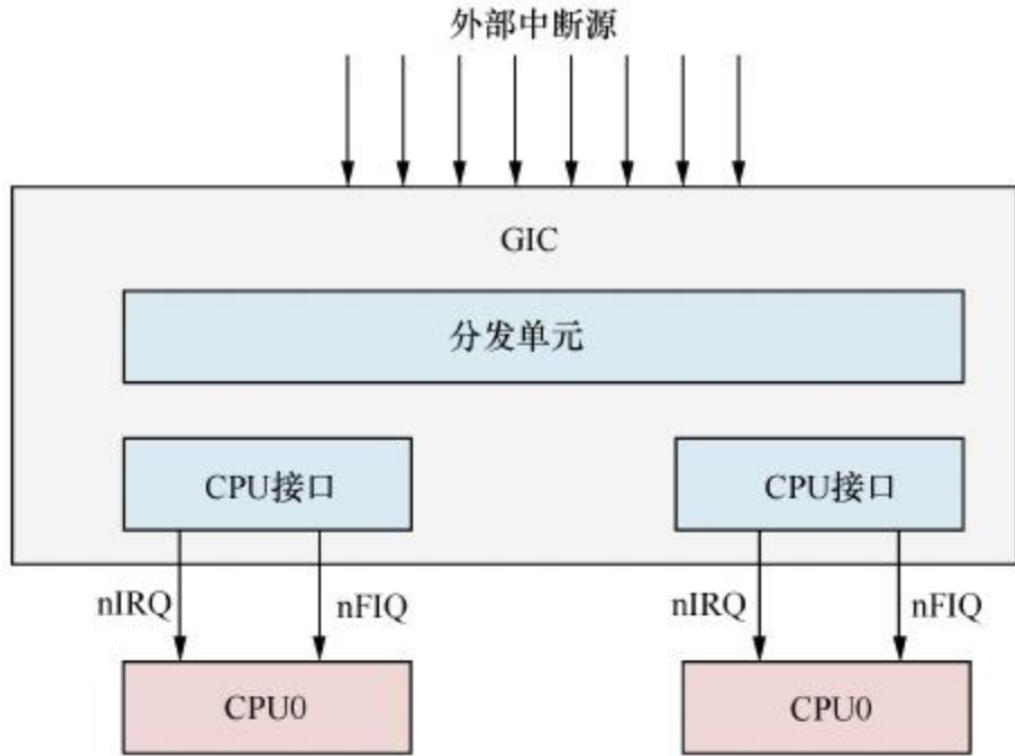
表 12.1 中断信号线

信 号	类 型	说 明
nIRQ	输入	IRQ 信号，每个 CPU 内核都有一根 nIRQ 信号线。它是一个低电平有效的信号线。低电平表示要激活这个 IRQ；高电平表示不要激活这个 IRQ。nIRQ 一直保持高电平直到触发 IRQ
nFIQ	输入	FIQ 信号，每个 CPU 内核都有一根 nFIQ 信号线。它是一个低电平有效的信号线。低电平表示要激活这个 FIQ；高电平表示不要激活这个 FIQ。nFIQ 一直保持高电平直到触发 FIQ



▲图 12.2 中断处理过程

NVIC GIC



▲图 12.1 中断控制器

## **NVIC Nested Vectored Interrupt Controller**

1.
    - 
    -
  2.
    - NVIC Cortex-M
    -
  3.
    - NVIC
  4.
    - VTOR
    - 240 Cortex-M

1.
    - GIC
    - Distributor CPU CPU Interface
  2.
    - 
    -
  3.
    - 1020
    - GICv3 GICv4
  - 4.

- GIC
- SoC
- 

- NVIC
  - Cortex-M
  - 
  -
- GIC
  - Cortex-A
  - 
  -

### 1.2.6 ARM GIC

[toc]

#### GIC

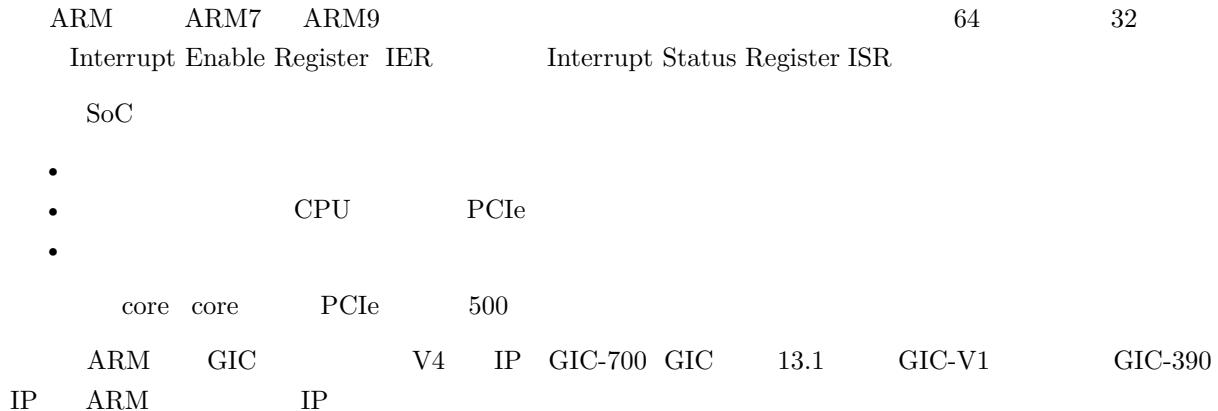
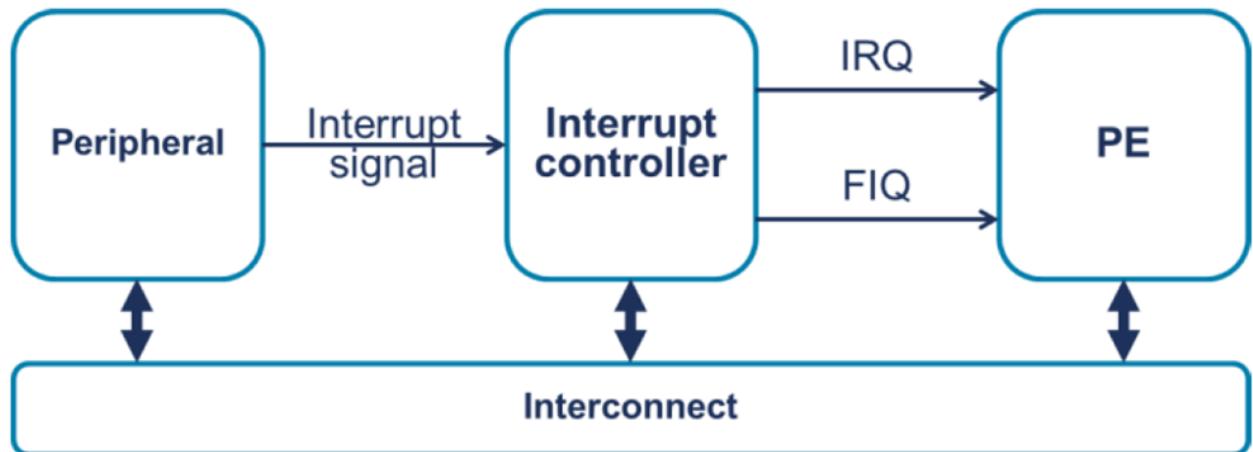


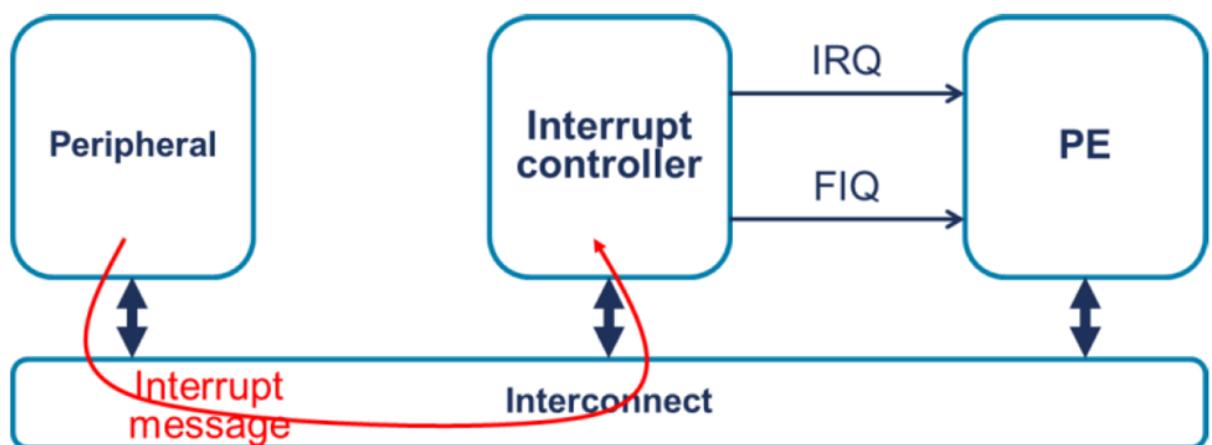
表 13.1 GIC 的发展史

版本	GIC-V1	GIC-V2	GIC-V3	GIC-V4
新增功能	<input type="checkbox"/> 支持 8 核, <input type="checkbox"/> 支持多达 1020 个中断源, <input type="checkbox"/> 支持用 8 位二进制数表示的优先级, <input type="checkbox"/> 支持软件触发中断, <input type="checkbox"/> 支持 TrustZone	<input type="checkbox"/> 支持虚拟化 <input type="checkbox"/> 改进对安全软件的支持	<input type="checkbox"/> 支持的 CPU 核数大于 8 <input type="checkbox"/> 支持基于消息的中断 <input type="checkbox"/> 支持更多的中断 ID	支持注入虚拟中断
IP 核心	GIC-390	GIC-400	GIC-500、GIC-600	GIC-700
应用场景	Cortex-A9 MPCore	Cortex-A7/A9 MPCore	Cortex-A76 MPCore	

C2K GIC V4.1.



MSI



**GIC-V3** GIC-V3

GIC-V3 CPU 8

PCIe MSIX

GIC-V3 IP GIC-500 GIC-600

PCIe PCIe

**GIC-V4** GICv4 GICv3

#### 1. vPE

GICv4 vPE

- GICv4 vPE Hypervisor

2.

GICv4 Direct Injection

- GICv4

3.

GICv4

- **vPE** GICv4 vPE

4.

GICv4

- **vLPI Control** GICv4

5.

GICv4

- GICv4

6.

GICv4

- GICv4

4	inactive	pending	CPU	active	CPU	active and
pending	CPU					
		edge-triggered		level-triggered		

GIC

GIC

GIC

表 13.2 GIC 分配的中断号范围

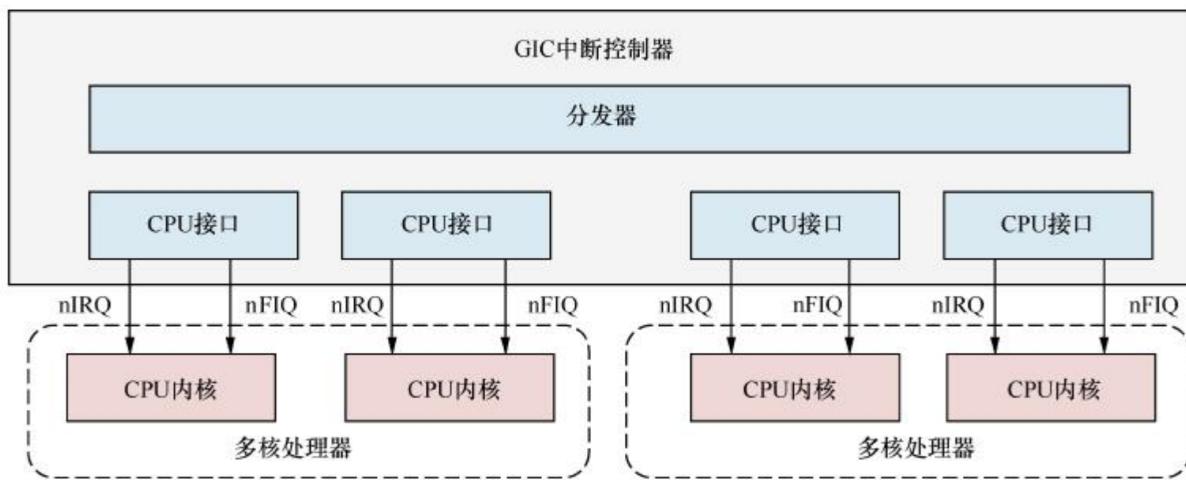
中断类型	中断号范围
软件触发中断 (SGI)	0~15
私有外设中断 (PPI)	16~31
共享外设中断 (SPI)	32~1019

SoC GIC-V2

- SGI              SGI Linux              Inter-Processor Interrupt IPI              CPU
  - PPI              PPI              CPU              CPU              local timer
  - SPI

SGI PPI CPU SPI CPU

GIC distributor CPU



▲图 13.2 GIC-V2 的内部结构

GIC-V2

## CPU CPU interface

13.2

CPU

nIRQ nFIQ

CPU

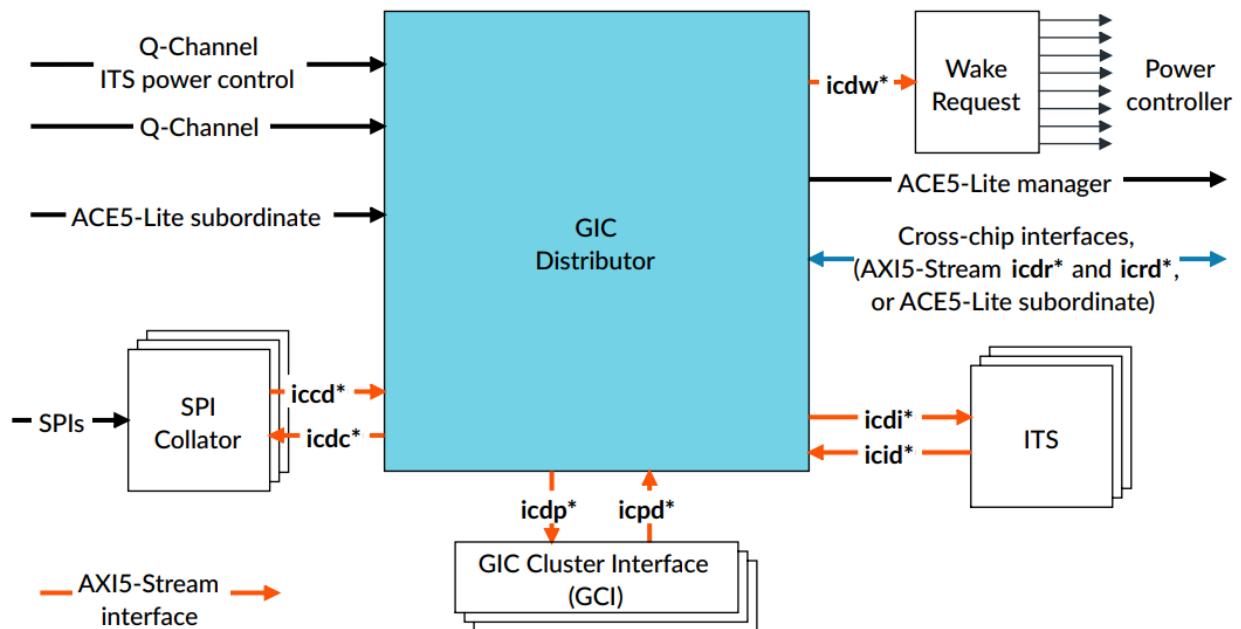
CPU CPU

CPU

GIC-700

GIC-700

**Figure 3-1: GIC-700 Distributor**



gic-v3 Distributor Redistributor CPU interface ITS Interrupt Translation Service

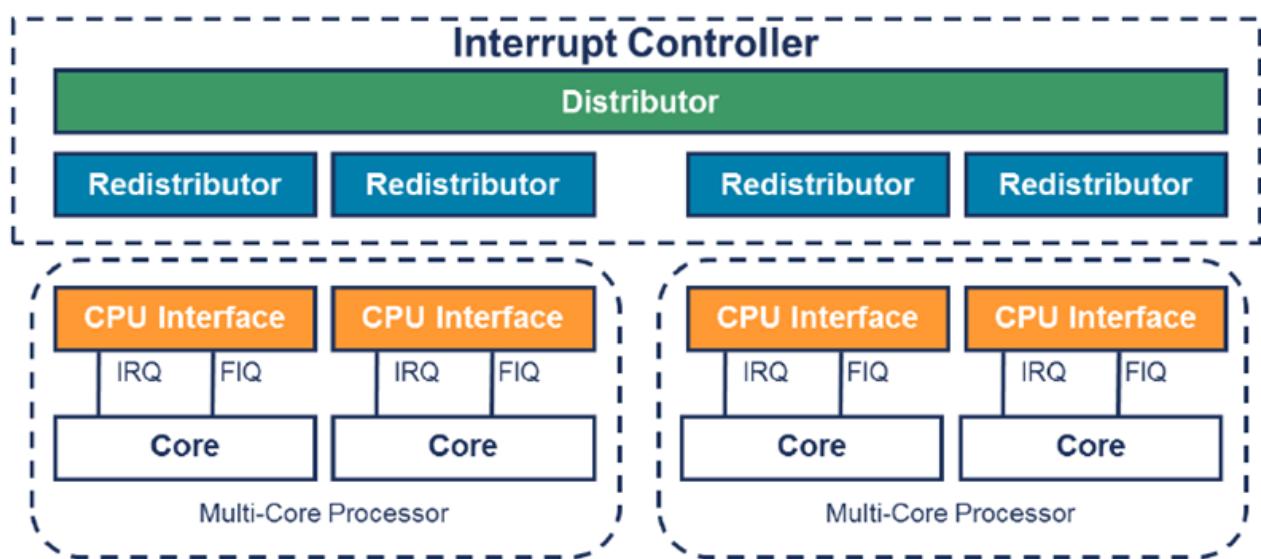
- Distributor SPI priority affinity trigger-type sgi spi Redistributor
  - Redistributor PPI LPI priority affinity trigger-type PPI LPI CPU interface
  - CPU interface Redistributor CPU CPU Distributor Redistributor
  - ITS msi msi LPI Redistributor

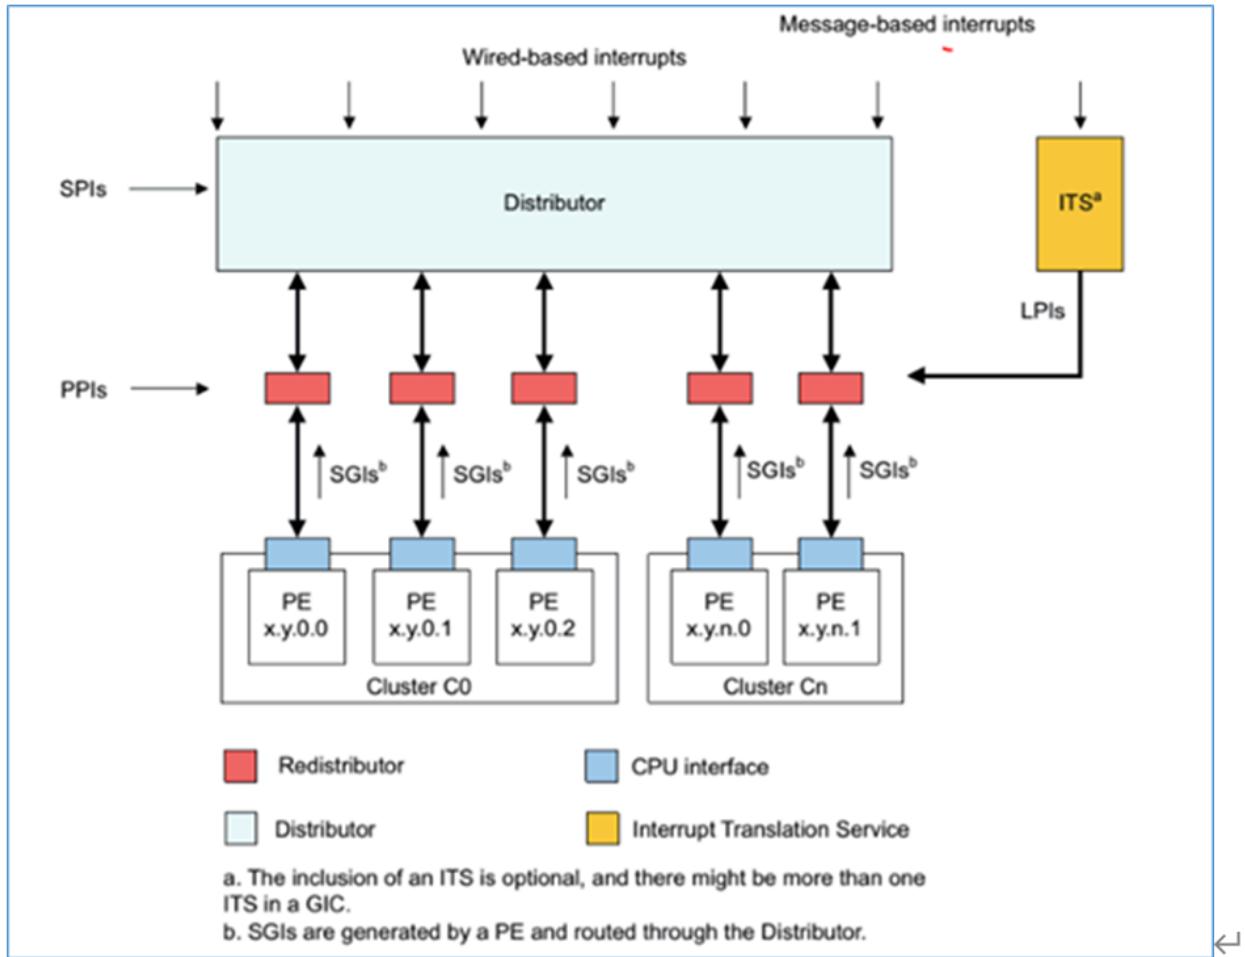
GIC	General Interrupt Controller
SGI	Software Generated Interrupt
PPI	Private Peripheral Interrupt
SPI	Shared Peripheral Interrupt
LPI	Locality-specific Peripheral Interrupt
PE	Process Element
ITS	Interrupt Translation Service

SGI	SGI	security	group	PE	PE / PE
PPI	PPI				
SPI			PE	PE /	PE
LPI	LPI	forwarding	/ITS		

SGI	0 - 15	PE
PPI	16 – 31	PE
SPI	32 - 1019	
	1020 - 1023	
	1024 - 1055	
PPI	1056 -1119	GICv3.1
	1120 - 4095	
SPI	4096 - 5119	GICv3.1
	5120 - 8191	
LPI	8192 -	

LPI      8192      LPI      GICD\_TYPER      IDbits field      LPI





GCI GIC Cluster Interface      cluster      Redistributor c2000      EXG      cluster      core 8 thread      GCI      8 Redistributor  
 c2k      ICI      ARM CPU interface

**SGI**      core A      CPU interface      GICD      threadde GICR      interface      thread

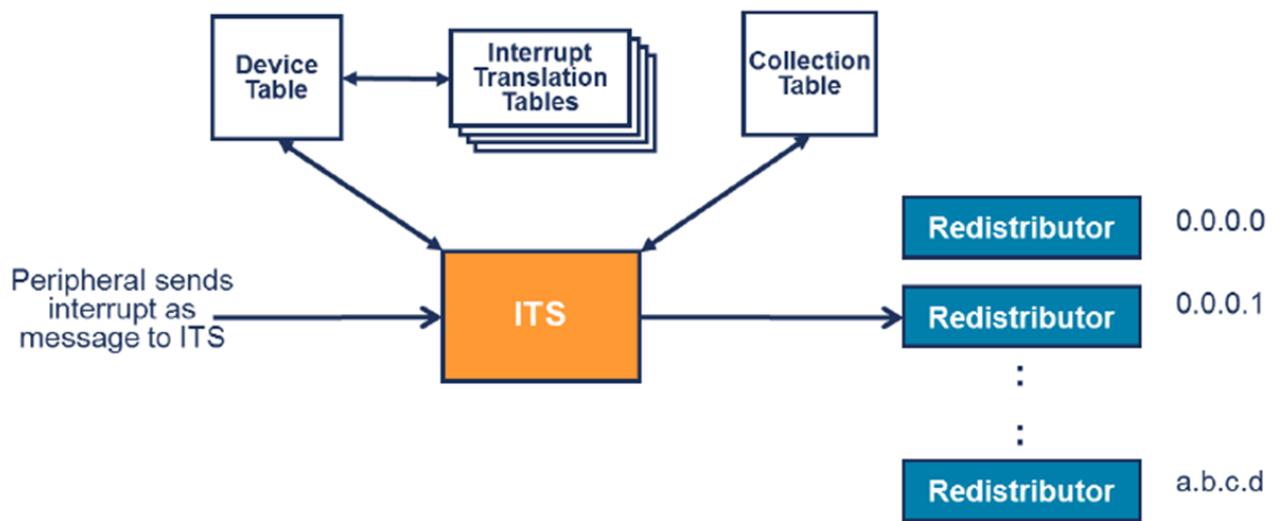
**PPI**      PPI      GICR interface      thread

EXG timer

**SPI**      SPI      GICD      GICR GICR interface      thread

PCIe

**LPI**      MSI      ITS      GICR GICR interface      thread



1. ITS device id event id device id event id
2. ITS GICR CPU
  - a) Device table device ITT
  - b) ITT event id collection id
  - c) Collection table collection id GICR
3. ITS ITS ITS CMDQ
 

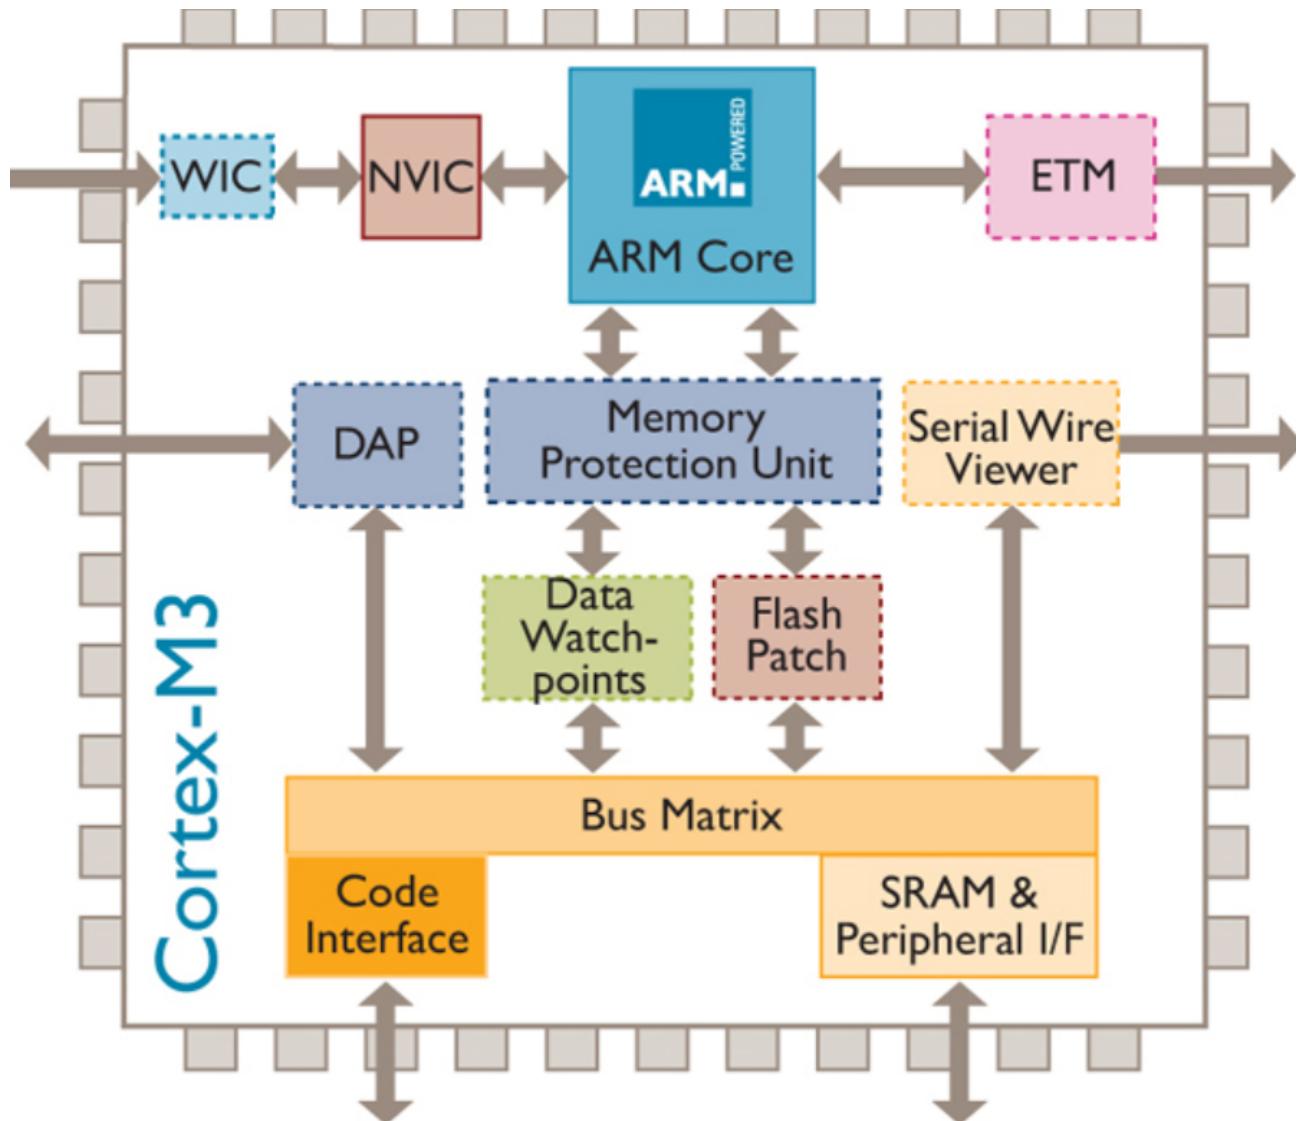
c2k	MSI	its->gicd->gicr	arm	its	redistributor	GICR	distributor(GICD)
GICD		ITS	GICD	GICD	ddr	GICR	c2k
							arm

### 1.2.7 RISCV AIA

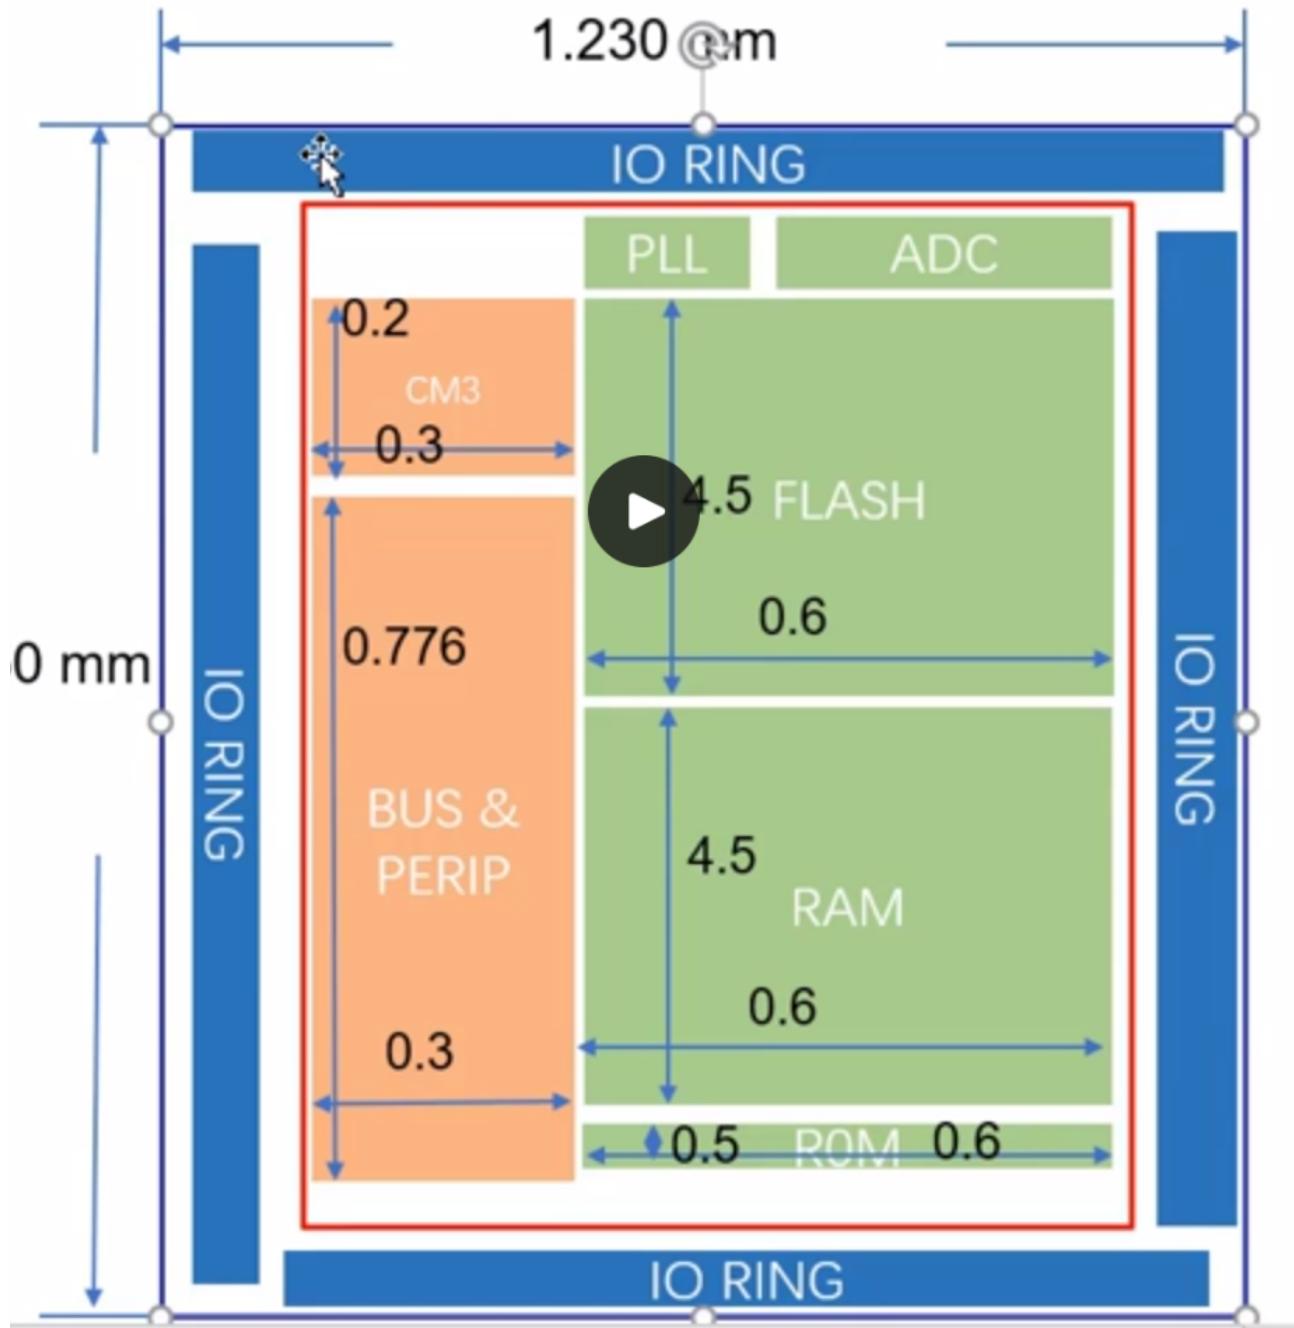
[toc]

### 1.2.8 Cortex M3 MCU

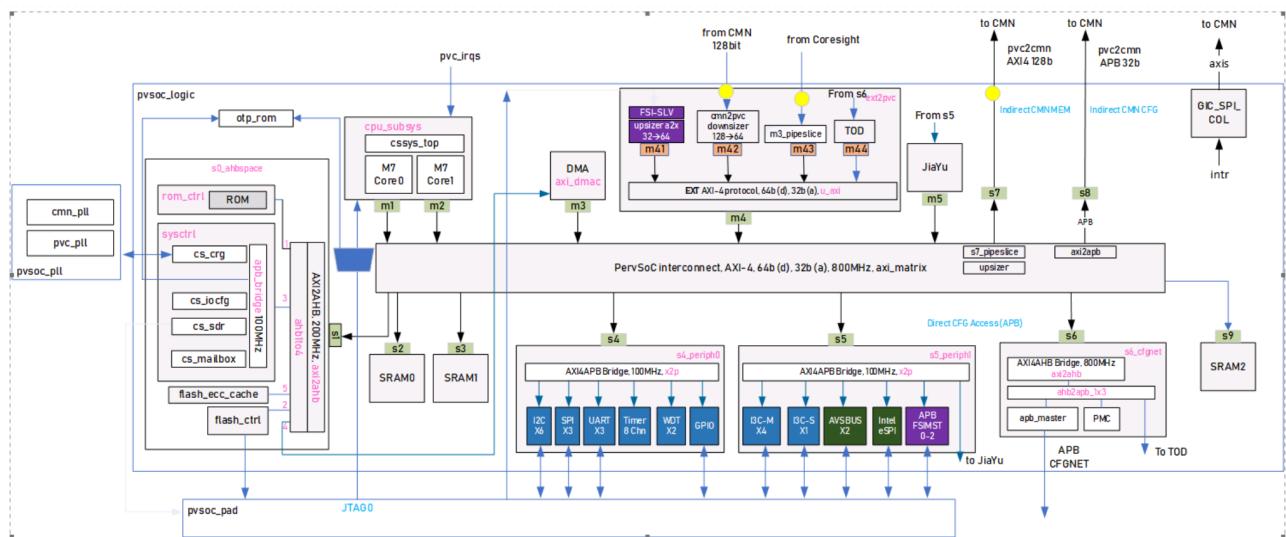
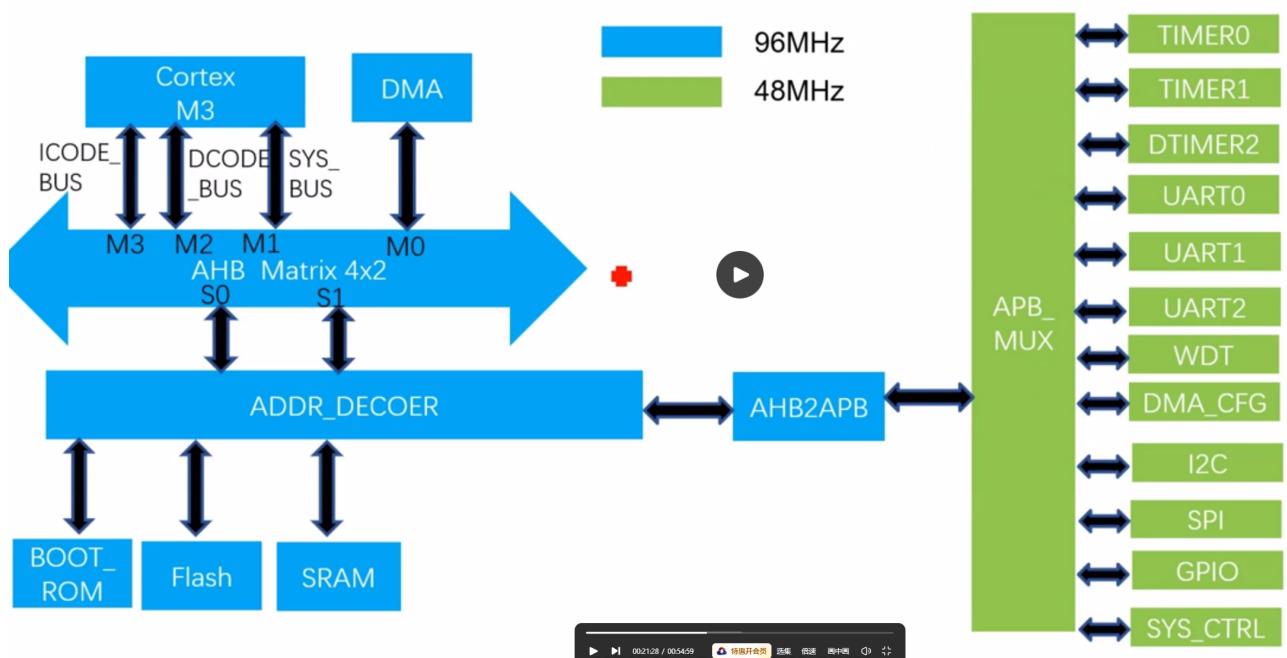
[toc]



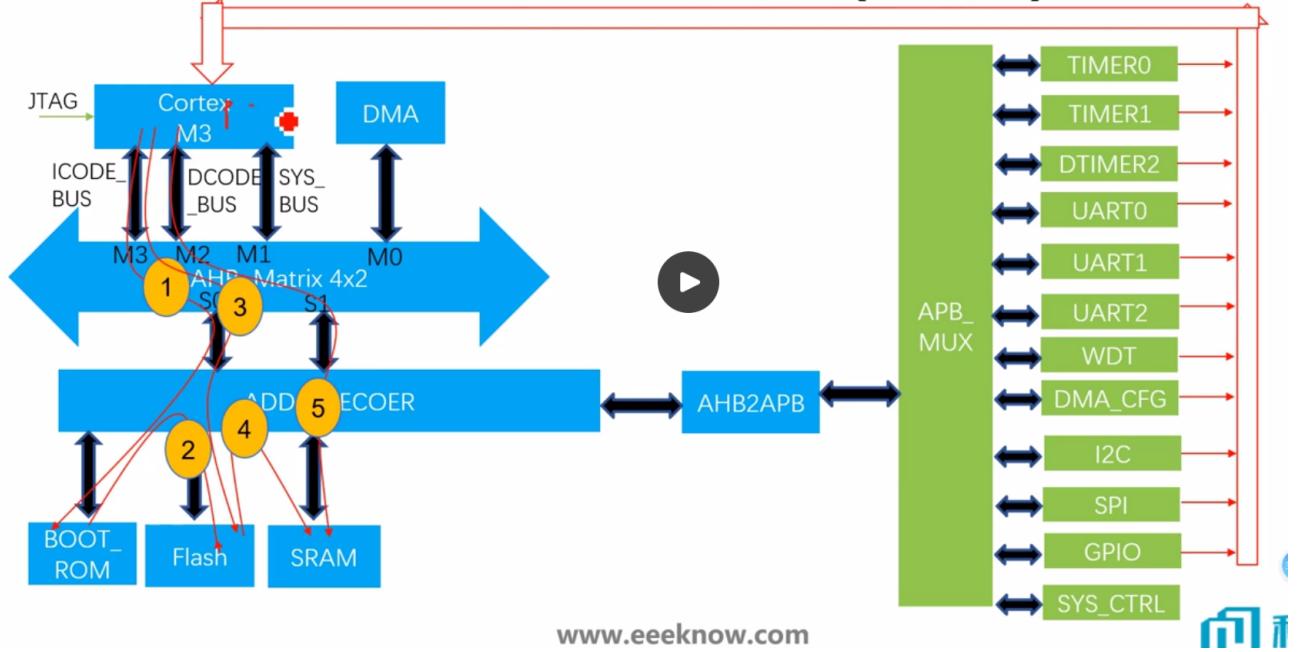
MCU



# Cortex-M3 MCU架构-整体架构



# Cortex-M3 MCU架构-整体架构(数据流)



## Cortex-M3 MCU架构-Memory Map

SUB	名字	结束地址	起始地址	大小 (KBYTE)
存储系统	Program	0000FFFF	00000000	64
	Reserved	00FFFFFF	00010000	16320
	Boot Loader	0100FFFF	01000000	64
	Reserved	1FFFFFFF	01010000	507840
	RAM	2000FFFF	20000000	64
	Reserved	3FFFFFFF	20010000	524224
APB 系统	TIMER0	40000FFF	40000000	4
	TIMER1	40001FFF	40001000	4
	DTIMER	40002FFF	40002000	4
	SYS_CTRL	40003FFF	40003000	4
	UART 0	40004FFF	40004000	4
	UART 1	40005FFF	40005000	4
	UART 2	40006FFF	40006000	4
	SPI	40007FFF	40007000	4
	WDT	40008FFF	40008000	4
	I2C0	40009FFF	40009000	4
	I2C1	4000AFFF	4000A000	4
	APB_TEST	4000BFFF	4000B000	4
	ADC	4000CFFF	4000C000	4
	RTC	4000DFFF	4000D000	4
	Reserved	4000EFFF	4000E000	4
保留	DMA	4000FFFF	4000F000	4
	Reserved	40010FFF	40010000	4
	Reserved	40011FFF	40011000	4
	Reserved	4001EFFF	40012000	52
CM3	Cortex3	FFFFFFFFFF	40020000	3145600

### 1.3 ARM N2

ARM N2    ARM Neoverse N2 reference design    ARM  
CPU

Neoverse N2    ARM

1. N2

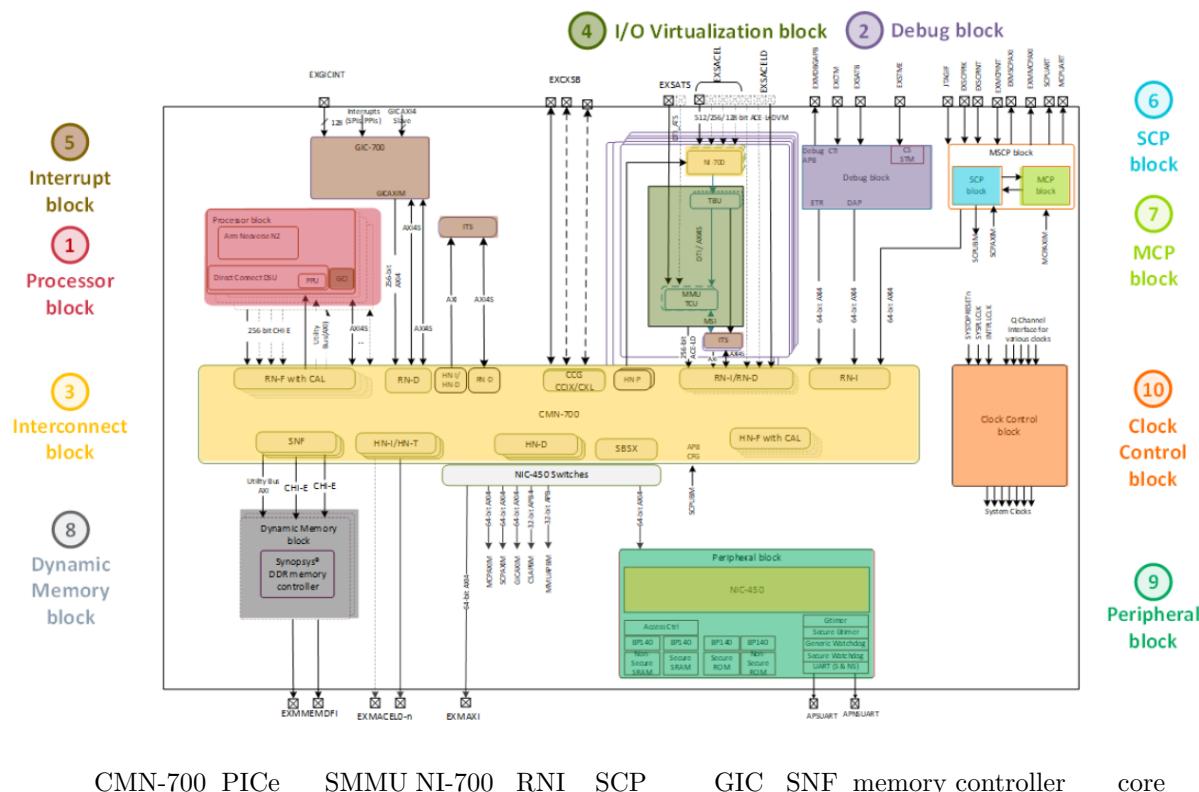
HPC

- 2. Neoverse N2
  - 3. ARM N2
  - 4. N2
  - 5. Neoverse N2      ARM      ARM  
                          ARM N2

ARM Neoverse N2	ARM IP	IP
1. <b>ARM Neoverse N2 CPU</b>	N2	N2 ARMv9
2. <b>System Level Cache (SLC)</b>	ARM N2	SLC
3. <b>CCIX CXL</b>	N2 CCIX Cache Coherent Interconnect for Accelerators	CXL Compute Express Link
4. <b>AMBA (Advanced Microcontroller Bus Architecture)</b>	AMBA	SoC IP N2
	AMBA	
5. <b>TrustZone</b>	ARM TrustZone	
6. <b>Mali GPU</b>	N2 CPU	ARM Mali GPU
7. <b>SMMU (System Memory Management Unit)</b>		
IP	ARM Neoverse N2	

Arm Neoverse N2 reference design Technical Overview

Figure 3-1: RD-N2 architecture block diagram



## 1.4

### 1.4.1 ARM

#### core

1. BMC C2000 CHIP POR\_N pervsoc
2. otprom valid SCP SRAM BISR PVO(SCP) ROM otprom
3. SCP SRAM
4. SCP ROM FW GPIO CHIP ID, CHIP ID CHIP;
5. I2C/FSI
6. CHIP SCP QSPI;
7. CHIP SCP CHIP CHIP
8. CHIP SCP CHIP SCP
9. CHIP SCP QSPI SCP RAM FW CHIP SRAM;
10. CHIP SCP 8 SCP RAM FW SCP, SCP SRAM:
11. CHIP SCP SRAM, SCP RAM FW:
12. CHIP SCP QSPI C2C SCP
13. CHIP SCP CHIP SCP C2C C2C;
14. SCP CHIP ID CMN OCM(SLC) SAM;
15. CHIP SCP MCP RAM FW, MCP( ) +
16. SCP CHIP L2 NCU SCOM/FIR;+
17. CHIP SCP Hostboot base CHIP OCM(SLC);+
18. CHIP SCP CHIP
19. srest 0x100

#### core-hostboot

1. kernel other thread spinlock, thread
2. hostboot kernel C/C++ main +
3. hostboot kernel
4. MMU(SLB/TLB), DIMM OCM(SLC, 28 \* 2MB=56MB 3/4
5. TB CPU
6. Scratch FSP/BMC;
7. CPU init main
8. kernel other thread spinlock 1 thread
9. hostboot task task user kernel hypersior
10. init main VFS;
11. VFS base modules

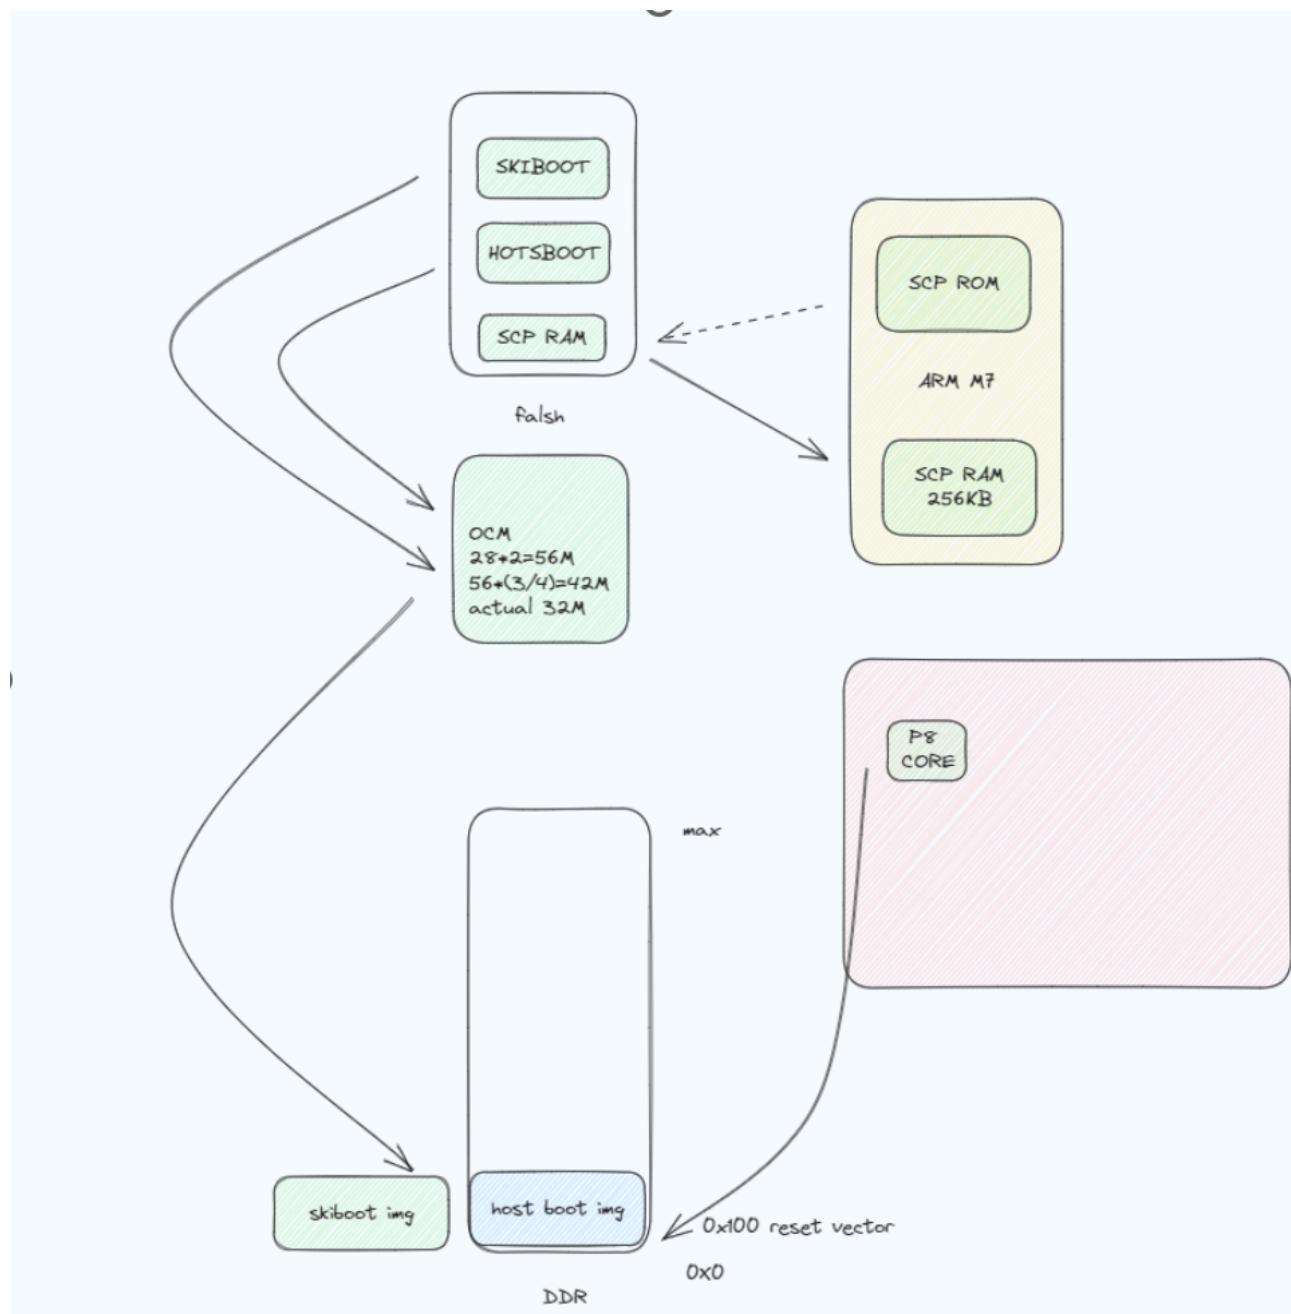
init

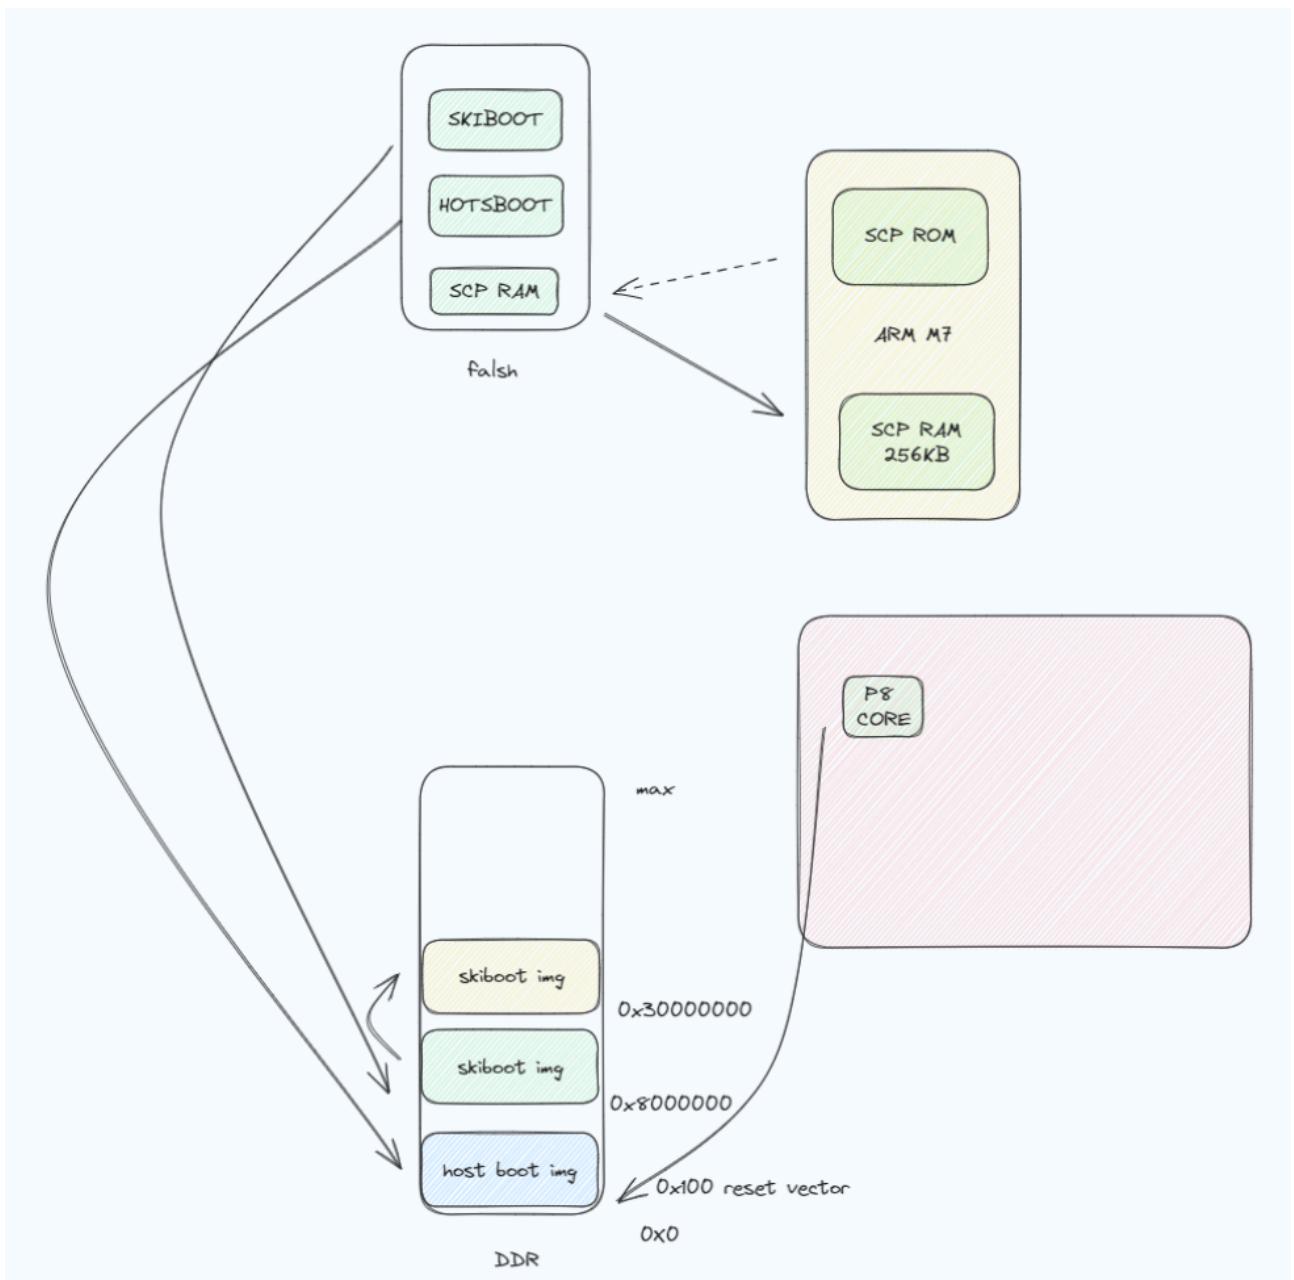
12. initservice
13. initservice g\_taskinfolist
14. secureboot\_base/ipmapbase/pnor/vfs/extinitsvc
15. hostboot extended image
16. extinitsvc g\_exttaskinfolist
17. istepdisp hostboot
18. istep hostboot
19. istep 21.03 payload(skiboot,OPAL);+
20. payload(skiboot,OPAL)

### **core-skiboot**

1. C2000 skiboot
2. thread skiboot relocate thread
3. thread boot flag=1, thread
4. thread main cpu entry, thread secondary cpu entry;
5. uart(console);
6. skiboot /
7. skiboot
8. opal call table;
9. hostboot device tree HDAT;
10. DT SCOM/eSPI(LPC):
11. DT C2000 DIMM
12. DT homer C2000
13. DT CPU(thread);
14. opal DT;
- 15.
16. BMC ipmi/BT/UART;
17. GIC
18. CPU thread
19. sreset CPU
20. TOD TOD
21. sensor OPAL
- 22.
- 23.NV

- 24. jiayu
- 25. TPM
- 26. opal console(UART)
- 27. PCIE HB PCIE (RC CXL PCIe device tree )
- 28. PCIE SLOT
- 29.
- 30.
- 31. zImage(boot loader petitboot);
- 32. boot loader





CMN OCM(on chip memory) L3 cache DDR Intel cache as ram

OCM memory CMN OCM(on chip memory) 28\*2=56M 3/4 32M

SCP ram DDR training OCM flash DDR

OCM SNF

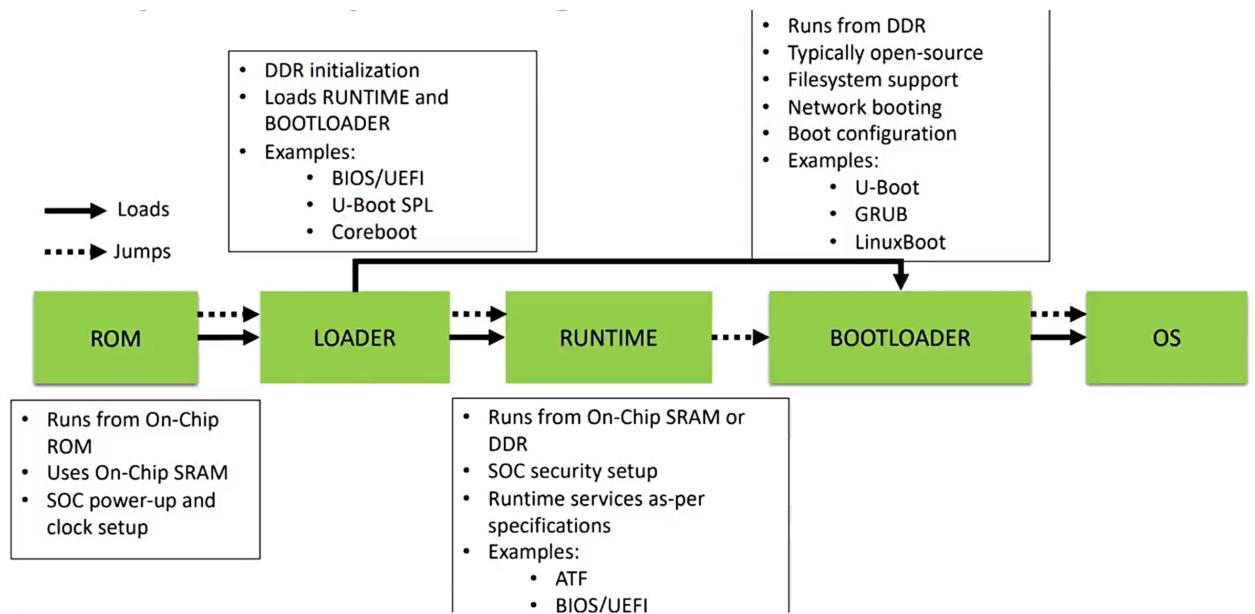
SCP OCM SCP RAM hostboot OCM skiboot OCM skiboot ddr training DDR OCM OCM cache

hostboot skiboot 0x800 0000

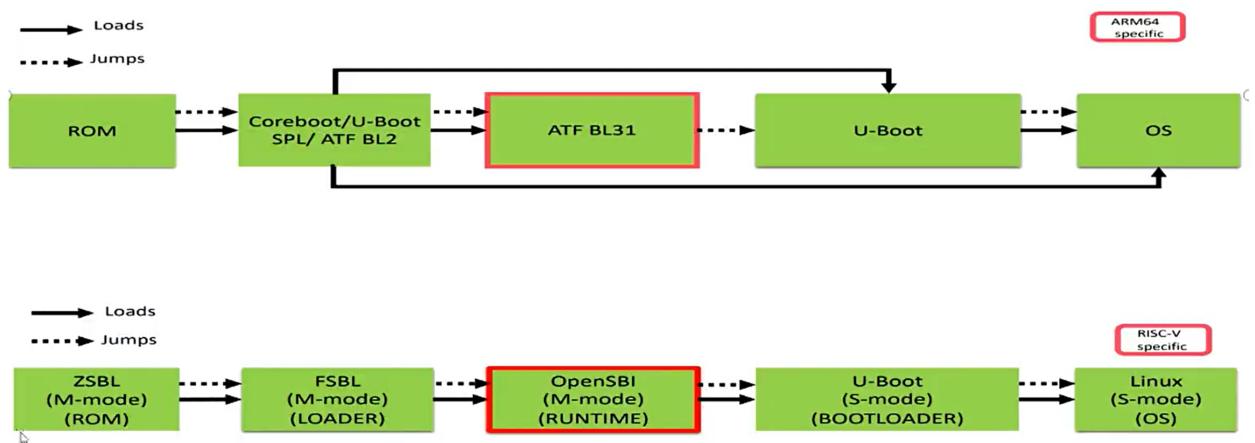
32M 0x200 0000, skiboot 0x3000 0000 OCM 0x0 0x0

## 1.4.2 RISC-V Firmware

Firmware

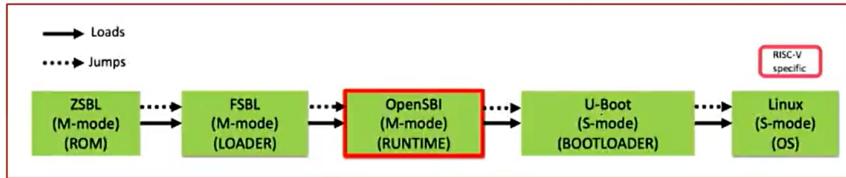


ARM64 vs RISC-V



uboot    EDK2    PEI    DXE

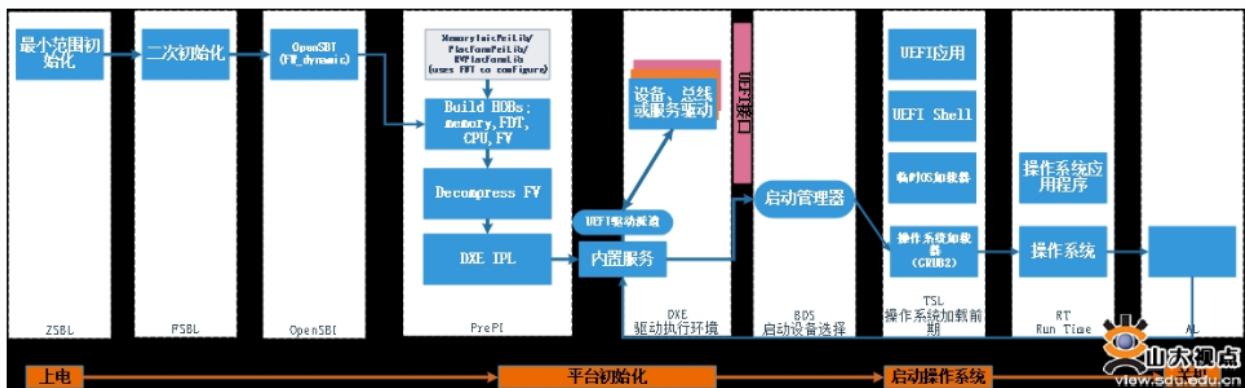
RISC-V linux启动流程如图，启动流程：ZSBL-->FSBL-->OpenSBI-->u-boot (UEFI)-->Linux



1. ZSBL (第0阶段Bootloader) at 0x00010000 (ROM in SOC)
2. FSBL (第1阶段Bootloader)
3. OpenSBI (runtime services) M模式
4. U-boot/UEFI (bootloader),
5. Linux启动linux需要以下文件：
  - a. /boot/extlinux/extlinux.conf: 包含启动菜单、默认启动等
  - b. Image: Linux镜像
  - c. Initrd: Linux执行需要的ramdisk, 执行其中的init文件并进行进一步处理
  - d. 设备树(.dtb) (可选) :未指定时会使用默认的dtb

ZSBL: SCP ROM

FSBL: SCP RAM



#### 1.4.3 Cache as RAM and On Chip Memory

x86 Cache as RAM, CAR;

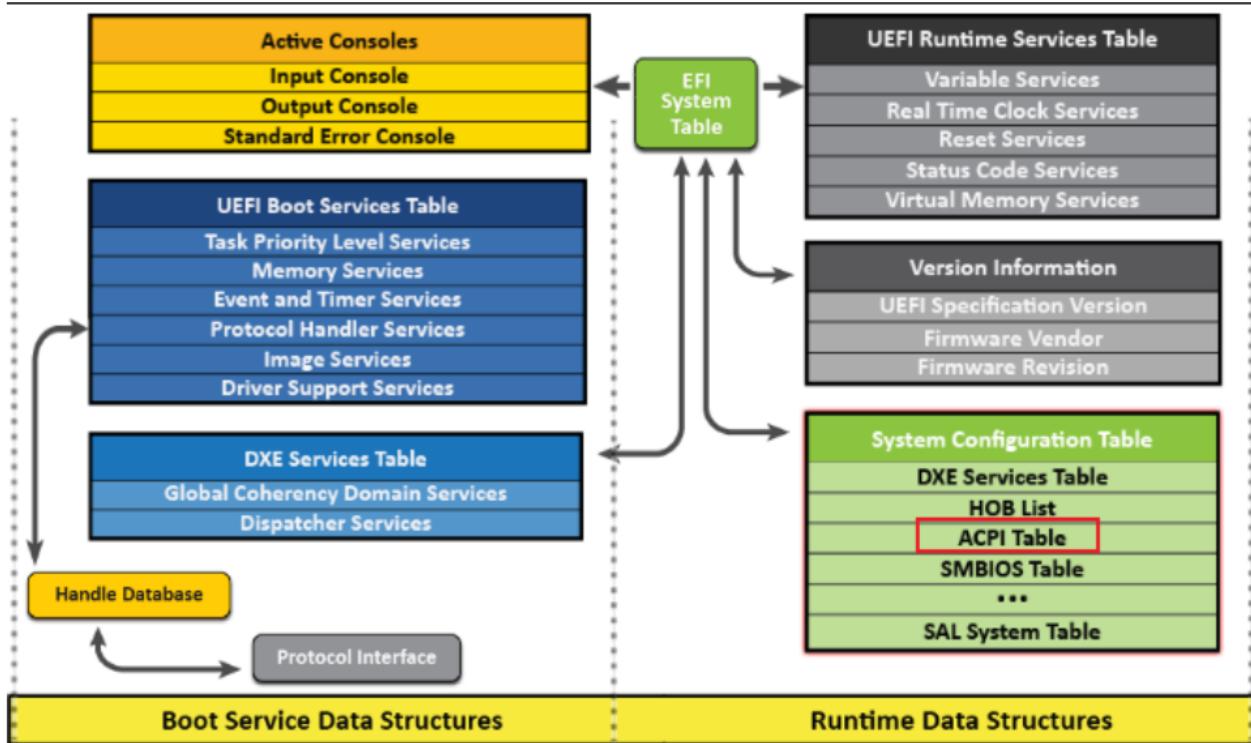
ARM On Chip Memory OCM

DDR cache RAM C

CAR

OCM

#### 1.4.4 System Table



MdeModulePkg\Core\DXe\DXeMain\DXeMain.c

```

//  

// Allocate the EFI System Table and EFI Runtime Service Table from EfiRuntimeServicesData  

// Use the templates to initialize the contents of the EFI System Table and EFI Runtime  

// Services Table  

//  

gDxeCoreST = AllocateRuntimeCopyPool (sizeof (EFI_SYSTEM_TABLE),  

    &mEfiSystemTableTemplate);  

ASSERT (gDxeCoreST != NULL);  

gDxeCoreRT = AllocateRuntimeCopyPool (sizeof (EFI_RUNTIME_SERVICES),  

    &mEfiRuntimeServicesTableTemplate);  

ASSERT (gDxeCoreRT != NULL);  

gDxeCoreST->RuntimeServices = gDxeCoreRT;  

///  

/// EFI System Table  

///  

typedef struct {  

    ///  

    /// The table header for the EFI System Table.  

    ///  

    EFI_TABLE_HEADER          Hdr;  

    ///  

    /// A pointer to a null terminated string that identifies the vendor

```

```

/// that produces the system firmware for the platform.

CHAR16 *FirmwareVendor;

/// A firmware vendor specific value that identifies the revision
/// of the system firmware for the platform.

UINT32 FirmwareRevision;

/// The handle for the active console input device. This handle must support
/// EFI_SIMPLE_TEXT_INPUT_PROTOCOL and EFI_SIMPLE_TEXT_INPUT_EX_PROTOCOL.

/// EFI_HANDLE ConsoleInHandle;

/// A pointer to the EFI_SIMPLE_TEXT_INPUT_PROTOCOL interface that is
/// associated with ConsoleInHandle.

/// EFI_SIMPLE_TEXT_INPUT_PROTOCOL *ConIn;

/// The handle for the active console output device.

/// EFI_HANDLE ConsoleOutHandle;

/// A pointer to the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL interface
/// that is associated with ConsoleOutHandle.

/// EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *ConOut;

/// The handle for the active standard error console device.
/// This handle must support the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL.

/// EFI_HANDLE StandardErrorHandler;

/// A pointer to the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL interface
/// that is associated with StandardErrorHandler.

/// EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *StdErr;

/// A pointer to the EFI Runtime Services Table.

/// EFI_RUNTIME_SERVICES *RuntimeServices;

/// A pointer to the EFI Boot Services Table.

/// EFI_BOOT_SERVICES *BootServices;

/// The number of system configuration tables in the buffer ConfigurationTable.

/// UINTN NumberOfTableEntries;

```

```

/// A pointer to the system configuration tables.
/// The number of entries in the table is NumberOfTableEntries.
///
EFI_CONFIGURATION_TABLE *ConfigurationTable;
} EFI_SYSTEM_TABLE;

```

## 1.4.5 EDK2 EVENT

[toc]

UEFI

```

/ intel 8253/8254      PC      IC      MSC8051,
Timer          cpu
PC/AT UEFI     8254   Timer mode 3   1.1931816MHz tick mTimerPeriod    1ms 10000 tick =10000*100ns
pin 8259 IRQ0   CPU     Timer
cpu archprotocol Legacy8259Protocol ,   cpu IDT   Timer TimerInterruptHandler

```

### Timer

Timer

\*\*1.\*\*

\*\*2.\*\*

\*\*3.\*\*

\*\*4.\*\* Timer CoreTimerTick(mTimerPeriod)/CoreTimerTick(100ms) TimerArchprotocol install EVT\_NOTIFY\_SIGNAL

\*\*5.\*\* System time mEfiSystemTime =mEfiSystemTime+ 1Tick(100ms)

**6.** timer event list event BS->SignalEvent signal

BS->SignalEvent-->CoreNotifyEvent-->CoreRestoreTpl-->CoreDispatchEventNotifies-->A. EVT\_NOTIFY\_SIGNAL  
>SignalCount = 0 Event->NotifyFunction (Event, Event->NotifyContext) Tpl gEventPending  
mask bit event

\*\*7.\*\*

\*\*8.\*\*

**1.** gEventQueue--- A list of event's to notify for each priority level

**2.** mEfiTimerList---- timer event

**3.** gEventSignalQueue --- A list of events to signal based on EventGroup type

**4.** EVT\_NOTIFY\_WAIT event BS-> **CoreCheckEvent**->CoreNotifyEvent-->CoreRestoreTpl-->CoreDispatchEventNotifies gEventQueue IEvent

gEventPending gEventQueue 1 0

gEventPending 32 1

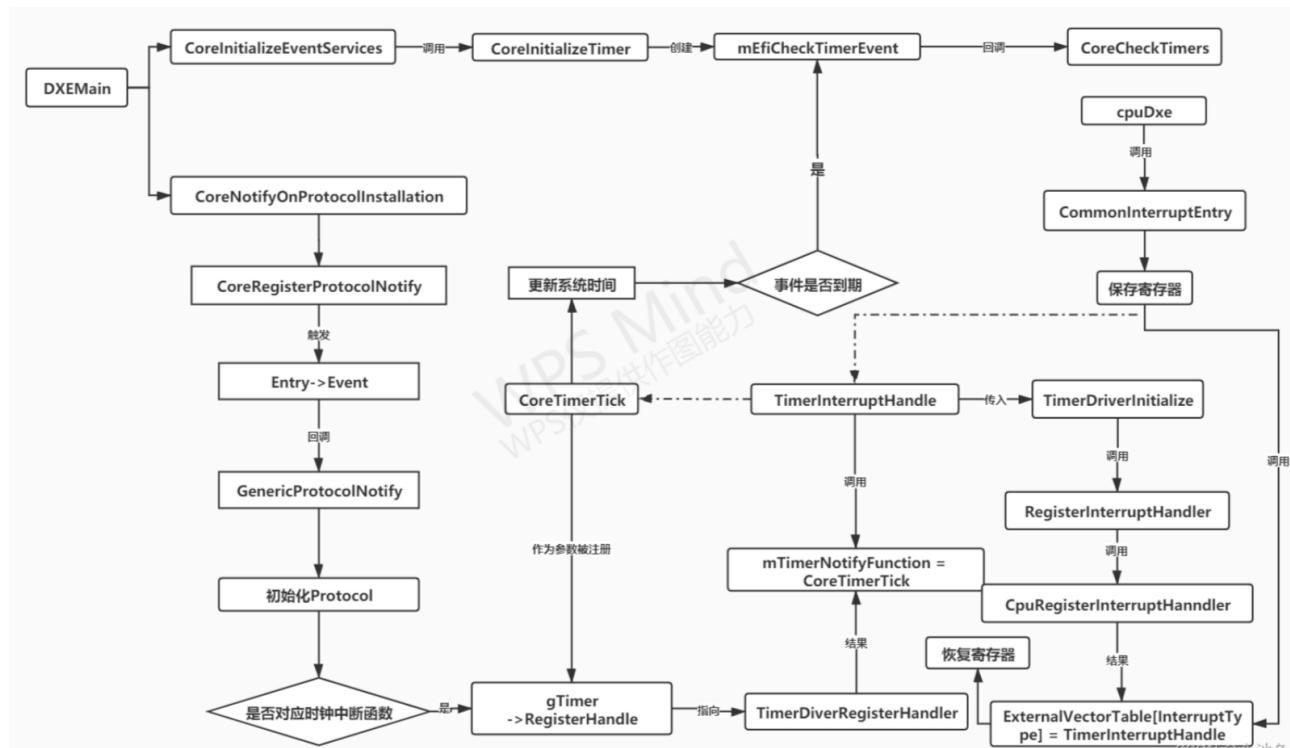
x86 IRQ

- IRQ0:
- IRQ1:
- IRQ2:
- IRQ3: 2 COM2
- IRQ4: 1 COM1
- IRQ5:
- IRQ6:
- IRQ7: 1 LPT1
- IRQ

event

```
//  
// Task priority level  
  
#define TPL_APPLICATION      4  
#define TPL_CALLBACK         8  
#define TPL_NOTIFY           16  
#define TPL_HIGH_LEVEL       31
```

31



DXE main

```
//  
// Initialize the Event Services
```

```

//  

Status = CoreInitializeEventServices ();  

...  

//  

// Register for the GUIDs of the Architectural Protocols, so the rest of the  

// EFI Boot Services and EFI Runtime Services tables can be filled in.  

// Also register for the GUIDs of optional protocols.  

//  

CoreNotifyOnProtocolInstallation ();  

event  

/**  

Called by the platform code to process a tick.  

@param Duration The number of 100ns elapsed since the last call  

to TimerTick  

**/  

VOID  

EFIAPI  

CoreTimerTick (  

    IN UINT64 Duration  

)  

{  

    IEVENT *Event;  

//  

// Check runtiem flag in case there are ticks while exiting boot services  

//  

CoreAcquireLock (&mEfiSystemTimeLock);  

//  

// Update the system time  

//  

mEfiSystemTime += Duration;  

//  

// If the head of the list is expired, fire the timer event  

// to process it  

//  

if (!IsListEmpty (&mEfiTimerList)) {  

    Event = CR (mEfiTimerList.ForwardLink, IEVENT, Timer.Link, EVENT_SIGNATURE);  

    if (Event->Timer.TriggerTime <= mEfiSystemTime) {  

        CoreSignalEvent (mEfiCheckTimerEvent);  

    }
}

```

```

        }

    }

    CoreReleaseLock (&mEfiSystemTimeLock);
}

timer event

/***
Initializes "event" support.

retval EFI_SUCCESS           Always return success

**/


EFI_STATUS
CoreInitializeEventServices (
    VOID
)
{
    UINTN       Index;

    for (Index=0; Index <= TPL_HIGH_LEVEL; Index++) {
        InitializeListHead (&gEventQueue[Index]);
    }

    CoreInitializeTimer ();

    CoreCreateEventEx (
        EVT_NOTIFY_SIGNAL,
        TPL_NOTIFY,
        EfiEventEmptyFunction,
        NULL,
        &gIdleLoopEventGuid,
        &gIdleLoopEvent
    );

    return EFI_SUCCESS;
}



- gEventQueue--- A list of event's to notify for each priority level
- CoreInitializeTimer



/***
Initializes timer support.

**/


VOID
CoreInitializeTimer (
    VOID
)
{

```

```

EFI_STATUS Status;

Status = CoreCreateEventInternal (
    EVT_NOTIFY_SIGNAL,
    TPL_HIGH_LEVEL - 1,
    CoreCheckTimers,
    NULL,
    NULL,
    &mEfiCheckTimerEvent
);

ASSERT_EFI_ERROR (Status);
}

```

```

#include <Uefi.h>
#include <Library/UefiBootServicesTableLib.h>
#include <Library/UefiLib.h>

```

```

EFI_EVENT TimerEvent;
EFI_EVENT ExitBootServicesEvent;
BOOLEAN TimerTriggered = FALSE;

VOID EFIAPI TimerEventHandler(
    IN EFI_EVENT Event,
    IN VOID *Context
) {
    TimerTriggered = TRUE;
    Print(L"Timer event triggered!\n");
}

```

```

EFI_STATUS
EFIAPI
UefiMain(
    IN EFI_HANDLE ImageHandle,
    IN EFI_SYSTEM_TABLE *SystemTable
) {
    EFI_STATUS Status;

    //          10000000 100ns (1)
    Status = gBS->CreateEvent(
        EVT_TIMER | EVT_NOTIFY_SIGNAL,
        TPL_CALLBACK,
        TimerEventHandler,
        NULL,
        &TimerEvent
    );
    if (EFI_ERROR(Status)) {

```

```

    Print(L"Failed to create event: %r\n", Status);
    return Status;
}

// 1
Status = gBS->SetTimer(
    TimerEvent,
    TimerRelative,
    10000000 // 100ns 10000000 = 1
);
if (EFI_ERROR(Status)) {
    Print(L"Failed to set timer: %r\n", Status);
    gBS->CloseEvent(TimerEvent);
    return Status;
}

// 
while (!TimerTriggered) {
    gBS->Stall(100000); // 100ms
}

// 
gBS->CloseEvent(TimerEvent);

return EFI_SUCCESS;
}

1

gBS->CreateEvent

typedef struct {
    UINT32 Type;
    EFI_TPL NotifyTpl;
    EFI_EVENT_NOTIFY NotifyFunction;
    VOID *NotifyContext;
    BOOLEAN Triggered;
} EFI_EVENT_INTERNAL;

EFI_STATUS
EFI API
CreateEvent (
    IN UINT32 Type,
    IN EFI_TPL NotifyTpl,
    IN EFI_EVENT_NOTIFY NotifyFunction,
    IN VOID *NotifyContext,
    OUT EFI_EVENT *Event
) {
    EFI_STATUS Status;
    EFI_EVENT_INTERNAL *NewEvent;
}

```

```

//  

if (Event == NULL || (Type & EVT_NOTIFY_SIGNAL && NotifyFunction == NULL)) {  

    return EFI_INVALID_PARAMETER;  

}  

//  

NewEvent = AllocateZeroPool(sizeof(EFI_EVENT_INTERNAL));  

if (NewEvent == NULL) {  

    return EFI_OUT_OF_RESOURCES;  

}  

//  

NewEvent->Type = Type;  

NewEvent->NotifyTpl = NotifyTpl;  

NewEvent->NotifyFunction = NotifyFunction;  

NewEvent->NotifyContext = NotifyContext;  

NewEvent->Triggered = FALSE;  

//  

Status = RegisterEvent(NewEvent);  

if (EFI_ERROR(Status)) {  

    FreePool(NewEvent);  

    return Status;  

}  

//  

*Event = (EFI_EVENT)NewEvent;  

return EFI_SUCCESS;
}  

event
event
if ((Type & EVT_NOTIFY_SIGNAL) != 0x00000000) {  

//  

// The Event's NotifyFunction must be queued whenever the event is signaled  

//  

InsertHeadList (&gEventSignalQueue, &IEvent->SignalLink);  

}  

32 bitmap 1

SignalEvent event 32
/**  

Signals the event. Queues the event to be notified if needed.  

@param UserEvent The event to signal .
```

```

@retval EFI_INVALID_PARAMETER Parameters are not valid.
@retval EFI_SUCCESS           The event was signaled.

*/
EFI_STATUS
EFIAPI
CoreSignalEvent (
    IN EFI_EVENT      UserEvent
)

/**
Queues the event's notification function to fire.

@param Event          The Event to notify

*/
VOID
CoreNotifyEvent (
    IN IEVENT        *Event
)
{

    /**
    // Event database must be locked
    //

    ASSERT_LOCKED (&gEventQueueLock);

    //
    // If the event is queued somewhere, remove it
    //

    if (Event->NotifyLink.ForwardLink != NULL) {
        RemoveEntryList (&Event->NotifyLink);
        Event->NotifyLink.ForwardLink = NULL;
    }

    //
    // Queue the event to the pending notification list
    //

    InsertTailList (&gEventQueue[Event->NotifyTpl], &Event->NotifyLink);
    gEventPending |= (UINTN)(1 << Event->NotifyTpl);
}

Duang          CoreTimerTick      TPL          restore TPL
timer  event mEfiTimerList
                    TPL
while (gEventPending != 0)

```

```

{
    PendingTpl = (UINTN) HighBitSet64 (gEventPending);
    if (PendingTpl <= NewTpl) {
        break;
    }
    gEfiCurrentTpl = PendingTpl;
    if (gEfiCurrentTpl < TPL_HIGH_LEVEL) {
        CoreSetInterruptState (TRUE);
    }
    CoreDispatchEventNotifies (gEfiCurrentTpl);
    // gEventQueue [gEfiCurrentTpl] event Notification
}

```

gEventPending 32 1

1.

2. TPL TPL 16 17 16 TPL 17

3. pending pending TPL 16 15 16

—

waitforevent 4 10ms 31 RestoreTpl 4, 31 4 30 5 4

## CoreTimerTick

/\*\*

*Lowers the task priority to the previous value. If the new priority unmasks events at a higher priority, they are dispatched.*

*@param NewTpl New, lower, task priority*

\*\*/

VOID

EFI API

CoreRestoreTpl (

IN EFI\_TPL NewTpl

)

{

EFI\_TPL OldTpl;

EFI\_TPL PendingTpl;

- OldTpl:

- PendingTpl:

OldTpl = gEfiCurrentTpl;

if (NewTpl > OldTpl) {

DEBUG ((EFI\_D\_ERROR, "FATAL ERROR - RestoreTpl with NewTpl(0x%x) > OldTpl(0x%x)\n",  
 ↵ NewTpl, OldTpl));

```

    ASSERT (FALSE);
}

ASSERT (VALID_TPL (NewTpl));


- OldTpl = gEfiCurrentTpl;
- NewTpl           NewTpl OldTpl

if (OldTpl >= TPL_HIGH_LEVEL && NewTpl < TPL_HIGH_LEVEL) {
    gEfiCurrentTpl = TPL_HIGH_LEVEL;
}


- TPL_HIGH_LEVEL      TPL_HIGH_LEVEL      TPL_HIGH_LEVEL

while (gEventPending != 0) {
    PendingTpl = (UINTN) HighBitSet64 (gEventPending);
    if (PendingTpl <= NewTpl) {
        break;
    }

    gEfiCurrentTpl = PendingTpl;
    if (gEfiCurrentTpl < TPL_HIGH_LEVEL) {
        CoreSetInterruptState (TRUE);
    }
    CoreDispatchEventNotifies (gEfiCurrentTpl);
}



- while (gEventPending != 0):
- PendingTpl = (UINTN) HighBitSet64 (gEventPending):
- PendingTpl  NewTpl
- PendingTpl      TPL_HIGH_LEVEL
- CoreDispatchEventNotifies



gEfiCurrentTpl = NewTpl;

if (gEfiCurrentTpl < TPL_HIGH_LEVEL) {
    CoreSetInterruptState (TRUE);
}



- NewTpl
- TPL_HIGH_LEVEL

CoreDispatchEventNotifies

/**
Dispatches all pending events.

@param Priority           The task priority level of event notifications

```

*to dispatch*

```
**/
VOID
CoreDispatchEventNotifies (
    IN EFI_TPL      Priority
)

{

    IEVENT          *Event;
    LIST_ENTRY       *Head;

    • Event:
    • Head:

    CoreAcquireEventLock ();
    ASSERT (gEventQueueLock.OwnerTpl == Priority);
    Head = &gEventQueue[Priority];

    • CoreAcquireEventLock():
    • ASSERT (gEventQueueLock.OwnerTpl == Priority):
    • Head = &gEventQueue[Priority]:

    while (!IsEmpty (Head)) {

        Event = CR (Head->ForwardLink, IEVENT, NotifyLink, EVENT_SIGNATURE);
        RemoveEntryList (&Event->NotifyLink);

        Event->NotifyLink.ForwardLink = NULL;
        • while (!IsEmpty (Head)):
        • Event = CR (Head->ForwardLink, IEVENT, NotifyLink, EVENT_SIGNATURE):
        • RemoveEntryList (&Event->NotifyLink):
        • Event->NotifyLink.ForwardLink = NULL:

    SIGNAL

    if ((Event->Type & EVT_NOTIFY_SIGNAL) != 0) {
        Event->SignalCount = 0;
    }

    • EVT_NOTIFY_SIGNAL   SignalCount

    CoreReleaseEventLock ();

    ASSERT (Event->NotifyFunction != NULL);
    Event->NotifyFunction (Event, Event->NotifyContext);

    • CoreReleaseEventLock():
    • ASSERT (Event->NotifyFunction != NULL):
```

- Event->NotifyFunction (Event, Event->NotifyContext):

```
    CoreAcquireEventLock ();
}
```

- CoreAcquireEventLock ():

```
gEventPending &= ~(UINTN)(1 << Priority);
CoreReleaseEventLock ();
}
```

- gEventPending &= ~(UINTN)(1 << Priority):
- CoreReleaseEventLock ():

- 31TPL
- 

[Blog](#)

[Blog](#)

## 1.4.6 EDK2 MEMORY MAP

GetMemoryMap

<https://blog.csdn.net/xiaopangzi313/article/details/109928878>

<https://www.lab-z.com/stu5/>

## 1.5

### 1.5.1

[toc]

Windows

#### 1. FAT32 File Allocation Table 32

- 2TB 4GB
- USB SD

#### 2. NTFS New Technology File System

- 
- Windows SSD

macOS

1. **HFS+ Hierarchical File System Plus**
    - 
    - macOS              APFS
  2. **APFS Apple File System**
    - SSD
    - macOS 10.13              Mac     iOS

## Linux

- ## 1. ext4 Fourth Extended File System

- - Linux

## 2. XFS

- 10

### 3. Btrfs B-tree File System

- 10

## 1. exFAT Extended File Allocation Table

- FAT32 NTFS
  - USB Windows macOS

## 1. ZFS Zettabyte File System

- RAID-Z
  -

### 1.5.2 RAID

[toc]

## RAID Redundant Array of Independent Disks

RAID

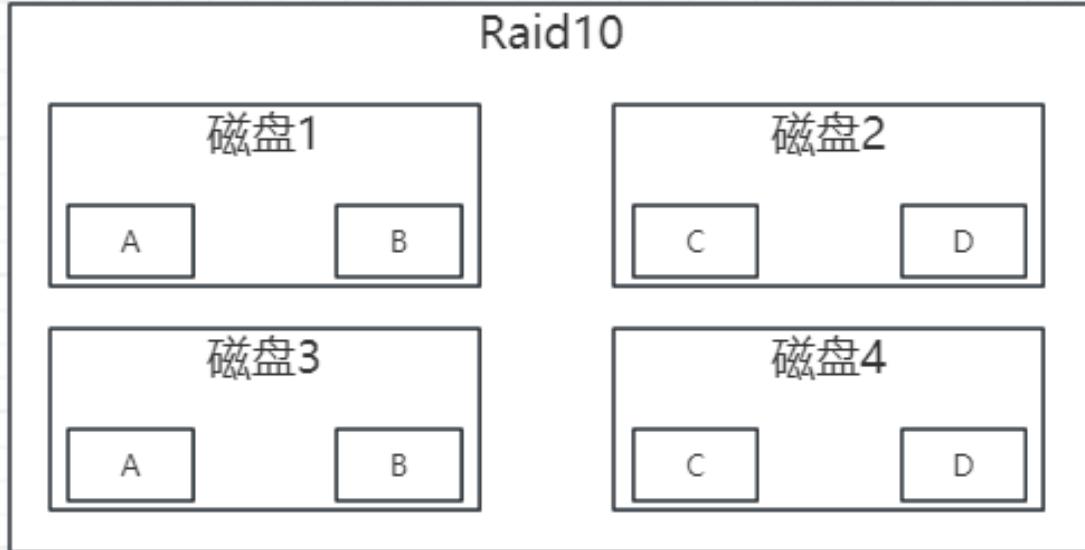
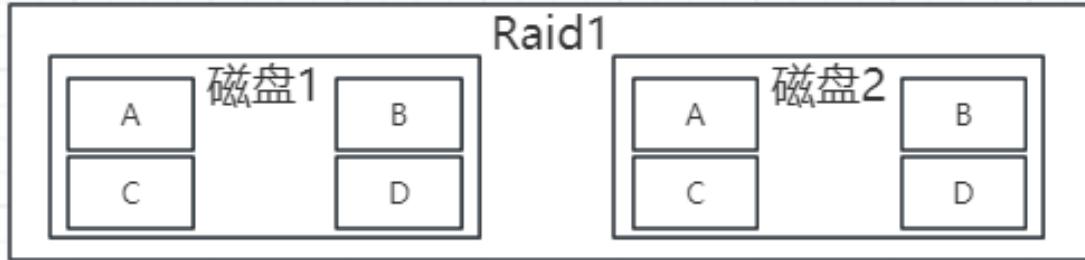
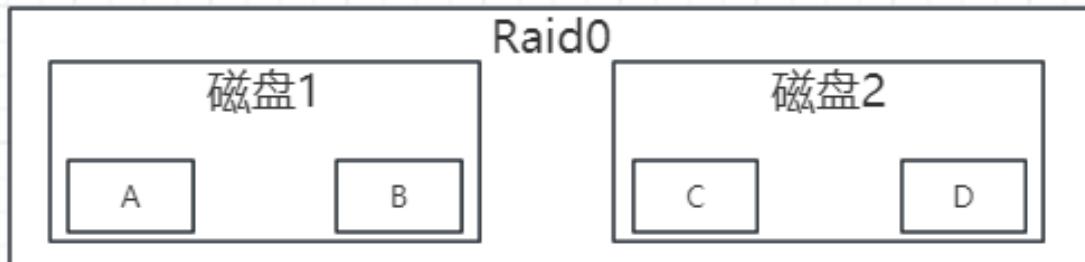
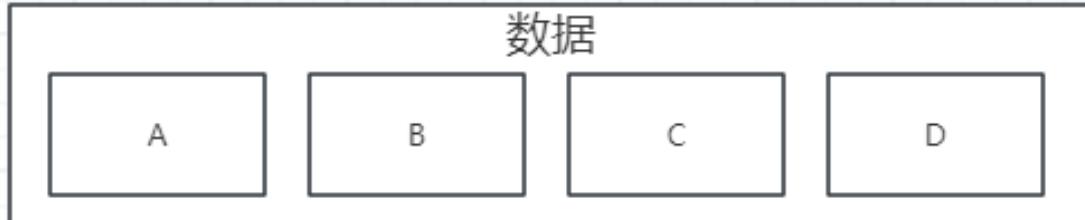
1. RAID 0
  2. RAID 1 RAID 5 RAID 6 RAID 1
  3. RAID 5

RAID

- RAID 0
  - RAID 1
  - RAID 5
  - RAID 6    RAID 5

- RAID 10 RAID 1+0 RAID 0 RAID 1

0 1 10



### RAID 3

- RAID 3
- RAID 3
-

## RAID 5

- RAID 5
- RAID 0
- $(\frac{(n-1)}{n})$
- RAID 1
- RAID 3

## RAID 6

- RAID 6 RAID 5
- RAID 6 RAID 5
- RAID 5
- RAID 5  $(\frac{(n-2)}{n})$
- RAID 5

- RAID 3
- RAID 5
- RAID 6

### 1.5.3 SSD

[toc]

SSD      SOC

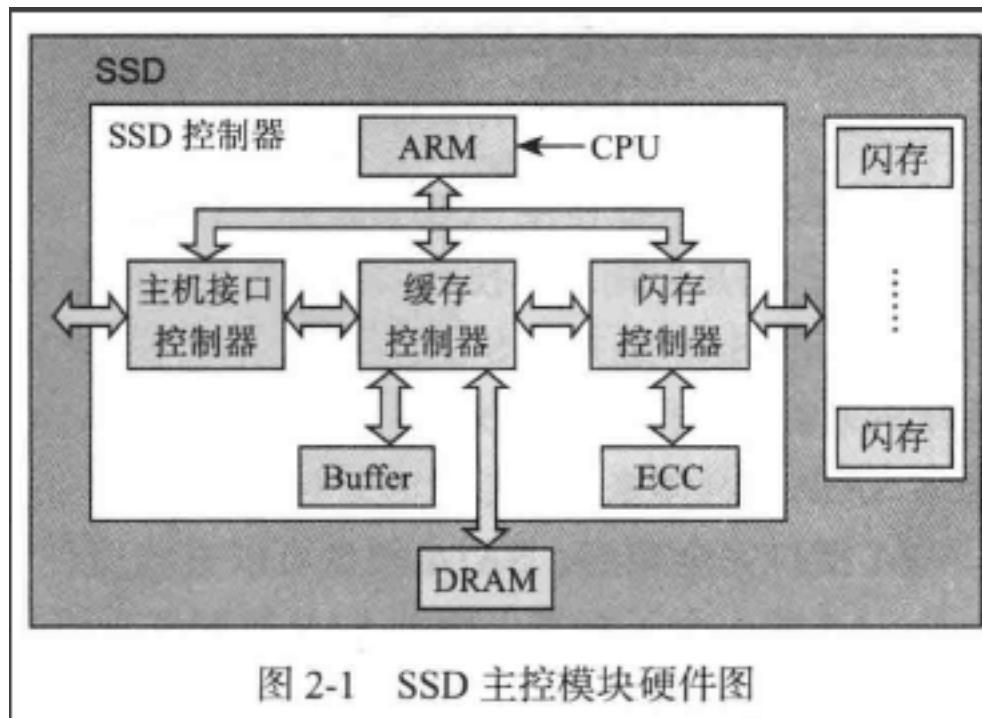


图 2-1 SSD 主控模块硬件图

SSD controller      FTL nand

PCIE/SATA host      FTL

FTL flash translation layer      LBA      SSD

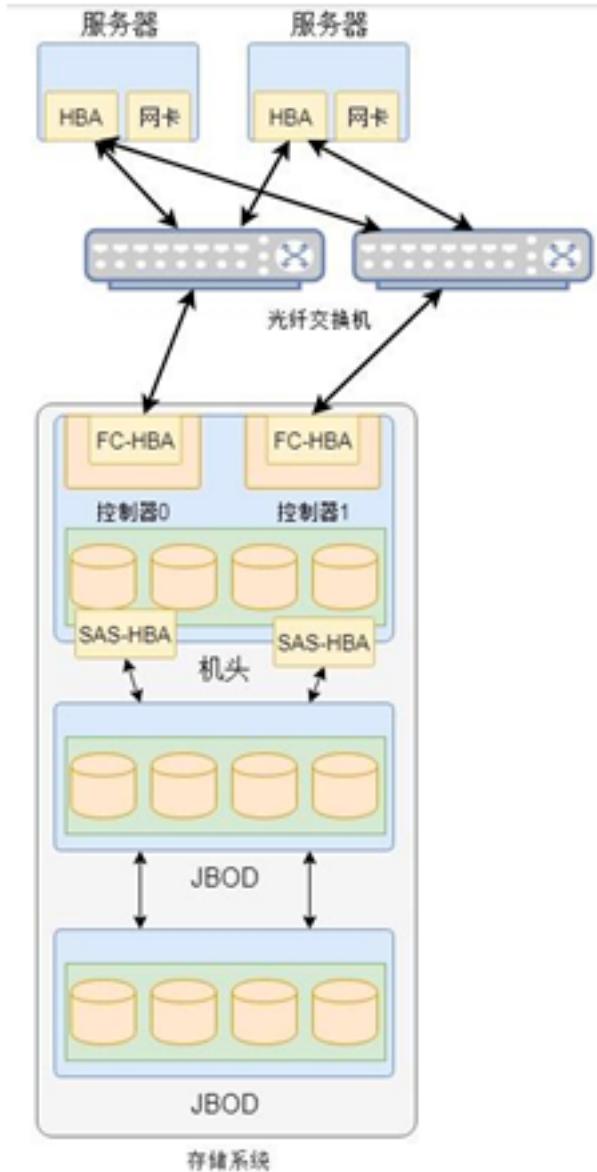
flash controller      ECC

#### 1.5.4

[toc]

( ) ( JBOD )

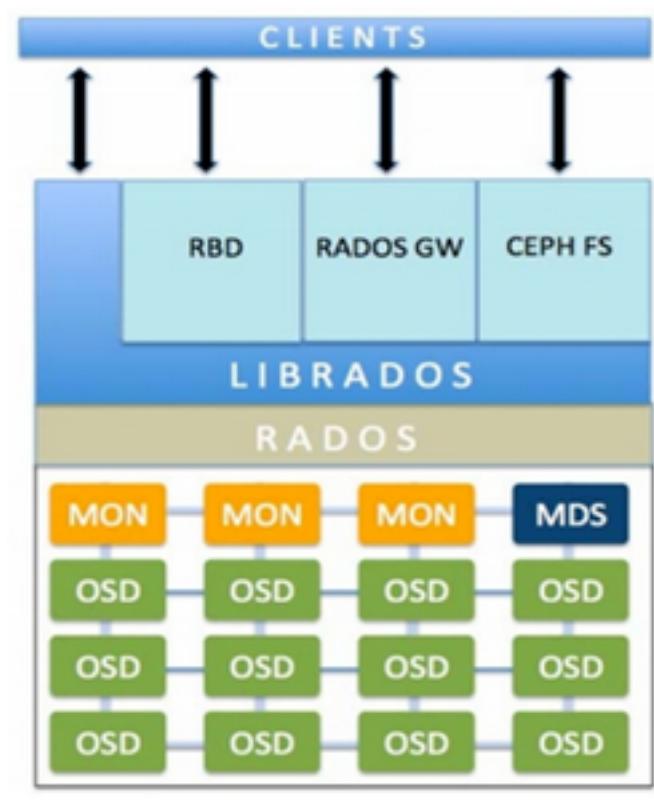
LUN



Web

—  
Ceph

HDFS



Ceph	MON	OSD	MDS
(1) MON		MON	
(2) OSD		OSD	
(3) MDS	CephFS	Ceph	RADOS
(4) RADOS	RADOS	ceph	Ceph
(5) RBD	( )	RADOS	
(6) CephFS	Ceph	Ceph	POSIX
(7) Librados	libRADOS	PHP	RUBY Java Python C++ RADOS
(8) RADOS	GW RGW	Ceph	RGW Amazon S3 openstack Swift RUSTFUL API
	RADOS GW	MON	( ) CephFS

Blog

1.6

### 1.6.1 Hypervisor

[toc]

## Introduction

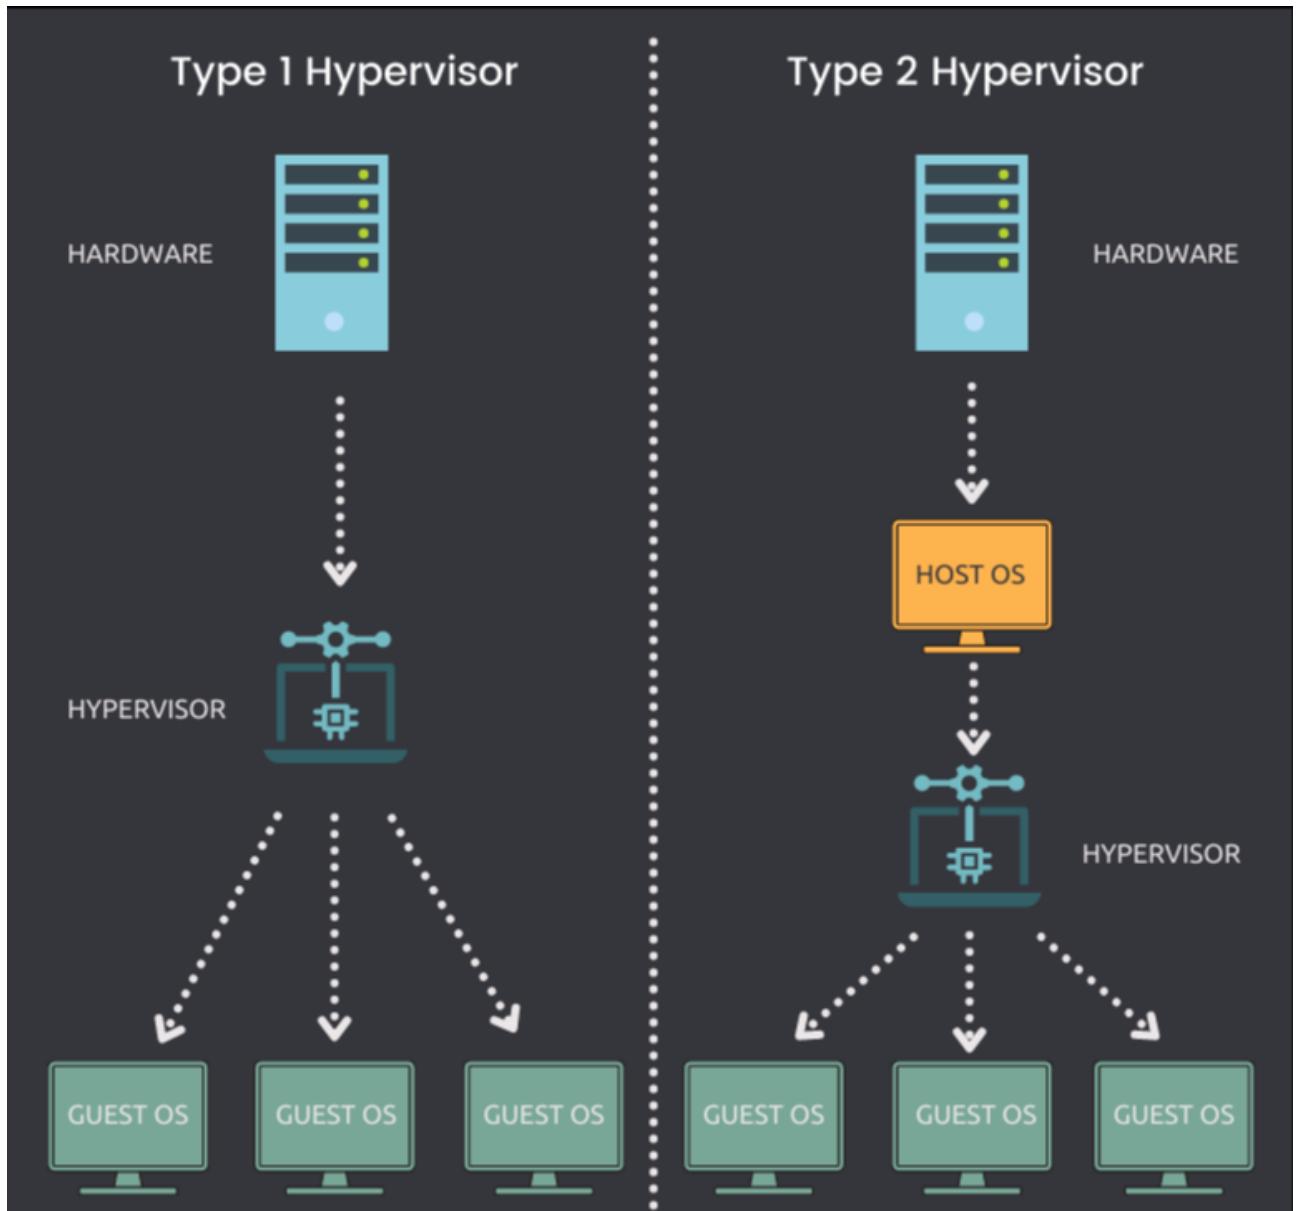
## Type 1 Hypervisor      Hypervisor Type 1 Hypervisor

- Type 1 Hypervisor
- (OS)
- CPU
- 
- 
- Type 1 Hypervisor
- 
- VMware ESXi Citrix Hypervisor Microsoft Hyper-V Type 1 Hypervisor

## Type 2 Hypervisor

2

- Type 2 hypervisor
- 
- CPU
- (OS)
- 
- Type 2 hypervisor
- 
- VMware Workstation Player VMware Workstation Pro VirtualBox Type 2 hypervisor



### RISC-V Type-1 Hypervisor

RVirt	RVirt	MIT PDOS	2018	Rust	<b>RISC-V Type-1 hypervisor</b>	RVirt	RISC-V
MIT PDOS		RVirt	S-mode	U-mode	CSR	hypervisor	hypervisor
QEMU virt	Hifive	Unleashed		fedora	RVirt	M S U	RISC-V
<b>Bao-hypervisor</b>	Bao	hypervisor			hypervisor		Bao
Bao							hypervisor
Bao	hypervisor			Bao	Bao		
Bao	RISC-V		Hypervisor Extension	rocket chip	CVA6		
<b>Xvisor</b>	Xvisor			Xvisor		VirtIO	
CPU	IO		Xvisor	Normal vCPUs			

2019	Xvisor	RISC-V	RISC-V	RISC-V Hypervisor Extension 1.0	Xvisor	Xvisor
<b>seL4</b>	seL4		NICTA	CSIRO data61	seL4	seL4
seL4						
•	seL4					
•	seL4					
•	seL4					
•	seL4	GPLv2				
•		seL4				
•	seL4	RISC-V				

<b>KVM</b>	KVM Kernel-based Virtual Machine	Linux	VM	VM	KVM	Intel
VT AMD-V	VM					
KVM		Red Hat IBM Intel KVM	Linux	Windows BSD	KVM	virt-manager virsh
KVM		Linux	Linux	Linux	Linux	

## RISC-V Type-1 Hypervisor

### 1.6.2 MIT 6.824 Distributed Systems

6.824 MIT                    Golang

B

<https://pdos.csail.mit.edu/6.82>

## 1.7 Memory consistency and cache coherence

### 1.7.1 Memory Consistency

[toc]

152	consistency model	152	252
	“ ”		
	consistency model		FENCE

```
static int x = 0, y = 0;
static int r1, r2;

static void int thread_cpu0(void)
{
    x = 1;      /* 1) */
    y = 2;      /* 2) */
    r1 = x;     /* 3) */
    r2 = y;     /* 4) */
}
```

```

r1 = y; /* 2) */

}

static void int thread_cpu1(void)
{
    y = 1; /* 3) */
    r2 = x; /* 4) */
}

static void check_after_assign(void)
{
    printk("r1 = %d, r2 = %d\n", r1, r2);
}

4           6

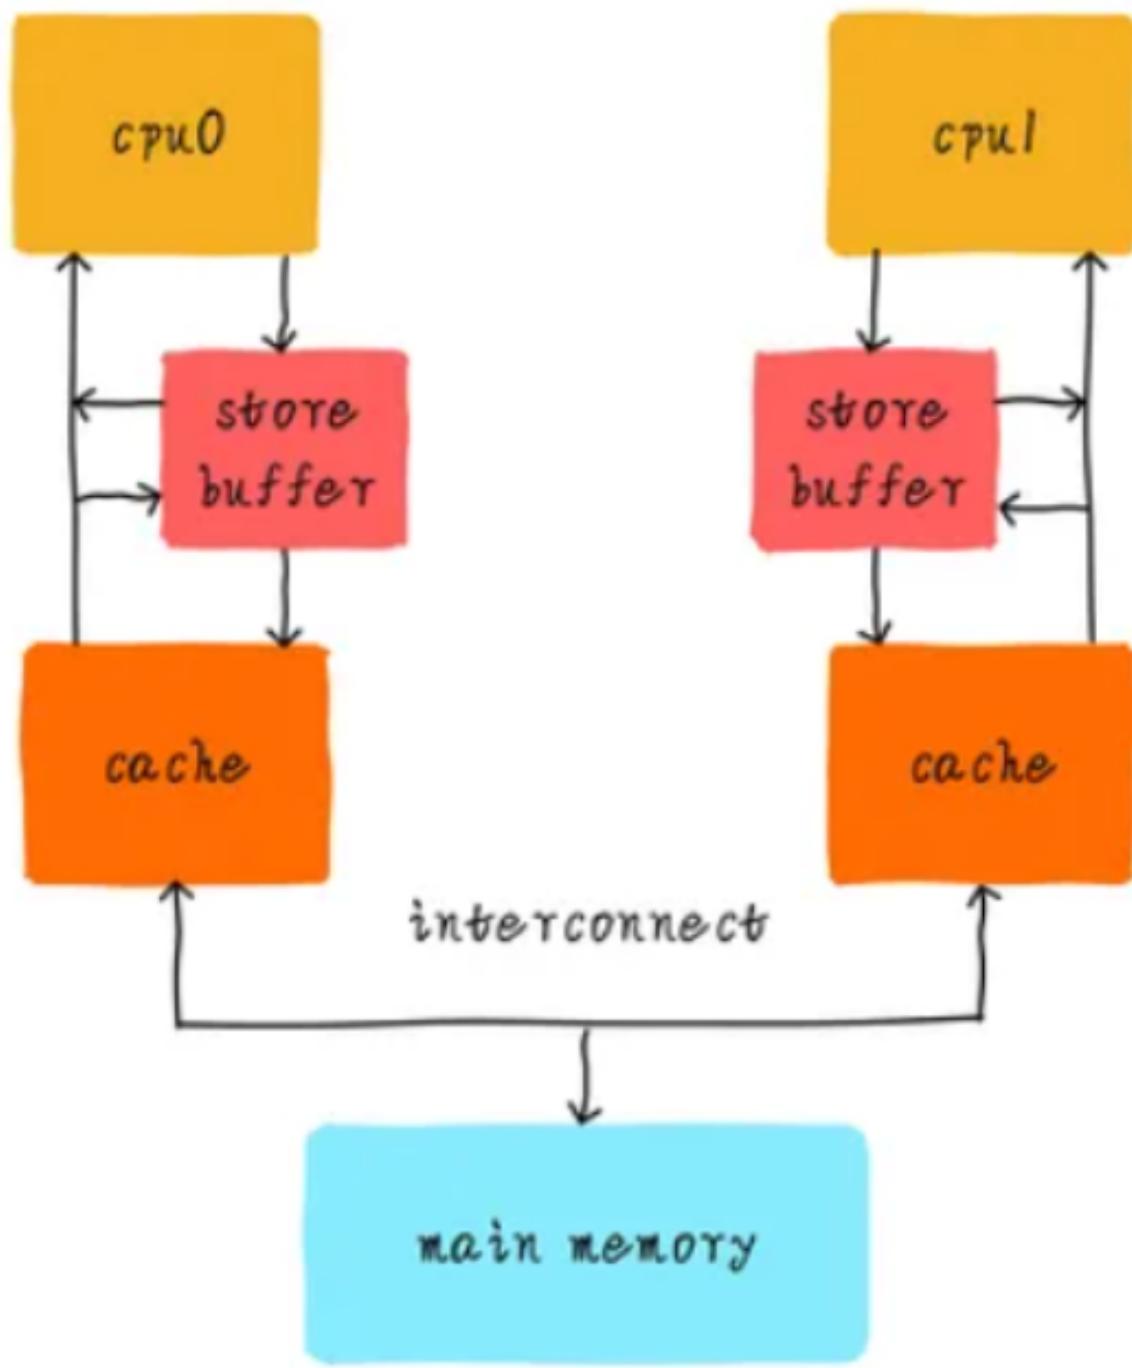
1) 2) 3) 4)
1) 3) 2) 4)
1) 3) 4) 2)
3) 4) 1) 2)
3) 1) 4) 2)
3) 1) 2) 4)

thread_cpu0 thread_cpu1      check_after_assign()      6      3

r1 = 1, r2 = 1
r1 = 0, r2 = 1 // 1) 2) 3) 4)
r1 = 1, r2 = 0 // 3) 4) 1) 2)

r1   r2          " "        write buffer prefetching  multiple cache banks
cache       write buffer      cache   cache

```



write buffer      1) 3) 2) 4) case      write buffer      core      0

$r1 = 0, r2 = 0$

3) 4) 1) 2)

\*\*

\*\* TSO

Memory Consistency Model

“ ” —— core 0 write buffer core1

memory consistency and cache coherence

cache CPU

CPU

4 Store-Load Store-Store Load-Store Load-Load

Sequential Consistency SC	Total	Store Order TSO	x86-64	Part	Store Order PSO
Relax Memory Order RMO					

### 1. Sequential Consistency, SC

- 
- 

1 → 3 → 2 → 4

3 → 1 → 4 → 2

1 → 3 → 4 → 2 //

3 → 1 → 2 → 4

- 
- 

1 → 3 → 4 → 2

1. 1: x = 1 CPU 0
2. 3: y = 1 CPU 1
3. 4: r2 = x CPU 1
4. 2: r1 = y CPU 0

CPU 0 CPU 1

- CPU 0 1 → 2
- CPU 1 3 → 4

1 → 3 → 4 → 2

- 4 x 1 r2 = 1
- 2 y 3 r1 = 1

2. Total Store Order, TSO TSO  
load load-store TSO store-load 3

store load 4 store-store store-load load-

•

```
1 -> 3 -> 2 -> 4  
3 -> 1 -> 4 -> 2  
1 -> 2 -> 3 -> 4  
3 -> 4 -> 1 -> 2
```

### 3. Partial Store Order, PSO PSO

## Processor Consistency, PC      Relaxed Consistency, RC

10

```
1 -> 3 -> 2 -> 4  
3 -> 1 -> 4 -> 2  
1 -> 2 -> 3 -> 4  
3 -> 4 -> 1 -> 2  
2 -> 1 -> 4 -> 3  
4 -> 3 -> 2 -> 1
```

#### 4. Relaxed Memory Order, RMO RMO

2

```
1 -> 3 -> 2 -> 4  
3 -> 1 -> 4 -> 2  
1 -> 2 -> 3 -> 4  
3 -> 4 -> 1 -> 2  
2 -> 1 -> 4 -> 3  
4 -> 3 -> 2 -> 1  
2 -> 4 -> 1 -> 3  
4 -> 2 -> 3 -> 1
```

	SC	TSO	PSO	RMO
Store-Store			Yes	Yes
Store-Load		Yes	Yes	Yes
Load-Load				Yes
Load-Store				Yes

`r1 = 0, r2 = 0`      Linux    `smp_mb()`      `smp_mb()`

- [A Primer on Memory Consistency and Cache Coherence \(2nd Edition\)](#)
- [https://blog.csdn.net/qq\\_29328443/article/details/107616795](https://blog.csdn.net/qq_29328443/article/details/107616795)
- <https://blog.csdn.net/anyegongjuezjd/article/details/125954805>
- [https://zhuanlan.zhihu.com/p/141655129?from\\_voters\\_page=true](https://zhuanlan.zhihu.com/p/141655129?from_voters_page=true)
-

# Chapter 2

## 2.1 ARM

### core

1. BMC C2000 CHIP POR\_N pervsoc
2. otprom valid SCP SRAM BISR PVO(SCP) ROM otprom
3. SCP SRAM
4. SCP ROM FW GPIO CHIP ID, CHIP ID CHIP;
5. I2C/FSI
6. CHIP SCP QSPI;
7. CHIP SCP CHIP CHIP
8. CHIP SCP CHIP SCP
9. CHIP SCP QSPI SCP RAM FW CHIP SRAM;
10. CHIP SCP 8 SCP RAM FW SCP, SCP SRAM:
11. CHIP SCP SRAM, SCP RAM FW:
12. CHIP SCP QSPI C2C SCP
13. CHIP SCP CHIP SCP C2C C2C;
14. SCP CHIP ID CMN OCM(SLC) SAM;
15. CHIP SCP MCP RAM FW, MCP( ) +
16. SCP CHIP L2 NCU SCOM/FIR;+
17. CHIP SCP Hostboot base CHIP OCM(SLC);+
18. CHIP SCP CHIP
19. srest 0x100

### core-hostboot

1. kernel other thread spinlock, thread
2. hostboot kernel C/C++ main +

3. hostboot kernel
4. MMU(SLB/TLB), DIMM OCM(SLC, 28 \* 2MB=56MB 3/4
5. TB CPU
6. Scratch FSP/BMC;
7. CPU init main
8. kernel other thread spinlock 1 thread
9. hostboot task task user kernel hypersior
10. init main VFS;
11. VFS base modules

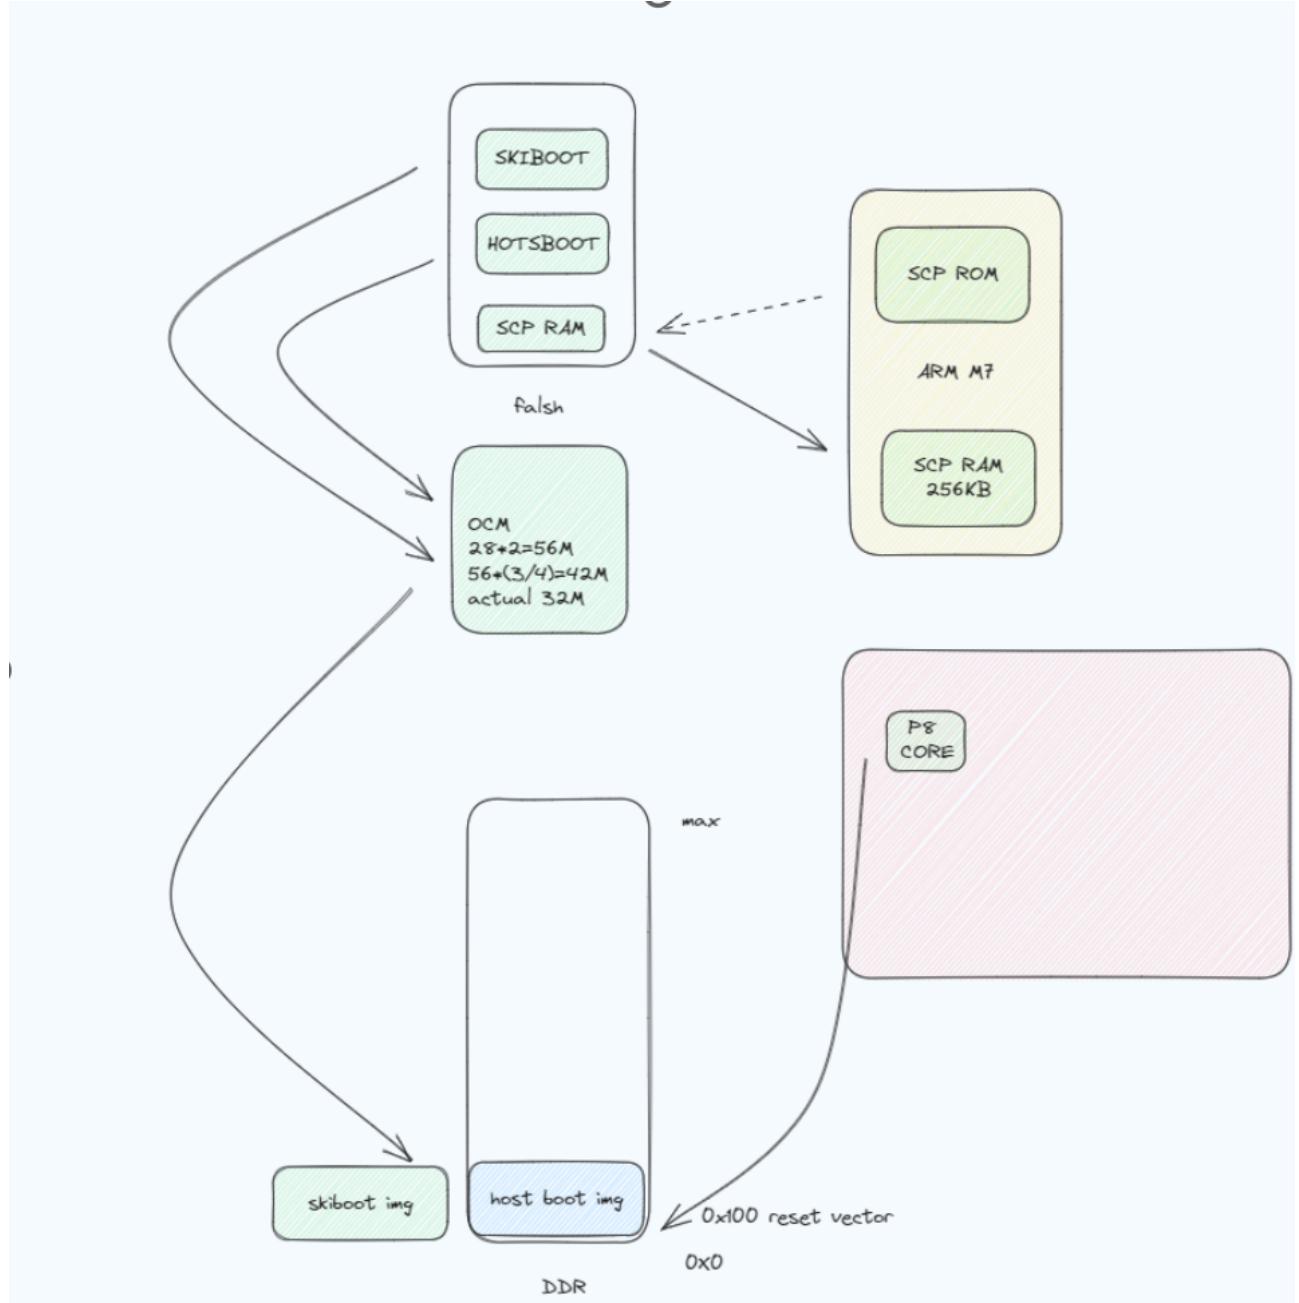
init

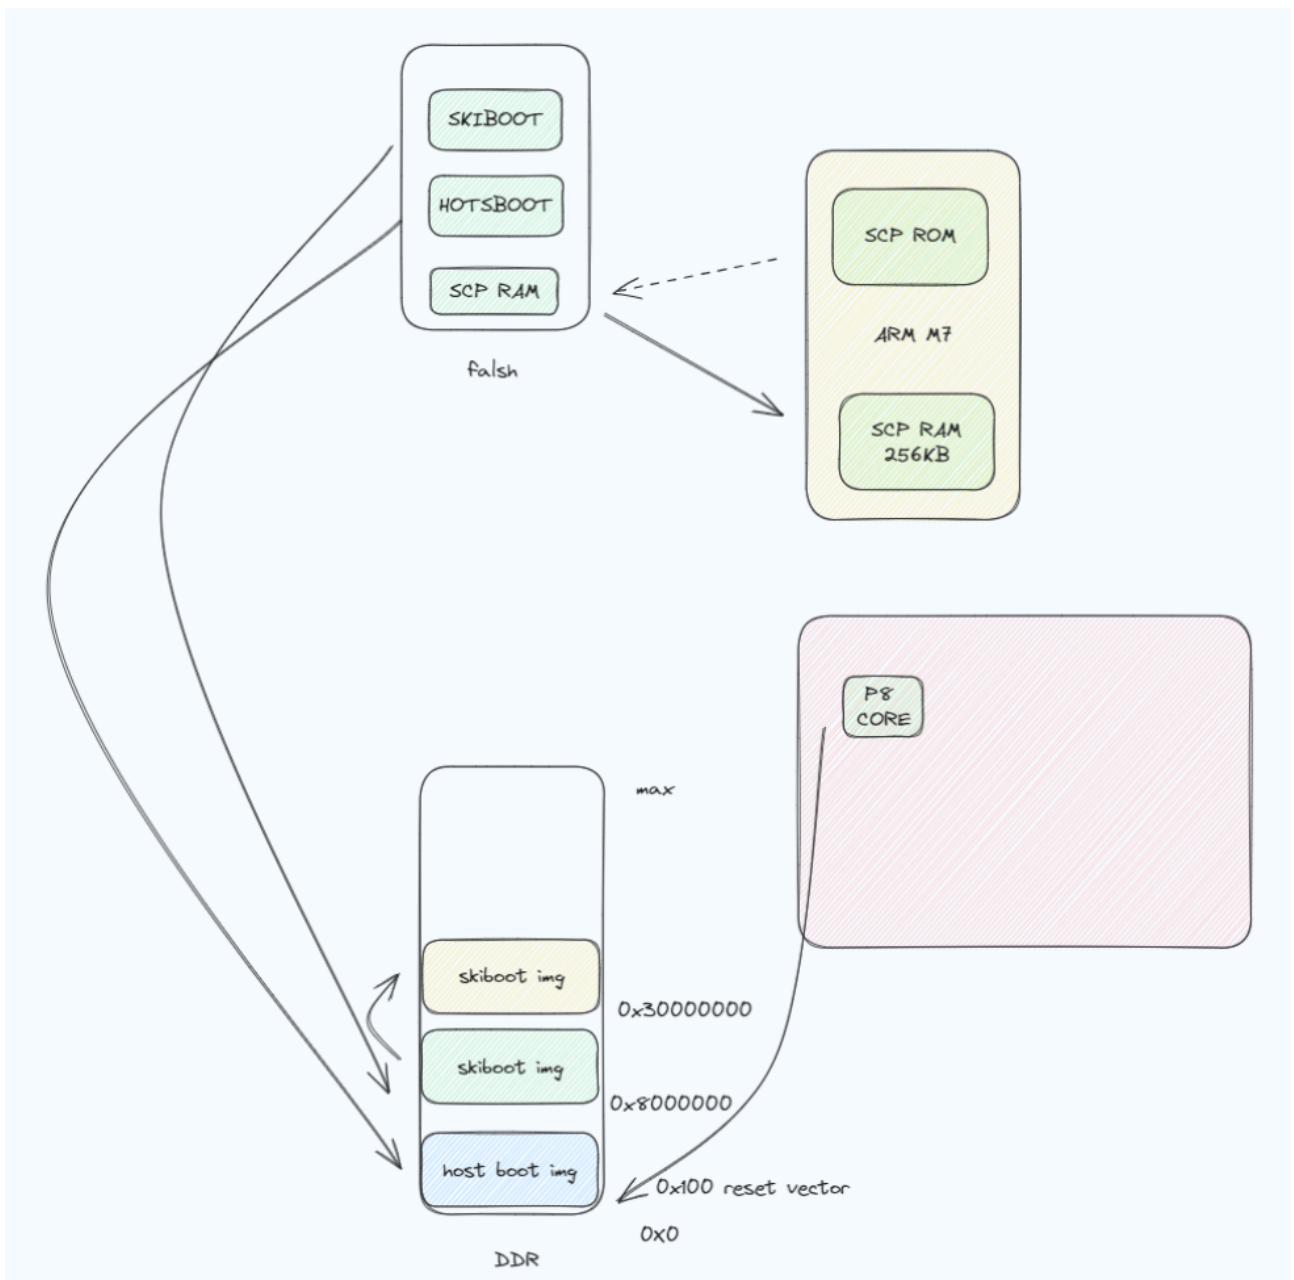
12. initservice
13. initservice g\_taskinfolist
14. secureboot\_base/ipmibase/pnor/vfs/extinitsvc
15. hostboot extended image
16. extinitsvc g\_exttaskinfolist
17. istepdisp hostboot
18. istep hostboot
19. istep 21.03 payload(skiboot,OPAL);+
20. payload(skiboot,OPAL)

## core-skiboot

1. C2000 skiboot
2. thread skiboot relocate thread
3. thread boot flag=1, thread
4. thread main cpu entry, thread secondary cpu entry;
5. uart(console);
6. skiboot /
7. skiboot
8. opal call table;
9. hostboot device tree HDAT;
10. DT SCOM/eSPI(LPC):
11. DT C2000 DIMM
12. DT homer C2000
13. DT CPU(thread);
14. opal DT;

- 15.
16.       BMC ipmi/BT/UART;
17.       GIC
18.      CPU thread
19. sreset CPU
20. TOD TOD
21. sensor OPAL
- 22.
- 23.NV
24.   jiayu
25.   TPM
26. opal console(UART)
27. PCIE HB   PCIE   (RC   CXL PCIe  device tree    )
- 28.PCIE SLOT
- 29.
- 30.
31.      zImage(boot loader  petitboot);
32. boot loader





CMN OCM(on chip memory) L3 cache DDR      Intel cache as ram

OCM memory      CMN OCM(on chip memory)      28\*2=56M      3/4      32M

SCP ram      DDR training      OCM      flash      DDR

OCM      SNF

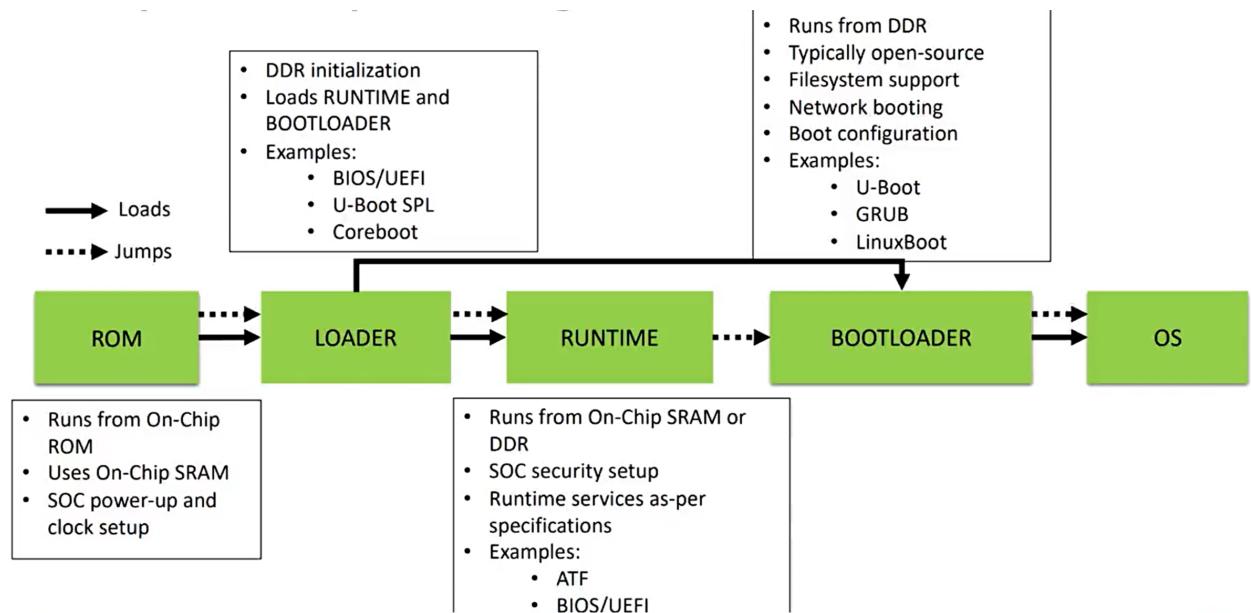
SCP      OCM      SCP RAM hostboot      OCM      skiboot      OCM      skiboot      ddr training      DDR      OCM      OCM      cache

hostboot skiboot 0x800 0000

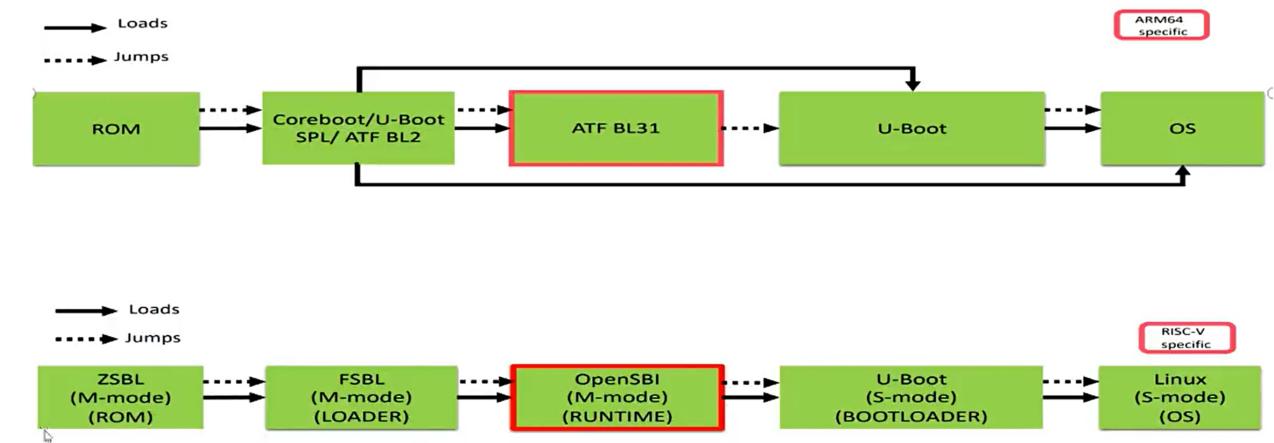
32M      0x200 0000, skiboot      0x3000 0000      OCM      0x0      0x0

## 2.2 RISC-V Firmware

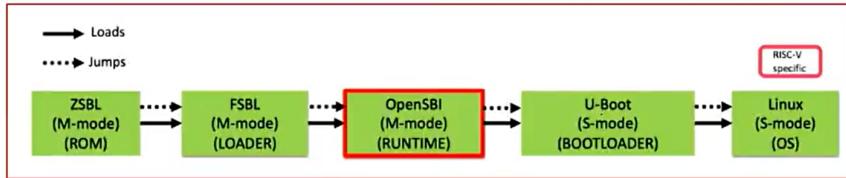
### Firmware



### ARM64 vs RISC-V



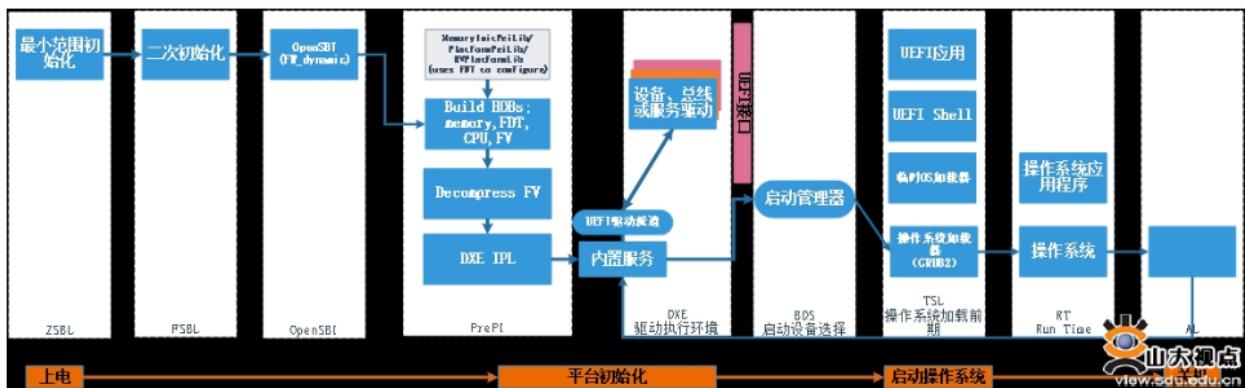
RISC-V linux启动流程如图，启动流程：ZSBL-->FSBL-->OpenSBI-->u-boot (UEFI)-->Linux



1. ZSBL (第0阶段Bootloader) at 0x00010000 (ROM in SOC)
2. FSBL (第1阶段Bootloader)
3. OpenSBI (runtime services) M模式
4. U-boot/UEFI (bootloader),
5. Linux启动linux需要以下文件：
  - a. /boot/extlinux/extlinux.conf: 包含启动菜单、默认启动等
  - b. Image: Linux镜像
  - c. Initrd: Linux执行需要的ramdisk, 执行其中的init文件并进行进一步处理
  - d. 设备树(.dtb) (可选) :未指定时会使用默认的dtb

ZSBL: SCP ROM

FSBL: SCP RAM



## 2.3 Cache as RAM and On Chip Memory

x86 Cache as RAM, CAR;

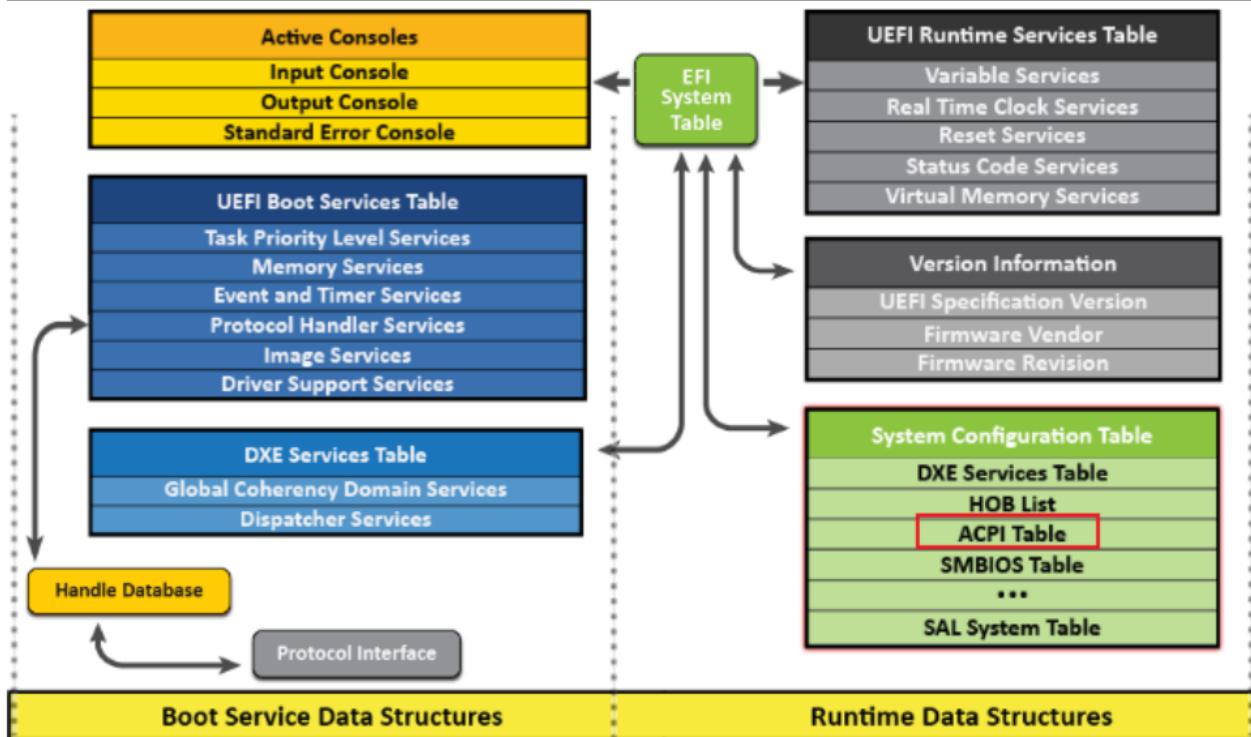
ARM On Chip Memory OCM

DDR cache RAM C

CAR

OCM

## 2.4 System Table



MdeModulePkg\Core\DXe\DXeMain\DXeMain.c

```
//  
// Allocate the EFI System Table and EFI Runtime Service Table from EfiRuntimeServicesData  
// Use the templates to initialize the contents of the EFI System Table and EFI Runtime  
// Services Table  
//  
gDxeCoreST = AllocateRuntimeCopyPool (sizeof (EFI_SYSTEM_TABLE),  
    &mEfiSystemTableTemplate);  
ASSERT (gDxeCoreST != NULL);  
  
gDxeCoreRT = AllocateRuntimeCopyPool (sizeof (EFI_RUNTIME_SERVICES),  
    &mEfiRuntimeServicesTableTemplate);  
ASSERT (gDxeCoreRT != NULL);  
  
gDxeCoreST->RuntimeServices = gDxeCoreRT;  
  
///  
/// EFI System Table  
///  
typedef struct {  
    ///  
    /// The table header for the EFI System Table.  
}
```

```

////
EFI_TABLE_HEADER          Hdr;
////
/// A pointer to a null terminated string that identifies the vendor
/// that produces the system firmware for the platform.
////
CHAR16                  *FirmwareVendor;
////
/// A firmware vendor specific value that identifies the revision
/// of the system firmware for the platform.
////
UINT32                  FirmwareRevision;
////
/// The handle for the active console input device. This handle must support
/// EFI_SIMPLE_TEXT_INPUT_PROTOCOL and EFI_SIMPLE_TEXT_INPUT_EX_PROTOCOL.
////
EFI_HANDLE               ConsoleInHandle;
////
/// A pointer to the EFI_SIMPLE_TEXT_INPUT_PROTOCOL interface that is
/// associated with ConsoleInHandle.
////
EFI_SIMPLE_TEXT_INPUT_PROTOCOL *ConIn;
////
/// The handle for the active console output device.
////
EFI_HANDLE               ConsoleOutHandle;
////
/// A pointer to the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL interface
/// that is associated with ConsoleOutHandle.
////
EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *ConOut;
////
/// The handle for the active standard error console device.
/// This handle must support the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL.
////
EFI_HANDLE               StandardErrorHandler;
////
/// A pointer to the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL interface
/// that is associated with StandardErrorHandler.
////
EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *StdErr;
////
/// A pointer to the EFI Runtime Services Table.
////
EFI_RUNTIME_SERVICES      *RuntimeServices;
////
/// A pointer to the EFI Boot Services Table.
////
EFI_BOOT_SERVICES         *BootServices;

```

```

/// 
/// The number of system configuration tables in the buffer ConfigurationTable.
///
UINTN NumberOfTableEntries;
///
/// A pointer to the system configuration tables.
/// The number of entries in the table is NumberOfTableEntries.
///
EFI_CONFIGURATION_TABLE *ConfigurationTable;
} EFI_SYSTEM_TABLE;

```

## 2.5 EDK2 EVENT

[toc]

UEFI

```

/ intel 8253/8254      PC      IC      MSC8051,
Timer          cpu
PC/AT UEFI     8254   Timer mode 3   1.1931816MHz tick mTimerPeriod    1ms 10000 tick =10000*100ns
pin 8259 IRQ0   CPU     Timer
cpu archprotocol Legacy8259Protocol ,   cpu IDT   Timer TimerInterruptHandler

```

### Timer

Timer

**\*\*1.\*\***

**\*\*2.\*\***

**\*\*3.\*\***

**\*\*4.\*\*** Timer CoreTimerTick(mTimerPeriod)/CoreTimerTick(100ms) TimerArchprotocol install EVT\_NOTIFY\_SIGNAL

**\*\*5.\*\*** System time mEfiSystemTime =mEfiSystemTime+ 1Tick(100ms)

**6.** timer event list event BS->SignalEvent signal

BS->SignalEvent-->CoreNotifyEvent-->CoreRestoreTpl-->CoreDispatchEventNotifies-->A. EVT\_NOTIFY\_SIGNAL  
>SignalCount = 0 Event->NotifyFunction (Event, Event->NotifyContext) Tpl gEventPending  
mask bit event

**\*\*7.\*\***

**\*\*8.\*\***

**1.** gEventQueue--- A list of event's to notify for each priority level

**2.** mEfiTimerList---- timer event

3. gEventSignalQueue --- A list of events to signal based on EventGroup type

4. EVT\_NOTIFY\_WAIT event BS-> CoreCheckEvent->CoreNotifyEvent-->CoreRestoreTpl->CoreDispatchEventNotifies gEventQueue IEvent

gEventPending gEventQueue 1 0

gEventPending 32 1

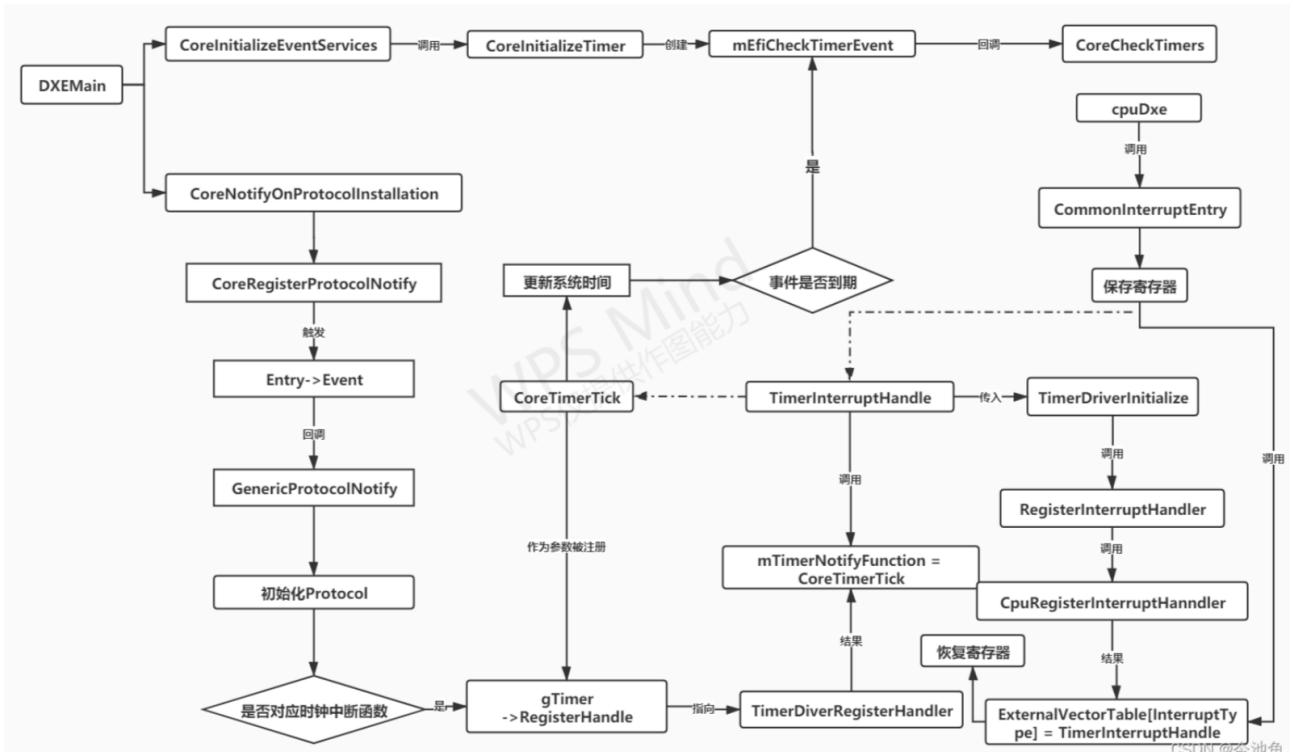
x86 IRQ

- IRQ0:
- IRQ1:
- IRQ2:
- IRQ3: 2 COM2
- IRQ4: 1 COM1
- IRQ5:
- IRQ6:
- IRQ7: 1 LPT1
- IRQ

event

```
//  
// Task priority level  
//  
#define TPL_APPLICATION 4  
#define TPL_CALLBACK 8  
#define TPL_NOTIFY 16  
#define TPL_HIGH_LEVEL 31
```

31



DXE main

```

// Initialize the Event Services
//
Status = CoreInitializeEventServices ();

...
// Register for the GIDs of the Architectural Protocols, so the rest of the
// EFI Boot Services and EFI Runtime Services tables can be filled in.
// Also register for the GIDs of optional protocols.
//
CoreNotifyOnProtocolInstallation ();

event

/**
Called by the platform code to process a tick.

@param Duration           The number of 100ns elapsed since the last call
                           to TimerTick
*/
VOID
EFIAPI
CoreTimerTick (

```

```

    IN UINT64    Duration
)
{
    IEVENT        *Event;

    //
    // Check runtiem flag in case there are ticks while exiting boot services
    //
    CoreAcquireLock (&mEfiSystemTimeLock);

    //
    // Update the system time
    //
    mEfiSystemTime += Duration;

    //
    // If the head of the list is expired, fire the timer event
    // to process it
    //

    if (!IsListEmpty (&mEfiTimerList)) {
        Event = CR (mEfiTimerList.ForwardLink, IEVENT, Timer.Link, EVENT_SIGNATURE);

        if (Event->Timer.TriggerTime <= mEfiSystemTime) {
            CoreSignalEvent (mEfiCheckTimerEvent);
        }
    }

    CoreReleaseLock (&mEfiSystemTimeLock);
}

timer event

/***
    Initializes "event" support.

    @retval EFI_SUCCESS           Always return success
*/
EFI_STATUS
CoreInitializeEventServices (
    VOID
)
{
    UINTN        Index;

    for (Index=0; Index <= TPL_HIGH_LEVEL; Index++) {
        InitializeListHead (&gEventQueue[Index]);
    }

    CoreInitializeTimer ();
}

```

```

CoreCreateEventEx (
    EVT_NOTIFY_SIGNAL,
    TPL_NOTIFY,
    EfiEventEmptyFunction,
    NULL,
    &gIdleLoopEventGuid,
    &gIdleLoopEvent
);

return EFI_SUCCESS;
}

```

- gEventQueue--- A list of event's to notify for each priority level
- CoreInitializeTimer

```

/***
Initializes timer support.
**/

```

```

VOID
CoreInitializeTimer (
    VOID
)
{
    EFI_STATUS Status;

    Status = CoreCreateEventInternal (
        EVT_NOTIFY_SIGNAL,
        TPL_HIGH_LEVEL - 1,
        CoreCheckTimers,
        NULL,
        NULL,
        &mEfiCheckTimerEvent
    );
    ASSERT_EFI_ERROR (Status);
}

```

```

#include <Uefi.h>
#include <Library/UefiBootServicesTableLib.h>
#include <Library/UefiLib.h>

```

```

EFI_EVENT TimerEvent;
EFI_EVENT ExitBootServicesEvent;
BOOLEAN TimerTriggered = FALSE;

```

```

VOID EFIAPI TimerEventHandler(

```

```

    IN EFI_EVENT Event,
    IN VOID *Context
) {
    TimerTriggered = TRUE;
    Print(L"Timer event triggered!\n");
}

EFI_STATUS
EFIAPI
UefiMain(
    IN EFI_HANDLE ImageHandle,
    IN EFI_SYSTEM_TABLE *SystemTable
) {
    EFI_STATUS Status;

    //      10000000 100ns (1)
    Status = gBS->CreateEvent(
        EVT_TIMER | EVT_NOTIFY_SIGNAL,
        TPL_CALLBACK,
        TimerEventHandler,
        NULL,
        &TimerEvent
    );
    if (EFI_ERROR(Status)) {
        Print(L"Failed to create event: %r\n", Status);
        return Status;
    }

    //      1
    Status = gBS->SetTimer(
        TimerEvent,
        TimerRelative,
        10000000 // 100ns 10000000 = 1
    );
    if (EFI_ERROR(Status)) {
        Print(L"Failed to set timer: %r\n", Status);
        gBS->CloseEvent(TimerEvent);
        return Status;
    }

    //
    while (!TimerTriggered) {
        gBS->Stall(100000); // 100ms
    }

    //
    gBS->CloseEvent(TimerEvent);

    return EFI_SUCCESS;
}

```

```

}

1

gBS->CreateEvent

typedef struct {
    UINT32 Type;
    EFI_TPL NotifyTpl;
    EFI_EVENT_NOTIFY NotifyFunction;
    VOID *NotifyContext;
    BOOLEAN Triggered;
} EFI_EVENT_INTERNAL;

EFI_STATUS
EFIAPI
CreateEvent (
    IN UINT32 Type,
    IN EFI_TPL NotifyTpl,
    IN EFI_EVENT_NOTIFY NotifyFunction,
    IN VOID *NotifyContext,
    OUT EFI_EVENT *Event
) {
    EFI_STATUS Status;
    EFI_EVENT_INTERNAL *NewEvent;

    //

    if (Event == NULL || (Type & EVT_NOTIFY_SIGNAL && NotifyFunction == NULL)) {
        return EFI_INVALID_PARAMETER;
    }

    //

    NewEvent = AllocateZeroPool(sizeof(EFI_EVENT_INTERNAL));
    if (NewEvent == NULL) {
        return EFI_OUT_OF_RESOURCES;
    }

    //

    NewEvent->Type = Type;
    NewEvent->NotifyTpl = NotifyTpl;
    NewEvent->NotifyFunction = NotifyFunction;
    NewEvent->NotifyContext = NotifyContext;
    NewEvent->Triggered = FALSE;

    //

    Status = RegisterEvent(NewEvent);
    if (EFI_ERROR(Status)) {
        FreePool(NewEvent);
        return Status;
    }
}

```

```

//  

*Event = (EFI_EVENT)NewEvent;  

return EFI_SUCCESS;  

}  

    event  

event  

if ((Type & EVT_NOTIFY_SIGNAL) != 0x00000000) {  

//  

// The Event's NotifyFunction must be queued whenever the event is signaled  

//  

InsertHeadList (&gEventSignalQueue, &IEvent->SignalLink);  

}  

32 bitmap 1

```

SignalEvent event 32

/\*\*  
*Signals the event. Queues the event to be notified if needed.*

**@param** UserEvent *The event to signal .*

**@retval** EFI\_INVALID\_PARAMETER *Parameters are not valid.*  
**@retval** EFI\_SUCCESS *The event was signaled.*

\*\*/  
EFI\_STATUS  
EFIAPI  
CoreSignalEvent (  
 IN EFI\_EVENT UserEvent  
)  
/\*\*  
*Queues the event's notification function to fire.*

**@param** Event *The Event to notify*

\*\*/  
VOID  
CoreNotifyEvent (  
 IN IEVENT \*Event  
)  
{  
//  
// Event database must be locked

```

//  

ASSERT_LOCKED (&gEventQueueLock);  
  

//  

// If the event is queued somewhere, remove it  

//  

if (Event->NotifyLink.ForwardLink != NULL) {  

    RemoveEntryList (&Event->NotifyLink);  

    Event->NotifyLink.ForwardLink = NULL;  

}  
  

//  

// Queue the event to the pending notification list  

//  

InsertTailList (&gEventQueue[Event->NotifyTpl], &Event->NotifyLink);  

gEventPending |= (UINTN)(1 << Event->NotifyTpl);  

}

```

Duang            CoreTimerTick            TPL            restore TPL

timer event mEfiTimerList

TPL

```

while (gEventPending != 0)
{
    PendingTpl = (UINTN) HighBitSet64 (gEventPending);
    if (PendingTpl <= NewTpl) {
        break;
    }
    gEfiCurrentTpl = PendingTpl;
    if (gEfiCurrentTpl < TPL_HIGH_LEVEL) {
        CoreSetInterruptState (TRUE);
    }
    CoreDispatchEventNotifies (gEfiCurrentTpl);
    // gEventQueue [gEfiCurrentTpl] event Notification
}

```

gEventPending 32      1

1.

2.        TPL    TPL                          16        17 16 TPL        17

3. pending      pending    TPL                  16        15        16

---

waitforevent    4        10ms                    31    RestoreTpl    4,        31 4    30 5        4

## CoreTimerTick

```
/**  
Lowers the task priority to the previous value. If the new  
priority unmasks events at a higher priority, they are dispatched.  
  
@param NewTpl New, lower, task priority  
  
**/  
VOID  
EFIAPI  
CoreRestoreTpl (  
    IN EFI_TPL NewTpl  
)  
  
{  
    EFI_TPL     OldTpl;  
    EFI_TPL     PendingTpl;  
  
    • OldTpl:  
    • PendingTpl:  
  
    OldTpl = gEfiCurrentTpl;  
    if (NewTpl > OldTpl) {  
        DEBUG ((EFI_D_ERROR, "FATAL ERROR - RestoreTpl with NewTpl(0x%x) > OldTpl(0x%x)\n",  
        ← NewTpl, OldTpl));  
        ASSERT (FALSE);  
    }  
    ASSERT (VALID_TPL (NewTpl));  
  
    • OldTpl = gEfiCurrentTpl:  
    • NewTpl           NewTpl OldTpl  
  
  
    if (OldTpl >= TPL_HIGH_LEVEL && NewTpl < TPL_HIGH_LEVEL) {  
        gEfiCurrentTpl = TPL_HIGH_LEVEL;  
    }  
  
    •           TPL_HIGH_LEVEL           TPL_HIGH_LEVEL           TPL_HIGH_LEVEL  
  
  
    while (gEventPending != 0) {  
        PendingTpl = (UINTN) HighBitSet64 (gEventPending);  
        if (PendingTpl <= NewTpl) {  
            break;  
        }  
    }  
}
```

```

gEfiCurrentTpl = PendingTpl;
if (gEfiCurrentTpl < TPL_HIGH_LEVEL) {
    CoreSetInterruptState (TRUE);
}
CoreDispatchEventNotifies (gEfiCurrentTpl);
}

• while (gEventPending != 0):
• PendingTpl = (UINTN) HighBitSet64 (gEventPending):
• PendingTpl NewTpl
• PendingTpl TPL_HIGH_LEVEL
• CoreDispatchEventNotifies

```

```

gEfiCurrentTpl = NewTpl;

if (gEfiCurrentTpl < TPL_HIGH_LEVEL) {
    CoreSetInterruptState (TRUE);
}

}

• NewTpl
• TPL_HIGH_LEVEL

```

### **CoreDispatchEventNotifies**

```

/***
Dispatches all pending events.

@param Priority           The task priority level of event notifications
                           to dispatch
**/>

```

```

VOID
CoreDispatchEventNotifies (
    IN EFI_TPL      Priority
)

```

```

{
    IEVENT        *Event;
    LIST_ENTRY    *Head;

    • Event:
    • Head:

```

```
CoreAcquireEventLock ();
```

```
ASSERT (gEventQueueLock.OwnerTpl == Priority);
```

```
Head = &gEventQueue[Priority];
```

- CoreAcquireEventLock ():
- ASSERT (gEventQueueLock.OwnerTpl == Priority):
- Head = &gEventQueue[Priority]:

```
while (!IsEmpty (Head)) {
```

```
    Event = CR (Head->ForwardLink, IEVENT, NotifyLink, EVENT_SIGNATURE);
```

```
    RemoveEntryList (&Event->NotifyLink);
```

```
    Event->NotifyLink.ForwardLink = NULL;
```

- while (!IsEmpty (Head)):
- Event = CR (Head->ForwardLink, IEVENT, NotifyLink, EVENT\_SIGNATURE):
- RemoveEntryList (&Event->NotifyLink):
- Event->NotifyLink.ForwardLink = NULL:

## SIGNAL

```
if ((Event->Type & EVT_NOTIFY_SIGNAL) != 0) {
```

```
    Event->SignalCount = 0;
```

```
}
```

- EVT\_NOTIFY\_SIGNAL SignalCount

```
CoreReleaseEventLock ();
```

```
ASSERT (Event->NotifyFunction != NULL);
```

```
Event->NotifyFunction (Event, Event->NotifyContext);
```

- CoreReleaseEventLock ():
- ASSERT (Event->NotifyFunction != NULL):
- Event->NotifyFunction (Event, Event->NotifyContext):

```
CoreAcquireEventLock ();
```

```
}
```

- CoreAcquireEventLock ():

```
gEventPending &= ~(UINTN)(1 << Priority);
```

```
CoreReleaseEventLock ();
```

```
}
```

- gEventPending &= ~(UINTN)(1 << Priority):
- CoreReleaseEventLock ():

- 31TPL
- 

[Blog](#)

[Blog](#)

## 2.6 EDK2 MEMORY MAP

GetMemoryMap

<https://blog.csdn.net/xiaopangzi313/article/details/109928878>

<https://www.lab-z.com/stu5/>

# Chapter 3

# OS

## 3.1 FreeRTOS

[TOC]

### 1. Determinism

RTOS                    RTOS

### 2. Low Latency

RTOS

### 3. Priority Scheduling

RTOS                    CPU                    Priority Preemptive Scheduling            Round-Robin Scheduling

### 4. Interrupt Handling

RTOS                    ISR

### 5. Time Management

RTOS

### 6. Resource Management

RTOS

## Linux

1.

- Linux
- Linux

2.

- Linux CPU Linux SCCHED\_FIFO SCCHED\_RR
- Linux

3.

- Linux
- **Swapping** Linux

4.

- Linux
- 

5.

- Linux CPU
- 

6.

- vs Linux Linux

## **Part II**

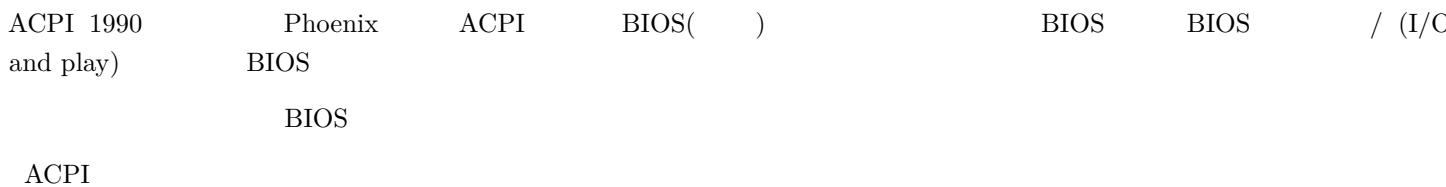
# Chapter 4

## ACPI

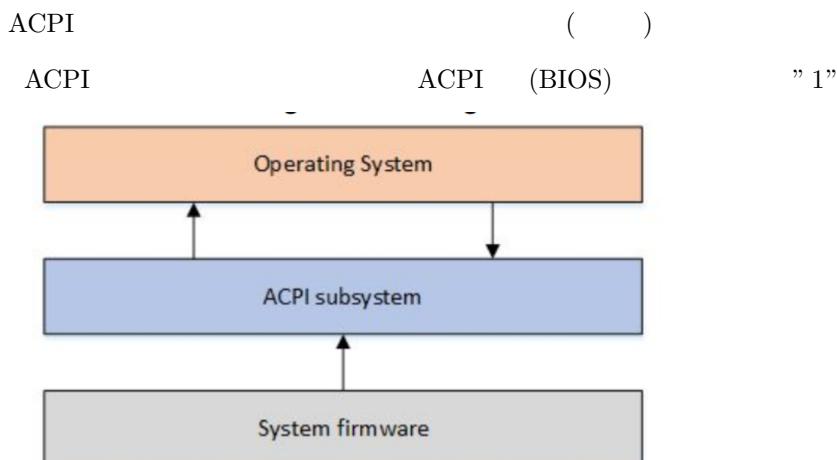
### 4.1 ACPI Introduction and Overview

[toc]

#### History of ACPI



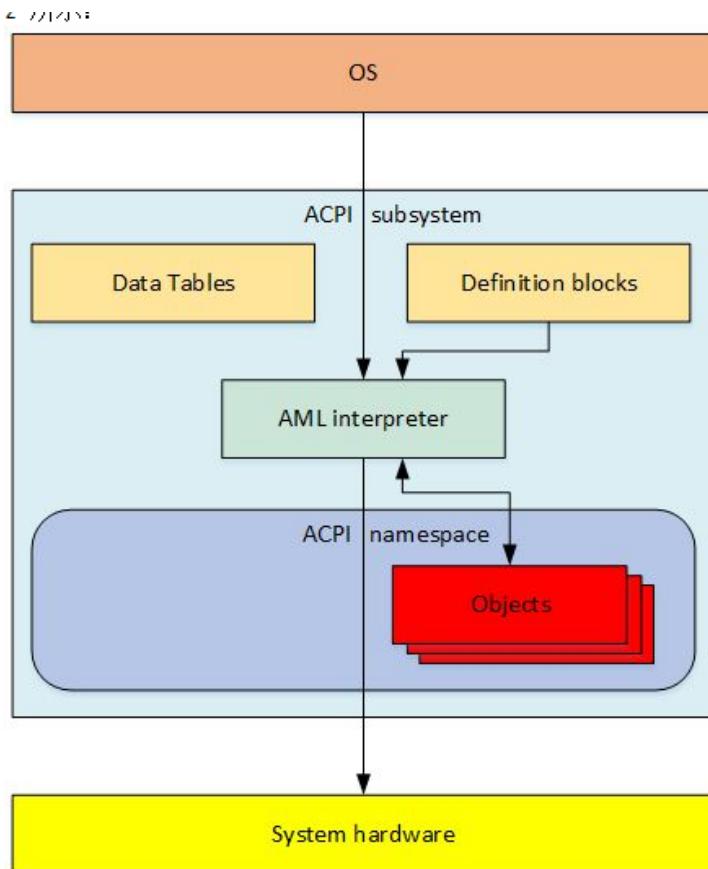
#### What is ACPI?



The *ACPI subsystem* is an interface layer between the *System firmware* and the *OS*. The arrows indicate data flow.

Figure 1: ACPI overview

ACPI (data table) (definition block) " 2 "



ACPI子系统由两类数据结构组成：  
数据表(data table)和定义块(definition block)

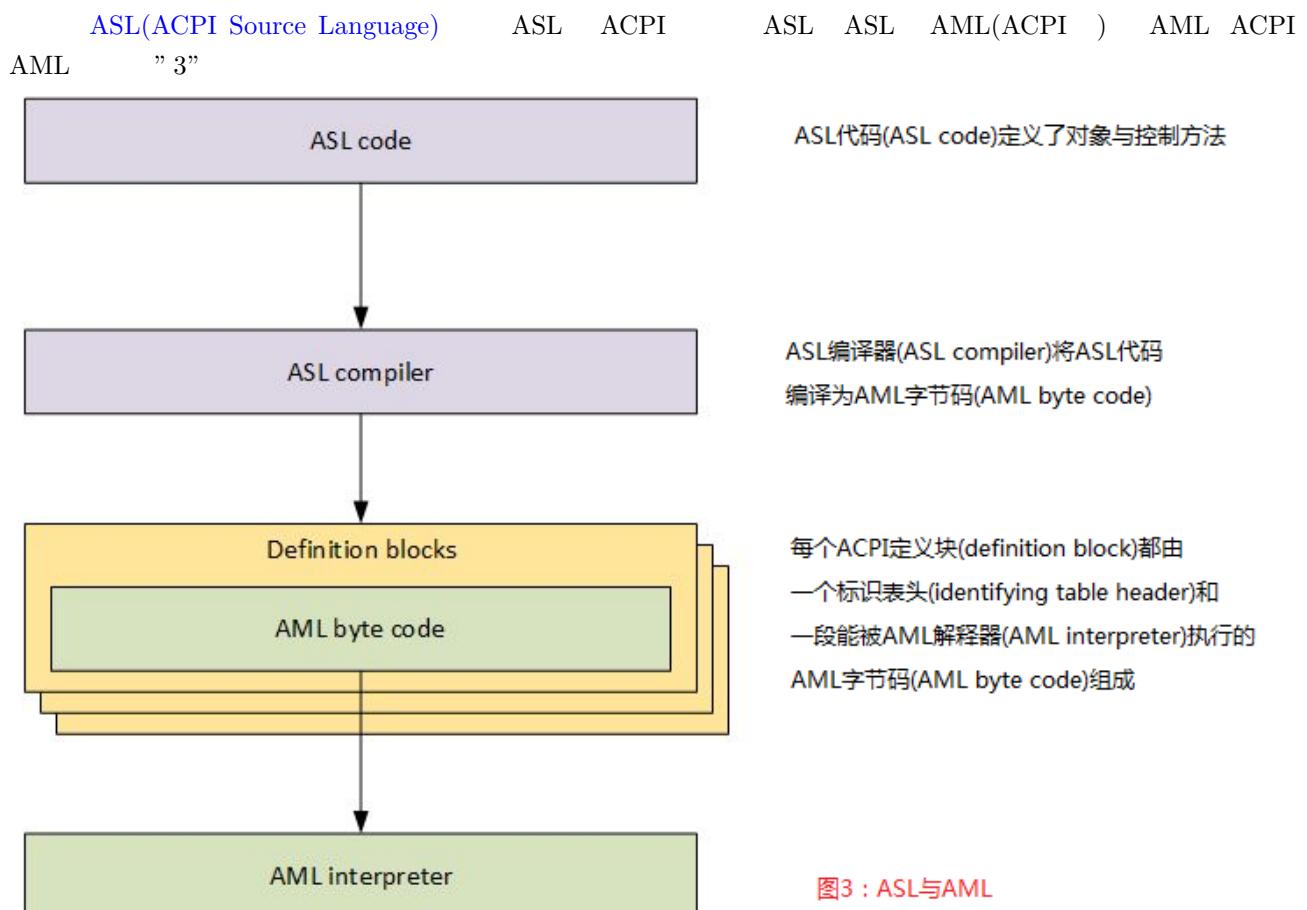
初始化时，AML解释器(AML interpreter)  
将定义块中的字节码提取为可枚举对象

可枚举对象的集合最终构成了操作系统内的  
“ACPI名字空间”(ACPI namespace)

可枚举对象(Object)  
要么拥有一个直接定义的值，  
要么必须由AML解释器执行

由操作系统控制的AML解释器，  
按照可枚举对象的指引与硬件交互，  
让硬件(System hardware)执行各种操作

图2：ACPI结构



ASL代码(ASL code)定义了对象与控制方法

ASL编译器(ASL compiler)将ASL代码  
编译为AML字节码(AML byte code)

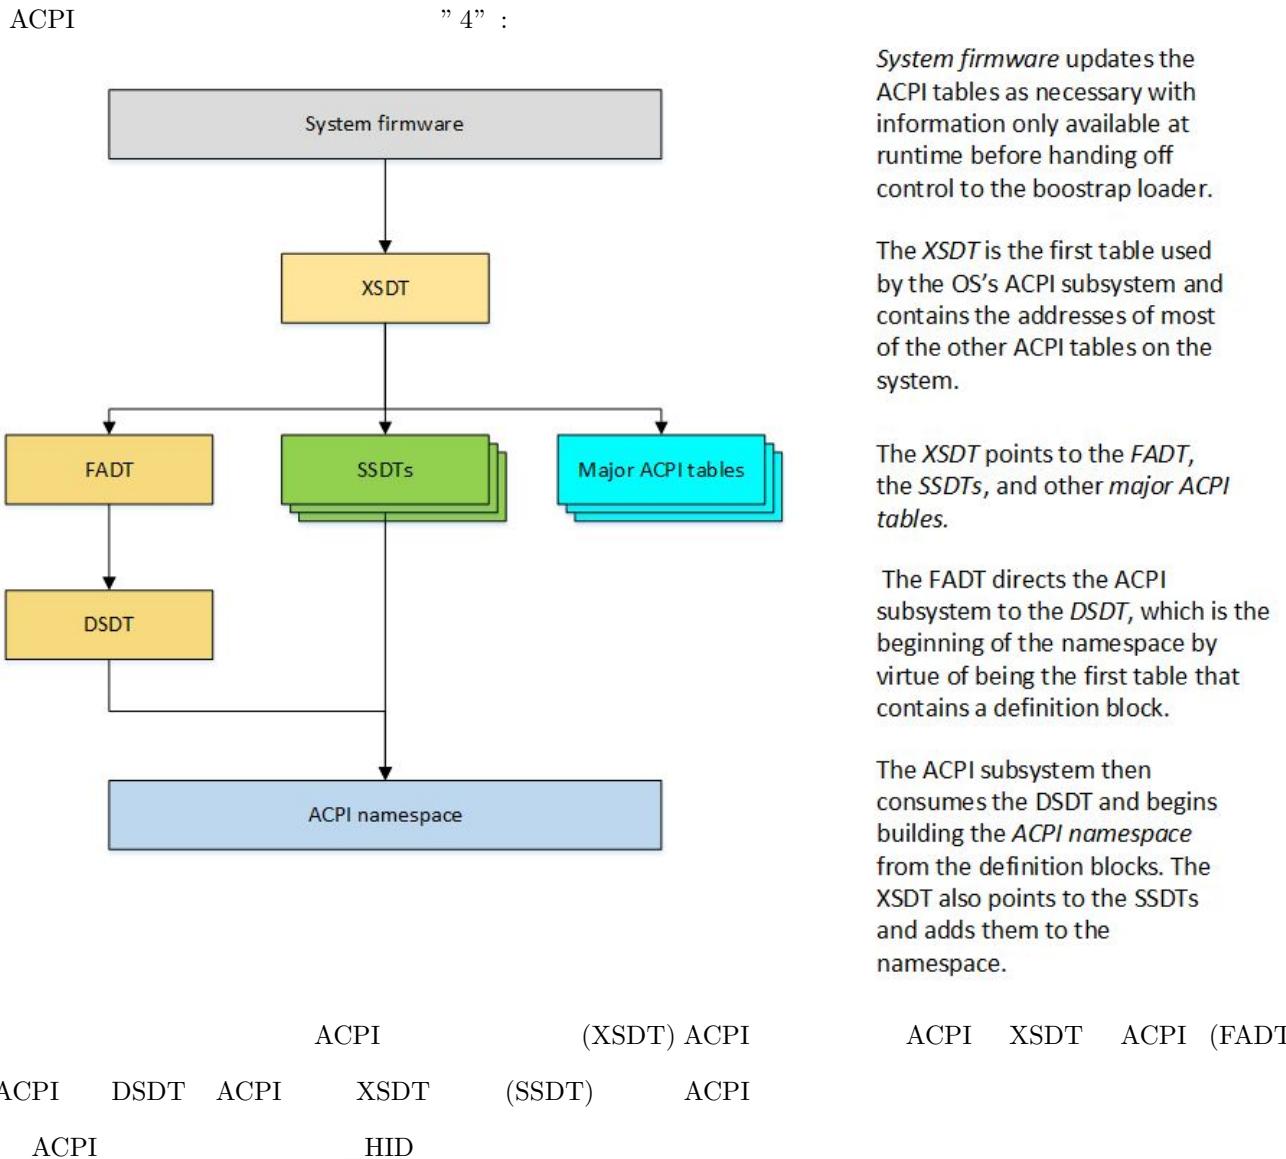
每个ACPI定义块(definition block)都由  
一个标识表头(identifying table header)和  
一段能被AML解释器(AML interpreter)执行的  
AML字节码(AML byte code)组成

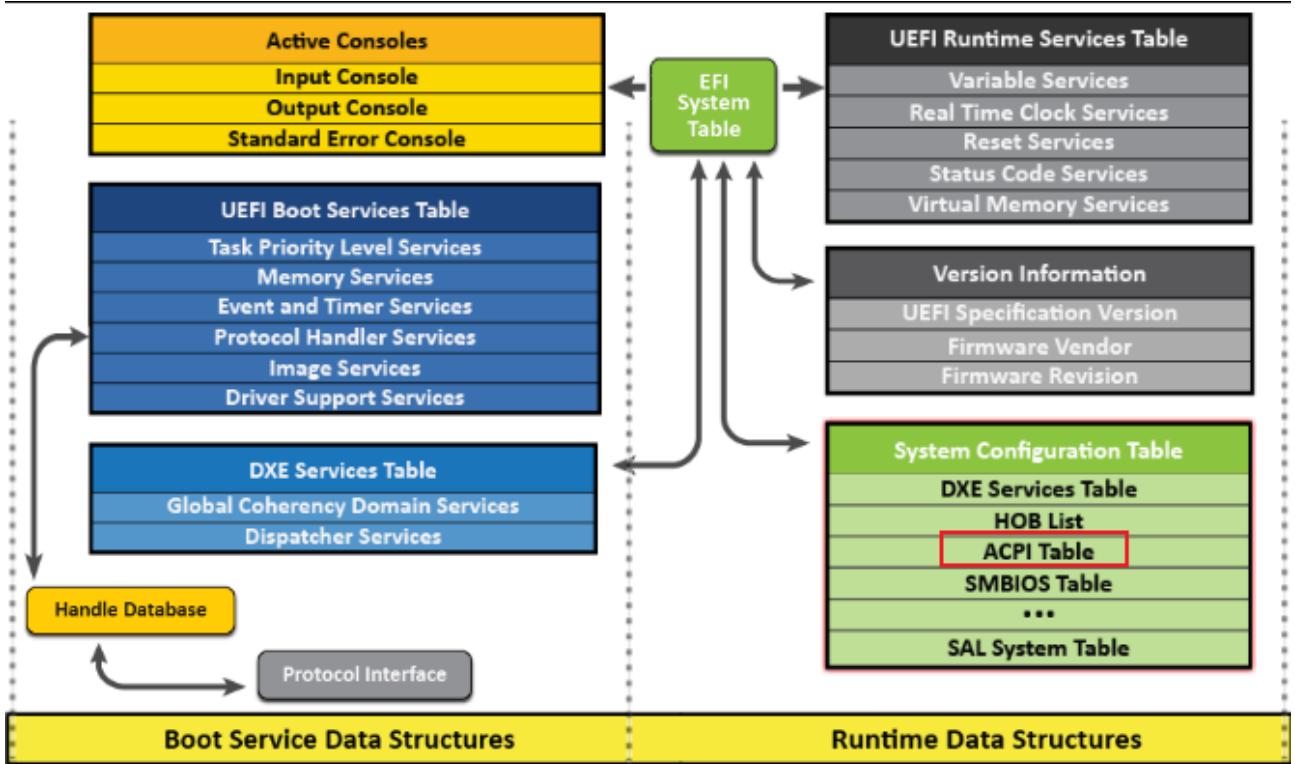
图3：ASL与AML

AML / I/O PCI / AML

"ACPI "	ACPI	ACPI	( PCI USB )	ACPI
_HID ( )	ACPI	_ADR ( )	ACPI	ACPI
PCI	PCI	ACPI	ACPI PCI	
ACPI		ACPI	ACPI	

## ACPI initialization





## ACPI EDK2

system table -> system configuration table -> ACPI table

system table      OS    OS

ACPI table    system configuration table    entry entry    GUID

///

/// Contains a set of GUID/pointer pairs comprised of the ConfigurationTable field in the  
/// EFI System Table.

///

**typedef struct {**

///

/// The 128-bit GUID value that uniquely identifies the system configuration table.

///

EFI\_GUID                                 VendorGuid;

///

/// A pointer to the table associated with VendorGuid.

///

VOID                                     \*VendorTable;

**} EFI\_CONFIGURATION\_TABLE;**

ACPI    GUID    MdePkg\Include\Guid\Acp.h

```
#define ACPI_TABLE_GUID \
{ \
    0xeb9d2d30, 0x2d88, 0x11d3, {0x9a, 0x16, 0x0, 0x90, 0x27, 0x3f, 0xc1, 0x4d} \
}
```

```
#define EFI_ACPI_TABLE_GUID \
{ \
```

```

0x8868e871, 0xe4f1, 0x11d3, {0xbc, 0x22, 0x0, 0x80, 0xc7, 0x3c, 0x88, 0x81 } \

}

#define ACPI_10_TABLE_GUID      ACPI_TABLE_GUID

//  

// ACPI 2.0 or newer tables should use EFI_ACPI_TABLE_GUID.  

//  

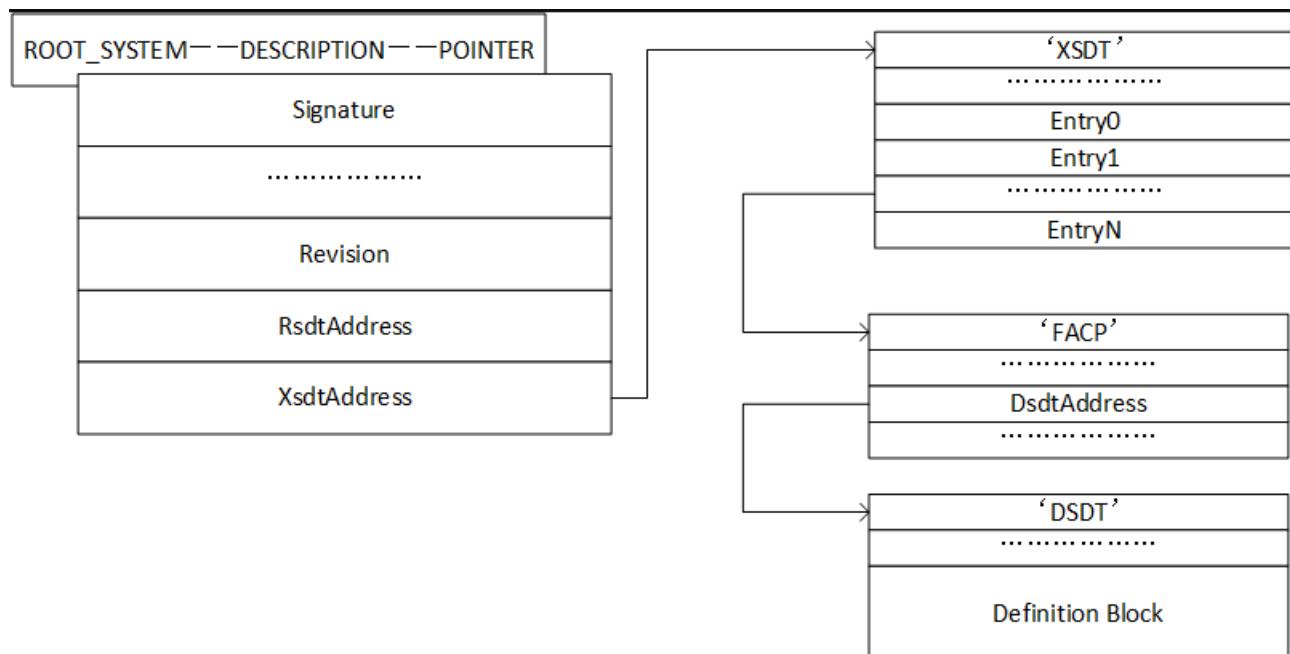
#define EFI_ACPI_20_TABLE_GUID  EFI_ACPI_TABLE_GUID

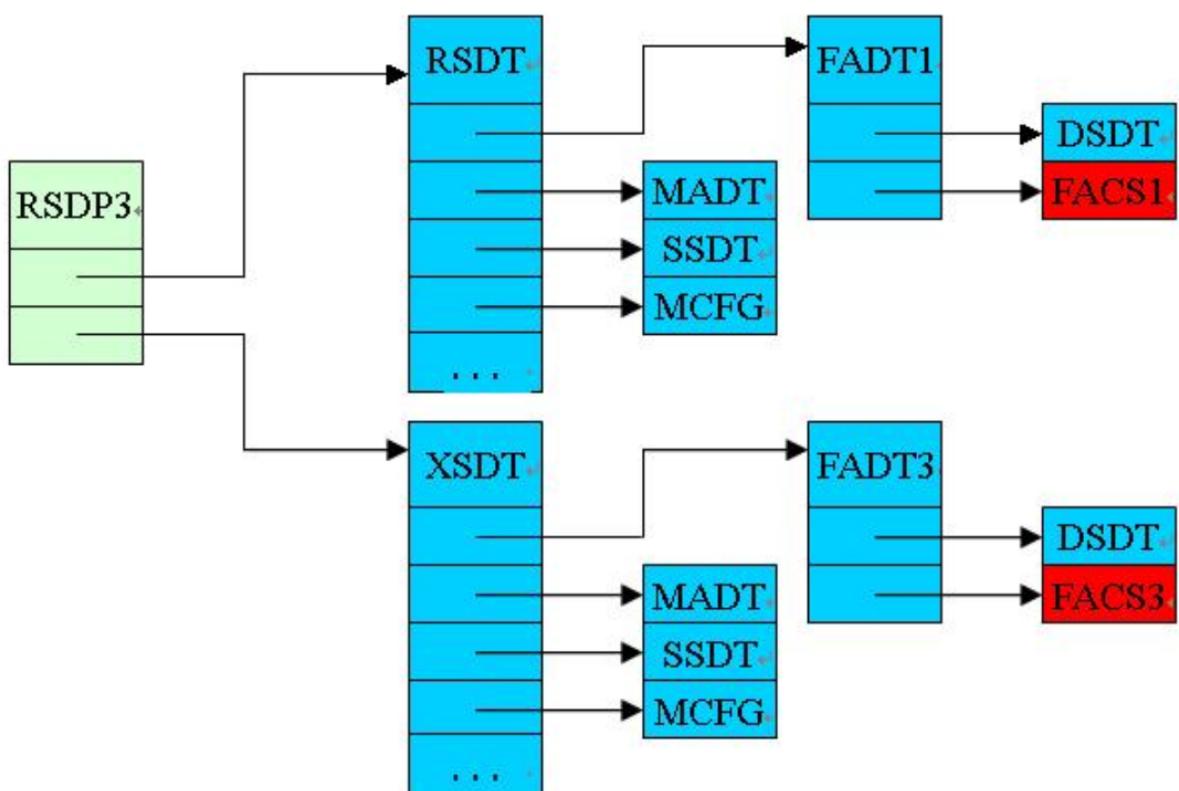
extern EFI_GUID gEfiAcpiTableGuid;
extern EFI_GUID gEfiAcpi10TableGuid;
extern EFI_GUID gEfiAcpi20TableGuid;

```

## ACPI table

1.0 2.0





Located in system's memory address space

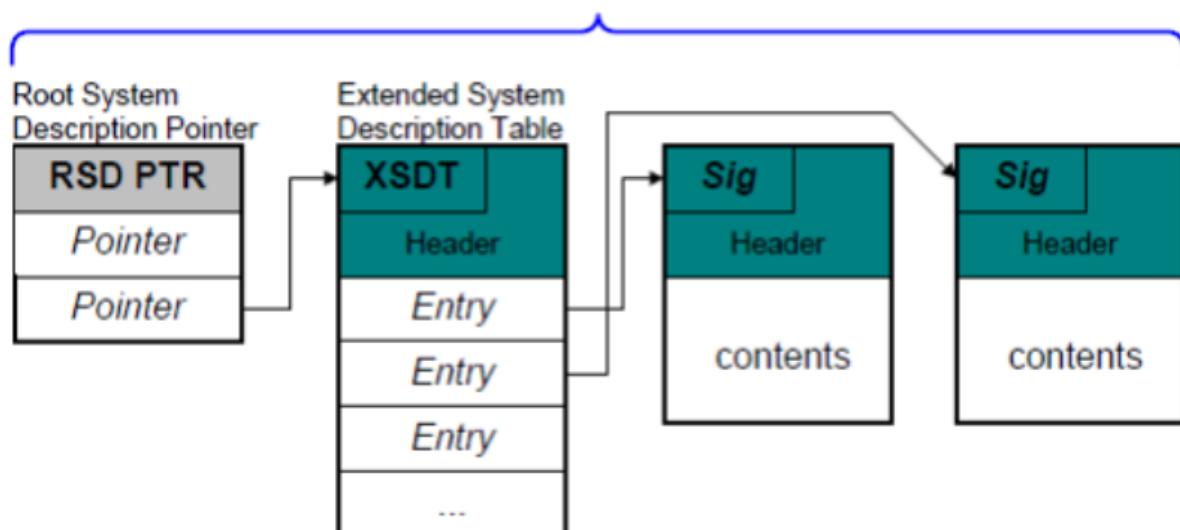


Fig. 5.1: Root System Description Pointer and Table

ACPI table entry      RSDP      RSDP

XSDT

#### RSDT XSDT

- RSDT 32      4GB      ACPI
- XSDT 64      64

- RSDT ACPI 1.0 32
- XSDT ACPI 2.0 64 32
- RSDT 32
- XSDT 64

## Linux Kernel ACPI

ACPI

### ACPI table

ACPI table XSDT 36 entry RSDP XSDT FADT Fixed ACPI Description Table DSDT Differentiated System Description Table FACS Firmware ACPI Control Structure

```

Signature "XSDT"
Length 0x000000144 (324)
Revision 0x01 (1)
Checksum 0x34 (52)
OEM ID "LENOVO"
OEM Table ID "CB-01 "
OEM Revision 0x00000001 (1)
Creator ID "
Creator Revision 0x01000013 (16777235)
Entry0 0x0000000044BC3000 (FACP)
Entry1 0x0000000044B2E000 (UEFI)
Entry2 0x0000000044BF3000 (SSDT)
Entry3 0x0000000044BF2000 (SSDT)
Entry4 0x0000000044BEC000 (SSDT)
Entry5 0x0000000044BE8000 (SSDT)
Entry6 0x0000000044BE4000 (SSDT)
Entry7 0x0000000044BD6000 (SSDT)
Entry8 0x0000000044BD5000 (SSDT)
Entry9 0x0000000044BD4000 (TPM2)
Entry10 0x0000000044BD3000 (SSDT)
Entry11 0x0000000044BD2000 (SSDT)
Entry12 0x0000000044BD1000 (MSDM)
Entry13 0x0000000044BCF000 (SSDT)
Entry14 0x0000000044BCE000 (LPIT)
Entry15 0x0000000044BCD000 (WSMT)
Entry16 0x0000000044BCC000 (SSDT)
Entry17 0x0000000044BC9000 (SSDT)
Entry18 0x0000000044BC8000 (DBG)
Entry19 0x0000000044BC7000 (DBG2)
Entry20 0x0000000044BC4000 (NHLT)
Entry21 0x0000000044BFD000 (ECDT)
Entry22 0x0000000044BC2000 (HPET)
Entry23 0x0000000044BC1000 (APIC)
Entry24 0x0000000044BC0000 (MCFG)

```

Entry25 0x0000000044B58000 (SSDT)  
Entry26 0x0000000044B56000 (SSDT)  
Entry27 0x0000000044B55000 (\$H20)  
Entry28 0x0000000044B54000 (DMAR)  
Entry29 0x0000000044B53000 (SSDT)  
Entry30 0x0000000044B4F000 (SSDT)  
Entry31 0x0000000044B4B000 (SSDT)  
Entry32 0x0000000044B4A000 (SSDT)  
Entry33 0x0000000044B49000 (FPDT)  
Entry34 0x0000000044B48000 (BGRT)  
Entry35 0x0000000044B47000 (PHAT)

### RSDP Root System Description Pointer

- RSDT XSDT
- 

### RSDT Root System Description Table XSDT Extended System Description Table

- ACPI
- RSDT 32 XSDT 64

### FADT Fixed ACPI Description Table

- 
- 

### DSDT Differentiated System Description Table

- AML
- 

### FACS Firmware ACPI Control Structure

- 
- 

## Memory ACPI table

ACPI SRAT System Resource Affinity Table SLIT System Locality Information Table

### MCFG Memory Configuration Table

- |                      |           |                      |                       |                |                      |
|----------------------|-----------|----------------------|-----------------------|----------------|----------------------|
| • MCFG<br>Mapped I/O | ACPI<br>/ | DRAM Dynamic<br>MCFG | Random Access<br>DRAM | Memory<br>MMIO | MMIO Memory-<br>MMIO |
|----------------------|-----------|----------------------|-----------------------|----------------|----------------------|

### SRAT System Resource Affinity Table

SRAT      CPU I/O      NUMA Non-Uniform Memory Access      SRAT

- 
- 
- CPU
- I/O

SRAT

```

struct ACPI_TABLE_SRAT {
    struct ACPI_TABLE_HEADER Header;           // ACPI
    uint32_t Reserved1;                      //
    uint64_t Reserved2;                      //
    struct ACPI_SRAT_ENTRY Entries[];         //
};

struct ACPI_SRAT_ENTRY {
    uint8_t Type;                           // CPU
    uint8_t Length;                         //
    // Type
};

//  

struct ACPI_SRAT_MEM_AFFINITY {
    uint8_t Type;                           // 0x01
    uint8_t Length;                         //
    uint32_t ProximityDomain;              //
    uint16_t Reserved1;                   //
    uint16_t Flags;                        //
    uint64_t BaseAddress;                 //
    uint64_t Length;                       //
    uint32_t Reserved2;                   //
    uint32_t Reserved3;                   //
};

};

```

## SLIT System Locality Information Table

SLIT	SLIT
------	------

```

struct ACPI_TABLE_SLIT {
    struct ACPI_TABLE_HEADER Header;           // ACPI
    uint64_t LocalityCount;                  //
    uint8_t Entry[];                        //
};


```

## CXL ACPI table

CEDT	ACPI	CXL Early Discovery Table CXL		
CEDT	CXL	CXL		
CEDT	CEDT	CXL	CXL	CXL
CEDT	CEDT	CXL	Entries	
1. <b>CXL Host Bridge Structure</b> CXL 2. <b>CXL Device Structure</b> CXL ID 3. <b>CXL Switch Structure</b> CXL Switch				
CXL	CXL			

## CEDT

```
struct cedt_header {
    uint32_t signature;          // 'CEDT'
    uint32_t length;             // Length of the entire table
    uint8_t revision;            // Revision of the CEDT table
    uint8_t checksum;             // Checksum of the entire table
    char oem_id[6];              // OEM ID
    char oem_table_id[8];         // OEM Table ID
    uint32_t oem_revision;        // OEM Revision
    uint32_t creator_id;          // Creator ID
    uint32_t creator_rev;         // Creator Revision
};

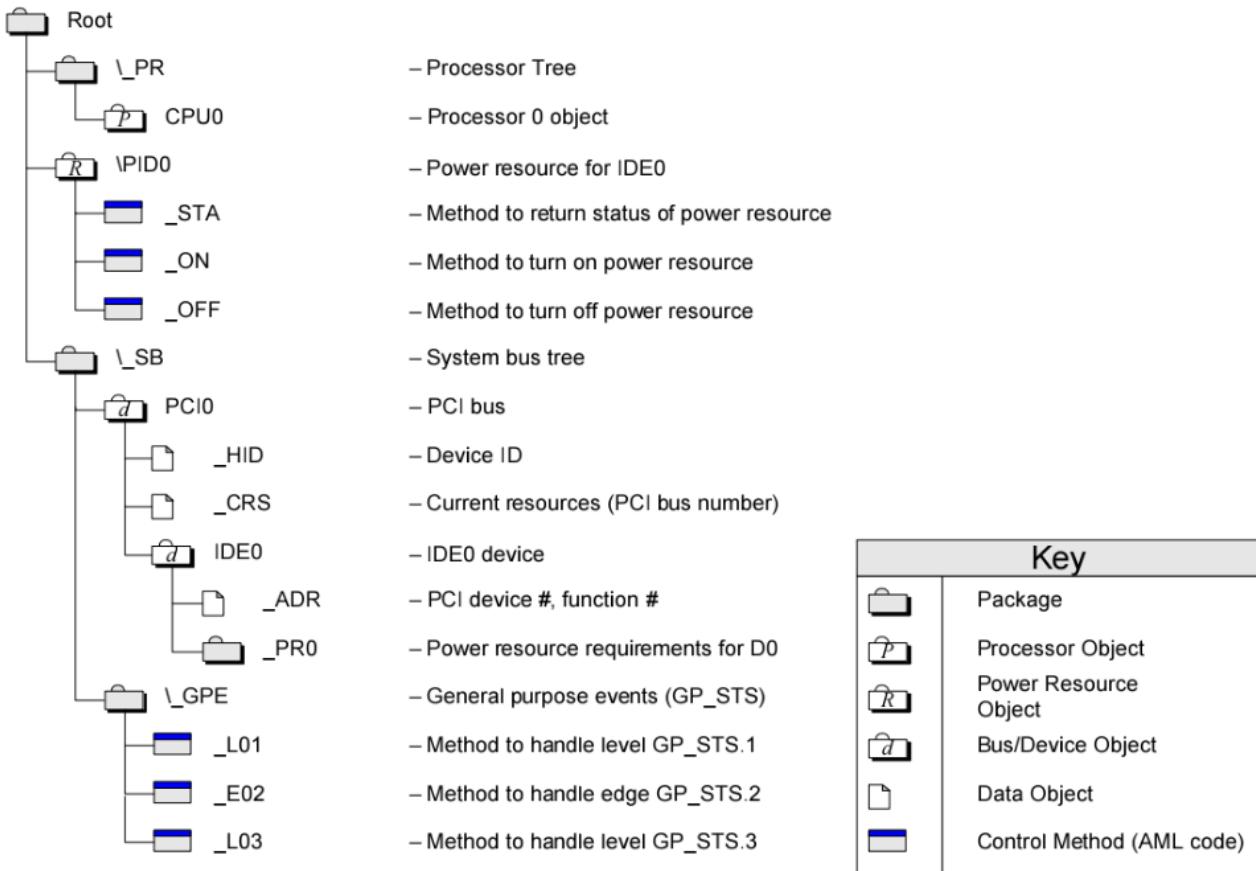
struct cxl_host_bridge {
    uint16_t type;               // Entry type (CXL Host Bridge)
    uint16_t length;              // Length of this entry
    uint32_t host_bridge_id;       // Host Bridge ID
    // Additional fields specific to the CXL Host Bridge
};

struct cxl_device {
    uint16_t type;               // Entry type (CXL Device)
    uint16_t length;              // Length of this entry
    uint32_t device_id;            // Device ID
    // Additional fields specific to the CXL Device
};

// Main CEDT Table containing multiple entries
struct cedt_table {
    struct cedt_header header;
    struct cxl_host_bridge host_bridge;
    struct cxl_device device;
    // Other entries as needed
};
```

## DSDT SSDT

ACPI namespace



Name	Description
\_GPE	General events in GPE register block.

Name	Description
\_PR	ACPI 1.0 Processor Namespace. ACPI 1.0 requires all Processor objects to be defined under this namespace. ACPI 2.0 and later allow Processor object definitions under the \_SB namespace. Platforms may maintain the \_PR namespace for compatibility with ACPI 1.0 operating systems, but it is otherwise deprecated. see the compatibility note in <a href="#">Section 5.2.12.12</a> . An ACPI-compatible namespace may define Processor objects in either the \_SB or \_PR scope but not both. For more information about defining Processor objects, see <a href="#">Section 8</a> , "Processor Configuration and Control."
\_SB	All Device/Bus Objects are defined under this namespace.
\_SI	System indicator objects are defined under this namespace. For more information about defining system indicators, see <a href="#">Section 9.2</a> , \_SI System Indicators."
\_TZ	ACPI 1.0 Thermal Zone namespace. ACPI 1.0 requires all Thermal Zone objects to be defined under this namespace. Thermal Zone object definitions may now be defined under the \_SB namespace. ACPI-compatible systems may maintain the \_TZ namespace for compatibility with ACPI 1.0 operating systems. An ACPI-compatible namespace may define Thermal Zone objects in either the \_SB or \_TZ scope but not both. For more information about defining Thermal Zone objects, see <a href="#">Section 11</a> , "Thermal Management."

## table

asl aml aml raw raw ffs

```

ffs aml memory protocol install

// Example SSDT: SSDT_MYCUSTOM.aml

DefinitionBlock ("SSDT_MYCUSTOM.aml", "SSDT", 1, "YOURID", "YOURID", 1)
{
    // Insert your ACPI table content here
    // Example:
    Method (_STA, 0, NotSerialized) // _STA: Status
    {
        Return (0x0F) // indicate device is present and working
    }
}

```

### 1. DefinitionBlock

- **DefinitionBlock**    ACPI              SSDT\_MYCUSTOM.aml    SSDT    Secondary System Description  
Table ACPI 1 OEM ID "YOURID" OEM ID "YOURID"

### 2. Method (\_STA)

- ACPI              Method (\_STA, 0, NotSerialized)    \_STA        0        NotSerialized
- \_STA              Return (0x0F)            0x0F

## EDK2 protocol

Status = gBS->LocateProtocol (&gEfiAcpTableProtocolGuid, NULL, (VOID \*\*) &AcpiTable);

Status = AcpiTable->InstallAcpiTable ( AcpiTable, Table, TableSize, &TableKey );    uninstall    ACPI table  
table              device tree

[https://www.nirsoft.net/utils/firmware\\_tables\\_view.html](https://www.nirsoft.net/utils/firmware_tables_view.html)

RW

UEFI Shell <https://acpica.org/downloads/uefi-support>

acpi windwos tool download link: <https://acpica.org/downloads/binary-tools>

UEFI Shell <https://github.com/andreiw/UefiToolsPkg/tree/master/Applications>

UEFI Shell acpiview from UefiShellAcpiViewCommandLib.inf

## Advanced Configuration and Power Interface (ACPI) Introduction and Overview

### (ACPI)

ACPI Spec

<https://www.lab-z.com/revmem/>

[ACPI Source Language \(ASL\) Tutorial](#)

## 4.2 CEDT

[toc]

CEDT ACPI CXL Early Discovery Table CXL

CEDT CXL CXL

## CEDT

CEDT CXL CXL CXL

## CEDT

CEDT CXL Entries

1. CXL Host Bridge Structure CXL
  2. CXL Device Structure CXL ID
  3. CXL Switch Structure CXL Switch
- CXL CXL

## CEDT

**Table 9-19.** CEDT Header

Field	Length in Bytes	Byte Offset	Description
<b>Header:</b>			
Signature	4	00h	'CEDT'. Signature for the CXL Early Discovery Table.
Length	4	04h	Length, in bytes, of the entire CEDT.
Revision	1	08h	Value is 1.
Checksum	1	09h	Entire table must sum to 0.
OEM ID	6	0Ah	OEM ID
OEM Table ID	8	10h	Manufacturer Model ID
OEM Revision	4	18h	OEM Revision
Creator ID	4	1Ch	Vendor ID of the utility that created the table.
Creator Revision	4	20h	Revision of the utility that created the table.
CEDT Structure[n]	Varies	24h	A list of CEDT structures for this implementation.

**Table 9-20.** CEDT Structure Types

Value	Description
0	CXL Host Bridge Structure (CHBS)
1	CXL Fixed Memory Window Structure (CFMWS)
2	CXL XOR Interleave Math Structure (CXIMS)
3	RCEC Downstream Port Association Structure (RDPAS)
4-255	Reserved

**Table 9-21. CHBS Structure**

Field	Length in Bytes	Byte Offset	Description
Type	1	00h	=0 to indicate that this is a CHBS entry
Reserved	1	01h	Reserved
Record Length	2	02h	Length of this record (20h).
UID	4	04h	CXL Host Bridge Unique ID. Used to associate a CHBS instance with a CXL Host Bridge instance. The value of this field shall match the output of _UID under the associated CXL Host Bridge in ACPI namespace.
CXL Version	4	08h	<ul style="list-style-type: none"> <li>• 0000 0000h: RCH</li> <li>• 0000 0001h: Host Bridge that is associated with one or more CXL root ports</li> </ul>
Reserved	4	0Ch	Reserved
Base	8	10h	<ul style="list-style-type: none"> <li>• If CXL Version = 0000 0000h, this represents the base address of the RCH Downstream Port RCRB</li> <li>• If CXL Version = 0000 0001h, this represents the base address of the CHBCR</li> </ul> <p>See <a href="#">Table 8-17</a> for more details.</p>
Length	8	18h	<ul style="list-style-type: none"> <li>• If CXL Version = 0000 0000h, this field must be set to 8 KB (2000h)</li> <li>• If CXL Version = 0000 0001h, this field must be set to 64 KB (1 0000h)</li> </ul>

```
#define INTERLEAVE_TARGETS 1

#pragma pack (1)

typedef struct {
    EFI_ACPI_CEDT_CFMWS_STRUCTURE  Cfmws;
    UINT32  InterleaveTarget[INTERLEAVE_TARGETS];
} EFI_ACPI_CEDT_CFMWS_AND_INTERLEAVE_TARGET_STRUCTURE;

typedef struct {
    EFI_ACPI_DESCRIPTION_HEADER  Header;
    EFI_ACPI_CEDT_CFMWS_AND_INTERLEAVE_TARGET_STRUCTURE CfmwsTarget;
    EFI_ACPI_CEDT_CHBS_STRUCTURE  Chbs;
} EFI_ACPI_CEDT_STRUCTURE_TABLE;

#pragma pack ()

STATIC EFI_ACPI_CEDT_STRUCTURE_TABLE Cedt = {
    ARM_ACPI_HEADER (
        EFI_ACPI_6_4_CEDT_SIGNATURE,
        EFI_ACPI_CEDT_STRUCTURE_TABLE,
        1
    ),
    {
        {
            // CFMWS
            {
                1,
                EFI_ACPI_RESERVED_BYTE,

```

```

        sizeof (EFI_ACPI_CEDT_CFMWS_AND_INTERLEAVE_TARGET_STRUCTURE)
    },
    EFI_ACPI_RESERVED_DWORD,
    0x3fe00000000,
    0x2000000000,
    0,
    0,
    EFI_ACPI_RESERVED_WORD,
    4,
    2,
    0
},
{
//Interleave target list
    1
}
},
{
// CHBS
{
    0,
    EFI_ACPI_RESERVED_BYTE,
    sizeof (EFI_ACPI_CEDT_CHBS_STRUCTURE)
},
1,
1,
1,
EFI_ACPI_RESERVED_DWORD,
0x10D0000000,
0x1000
}
};

//  

// Reference the table being generated to prevent the optimizer from removing  

// the data structure from the executable  

//  

VOID* CONST ReferenceAcpiTable = &Cedt;

```

[CXL Type-3 device discovery, configuration in firmware and prepare ACPI tables for kernel usage](#)

[EDK2 CEDT Table](#)

CXL Spec3.0 Chapter 9.17

# **Chapter 5**

## **AMBA**

### **5.1 PCIe**

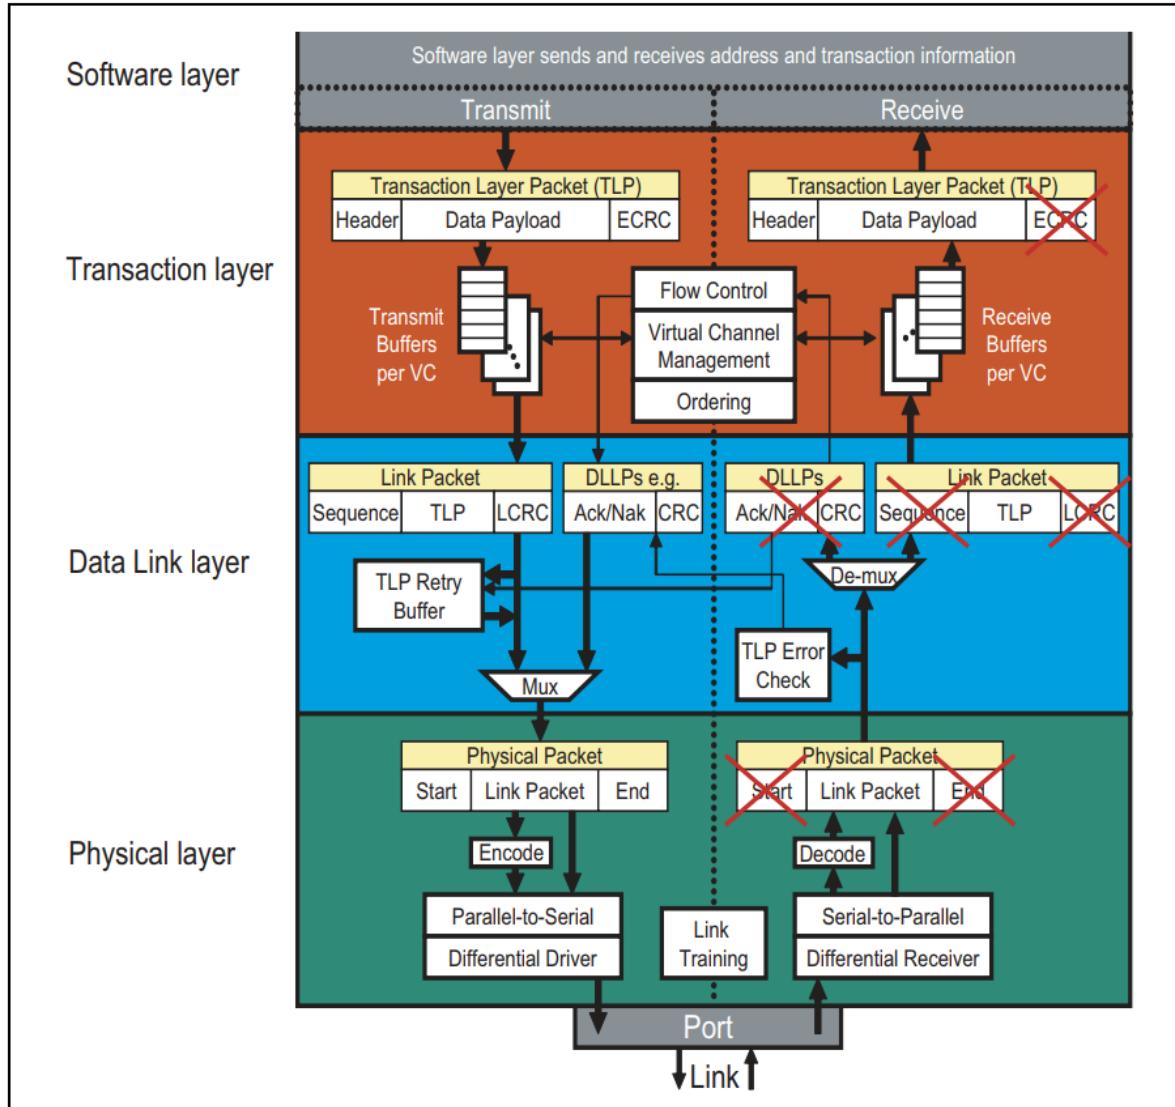
### **5.2 PCIe**

#### **5.2.1**

Link Lane      EQ LTSSM

#### **Introduction**

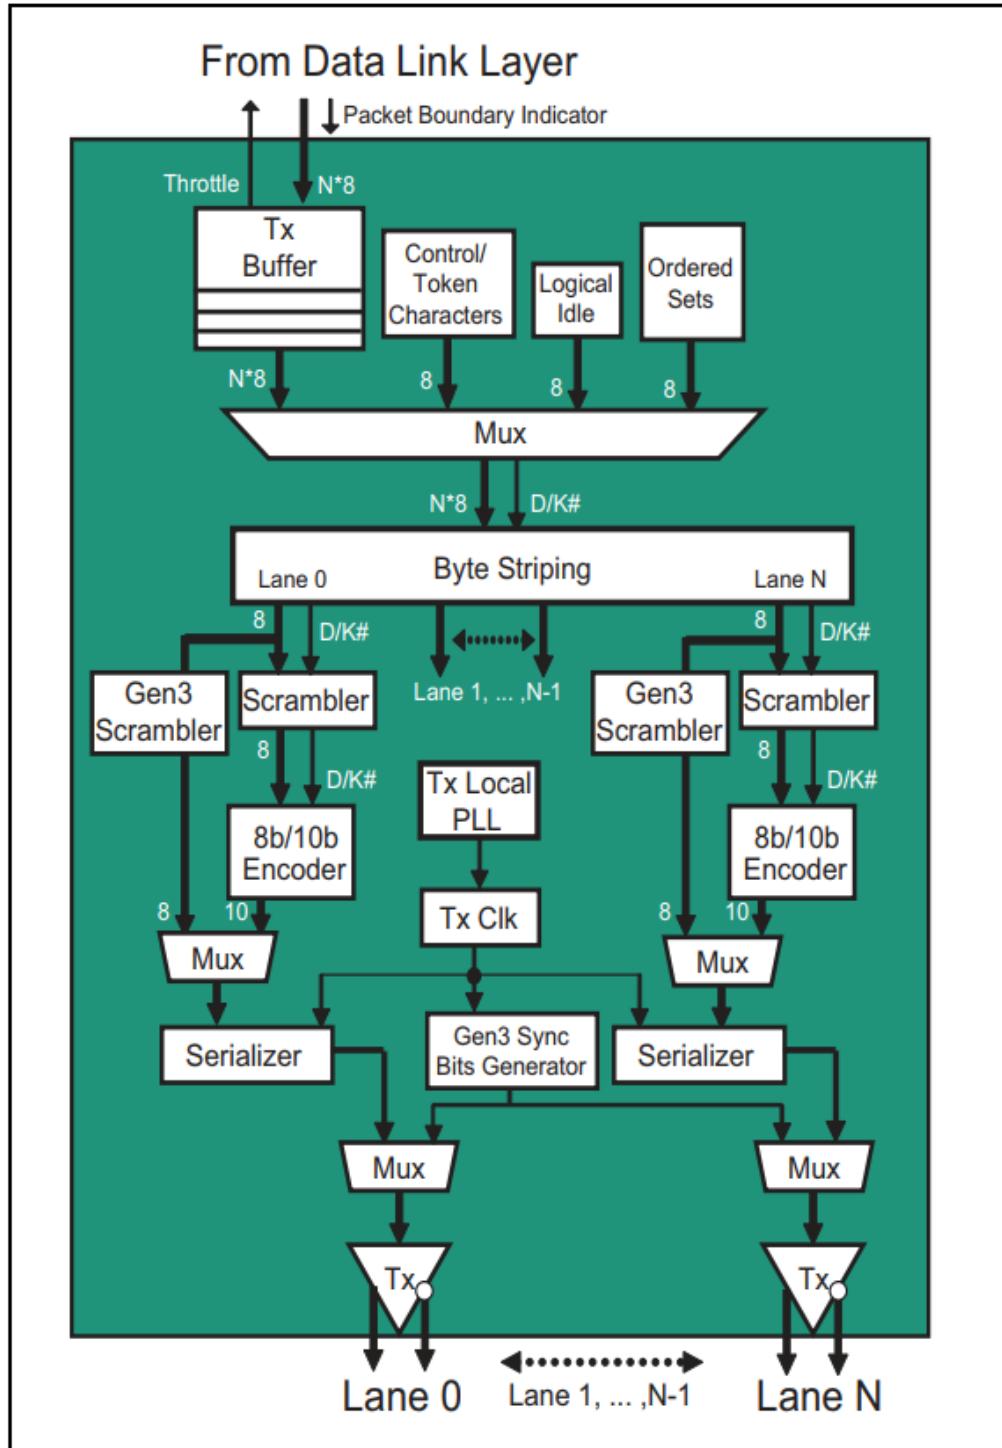
Figure 11-1: PCIe Port Layers



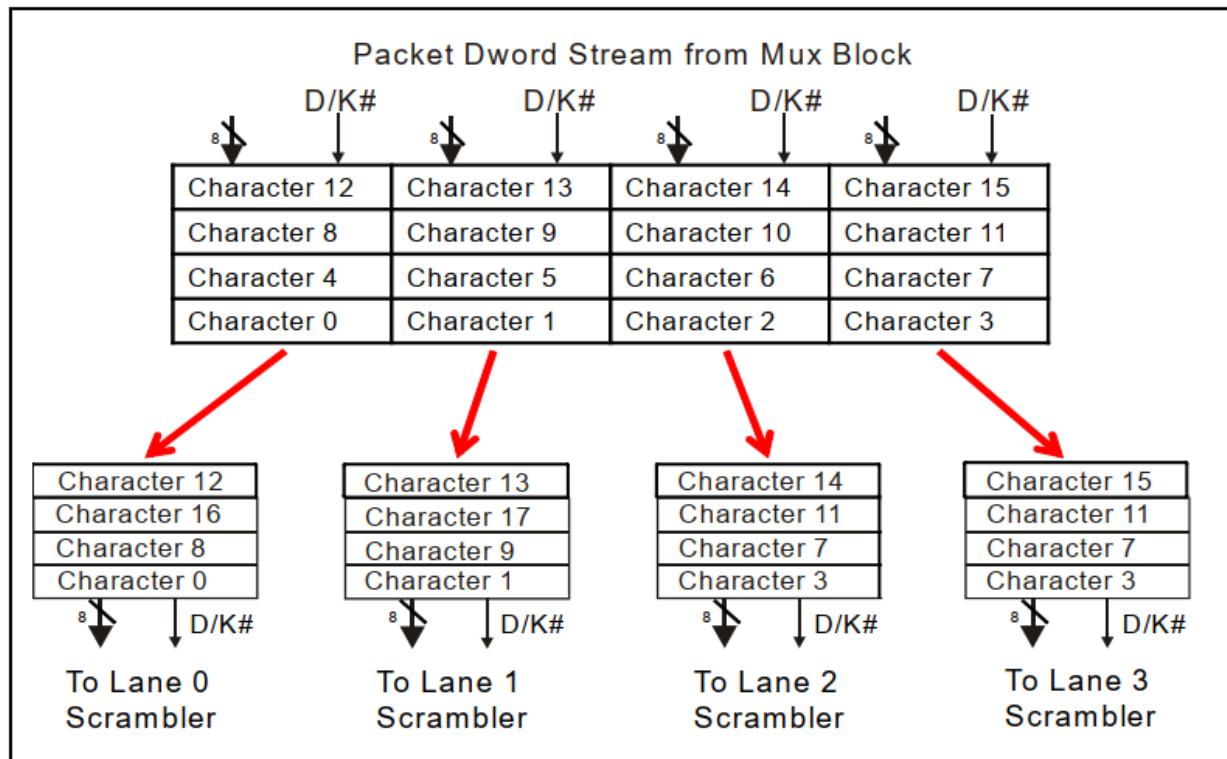
- Logical Sub-Block
- Electrical Sub-Block

## Logical Sub-Block

Figure 11-3: Physical Layer Transmit Details



<b>Tx Buffer &amp; Mux</b> buffer DL Mux DL skip buffer	buffer DL Mux DL skip buffer
<b>Byte Striping</b> lane	



## Data Scrambling

## Scrambling

pseudo-random number generator

PCIo

EMI

0 1

01

### LFSR Linear-feedback shift register

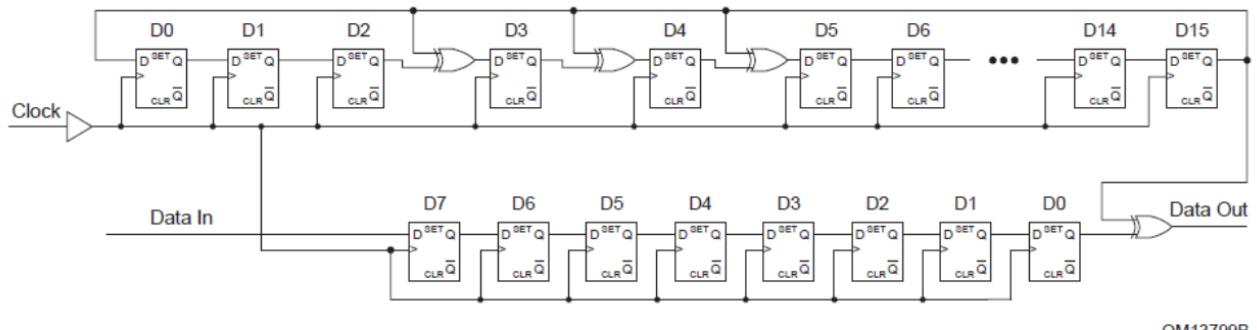
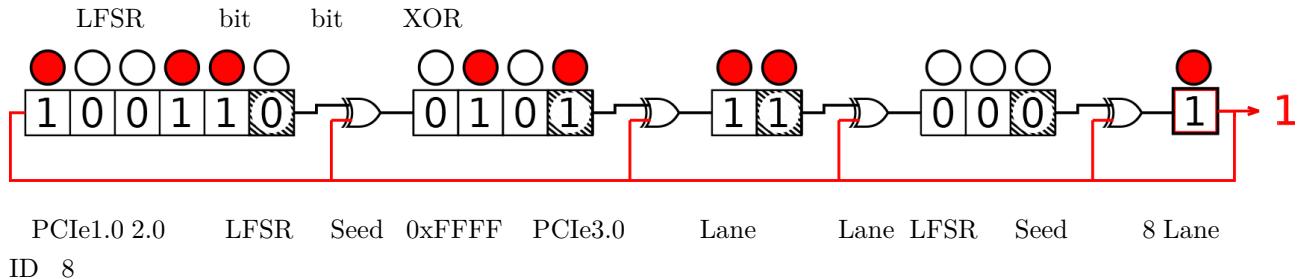


Figure 4-10 LFSR with 8b/10b Scrambling Polynomial

PCIe Galois LFSR PCIe 1.0 2.0 16 LFSR

$$G(X) = X^{16} + X^5 + X^4 + X^3 + 1$$

$$G(X) = X^{23} + X^{21} + X^{16} + X^8 + X^5 + X^2 + 1$$



PCIe1.0 2.0      LFSR      Seed 0xFFFF      PCIe3.0      Lane      Lane LFSR      Seed      8 Lane  
ID 8

Lane	Seed
0	1DBFBCh
1	0607BBh
2	1EC760h
3	18C0DBh
4	010F12h
5	19CFC9h
6	0277CEh
7	1BB807h

XOR                    XOR                    Seed                    LFSR

Encoding	Encoding	0 1	DC	DC	Balance	PCIe	Clock	Recov-
PCIe	8b/10b    128b/130b    242B/256B FLIT					PCIe	Lane	PCIe 1.0 8b/10b    8bit    10b

$$\frac{2.5GT/s \times \frac{8bits}{10bits}}{8bits} = 250MB/s$$

PCIe

PCIe Version	Line code	Transfer rate per lane	Throughput x1	Throughput x16
1.0	8b/10b	2.5 GT/s	250 MB/s	4 GB/s
2.0	8b/10b	5 GT/s	500 MB/s	8 GB/s
3.0	128b/130b	8 GT/s	984.6 MB/s	15.75 GB/s
4.0	128b/130b	16 GT/s	1.969 GB/s	31.51 GB/s
5.0	128b/130b	32 GT/s	3.938 GB/s	63.02 GB/s
6.0	1b/1b 242B/256B FLIT	64 GT/s	7.564 GB/s	121.00 GB/s

242B/256B B bit Byte

8b/10b 8b/10b PCIe 1.0 2.0 2.5GT/s 5GT/s [4] 8bits 5bits 3bits 6bits 4bits

8b/10b K D ABCDE FGH D.<ABCDE>.<FGH> K.<ABCDE>.<FGH>

### 5b/6b code (abcdei) [edit]

Input		RD = -1	RD = +1
Code	EDCBA	a b c d e i	
D.00	00000	100111	011000
D.01	00001	011101	100010
D.02	00010	101101	010010
D.03	00011	110001	
D.04	00100	110101	001010
D.05	00101	101001	
D.06	00110	011001	
D.07	00111	111000	000111
D.08	01000	111001	000110
D.09	01001	100101	
D.10	01010	010101	
D.11	01011	110100	
D.12	01100	001101	
D.13	01101	101100	
D.14	01110	011100	
D.15	01111	010111	101000
not used		111100	000011

Input		RD = -1	RD = +1
Code	EDCBA	a b c d e i	
D.16	10000	011011	100100
D.17	10001	100011	
D.18	10010	010011	
D.19	10011	110010	
D.20	10100	001011	
D.21	10101	101010	
D.22	10110	011010	
D.23 †	10111	111010	000101
D.24	11000	110011	001100
D.25	11001	100110	
D.26	11010	010110	
D.27 †	11011	110110	001001
D.28	11100	001110	
D.29 †	11101	101110	010001
D.30 †	11110	011110	100001
D.31	11111	101011	010100
K.28 ‡	11100	001111	110000

## **3b/4b code (fghj) [ edit ]**

Input		RD = -1	RD = +1
Code	HGF	f g h j	
D.x.0	000	1011	0100
D.x.1	001		1001
D.x.2	010		0101
D.x.3	011	1100	0011
D.x.4	100	1101	0010
D.x.5	101		1010
D.x.6	110		0110
D.x.P7 †	111	1110	0001
D.x.A7 †		0111	1000

Input		RD = -1	RD = +1
Code	HGF	f g h j	
K.x.0	000	1011	0100
K.x.1 ≠	001	0110	1001
K.x.2	010	1010	0101
K.x.3	011	1100	0011
K.x.4	100	1101	0010
K.x.5 ≠	101	0101	1010
K.x.6	110	1001	0110
K.x.7 ≠	111	0111	1000

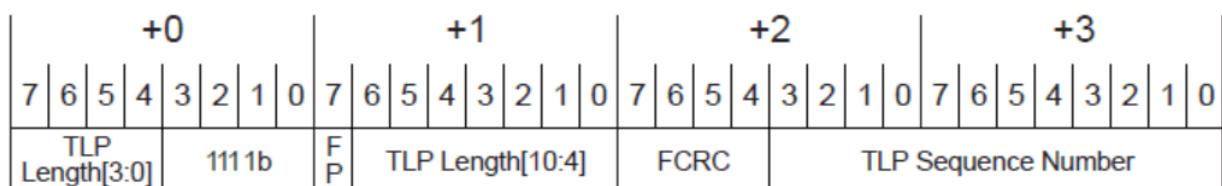
\*\*RD Running Disparity \*\* 1 0 01 8b/10b RD ±1 RD ±2  
8b/10b 01 DC 20% PCIe 3.0 8b/10b

<b>128b/130b</b>	PCIe 3.0	PCIe	128b/130b	130 bits	2 bits	98.46%		
128b/130b	64b/66b	[5]	payload	64b -> 128b	8b/10b		3.0 LFSR	DC
LFSR	128b/130b	Framing Encoding						

Framing

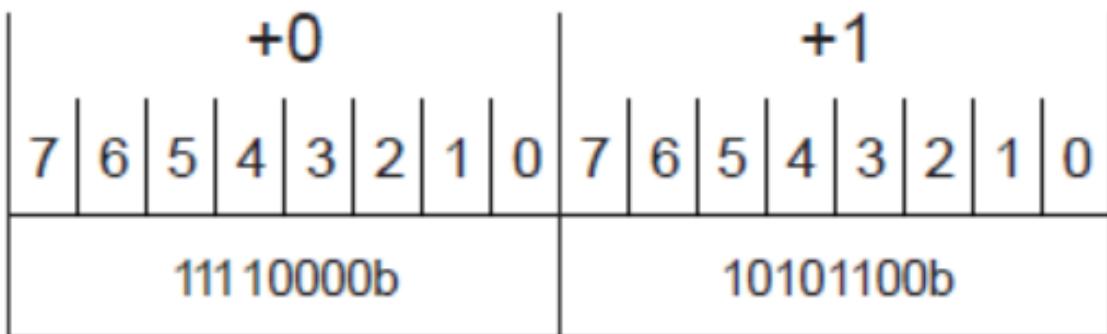
128b/130b                  block                  Token

- TLP 2 4 bits FCRC Seq 4 token — STP Start of TLP



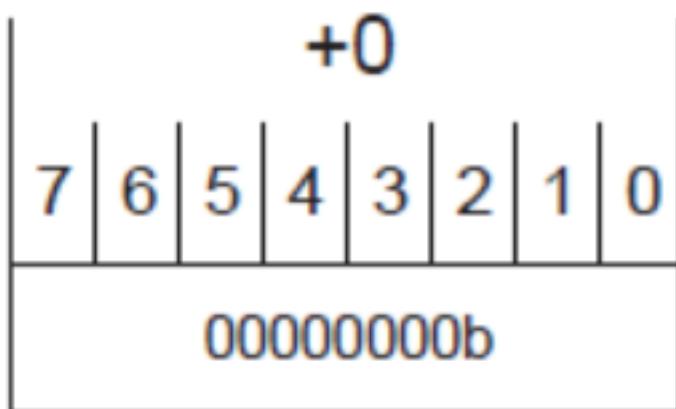
## STP Token

- DLLP 2 Token - SDP Start of DLLP



## SDP Token

Token	Lane 0	Byte 0	IDL	Logical	Idle	Nullify	TLP	EDB	EnD	Bad	EDS	End of
Data Stream												



## IDL Token

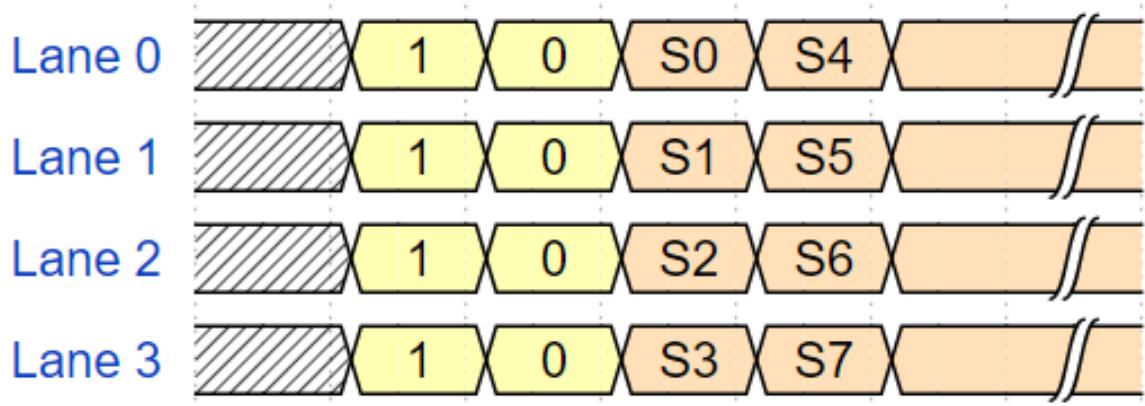
Encoding

128b/130b	128bit block	128bits payload	2bits	SyncHeader	payload
-----------	--------------	-----------------	-------	------------	---------

- 01b payload Data Block 128bits
- 10b payload Ordered Set Block 128bits

16	S0-S15	bit .0	.7	01b	10b
----	--------	--------	----	-----	-----





**242B/256B FLIT**    128b/130b

Amplitude Modulation 4

PCIe6.0

PCIe1.0

NRZ Non-Return-To-Zero

PAM4 Pulse

242B/256B FLIT    FLIT    256    bits

x8 Lanes	0	1	2	3	4	5	6	7
256 UI								
TLP Bytes	0	1	2	3	4	5	6	7
(0-299)	8	9	10	11	12	13	14	15
	16	17	18	19	20	21	22	23
	24	25	26	27	28	29	30	31
	32	33	34	35	36	37	38	39
	40	41	42	43	44	45	46	47
	48	49	50	51	52	53	54	55
	56	57	58	59	60	61	62	63
	64	65	66	67	68	69	70	71
	72	73	74	75	76	77	78	79
	80	81	82	83	84	85	86	87
	88	89	90	91	92	93	94	95
	96	97	98	99	100	101	102	103
	104	105	106	107	108	109	110	111
	112	113	114	115	116	117	118	119
	120	121	122	123	124	125	126	127
	128	129	130	131	132	133	134	135
	136	137	138	139	140	141	142	143
	144	145	146	147	148	149	150	151
	152	153	154	155	156	157	158	159
	160	161	162	163	164	165	166	167
	168	169	170	171	172	173	174	175
	176	177	178	179	180	181	182	183
	184	185	186	187	188	189	190	191
	192	193	194	195	196	197	198	199
	200	201	202	203	204	205	206	207
	208	209	210	211	212	213	214	215
	216	217	218	219	220	221	222	223
	224	225	226	227	228	229	230	231
	232	233	234	235	dip0	dip1	dip2	dip3
	dip4	dip5	crc0	crc1	crc2	crc3	crc4	crc5
	crc6	crc7	ecc0	ecc0	ecc0	ecc1	ecc1	ecc1

236 TLP 6 DLLP Data Link Layer Packet 8 CRC 6 FEC Forward Error Correction

tion	FLIT	CRC	DLLP	TLP	CRC	FEC	FLIT				
	FLIT	PCIe	FLIT		NRZ	2.5 GT/s	5.0 GT/s	8.0 GT/s	16.0 GT/s	32.0 GT/s	NRZ

### Electrical Sub-block

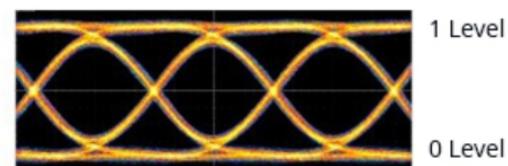
Parallel to Serial

Parallel to Serial

### Modulation

<b>NRZ Non-Return-to-Zero</b>	PCIe1.0 5.0 PCIe	NRZ Non-Return-to-Zero [10]	RZ Return-
to-Zero	1	Bipolar NRZ level	0    -V    1    +V

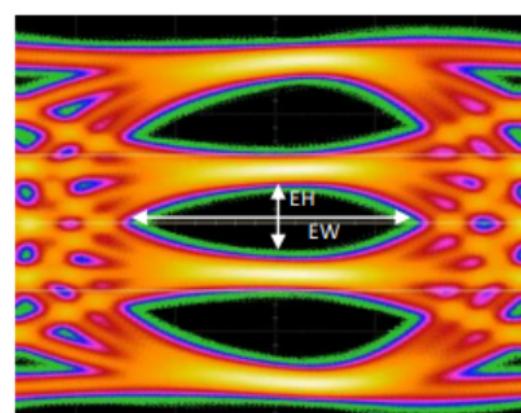
#### NRZ



Encoding   Scrambling

<b>PAM4</b>	PCIe6.0	PAM4 Pulse Amplitude Modulation 4	00   -V   10   +V   01   -V/3   11   +V/3
-------------	---------	-----------------------------------	---

#### Voltage Level



#### 2-Bit Encoding

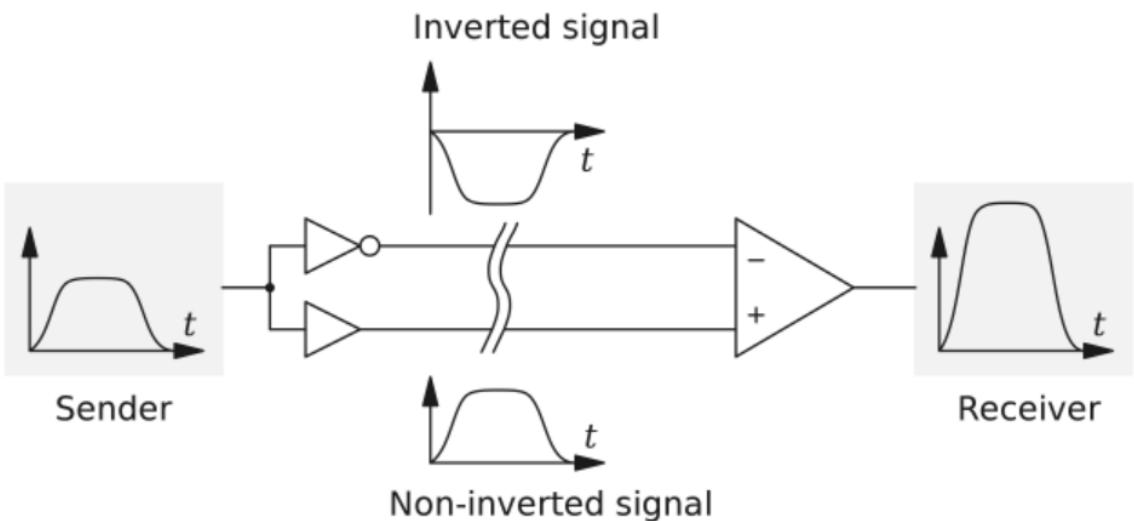
#### DC Balance Values

10	+3
11	+1
01	-1
00	-3

Differential Signal  
Pair

PCIe   Differential Signaling

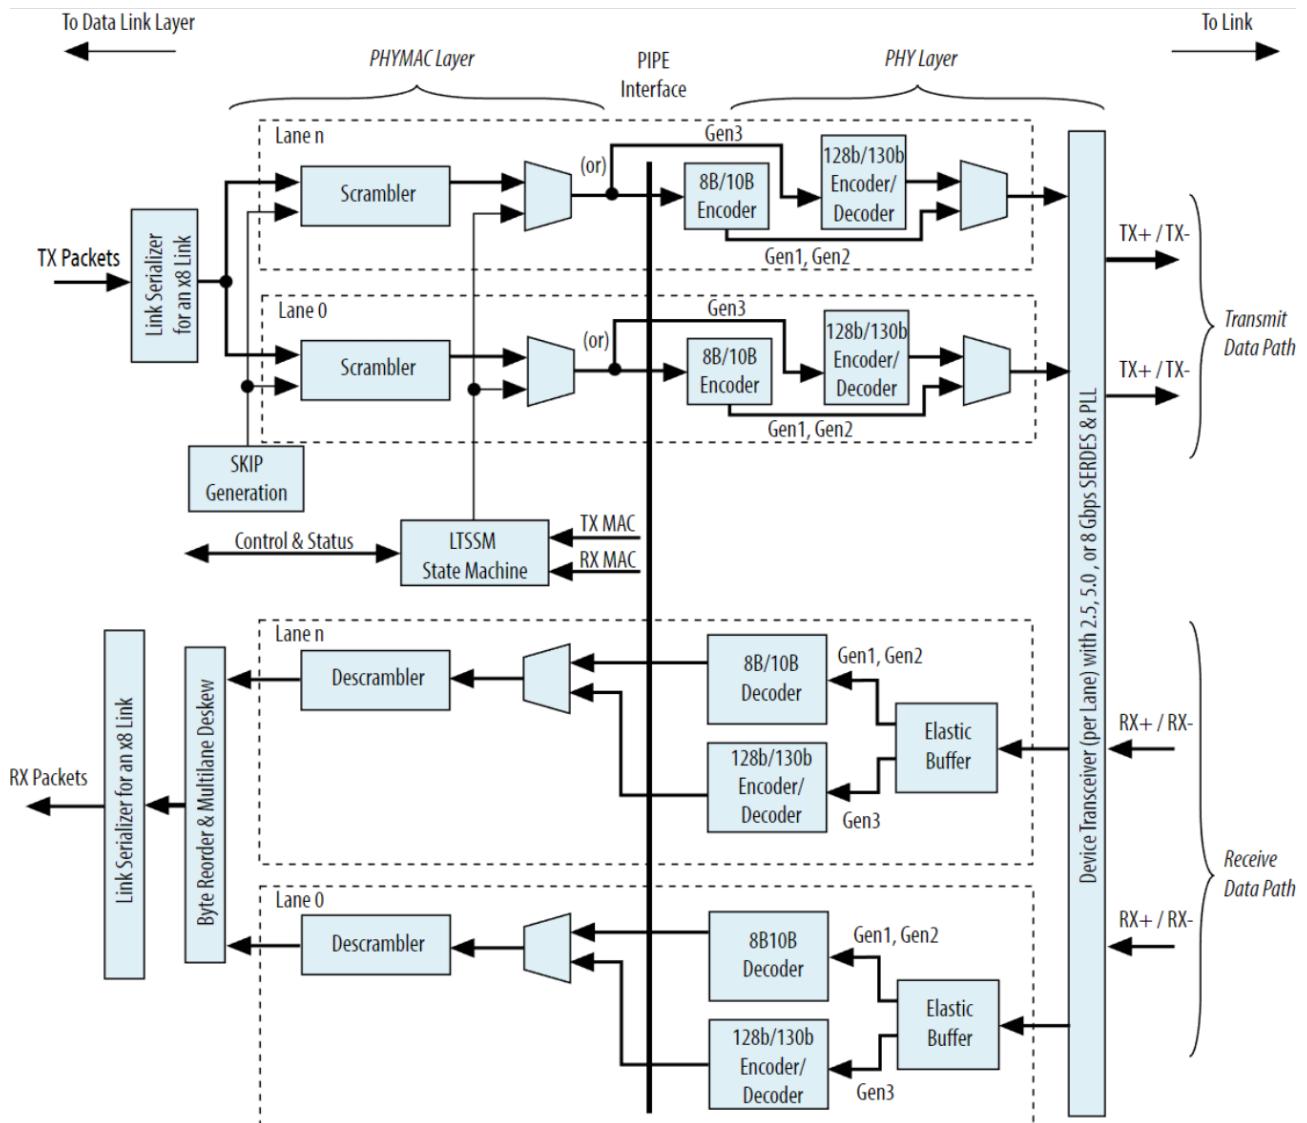
Pinout



我们假设原始信号中的电压为 $V_{Tx}$ , 经过差分处理的两路信号为 $+\frac{V_{Tx}}{2}$ 和 $-\frac{V_{Tx}}{2}$ , 干扰导致的电压变化为 $\Delta V_{noise}$ , 那么忽略传输的损耗, 在接收端收到的电压就是:

$$V_{Rx} = \left( +\frac{V_{Tx}}{2} + \Delta V_{noise} \right) - \left( -\frac{V_{Tx}}{2} + \Delta V_{noise} \right) = V_{Tx}$$

这样, 通过差分信号, 我们就抵消了信道上的干扰。



[Blog](#)

### 5.2.2 Data Link Layer

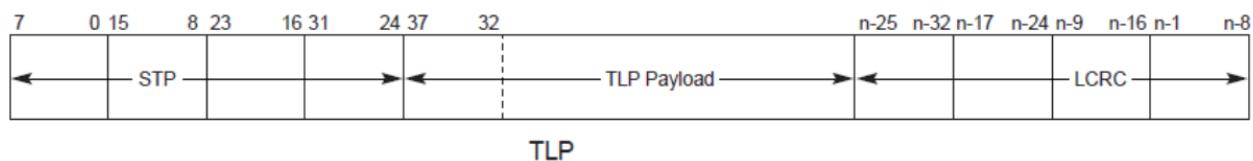
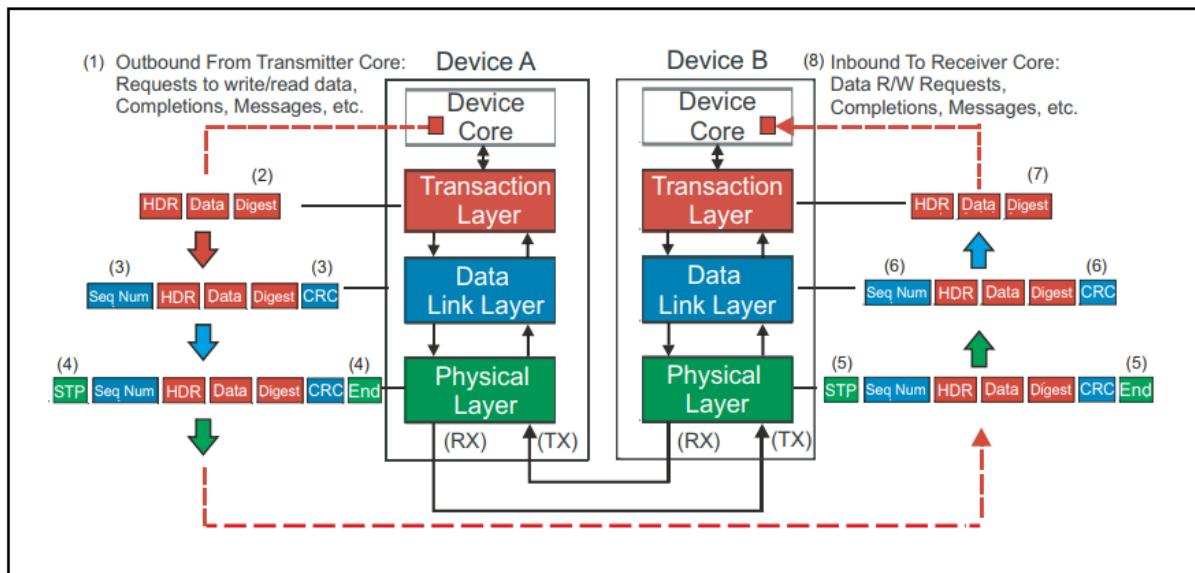
#### Introduction

##### TLP Ack/Nack

Data Link Layer

TLP	Feature	ack/nack	Token	STP Start of TLP	TLP	SDP Start
of DLLP	DLLP Data Link Layer Packet					
Transaction layer	TLPI	TLPI	buffer	TLPI		

Figure 5-2: PCIe TLP Assembly/Disassembly

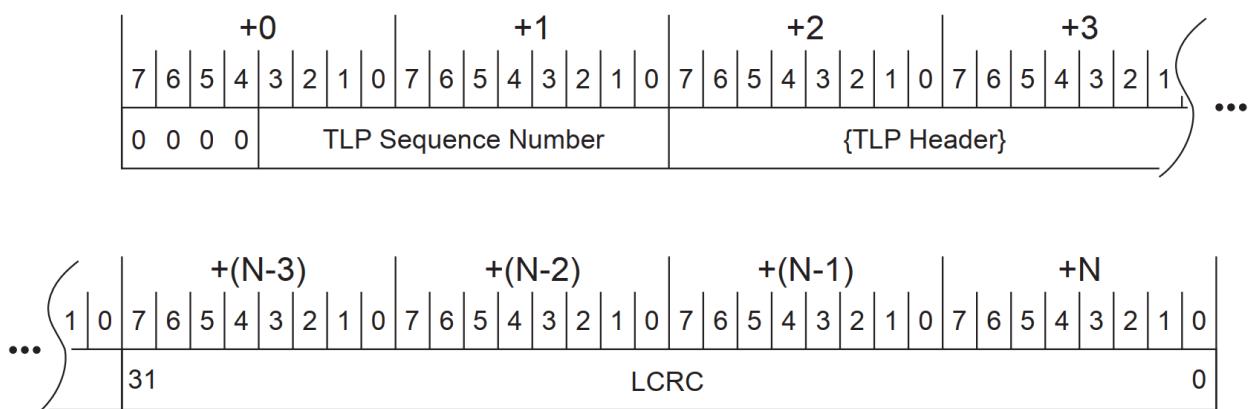


A-0803

Figure 4-14 TLP and DLLP Layout

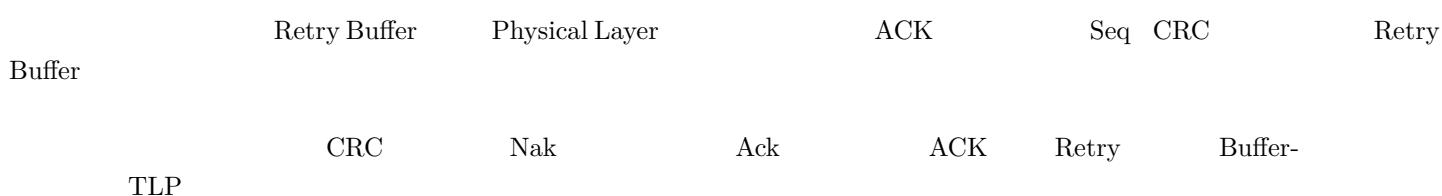
## TLP

- |    |                 |      |          |              |     |
|----|-----------------|------|----------|--------------|-----|
| 1. | Sequence Number | 2    | Link     | NEXT_RCV_SEQ |     |
| 2. | CRC             | LCRC | Link CRC | 4            | CRC |



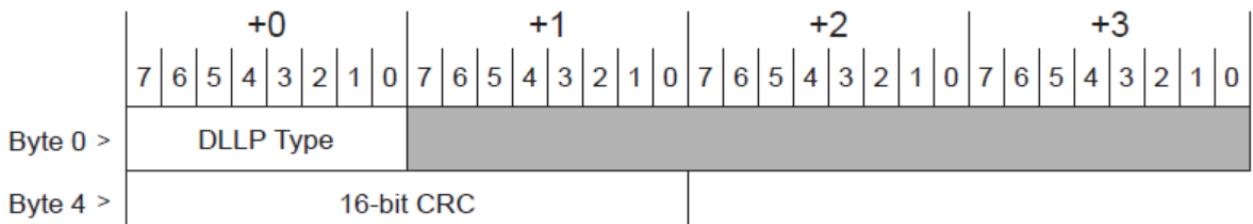
OM13786A

Figure 3-14 TLP with LCRC and TLP Sequence Number Applied



#### DLLP

TLP SDP end gen3	2B SDP	Ack Nak	DLLP Data Link Layer Packet	8B DLLP	6B	2B	(gen1&2)
			8B				



OM14303A

Figure 3-4 DLLP Type and CRC Fields

Table 4-1 shows the Special Symbols used for PCI Express and provides a brief description for each. These Symbols will be discussed in greater detail in following sections. Each of these Special Symbols, as well as the data Symbols, must be interpreted by looking at the 10-bit Symbol in its entirety.

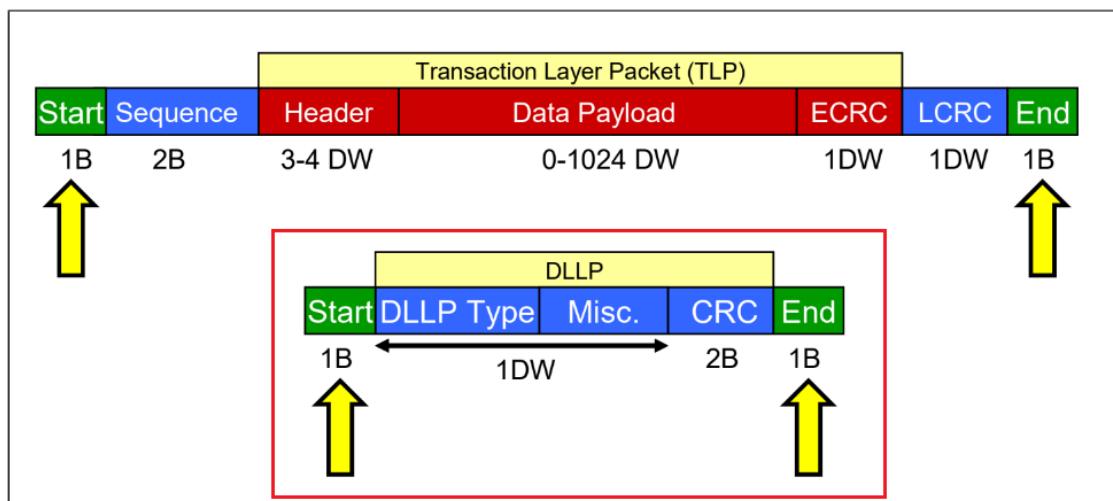
Table 4-1 Special Symbols

Encoding	Symbol	Name	Description
K28.5	COM	Comma	Used for Lane and Link initialization and management
K27.7	STP	Start TLP	Marks the start of a Transaction Layer Packet
K28.2	SDP	Start DLLP	Marks the start of a Data Link Layer Packet
K29.7	END	End	Marks the end of a Transaction Layer Packet or a Data Link Layer Packet
K30.7	EDB	EnD Bad	Marks the end of a nullified TLP
K23.7	PAD	Pad	Used in Framing and Link Width and Lane ordering negotiations
K28.0	SKP	Skip	Used for compensating for different bit rates for two communicating Ports
K28.1	FTS	Fast Training Sequence	Used within an Ordered Set to exit from L0s to L0
K28.3	IDL	Idle	Used in the Electrical Idle Ordered Set (EIOS)
K28.4			Reserved
K28.6			Reserved

## Physical Layer - Logical

TLPs and DLLPs from the Data Link Layer are clocked into a buffer in the Physical Layer, where Start and End characters are added to facilitate detection of the packet boundaries at the receiver. Since the Start and End characters appear on both ends of a packet they are also called “framing” characters. The framing characters are shown appended to a TLP and DLLP in Figure 2-28 on page 77, which also shows the size of each field.

Figure 2-28: TLP and DLLP Structure at the Physical Layer

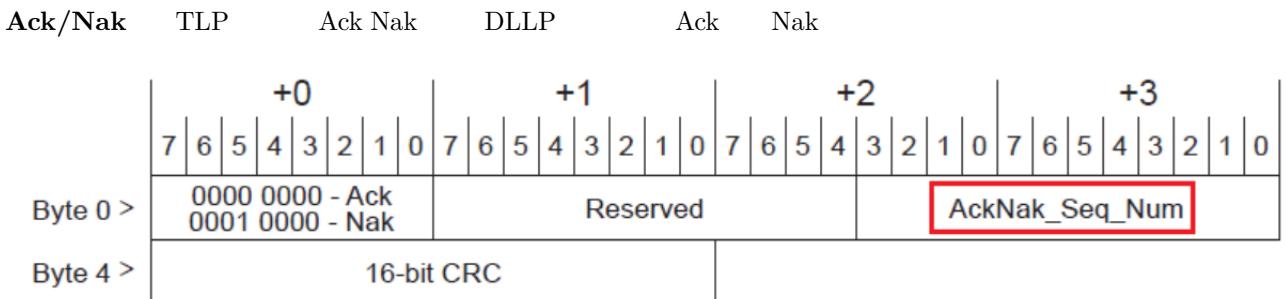


Within this layer, each byte of a packet is split out across all of the lanes in use

TLP    DLL    CRC 32bit    DLLP    CRC 16bit

DLLP DLLP Type              16 CRC

	Type	
Ack	00000000b	TLP
Nak	00010000	TLP
<InitFC1/InitFC2/UpdateFC>-<P/NP/Cpl>	Type	P/NP/Cpl
MRInitFC1/MRInitFC2/MRUpdateFC<0111/1111/1011>0xxxb		P/NP/Cpl
PM_*	00100xxxb	
NOP	00110001b	
Data_Link_Feature	00000010b	Scaled Flow Control
Vendor-specific	00110000b	DLLP



OM13781A

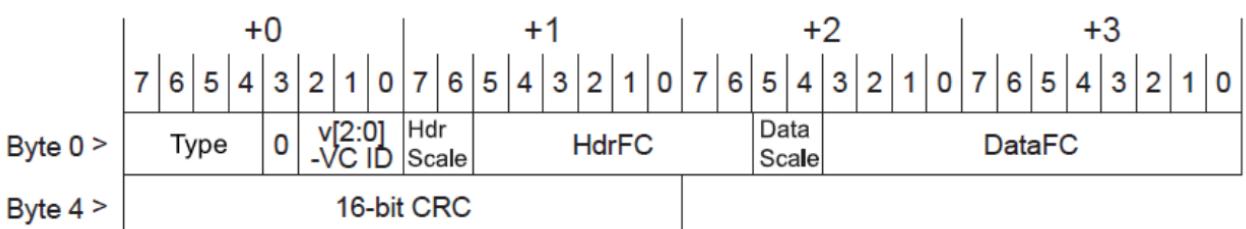
Figure 3-5 Data Link Layer Packet Format for Ack and Nak

AckNak_Seq_Num	TCP	PCIe	Ack Nak	Ack Nak	Retry Buffer	Ack/Nak
Buffer						
DDLP	4		retrain			

VC Virtual Channel	TC	TLP	PCIe	Traffic Class	VC Virtual Channel	VC
--------------------	----	-----	------	---------------	--------------------	----

1. TLP Posted P Non-Posted NP Completion Cpl
  2. VC Link Link VC VC
- N/NP/Cpl

- **InitFC1-P/NP/Cpl**
- **InitFC2-P/NP/Cpl** InitFC1
- **UpdateFC-P/NP/Cpl**



OM13782B

- Type ID

Type	Id
InitFC1-P	0100b
InitFC1-NP	0101b
InitFC1-Cpl	0110b
InitFC2-P	1100b
InitFC2-NP	1101b
InitFC2-Cpl	1110b
UpdateFC-P	1000b
UpdateFC-NP	1001b
UpdateFC-Cpl	1010b

- VC ID v[2:0] Virtual Channel Id Id 3 8 VC
  - HdrFC TLP Credit TLP Header Credit TLP
  - DataFC TLP Credit DW Double Word 4 Data Credit
- 64 128 4 DW TLP 128 Payload 1 DW TLP Digest 1 Header Credit (128 + 4) / 4 = 33 Data Credit
- |                       |      |     |
|-----------------------|------|-----|
| UpdateFC              | 30us | CRC |
| Scaled Flow Control 2 |      |     |

*Table 3-2 Scaled Flow Control Scaling Factors*

Scale Factor	Scaled Flow Control Supported	Credit Type	FC DLLP field			
			Transmitted	Received		
00b	No	Hdr	1	127	8 bits	HdrFC HdrFC
		Data	1	2,047	12 bits	DataFC DataFC
01b	Yes	Hdr	1	127	8 bits	HdrFC HdrFC
		Data	1	2,047	12 bits	DataFC DataFC
10b	Yes	Hdr	4	508	10 bits	HdrFC >> 2 HdrFC << 2
		Data	4	8,188	14 bits	DataFC >> 2 DataFC << 2
11b	Yes	Hdr	16	2,032	12 bits	HdrFC >> 4 HdrFC << 4
		Data	16	32,752	16 bits	DataFC >> 4 DataFC << 4

AVIP log credit 6 stack post non-post cpl head and data

	PH	PD	NPH	NPD	CPL	CPLD	
VC:0	0x390	0x5dd	0x390	0xe4	0	0	
VC:0x1	0	0	0	0	0	0	

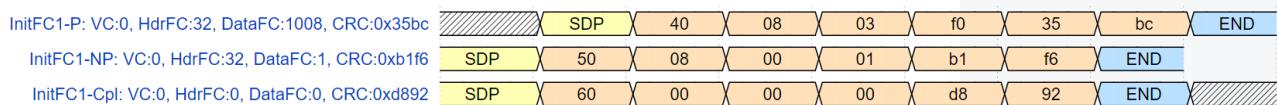
TC->VC map

TC: 0 1 2 3 4 5 6 7

VC: 0 0 0 0 0 0 0 0

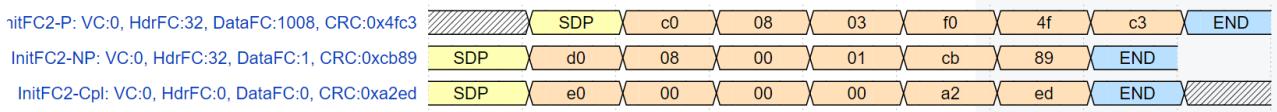
### HdrFC DataFC

#### 1. PCIe Endpoint Switch



#### 2. Switch      Endpoint      Switch Endpoint

#### 3. Switch    InitFC1 DLLP    InitFC2 DLLP



Blog

### 5.2.3 Transaction Layer

[toc]

#### VC Ordering

##### Introduction

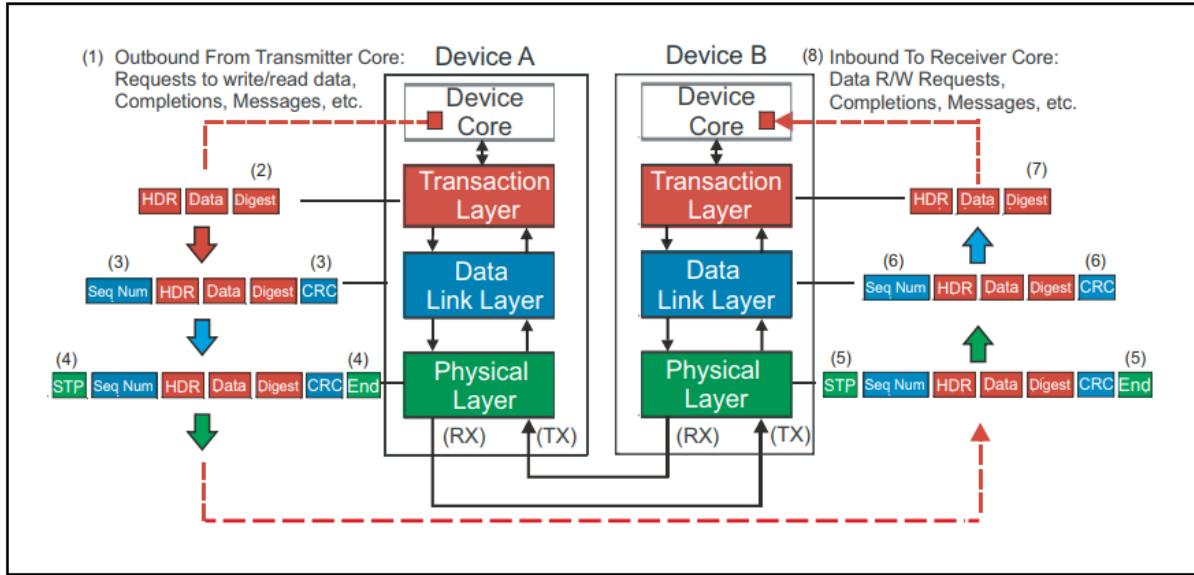
PCIe              Transaction

- Memory Transaction
- IO IO Transaction
- Configuration Transaction
- Message Transaction

- **Non-Posted**              Completion
- **Posted**              Fire and forget

Non-Posted NP Posted P Completion Cpl

Figure 5-2: PCIe TLP Assembly/Disassembly



### TLP Transaction Layer Packet

PCIe      TLP Transaction Layer Packet

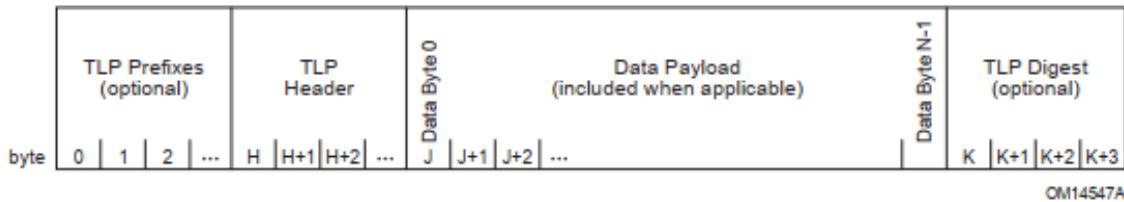


Figure 2-2 Serial View of a TLP

- **TLP Prefix**                      Precision Time Measurement  
SMMU    substreamID    PASID
- **TLP Digest** 4                  CRC  
      DLLP    CRC    CRC    switch        DLLP    CRC    TLP        switch    TLP        TLP
- **TLP Header**                  TLP  
      TLP                              Payload    Payload              Max\_Payload\_Size    4096
- **TLP Payload**                  TLP  
      buffer                      buffer              Max\_Payload\_Size

**TLP**      TLP                      12    3DW    16    4DW    4    DW                      8    DW

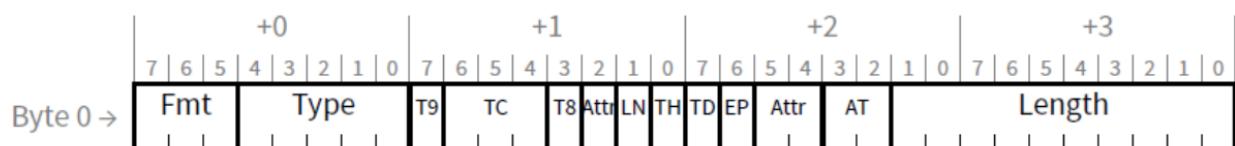


Figure 2-5 Fields Present in All TLP Headers

- **Fmt**

: TLP

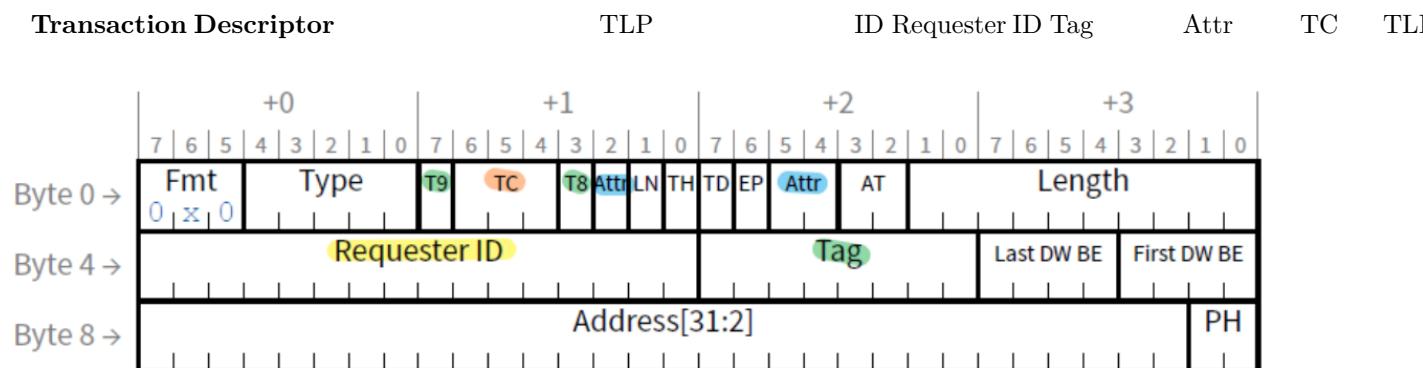
- **Bit 7** 1 Fmt 100 TLP Prefix
- **Bit 6** 1 = TLP Payload 0 = TLP Payload
- **Bit 5** 1 = 32 12 3DW Header 0 = 64 16 4DW Header

- **Type** IO

TLP

- **LN Lightweight Notification**

- **TH TLP Hints** TPH TLP Processing Hint TPH TLP Prefix
- **TD TLP Digest** 1 = TLP Digest 0 = TLP Digest
- **EP Error Poisoning** 1 = 0 =
- **AT Address Type** ATS 00 = 01 = 10 = 11 =
- **Length Payload DW Double Word** 1DW = 4



*Figure 2-18 Request Header Format for 32-bit Addressing of Memory*

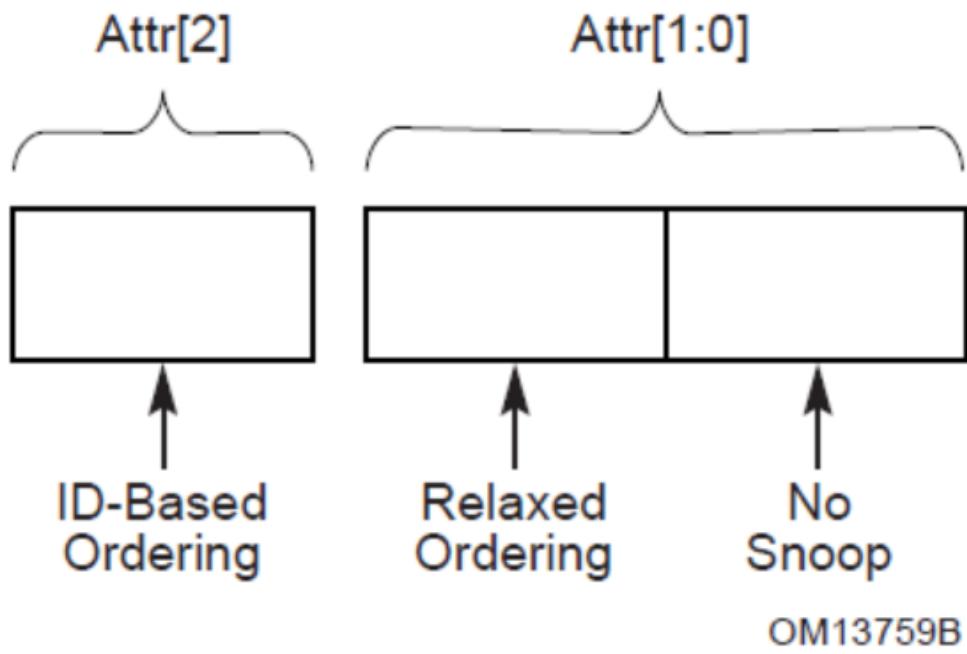
ID Transaction ID

ID Requester ID Tag

- **Requester ID** 16 bit BDF 8:5:3.
- **Tag** 10 bit TLP PCIe T8 T9 bits tag bits 10-Bit Tag Requester Enable

Attributes

bits Attr[2:1] Byte 1 - Bit 2 Byte 2 - Bit 5 Attr[0] Byte 2 - Bit 4 Coherency

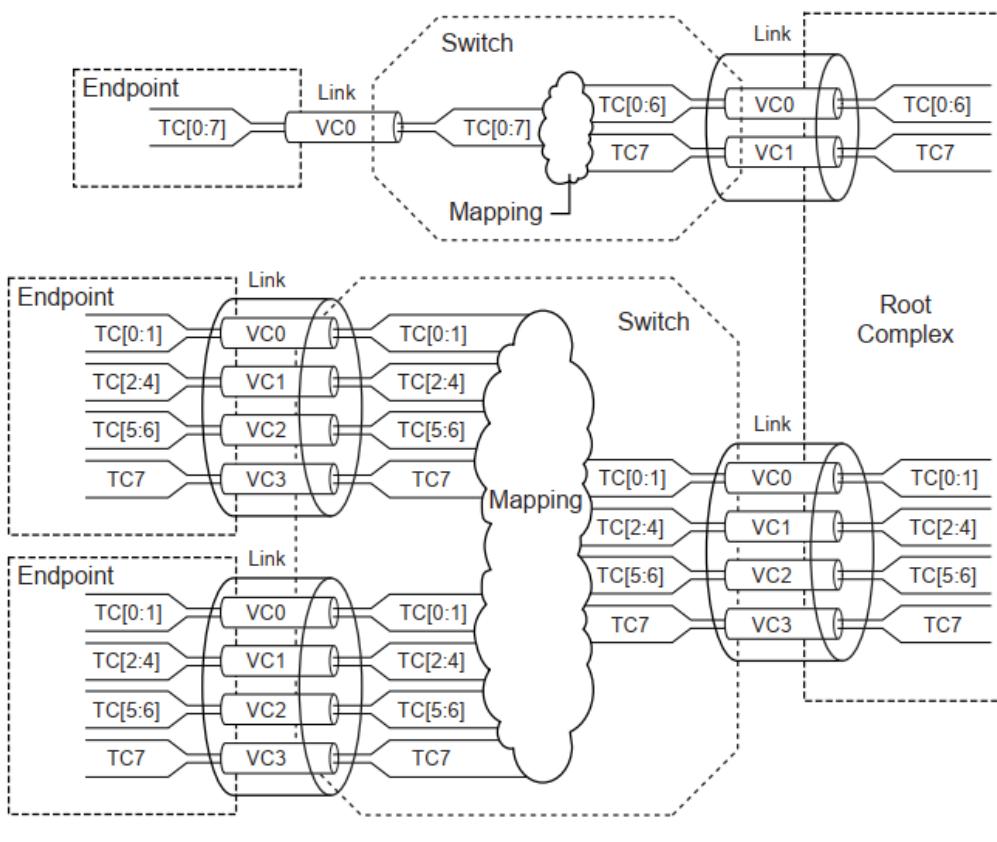


*Figure 2-16 Attributes Field of Transaction Descriptor*

Ordering

PCIe      virtual channel    TLP      TC      buffer      VC      VC    TC TLP      ordering

Figure 2-47 provides a graphical illustration of TC to VC mapping in several different Link configurations. For additional considerations on TC/VC, refer to [Section 6.3](#).



OM13762

[Figure 2-47 An Example of TC/VC Configurations](#)

### 2.5.3 VC and TC Rules

Attr[2:1] Bits

Attr[2]	Attr[1]	
0	0	
0	1	Relaxed Ordering
1	0	ID-based Ordering
1	1	Relaxed Ordering ID-based Ordering

No Snoop

NoSnoop Attr[0]	0	PCIe	Cache	1 PCIe	Cache	Cache
flag		flag				
	IO	MSI	DMA			

Traffic Class

Traffic Class 3 bit 8

TC VC Virtual Channel

1. PCIe Link VC Virtual Channel VC

2. TC VC TLP TC TLP VC  
 3. VC VC Credit Pool VC Credit 0 TLP Credit VC Credit  
 4. TC 0 Hardcoded VC0 TC TLP VC0  
 5. VC TC

**TLP ID**

PCIe

**Address-Based Routing** Memory Transaction IO IO Transaction

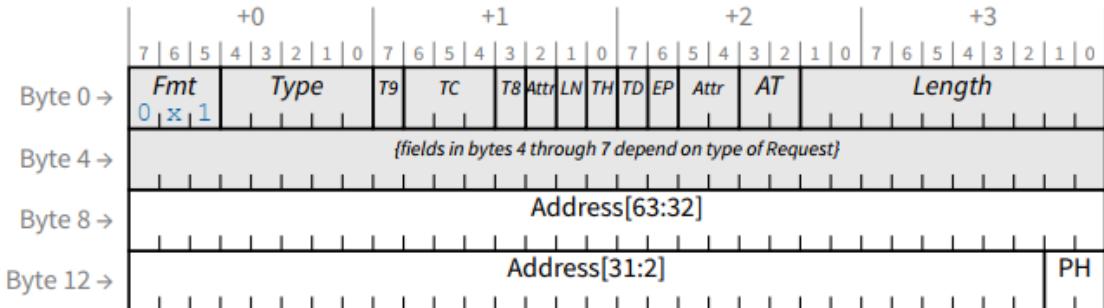


Figure 2-7 64-bit Address Routing

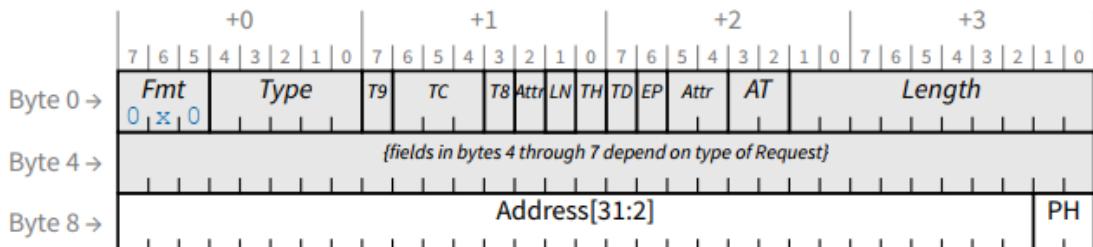


Figure 2-8 32-bit Address Routing

**BDF ID ID Based Routing** Configuration Transaction Message Transaction  
Completion

Implicit routing is used only with **Message Requests**, and is covered in Section 2.2.8

### 5.2.4 PCIe

[toc]

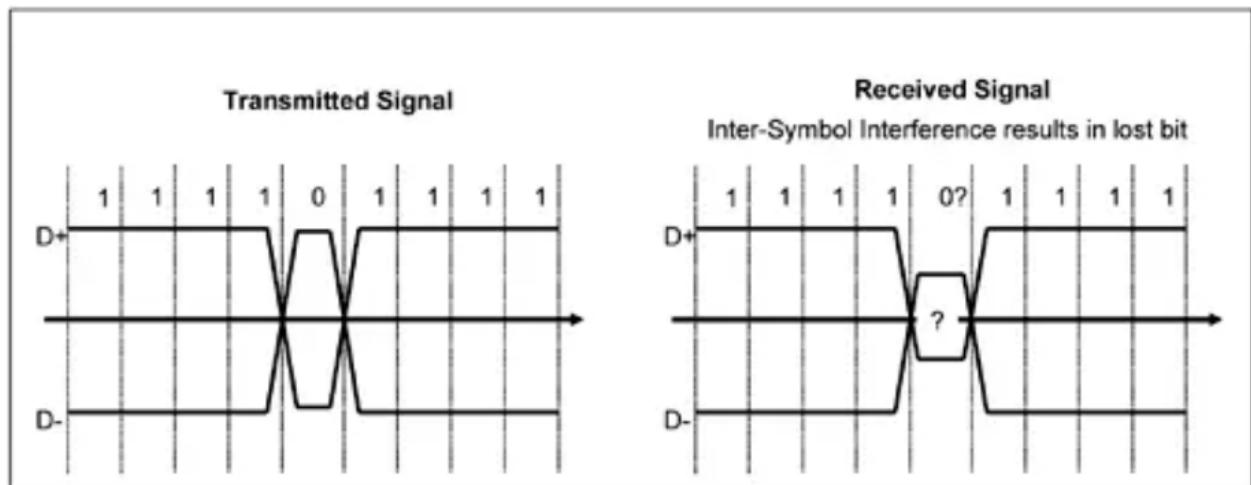
#### Overview

( )

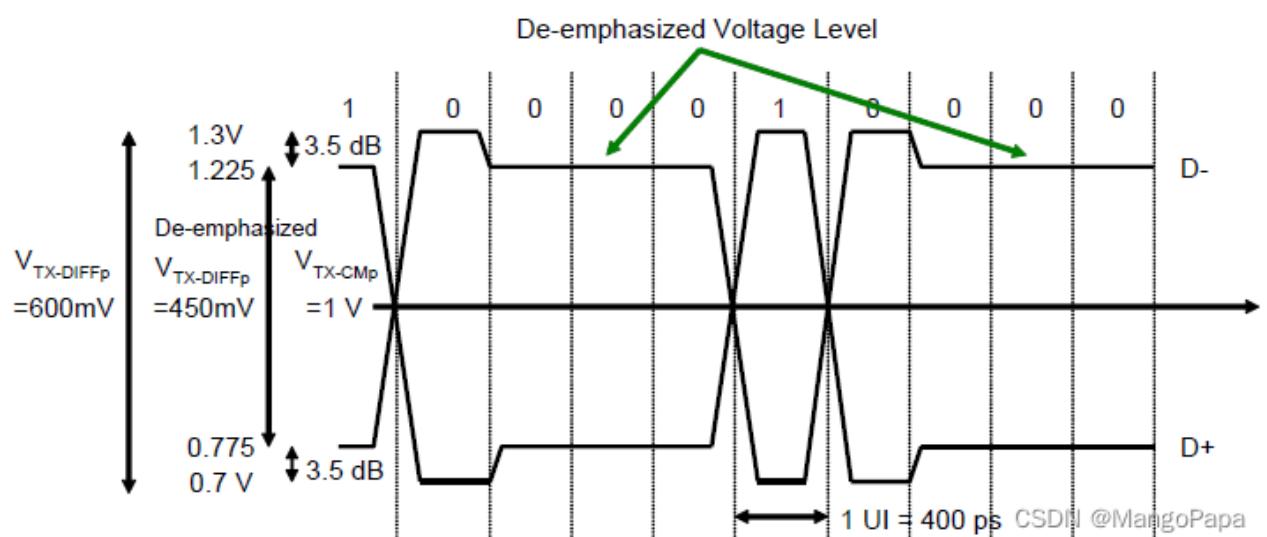
De-emphasis Pre-shoot

Feed forward Equalizer FFE

111101111      0      0      1      0



1      0

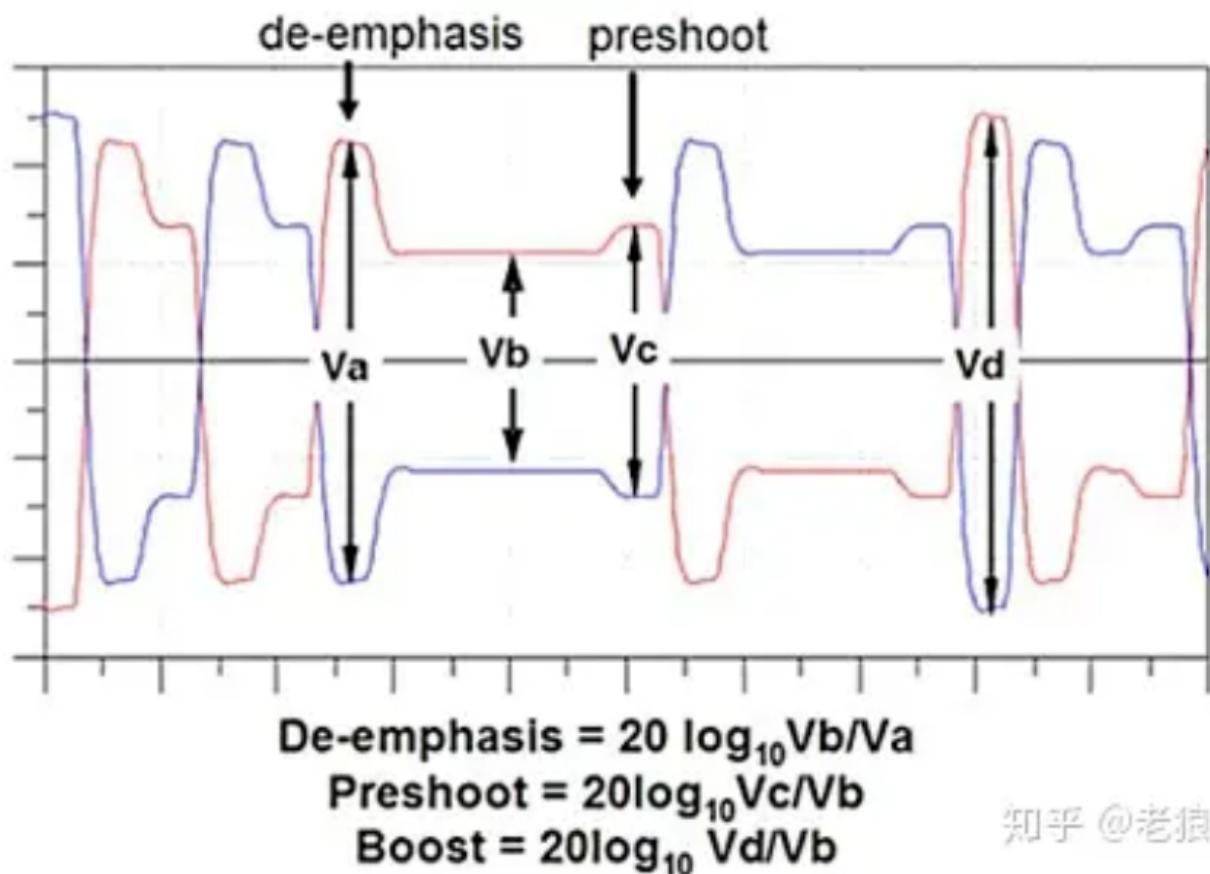


▲ 图 4: 去加重规则示意图

- 1.
- 2.

Preshoot

De-emphasis



FIR 8 GT/s

## Finite Impulse Response, FIR

FIR 5

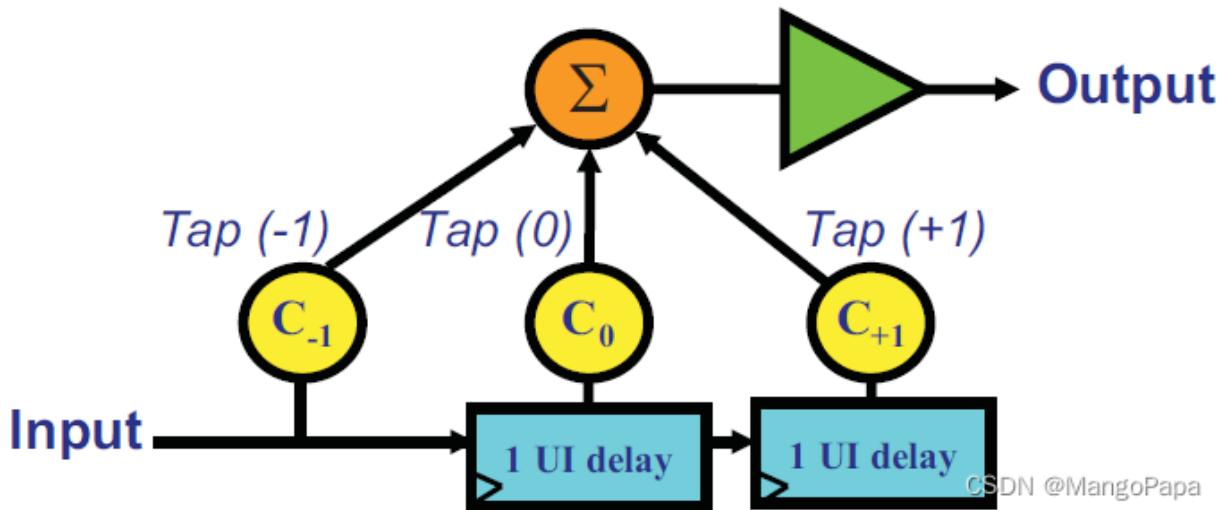
$$V_{out}(n) = C_{-1} V_{in}(n-1) + C_0 V_{in}(n) + C_{+1} V_{in}(n+1)$$

C-1 C0 C+1      FIR      Pre-cursor Cusor      Post-cursor      Pre-cursor      Pre-shoot      Post-cursor

$$|C_{-1}| + |C_0| + |C_{+1}| = 1, \quad C_{-1}$$

Pre-cursor Cusor Post-cursor

Pre-cursor Post-cursor Cursor



de-emphasize preshoot

**preset**

▼ 表 1: Tx Preset Encoding, Ratios and Corresponding Coefficient Values

Encoding	Preset #	Pre-shoot (dB)	De-emphasis (dB)	C <sub>-1</sub>	C <sub>+1</sub>	V <sub>a</sub> /V <sub>d</sub>	V <sub>b</sub> /V <sub>d</sub>	V <sub>c</sub> /V <sub>d</sub>
0000b	P0	0.0	-6.0 ± 1.5 dB	0.000	-0.250	1.000	0.500	0.500
0001b	P1	0.0	-3.5 ± 1 dB	0.000	-0.167	1.000	0.668	0.668
0010b	P2	0.0	-4.4 ± 1.5 dB	0.000	-0.200	1.000	0.600	0.600
0011b	P3	0.0	-2.5 ± 1 dB	0.000	-0.125	1.000	0.750	0.750
0100b	P4	0.0	0.0	0.000	0.000	1.000	1.000	1.000
0101b	P5	1.9 ± 1 dB	0.0	-0.100	0.000	0.800	0.800	1.000
0110b	P6	2.5 ± 1 dB	0.0	-0.125	0.000	0.750	0.750	1.000
0111b	P7	3.5 ± 1 dB	-6.0 ± 1.5 dB	-0.100	-0.200	0.800	0.400	0.600
1000b	P8	3.5 ± 1 dB	-3.5 ± 1 dB	-0.125	-0.125	0.750	0.500	0.750
1001b	P9	3.5 ± 1 dB	0.0	-0.166	0.000	0.668	0.668	1.000
1010b	P10	0.0	不定	0.000	不定	1.000	不定	不定
1011b~1111b	Reserved							

Continuous-Time Linear Equalization, CTLE

Decision Feedback Equalization DFE

**CTLE**

**DFE** DFE

DFE

DFE

**2.5 GT/s** -3.5dB

**5 GT/s** -3.5dB -6dB

**8 GT/s 16 GT/s 32 GT/s** gen1, gen2

gen3 FIR CTLE DFE

PCIe

PCIe

8 GT/s 16 GT/s 32 GT/s

**L0 PCIe EQ EQ DLLP**

- Link Control 3 Register Perform Equalization
- Link Control 2 Register Target Link Speed 8 GT/s
- Link Control Register Retrain Link

DLLP Blocking

Full equalization mode PCIe

training

PCIe EQ 100 ms EQ PCIe Gen5 PCIe

LTSSM Configuration TS Equalization bypass to highest rate

32GT Capability Reg Equalization bypass to highest rate

Bypass equalization to highest rate 2.5 GT/s L0 Recovery 32 GT/s 1  
EQ

32 GT/s No equalization needed 5 GT/s EQ EQ EQ  
Gen1 gen5

Phase PCIe Controller link status Equalization Phase Successful

Phase 3 Equalization Complete

8GT link status

16 32GT 16 32GT status

## EQ

1. EQ

2. EQ

3.

Port EQ      Port EQ

8 GT/s Link Status 2 Register    Link Equalization Request 8 GT/s      16 GT/s 16 GT/s Status Register  
Link Equalization Request 16 GT/s      32 GT/s 32 GT/s Status Register    Link Equalization Request 32  
GT/s

## EQ

1. gen1      DSP      gen3      preset

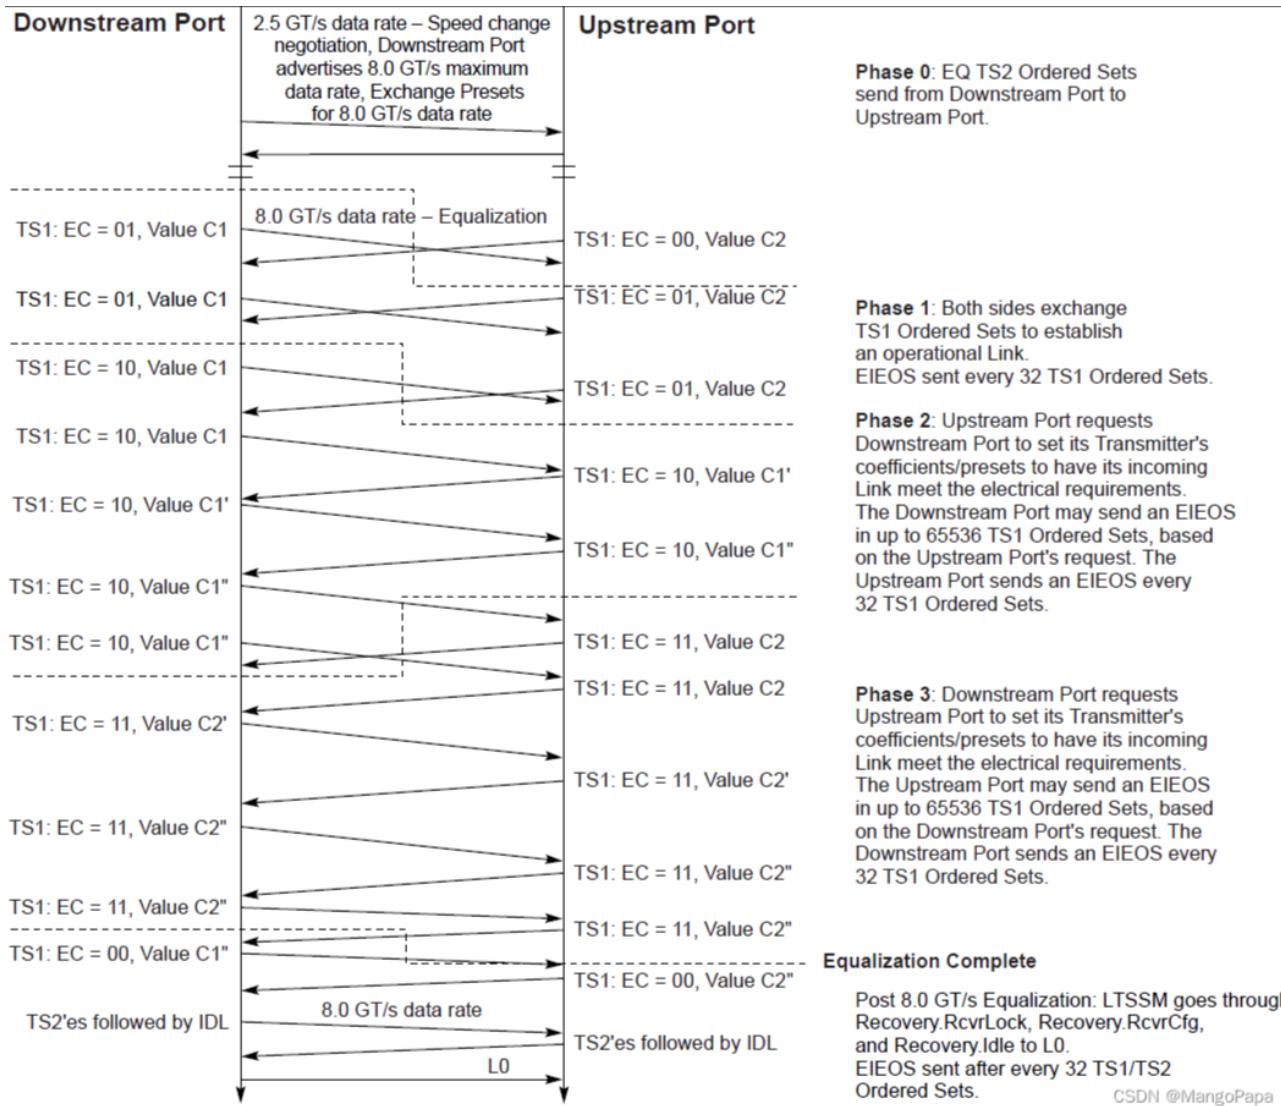
2. EQ Phase 0 EC=0    8 GT/s USP    Preset Coefficients    DSP    DSP    **Phase 0**

3. EQ Phase 1 EC=1 DSP    USP    LF (Low Frequency symbol7) FS (Full Swing symbol8) Post-cursor  
(symbol9)    Phase 2/3    32 TS1    EIEOS

4. EQ Phase 2 EC=2 USP    Master    DSP    Tx Preset    USP    USP      Preset    coefficient  
USP    32 TS1    EIEOS    DSP    USP    65536 TS1    EIEOS

5. EQ Phase 3 EC=3 DSP    Master    USP    Tx Preset    DSP    DSP      Preset    coefficient  
DSP    32 TS1    EIEOS    USP    DSP    65536 TS1    EIEOS

6. EQ    EC=0 LTSSM    Recovery.RcvrLock -> Recovery.RcvrCfg -> Recovery.Idle -> L0    L0    32  
TS1/TS2    EIEOS



PCIe EQ DSP phase1

phase0

<https://mangopapa.blog.csdn.net/article/details/124539607>

### 5.3 ARM N2

ARM N2 ARM Neoverse N2 reference design ARM Neoverse N2 ARM  
CPU

1. N2 HPC
2. Neoverse N2
3. ARM N2
4. N2
5. Neoverse N2 ARM ARM

ARM N2

ARM Neoverse N2 ARM IP IP

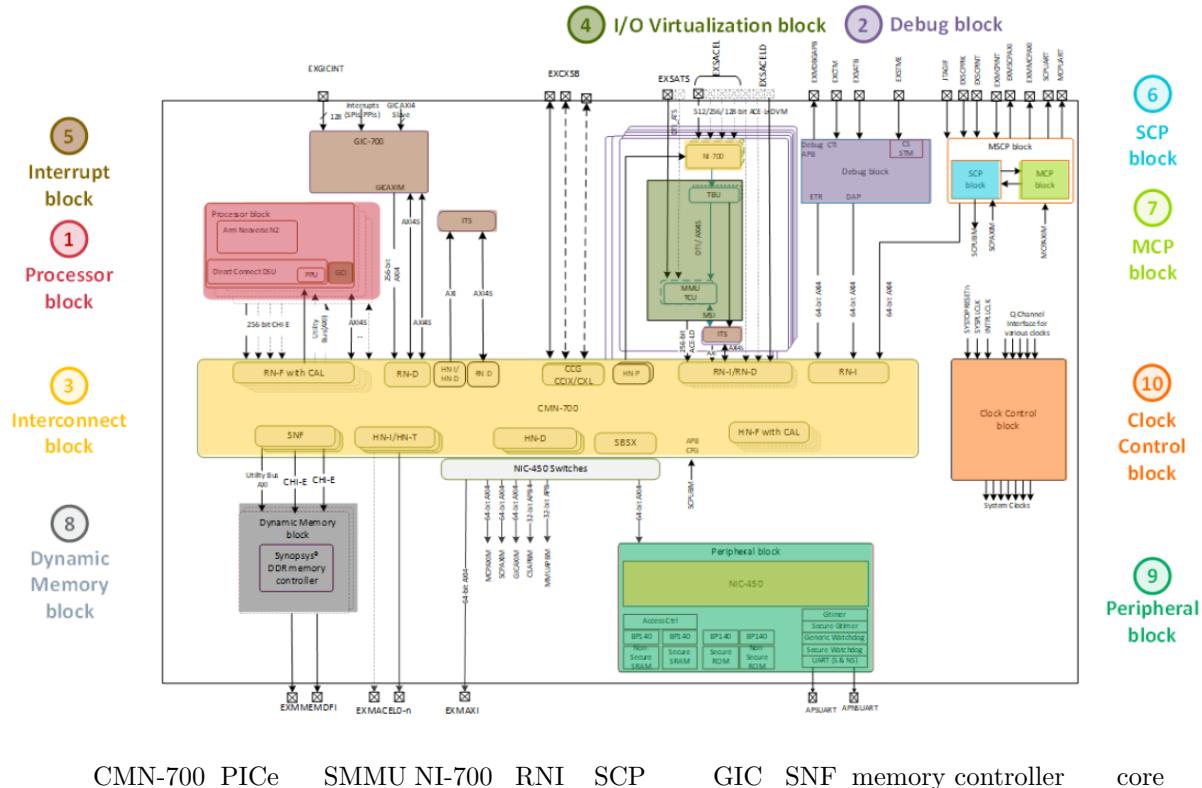
1. ARM Neoverse N2 CPU N2 N2 ARMv9
2. System Level Cache (SLC) ARM N2 SLC
3. CCIX CXL N2 CCIX Cache Coherent Interconnect for Accelerators CXL Compute Express Link
4. AMBA (Advanced Microcontroller Bus Architecture) AMBA SoC IP N2 AMBA
5. TrustZone ARM TrustZone
6. Mali GPU N2 CPU ARM Mali GPU
7. SMMU (System Memory Management Unit)

IP ARM Neoverse N2

High Level

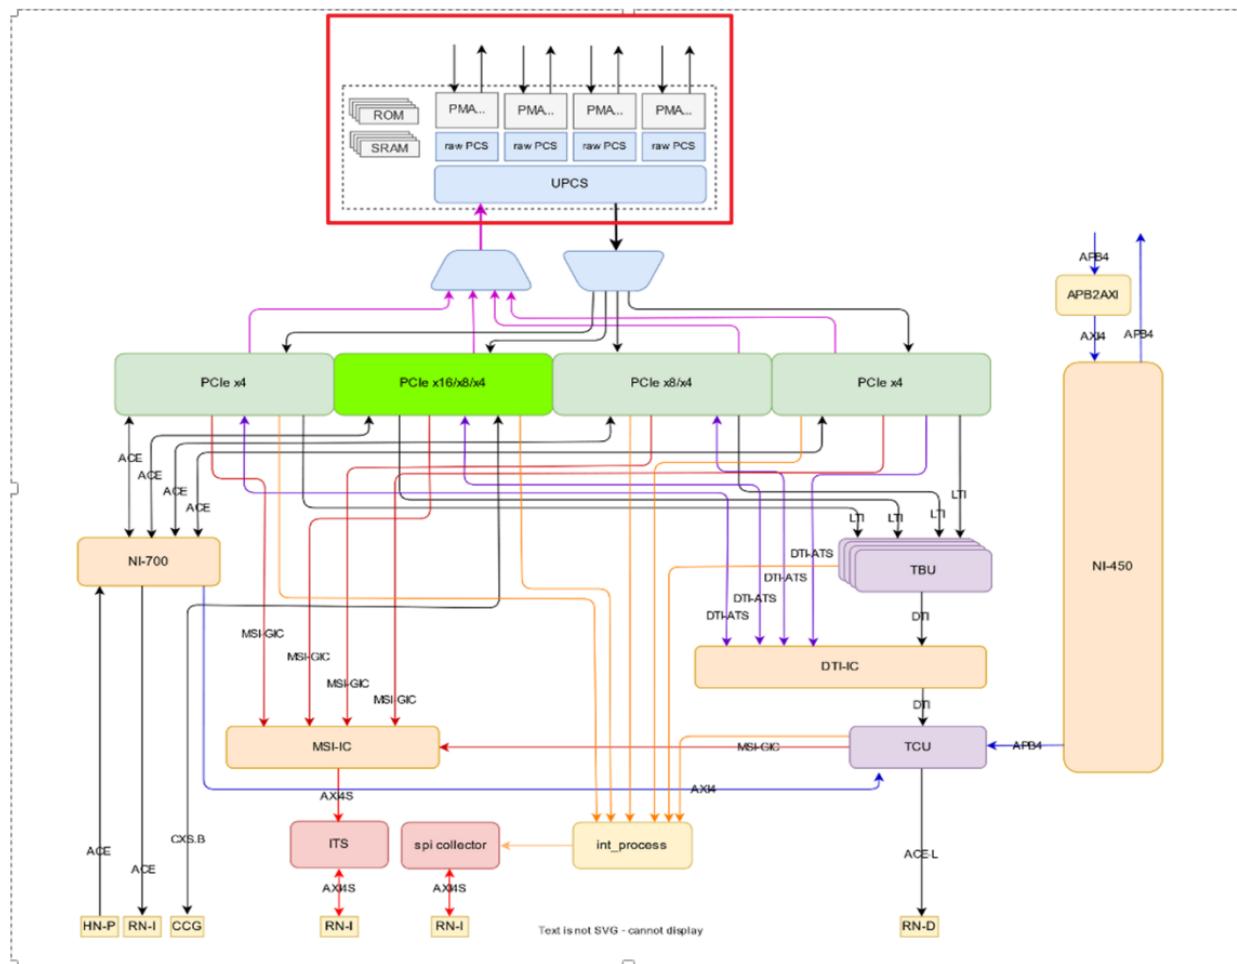
[Arm Neoverse N2 reference design Technical Overview](#)

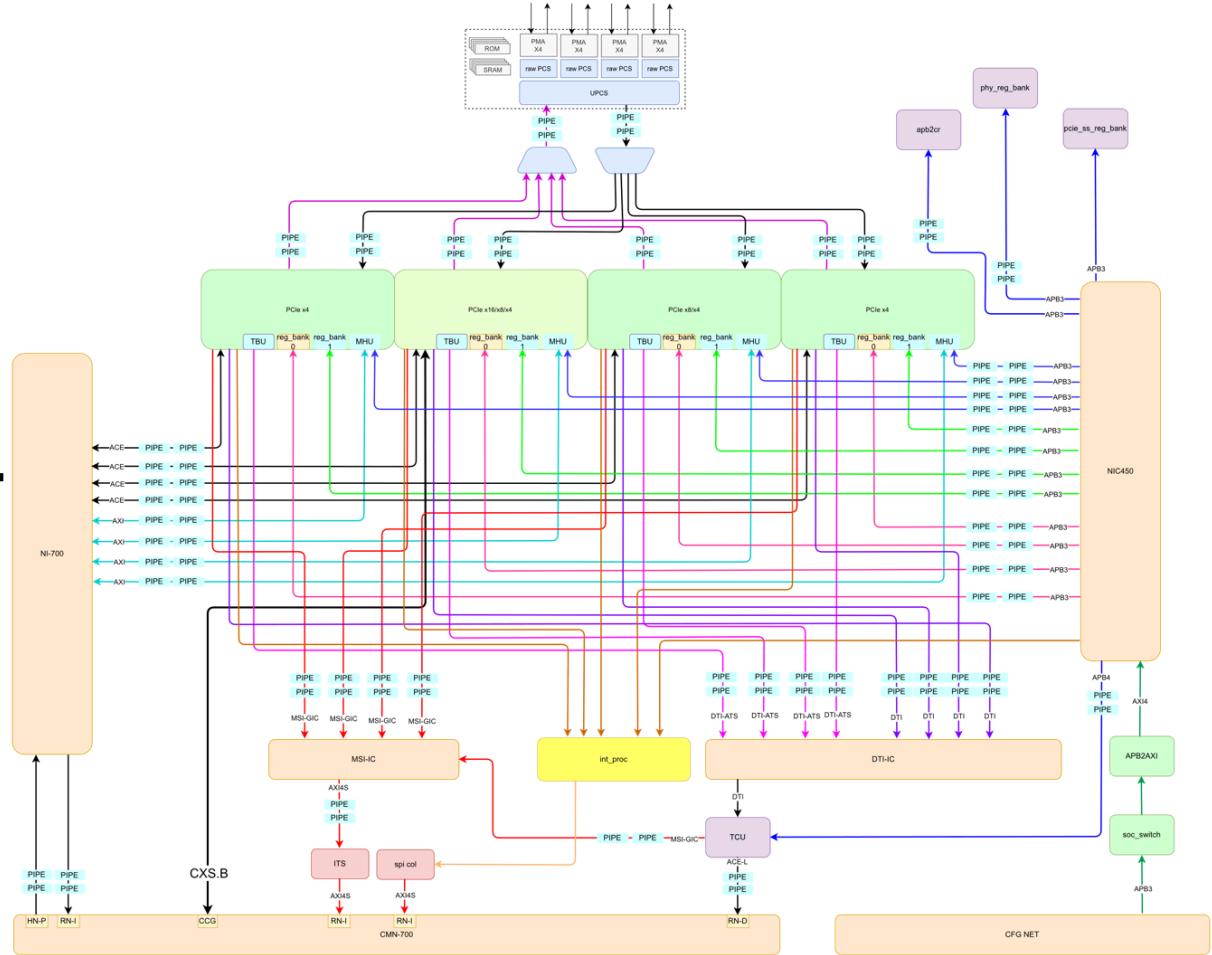
**Figure 3-1: RD-N2 architecture block diagram**



CMN-700 PICe SMMU NI-700 RNI SCP GIC SNF memory controller core

### 5.3.1 PCIe





4 S CXL CXL PCIe x16 x4 4 controller bifircation

CMN RNI HNP CCG NI-700 NI-700 controller X16 CCG CXL CCIX

SMMU-700 MMU TBU TCU DTIIC

MSI-IC MSI switch controller MSI MSI

ITS PCIe MSI GIC ITS GIC

int process

NIC-450 NI-700 core cfg controller NI-700 PCIe AMBA ACE

## 5.4 PCIe

### 5.4.1 Resource Degration

[toc]

BAR 5 resource pool

图1是Memory BAR寄存器的结构，图2是IO类型的BAR寄存器

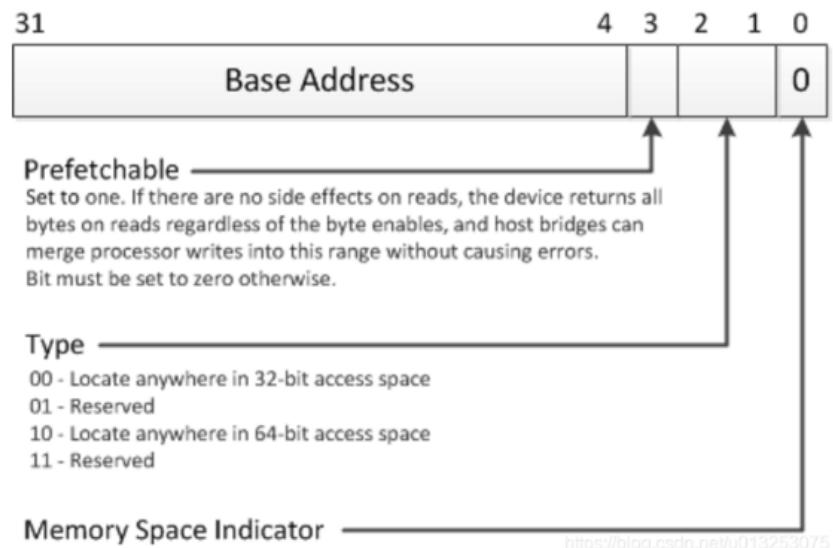


图1 Base Address Register for Memory

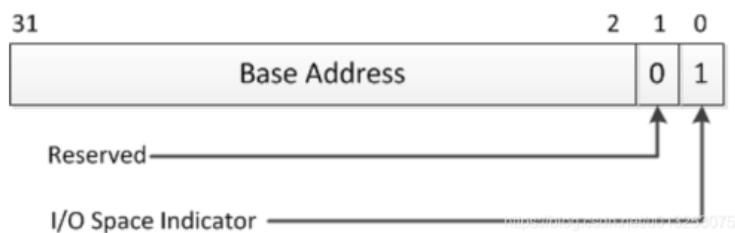


图2 Base Address Register for I/O

**bit0:**表示设备寄存器是映射到memory (0) 还是IO (1)空间。

**bit1:** reserved 0

**bit2:** 在base address register for Memory 中0表示32位地址空间，1表示64位地址空间。

**bit3:**在memory BAR中用来表示该设备是否允许prefetch, 1表示可以预取, 0表示不可以预区。

**bit4~31:**用来表示设备需要占用的地址空间大小。

```
InitializeResourcePool (&IoPool, PciBarTypeIo16);  
InitializeResourcePool (&Mem32Pool, PciBarTypeMem32);  
InitializeResourcePool (&PMem32Pool, PciBarTypePMem32);  
InitializeResourcePool (&Mem64Pool, PciBarTypeMem64);  
InitializeResourcePool (&PMem64Pool, PciBarTypePMem64);
```

RC                   prefetchable.

```

//  

// The base address of MMIO Registers  

//  

#define AC01_PCIE_MMIO_BASE_LIST      0x300000000000, 0x340000000000, 0x380000000000, 0x3C0000000000  

//  

// The size of MMIO space  

//  

#define AC01_PCIE_MMIO_SIZE_LIST     0x3FFE00000000, 0x3FFE00000000, 0x3FFE00000000, 0x3FFE00000000  

//  

// The base address of MMIO32 Registers  

//  

#define AC01_PCIE_MMIO32_BASE_LIST   0x000020000000, 0x000028000000, 0x000030000000, 0x000038000000  

//  

// The size of MMIO32 space  

//  

#define AC01_PCIE_MMIO32_SIZE_LIST   0x8000000, 0x8000000, 0x8000000, 0x8000000, 0x4000000

```

BAR 11              Prefetchable vs. Non-Prefetchable Memory Space

- Reads do not have side effects
- Write merging is allowed

PCIe Technology 3.0 Chapter 4.1.2.2

## 1. DegradeResource()

### 1.1 Degrade Resource For OptionRom

If any child device has both option ROM and 64-bit BAR, degrade its PMEM64/MEM64 requests in case that if a legacy option ROM image can not access 64-bit resources.

PCD PcdPciDegradeResourceForOptionRom

- X64 ---- TRUE
- X32 ARM AARCH64 ---- FALSE

root bridge bar0 mem32 bar0 bar 64bit

### 1.2 Degrade Resource For Others

```

//  

// Degrade resource if necessary  

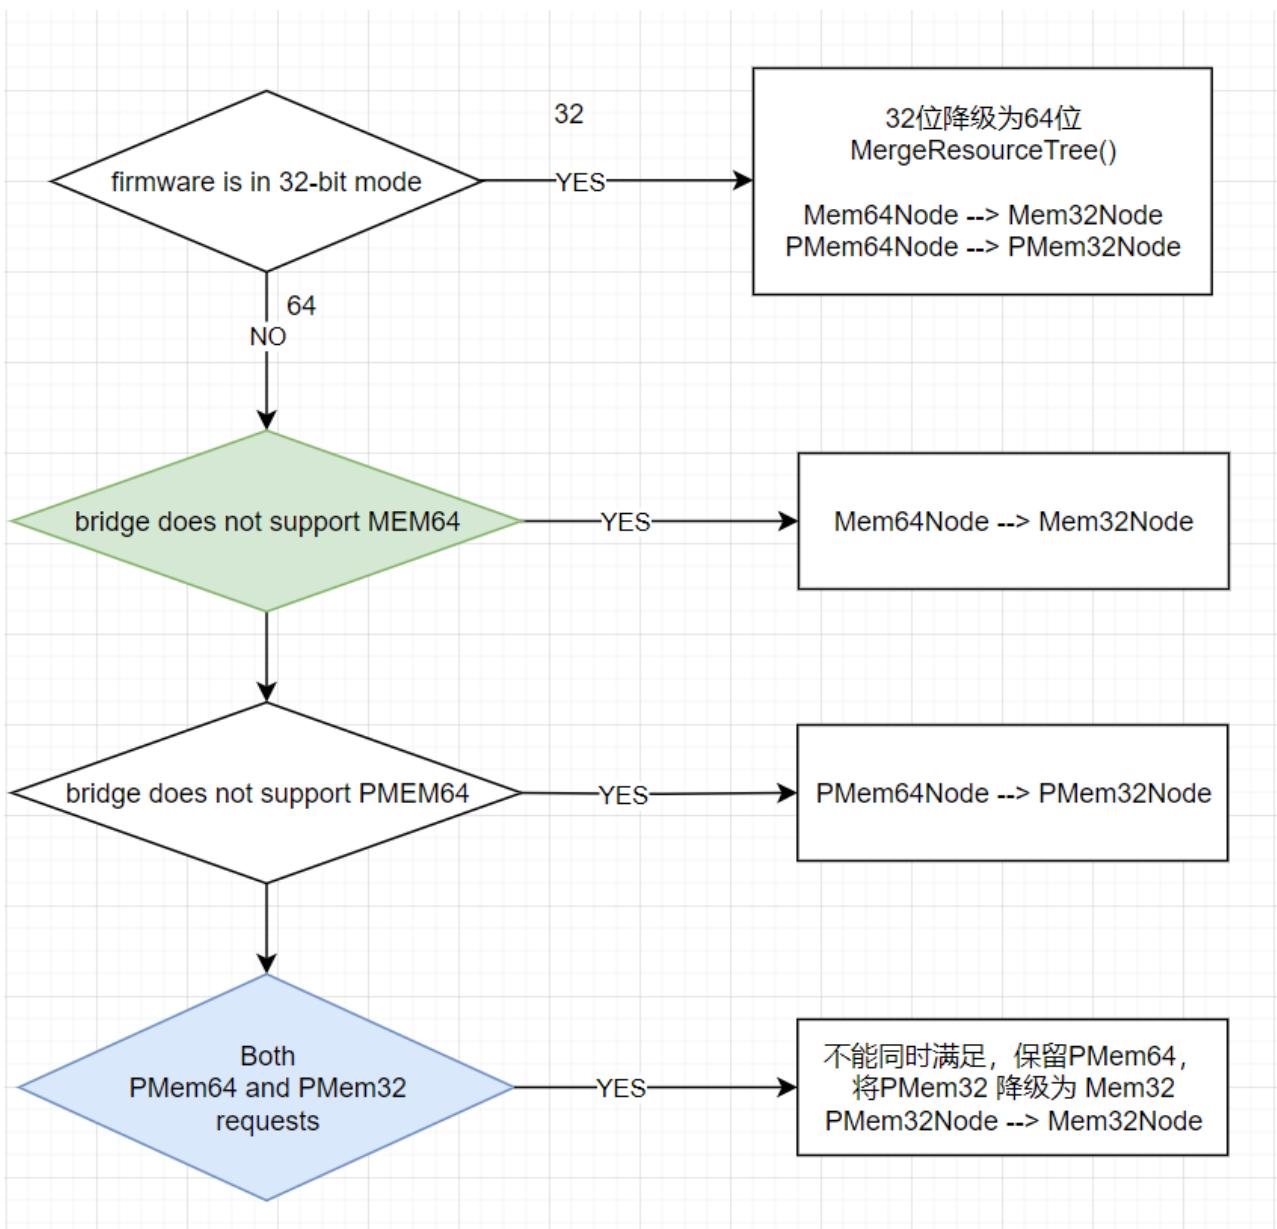
//  

DegradeResource (Bridge, Mem32Node, PMem32Node, Mem64Node, PMem64Node);  

BridgeSupportResourceDecode() bridge  

MergeResourceTree() resoure tree

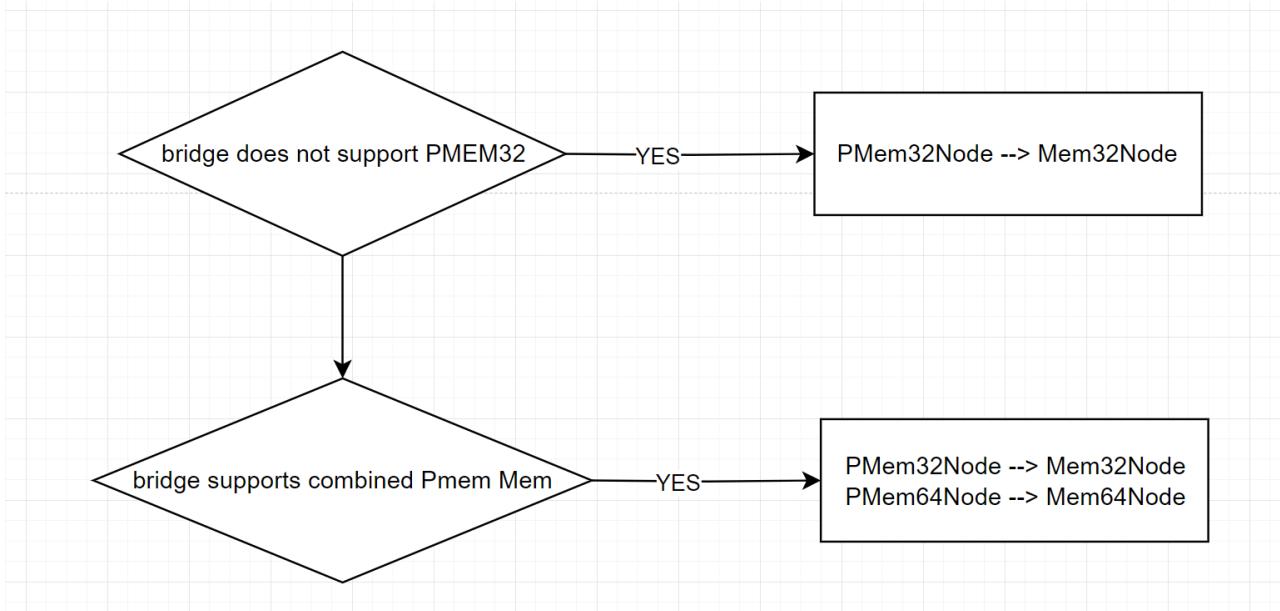
```



PMem64 PMem32 bridge

NP-Mem64 32 32

Secondary Status	I/O Limit	I/O Base	
Memory Limit	Memory Base		20h
Prefetchable Memory Limit	Prefetchable Memory Base		24h
Prefetchable Base Upper 32 Bits			28h
Prefetchable Base Limit 32 Bits			2Ch
I/O Limit Upper 16 Bits	I/O Base Upper 16 Bits		30h
		Capabilities	



2.

```

0004:03:00.0 USB controller: Renesas Technology Corp. uPD720201 USB 3.0 Host Controller (rev 03) (prog-
Control: I/O- Mem+ BusMaster+ SpecCycle- MemWINV- VGASnoop- ParErr- Stepping- SERR+ FastB2B- DisINT-
Status: Cap+ 66MHz- UDF- FastB2B- ParErr- DEVSEL=fast >TAbort- <TAbort- <MAbort- >SERR- <PERR- INTx
Latency: 0, Cache Line Size: 64 bytes
Interrupt: pin A routed to IRQ 198
NUMA node: 0
IOMMU group: 0
Region 0: Memory at 11800000 (64-bit, non-prefetchable) [size=8K]
Capabilities: [50] Power Management version 3
  Flags: PMEClk- DS1- D1- D2- AuxCurrent=375mA PME(D0+,D1-,D2-,D3hot+,D3cold+)
  Status: D0 NoSoftRst+ PME-Enable- DSel=0 DScale=0 PME-
Capabilities: [70] MSI: Enable- Count=1/8 Maskable- 64bit+
  
```

## 2.1 Ampere PCI log

```

0004:03:00.0 USB controller: Renesas Technology Corp. 1
00: 12 19 14 00 06 05 10 00 03 30 03 0c 10 00 00 00
10: 04 00 80 11 00 00 00 00 00 00 00 00 00 00 00 00
20: 00 00 00 00 00 00 00 00 00 00 00 00 ff ff ff ff
30: 00 00 00 00 50 00 00 00 00 00 00 00 ff 01 00 00
40: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  
```

# 1180 0004

HEX 1180 0004

DEC 293,601,284

OCT 2 140 000 004

BIN 0001 0001 1000 0000 0000 0000 0000 0100



bit0 - 0 memory

bit 2:1 memory 00 32bit memory 10- 64 bit memory B

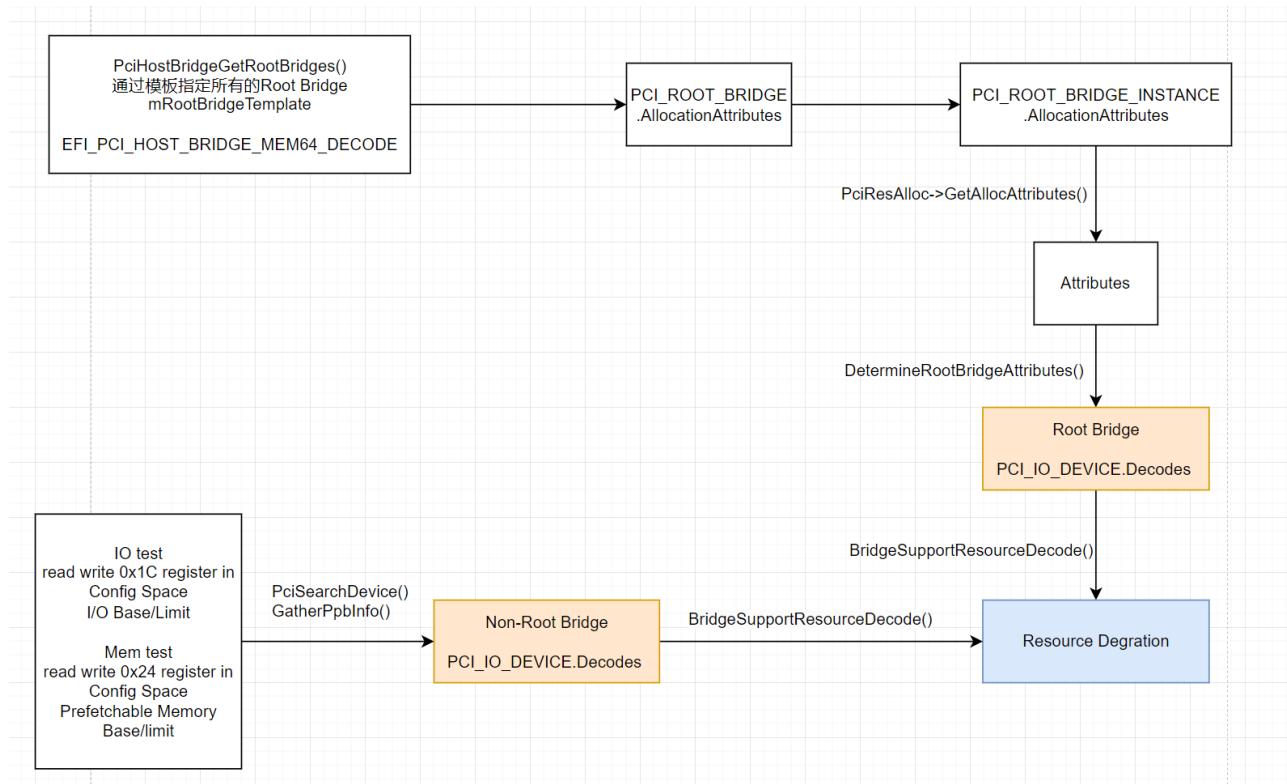
bit3 - 0 Non-prefetchable

## NP-Mem64

### 2.2 Degrade Resource

BridgeSupportResourceDecode()

PCI\_IO\_DEVICE.Decodes



Secondary Status	I/O Limit	I/O Base
Memory Limit	Memory Base	20h
Prefetchable Memory Limit	Prefetchable Memory Base	24h
	Prefetchable Base Upper 32 Bits	28h
	Prefetchable Base Limit 32 Bits	2Ch
I/O Limit Upper 16 Bits	I/O Base Upper 16 Bits	30h
		Capabilities

bridge

```
#define EFI_BRIDGE_IO32_DECODE_SUPPORTED      0x0001
#define EFI_BRIDGE_PMEM32_DECODE_SUPPORTED     0x0002
#define EFI_BRIDGE_PMEM64_DECODE_SUPPORTED     0x0004
#define EFI_BRIDGE_IO16_DECODE_SUPPORTED       0x0008
#define EFI_BRIDGE_PMEM_MEM_COMBINE_SUPPORTED 0x0010
#define EFI_BRIDGE_MEM64_DECODE_SUPPORTED      0x0020
#define EFI_BRIDGE_MEM32_DECODE_SUPPORTED      0x0040

struct _PCI_IO_DEVICE {
    ...
    // 
    // The resource decode the bridge supports
    //
    UINT32 Decodes;
    ...
}

PCI_IO_DEVICE.Decodes = DetermineRootBridgeAttributes() protocol

PciResAlloc->GetAllocAttributes (
    PciResAlloc,
    RootBridgeHandle,
    &Attributes
);
```

- root bridge            RC    root bridge
- bridge      bridge

*Table 7-10 I/O Addressing Capability*

Bits 3:0	I/O Addressing Capability
0h	16-bit I/O addressing
1h	32-bit I/O addressing
2h-Fh	Reserved

- IO capability -- ACPI spec 7.5.1.3.6
- Mem capability -- PCIe spec 7.5.1.3.9

The bottom 4 bits of both the Prefetchable Memory Base 0x24 and Prefetchable Memory Limit registers are read-only, contain the same value, and encode whether or not the bridge supports 64-bit addresses

- \* 0 - the bridge supports only 32 bit addresses.
- \* 1 - the bridge supports 64-bit addresses. bridge

```

0004:00:02.0 PCI bridge: Ampere Computing, LLC Altra PCI Express Root Port a1 (rev 04)
00: ef 1d 02 e1 07 01 10 00 04 00 04 06 10 00 01 00
10: 00 00 00 00 00 00 00 00 03 03 00 f0 00 00 00
20: 80 11 90 11 21 00 31 00 00 28 00 00 00 28 00 00
30: 00 00 00 00 40 00 00 00 00 00 00 ff 01 02 00
40: 01 70 c3 5f 08 00 00 00 00 00 00 00 00 00 00 00
50: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
60: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
70: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
    
```

Prefetchable Memory Base 0x0021

Prefetchable Memory Limit 0x0031

0x28	32	PMEM32	PMEM64 flag	PMEM32	PMEM32	MEM32	PMEM32	PMEM64	PMEM32
PCIe Spec		EDK2	Root Bridge	bridge	MEM64	Flag	bridge	NP-MEM64	NP-MEM64
NP-MEM64 tree		merge	NP-MEM32 tree						

RC

```

InitializeResourcePool (&IoPool, PciBarTypeIo16);
InitializeResourcePool (&Mem32Pool, PciBarTypeMem32);
// InitializeResourcePool (&Pmem32Pool, PciBarTypePmem32);
// InitializeResourcePool (&Mem64Pool, PciBarTypeMem64);
InitializeResourcePool (&Pmem64Pool, PciBarTypePmem64);
    
```

## 2.3 RC RC Root Bridge

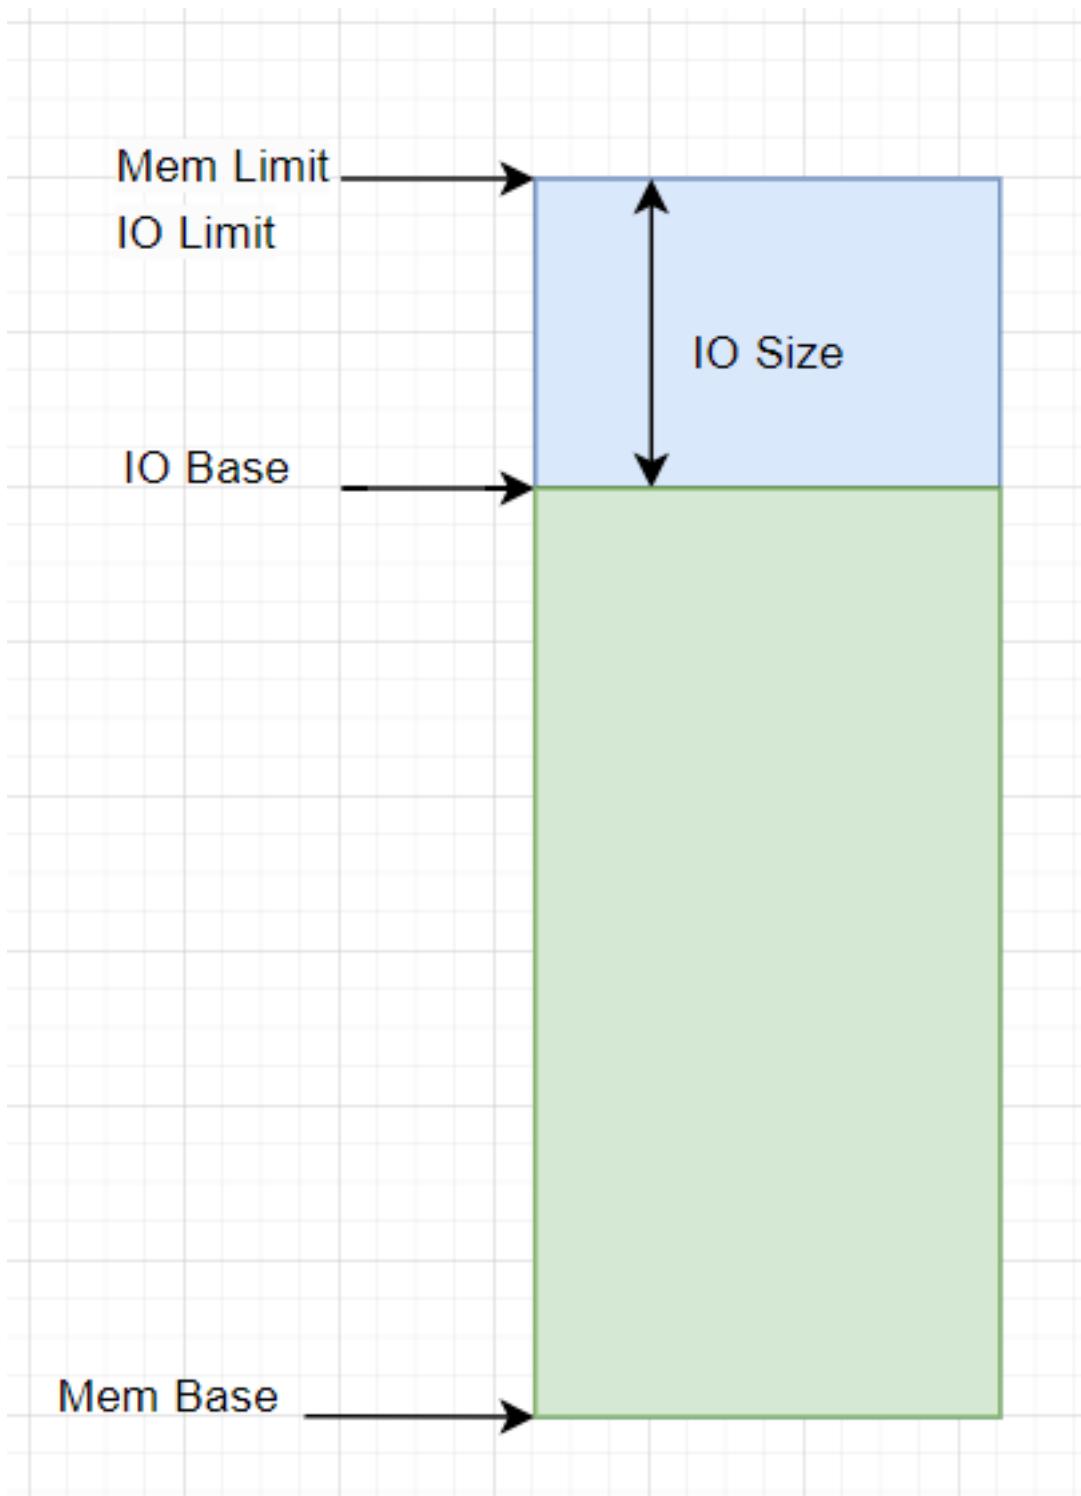
Mem PMem =Mmio32Base

Io.Base = Mem Limit IO\_SIZE

PMemAbove4G = MmioBase

MemAbove4G

```
{ // PMem  
MAX_UINT64,  
0,  
0  
}
```



```

CopyMem (&RootBridge->Bus, &Bridge->Bus, sizeof (PCI_ROOT_BRIDGE_APERTURE));
CopyMem (&RootBridge->Io, &Bridge->Io, sizeof (PCI_ROOT_BRIDGE_APERTURE));
CopyMem (&RootBridge->Mem, &Bridge->Mem, sizeof (PCI_ROOT_BRIDGE_APERTURE));
CopyMem (&RootBridge->MemAbove4G, &Bridge->MemAbove4G, sizeof (PCI_ROOT_BRIDGE_APERTURE));
CopyMem (&RootBridge->PMem, &Bridge->PMem, sizeof (PCI_ROOT_BRIDGE_APERTURE));
CopyMem (&RootBridge->PMemAbove4G, &Bridge->PMemAbove4G, sizeof (PCI_ROOT_BRIDGE_APERTURE));

```

```

if (RootComplex->Mmio32Base != 0) {
    RootBridge->Mem.Base = RootComplex->Mmio32Base;
    RootBridge->Mem.Limit = RootComplex->Mmio32Base + RootComplex->Mmio32Size - 1;
    RootBridge->PMem.Base = RootBridge->Mem.Base;
    RootBridge->PMem.Limit = RootBridge->Mem.Limit;
    RootBridge->Io.Base = RootComplex->Mmio32Base + RootComplex->Mmio32Size -
    ← AC01_PCIE_IO_SIZE;
    RootBridge->Io.Limit = RootBridge->Mem.Limit;
}

if (RootComplex->MmioBase != 0) {
    RootBridge->PMemAbove4G.Base = RootComplex->MmioBase;
    RootBridge->PMemAbove4G.Limit = RootComplex->MmioBase + RootComplex->MmioSize - 1;
}

if ((Attributes & EFI_PCI_HOST_BRIDGE_COMBINE_MEM_PMEM) != 0) {
    RootBridgeDev->Decodes |= EFI_BRIDGE_PMEM_MEM_COMBINE_SUPPORTED;
}

if ((Attributes & EFI_PCI_HOST_BRIDGE_MEM64_DECODE) != 0) {
    RootBridgeDev->Decodes |= EFI_BRIDGE_MEM64_DECODE_SUPPORTED;
    RootBridgeDev->Decodes |= EFI_BRIDGE_PMEM64_DECODE_SUPPORTED;
}

RootBridgeDev->Decodes |= EFI_BRIDGE_MEM32_DECODE_SUPPORTED;
RootBridgeDev->Decodes |= EFI_BRIDGE_PMEM32_DECODE_SUPPORTED;
RootBridgeDev->Decodes |= EFI_BRIDGE_IO16_DECODE_SUPPORTED;

```

NP-MEM64 NP-MEM32 bar NP-MEM32

Linux kernel 5.18.0 EDK2

PCIe Spec NP-MEM64

### 1.3.2.2 PCI Express Endpoint Rules

- A PCI Express Endpoint must be a Function with a Type 00h Configuration Space header.
- A PCI Express Endpoint must support Configuration Requests as a Completer.
- A PCI Express Endpoint must not depend on operating system allocation of I/O resources claimed through BAR(s).
- A PCI Express Endpoint must not generate I/O Requests.
- A PCI Express Endpoint must not support Locked Requests as a Completer or generate them as a Requester. PCI Express-compliant software drivers and applications must be written to prevent the use of lock semantics when accessing a PCI Express Endpoint.
- A PCI Express Endpoint operating as the Requester of a Memory Transaction is required to be capable of generating addresses greater than 4 GB.
- A PCI Express Endpoint is required to support MSI or MSI-X or both if an interrupt resource is requested., If MSI is implemented, a PCI Express Endpoint must support the 64-bit Message Address version of the MSI Capability structure.
- A PCI Express Endpoint requesting memory resources through a BAR must set the BAR's Prefetchable bit unless the range contains locations with read side-effects or locations in which the Function does not tolerate write merging. See Section 7.5.1.2.1 for additional guidance on having the Prefetchable bit Set.
  - For a PCI Express Endpoint, 64-bit addressing must be supported for all BARs that have the Prefetchable bit Set. 32-bit addressing is permitted for all BARs that do not have the Prefetchable bit Set.
  - The minimum memory address range requested by a BAR is 128 bytes.
  - A PCI Express Endpoint must appear within one of the hierarchy domains originated by the Root Complex.

### PCIe 3.0&4.0 Spec

- For a PCI Express Endpoint, 64-bit addressing must be supported for all BARs that have the Prefetchable bit Set. 32-bit addressing is permitted for all BARs that do not have the Prefetchable bit Set.

### 5.4.2 EDK2 PCIe OpRom

[toc]

#### 1. OpRom

1.1	0x30	Expansion ROM Base Address Register	option rom	Bridge	0x38
		Reserved	Capabilities Pointer		34h
		Expansion ROM Base Address Register (XROMBAR)			38h
		Bridge Control	Interrupt Pin	Interrupt	3Ch

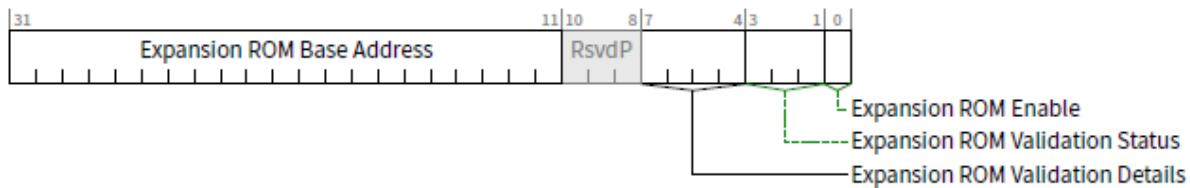


Figure 7-13 Expansion ROM Base Address Register

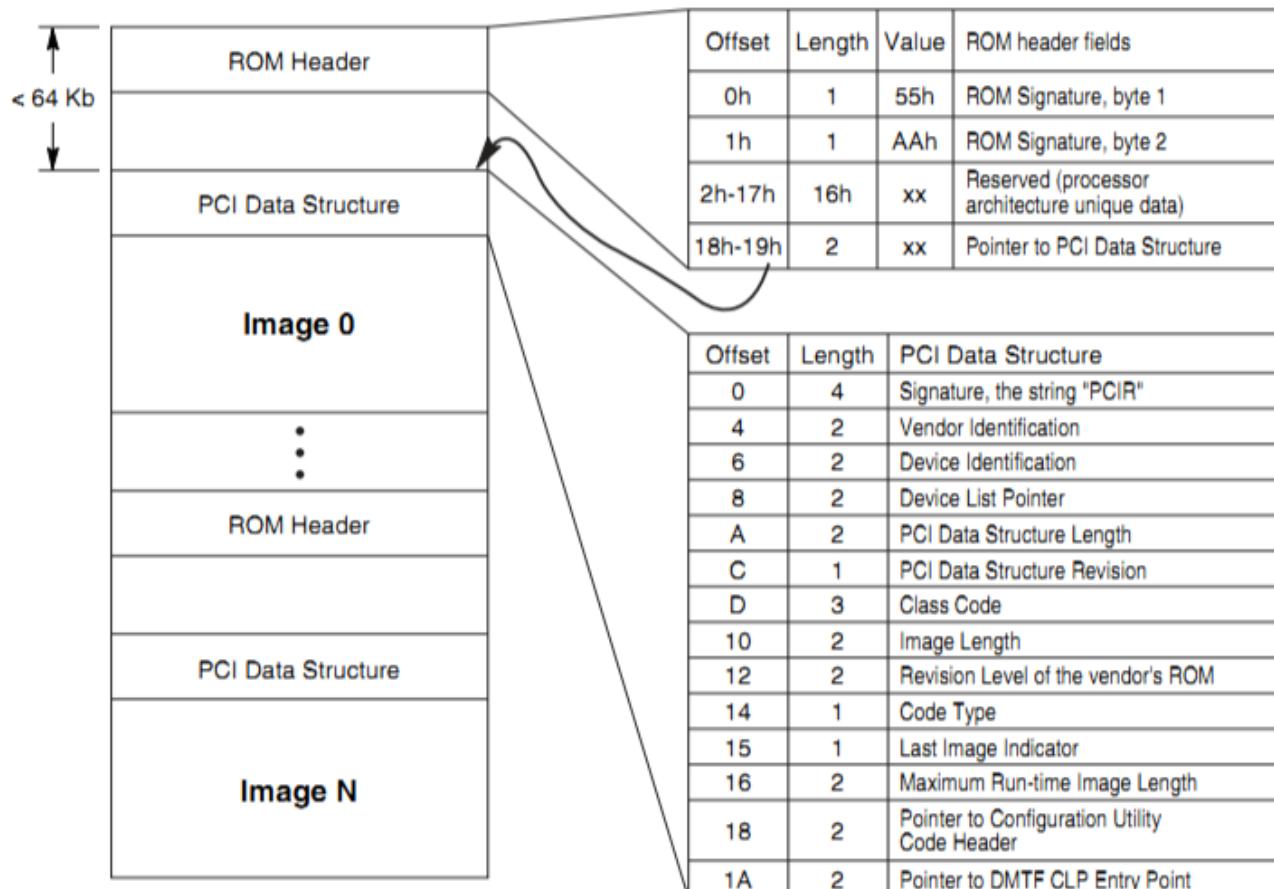
21bit      8bit

## 1.2 OpRom      Plug and Play BIOS Specification

EFI takes advantage of both the **PCI Firmware Specification** and the PE/COFF Specification to store EFI images in a PCI Option ROM. There are several rules that must be followed when constructing a PCI Option ROM:

*From UEFI Spec 14.4.21 PCI Option ROMs*

*PCI Firmware Specification 3.0*



A-0521

Figure 5-2: Image and Header Organization

## 2. OpRomImage

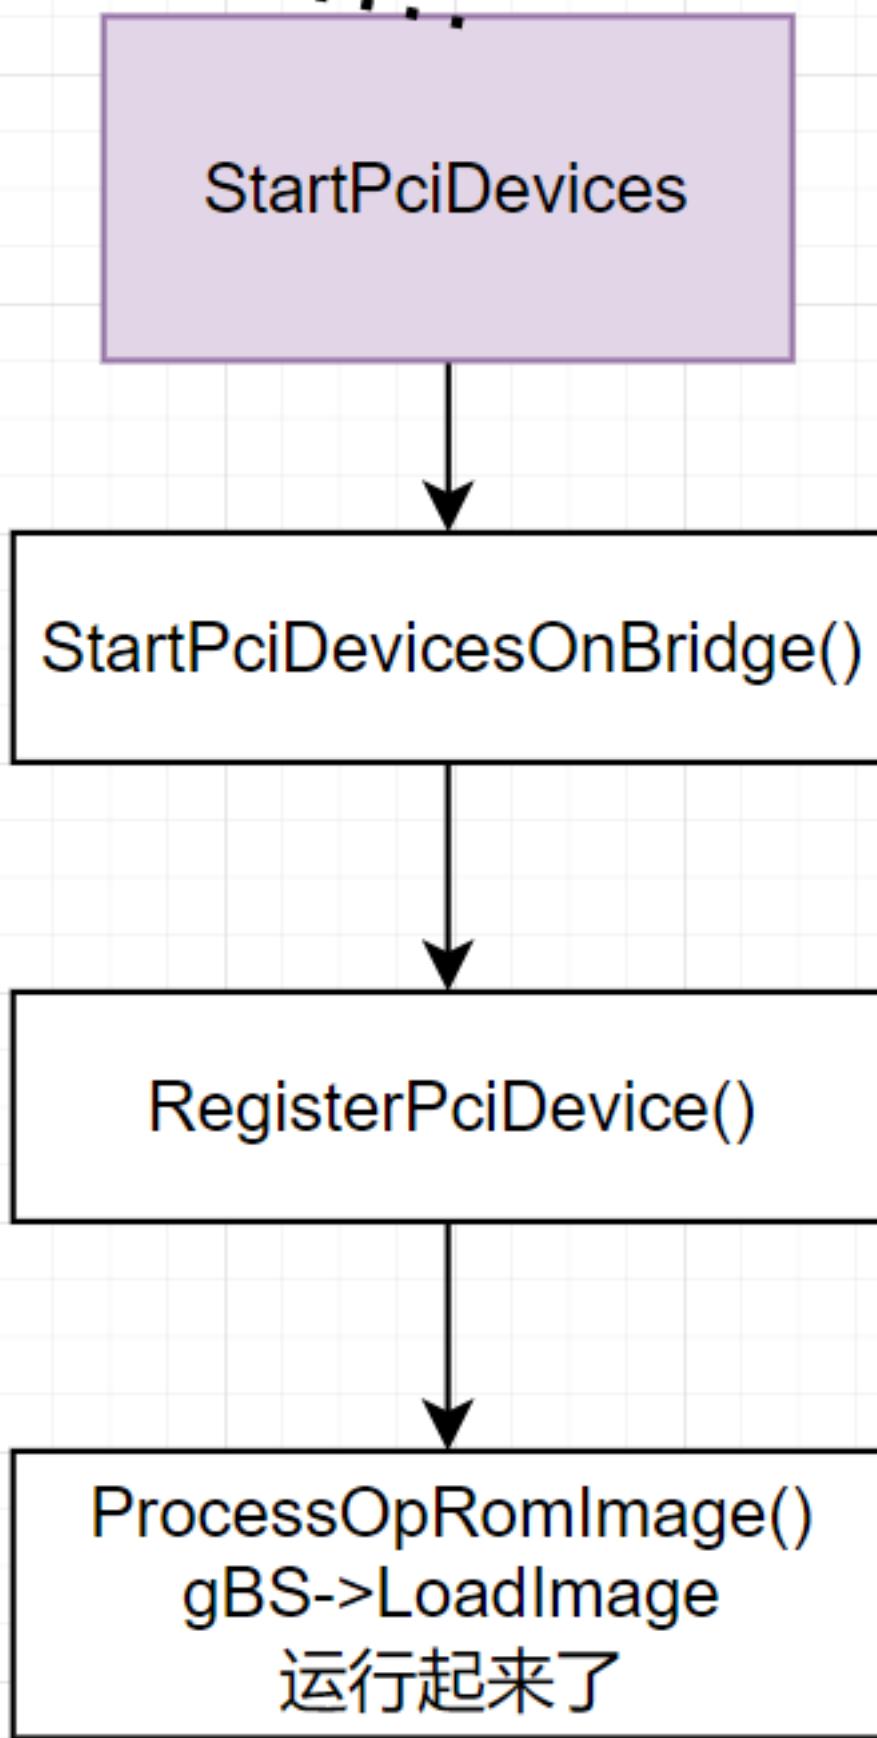
```
OpRom PciSearchDevice() --> GetOpRomInfo() OpRom 1 RomSize PCI_IO_DEVICE
0 1 NOT FOUND.

OpRom root bridge BAR BAR ProcessOptionRom() , root optionRom
LoadOpRomImage pci

LoadOpRomImage Rom PCIR rom root bridge root bridge option rom size root
bridge bar 0 root bridge BAR0 OptionRom Addr option rom shadow
shadow oprom program bar oprom bar oprom root bridge bar0 shadow
oprom

OpRom bar oprom bar
rombar 0x55aa 512 RomSize 3.1
optionRom FirstCheck FALSE Pcir 0 4
PCIr image legacy LegacyImageLength
Legacy Rom size rom 0 rom
rom PciDevice->PciIo optionRom

OpRom Image StartPciDevices image
```



\edk2\MdeModulePkg\Bus\Pci\PciBusDxe\PciOptionRomSupport.c

Resources

PCIe 5.0 Spec Chapter 7.5.1.2.4

<https://blog.csdn.net/robinsongsog/article/details/51785335>

<https://zhuanlan.zhihu.com/p/343464819>

<https://www.cnblogs.com/free-1122/p/16611254.html>

## 5.5 PCIe

### 5.5.1 PCIe AER

### 5.5.2 PCIe Interrupt

### 5.5.3 PCIe Hot-Plug

### 5.5.4 PCIe Power Management

### 5.5.5 *DPC*

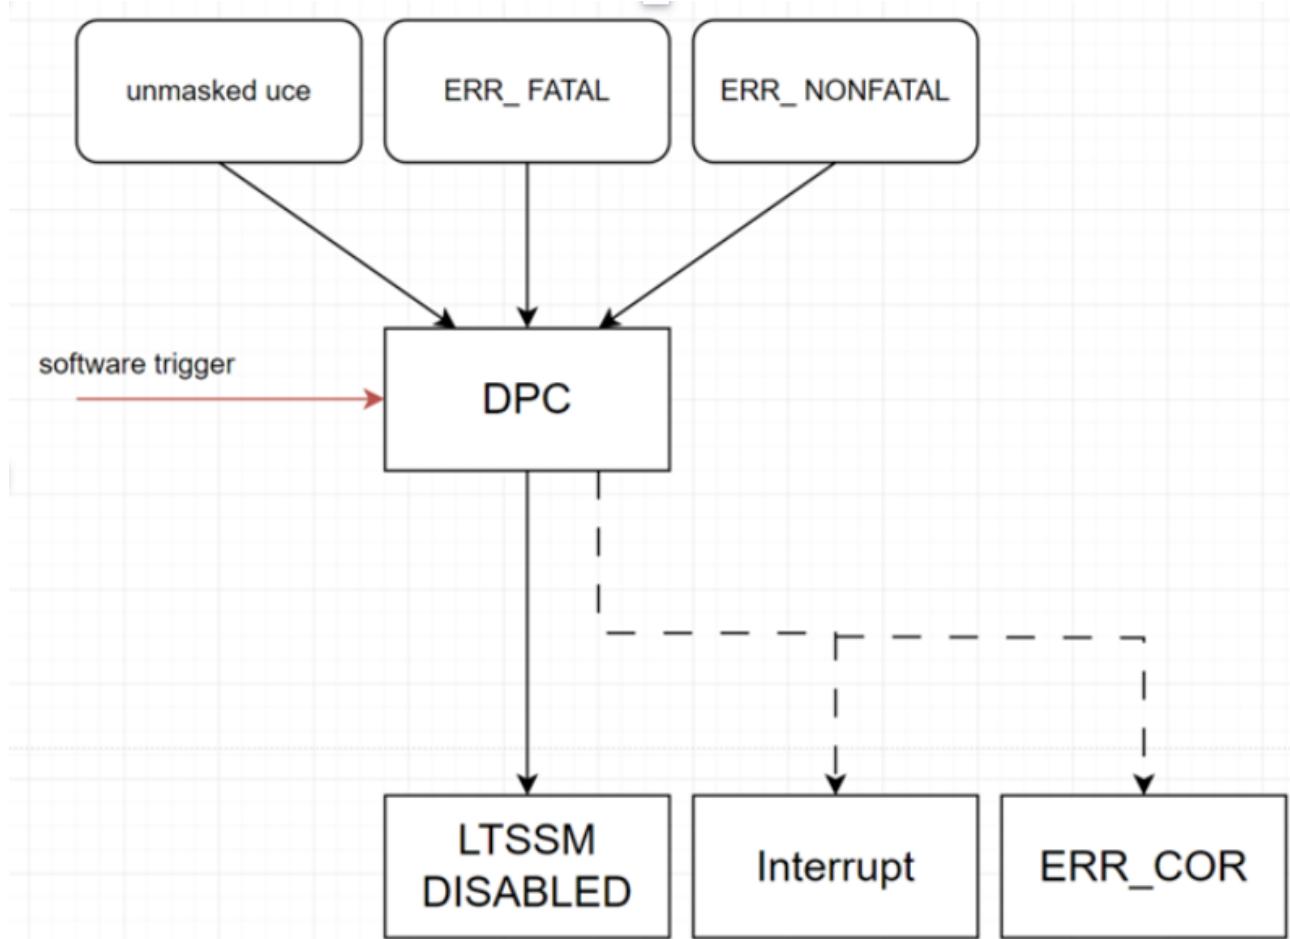
DPC downstream port containment PCIe Spec 3.1 root port pcie switch halt DPC

DPC unmasked uncorrectable error ERR\_FATAL ERR\_NONFATAL DPC

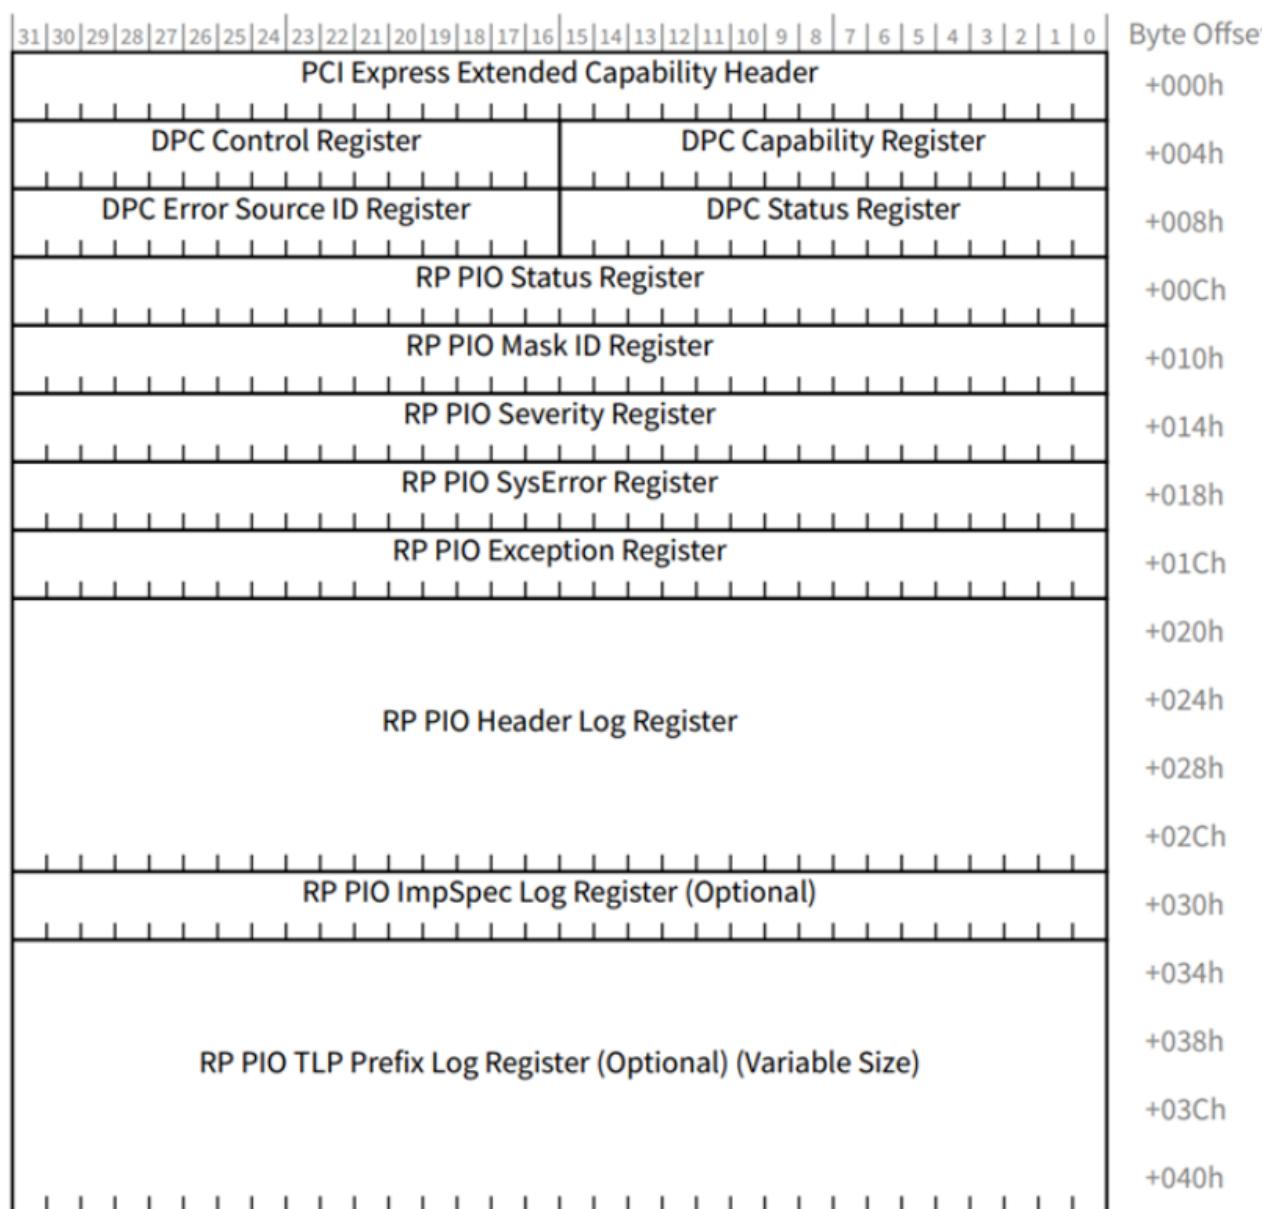
DPC DP DPC Trigger Status bit and DPC Trigger Reason field LTSSM disable link LTSSM disable  
Trigger Status bit

DPC LTSSM detect

DPC DPC ERR\_COR message



DPC Extended Capability structure



*Figure 7-249 DPC Extended Capability*

DPC extended capability    ID 1Dh

RP PIO

## DPC

DPC              DPC control register

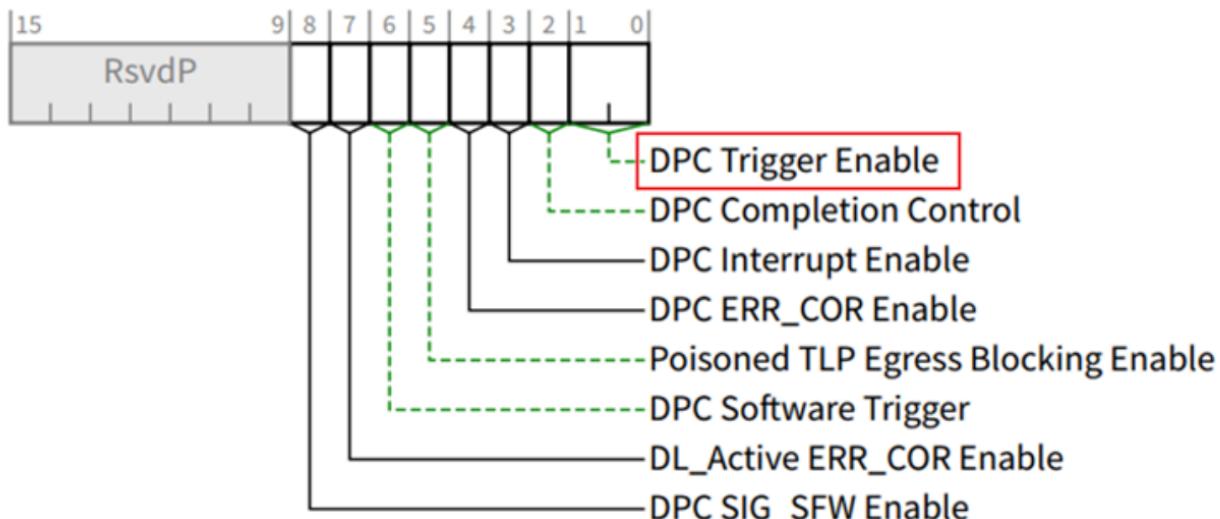


Figure 7-252 DPC Control Register

DPC Trigger Enable      DPC

Bit Location	Register Description	Attributes
1:0	<p><b>DPC Trigger Enable</b> - This field enables DPC and controls the conditions that cause DPC to be triggered.</p> <p>Defined encodings are:</p> <ul style="list-style-type: none"> <li><b>00b</b>      DPC is disabled</li> <li><b>01b</b>      DPC is enabled and is triggered when the Downstream Port detects an unmasked uncorrectable error or when the Downstream Port receives an ERR_FATAL Message</li> </ul>	RW

## 5.0-1.0-PUB — PCI Express® Base Specification Revision 5.0 Version 1.0

Bit Location	Register Description	Attributes
	<ul style="list-style-type: none"> <li><b>10b</b>      DPC is enabled and is triggered when the Downstream Port detects an unmasked uncorrectable error or when the Downstream Port receives an <u>ERR_NONFATAL</u> or <u>ERR_FATAL</u> Message</li> <li><b>11b</b>      Reserved</li> </ul> <p>Default value of this field is 00b.</p>	

DPC Trigger Enable field 01b DPC

ERR\_FATAL Message DPC

DPC Trigger Enable field 10b DPC

ERR\_NONFATAL or ERR\_FATAL Message DPC

## DPC

DPC      DPC

DPC enable

Port DPC

DPC Software Triggering Supported bit in DPC Capability register 0x1

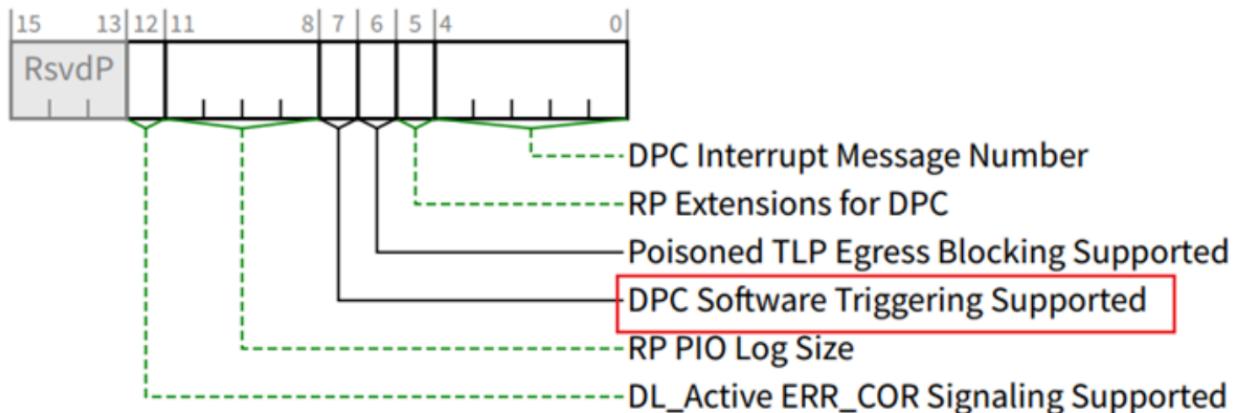


Figure 7-251 DPC Capability Register

DPC

1b DPC Software Trigger bit in the DPC Control Register Port DPC

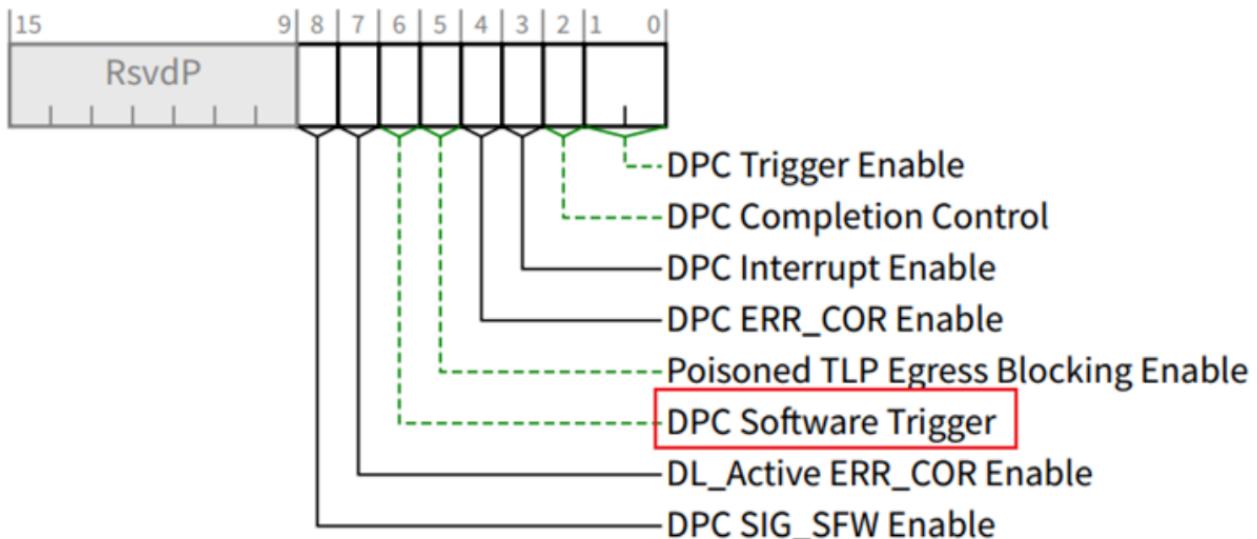
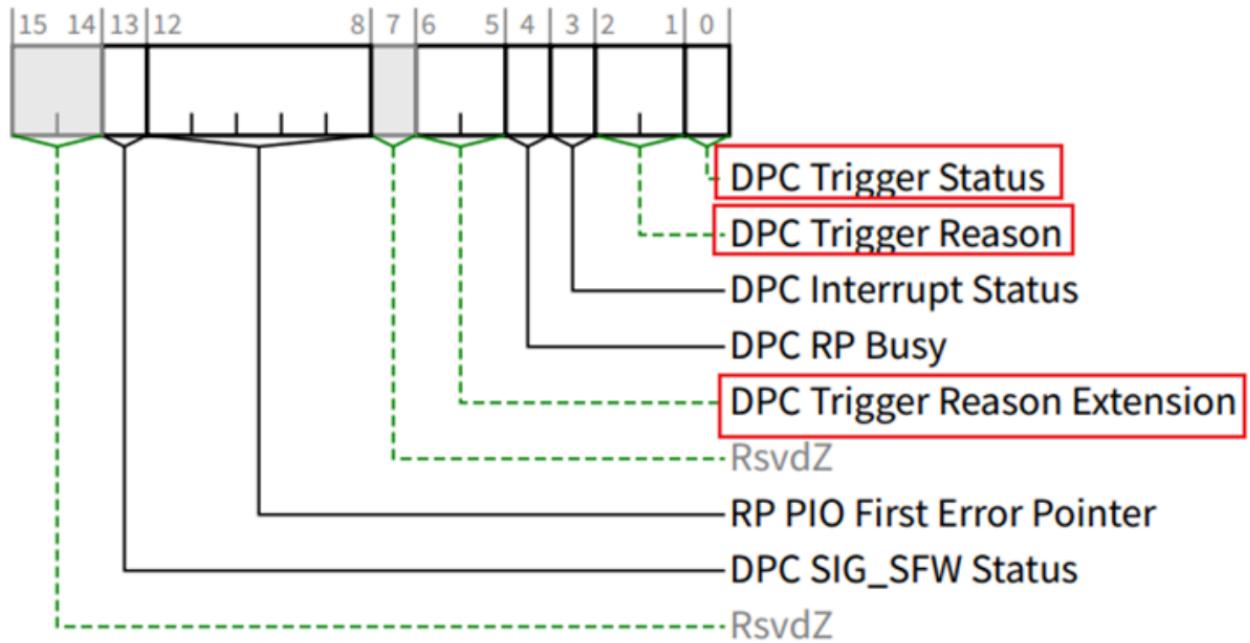


Figure 7-252 DPC Control Register

DPC Status Register



*Figure 7-253 DPC Status Register*

DPC Trigger Status 1      DPC

DPC Trigger Reason 0x11      DPC Trigger Reason Extension

DPC Trigger Reason Extension 0x01      DPC software Trigger

#### **DPC**

DPC DSP      DPC

DPC      DPC Control Register DPC Interrupt Enable bit

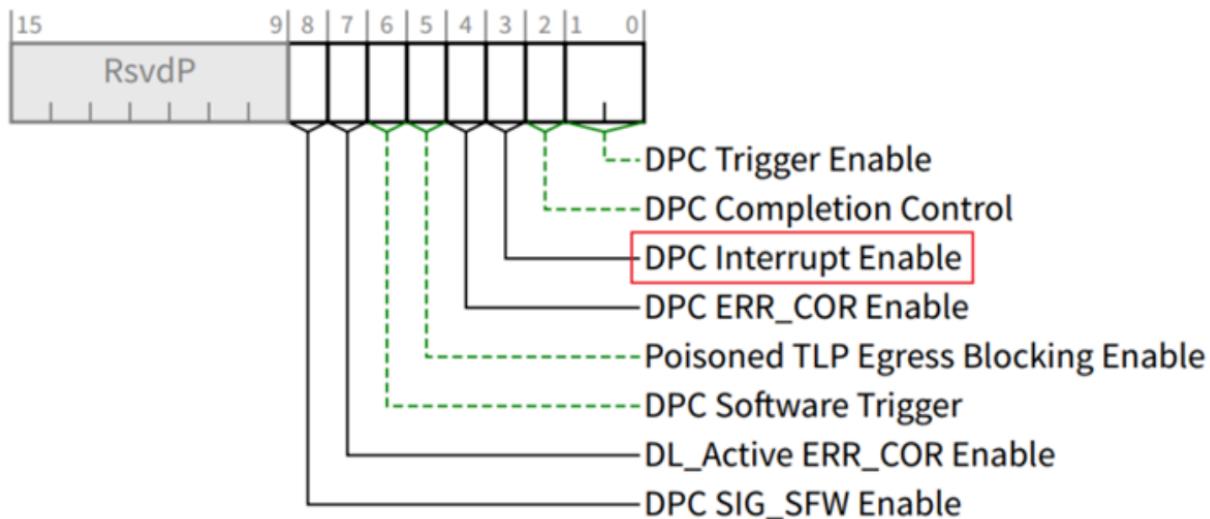
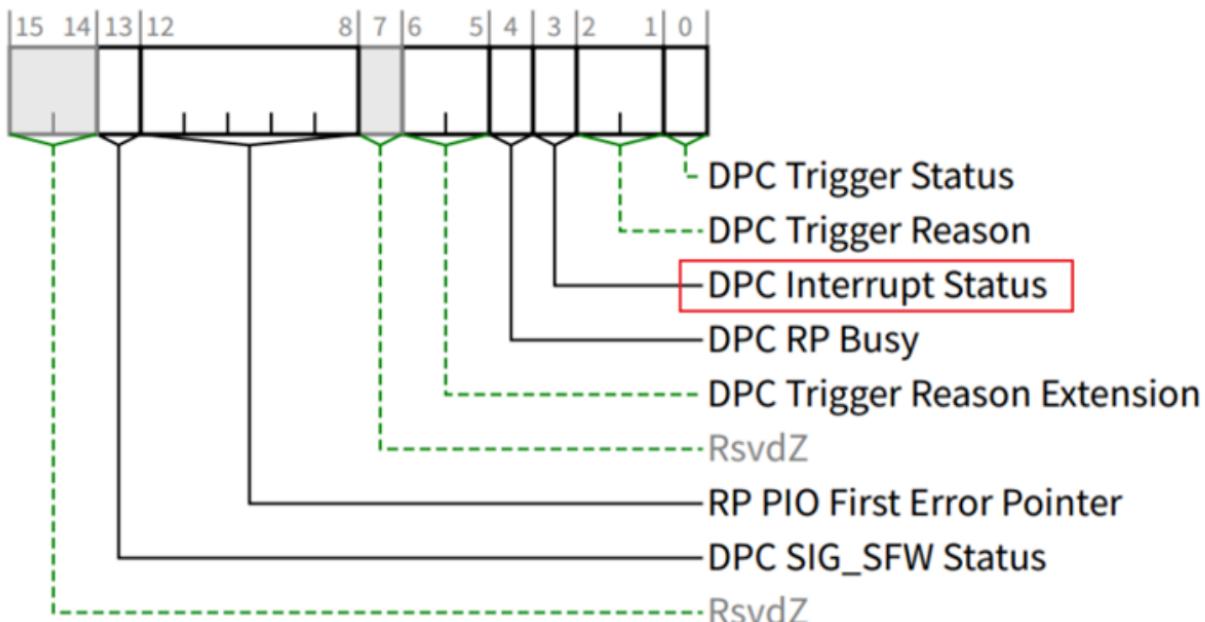


Figure 7-252 DPC Control Register

3	<b>DPC Interrupt Enable</b> - When Set, this bit enables the generation of an interrupt to indicate that DPC has been triggered. See <a href="#">Section 6.2.10.1</a> .	<u>RW</u>
Default value of this bit is 0b.		

DPC      DPC Status Register DPC Interrupt Status bit



3	<b>DPC Interrupt Status</b> - This bit is Set if DPC is triggered while the <a href="#">DPC Interrupt Enable</a> bit is Set. This may cause the generation of an interrupt. See <a href="#">Section 6.2.10.1</a> .	<u>RW1CS</u>
Default value of this bit is 0b.		

#### DPC ERR\_COR signaling

DPC DSP    ERR\_COR       Advanced Error Reporting (AER)

DPC ERR\_COR DPC

DPC ERR\_COR signaling is enabled by the DPC ERR\_COR Enable bit in the DPC Control Register.

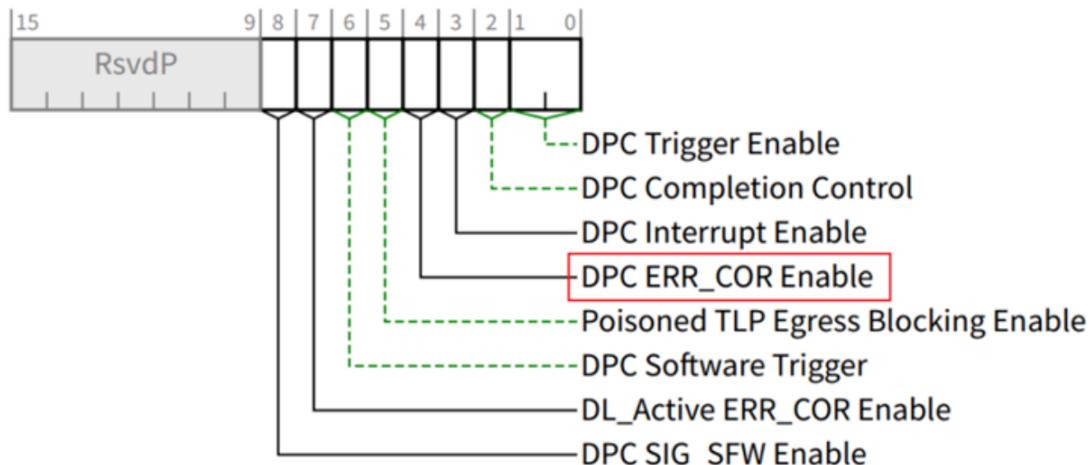


Figure 7-252 DPC Control Register

the Correctable Error Reporting Enable bit in the Device Control Register

or the DPC SIG\_SFW Enable bit in the DPC Control Register is Set

PCIe Spec OS DPC FW ERR\_COR signaling

## DPC

DPC Trigger Status bit in DPC Status Register Root port DPC

LTSSM disable Data Link Layer Link Active bit in the Link Status Register reads 0b DPC

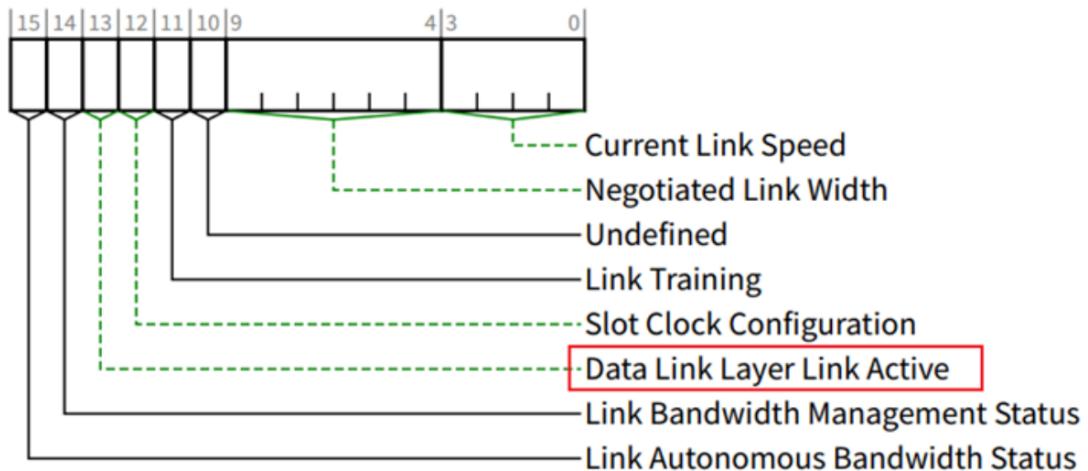
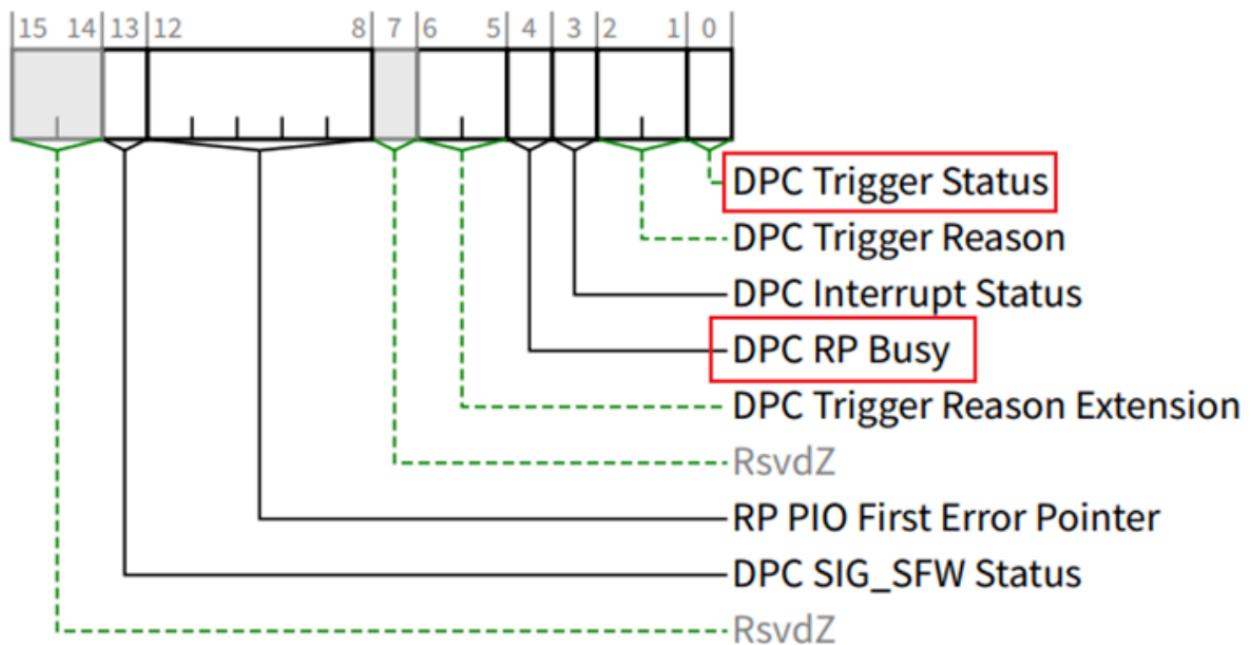


Figure 7-29 Link Status Register

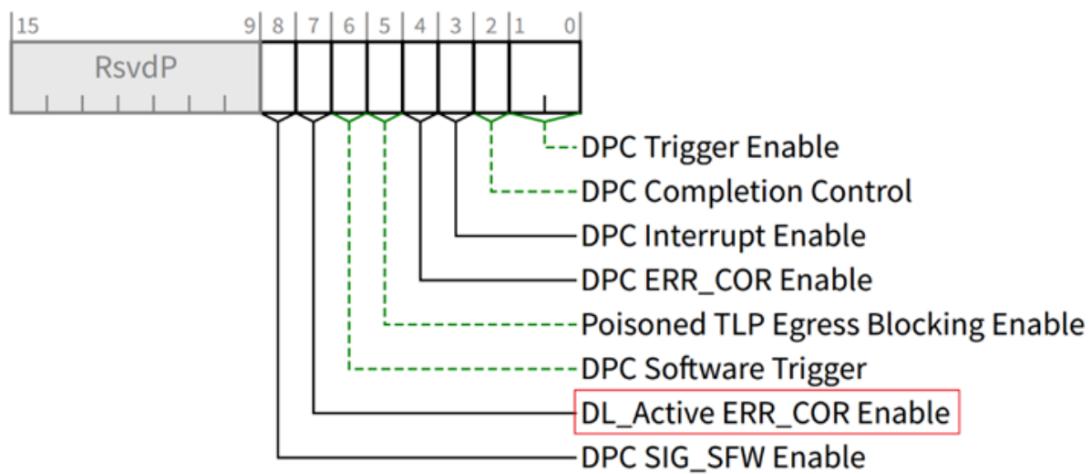
Root Port RP DPC DPC RP Busy bit reads 0b.



*Figure 7-253 DPC Status Register*

#### DL\_Active ERR\_COR signaling

DPC active DL\_Active ERR\_COR  
 DL\_Active the Data Link Layer Link Active bit in the Link Status Register DPC bit  
 ERR\_COR ERR\_COR active  
 DL\_Active ERR\_COR signaling DPC control :



*Figure 7-252 DPC Control Register*

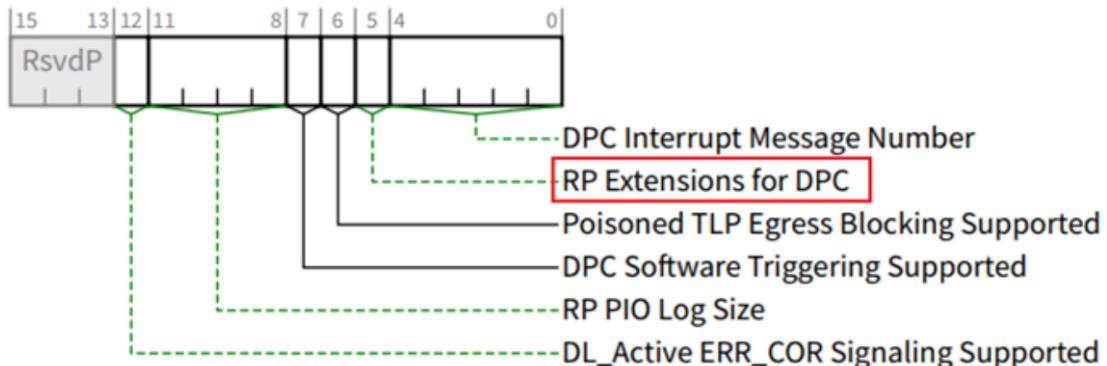
#### DPC ERR\_COR

the Correctable Error Reporting Enable bit in the Device Control Register  
 or the DPC SIG\_SFW Enable bit in the DPC Control Register is Set

DL_ACTIVE	ERR_COR	MSI/MSI-X	MSI/MSI-X	ERR_COR
INTx	DL_ACTIVE	INTx		ERR_COR

## eDPC

eDPC	DPC	PCIe Spec 3.1	DPC	DPC	RP PIO Root Port Programmed I/O
RP PIO	PCIe	RP PIO			
DPC capability	bit5	RP	eDPC		

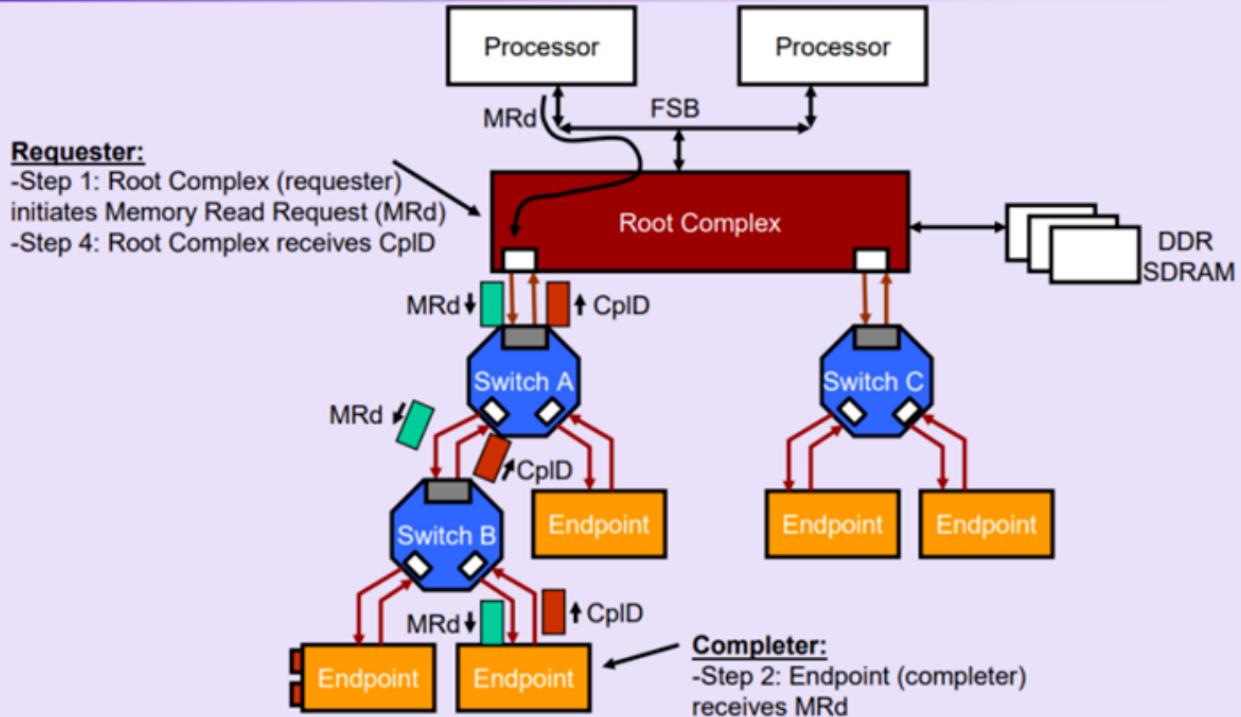


*Figure 7-251 DPC Capability Register*

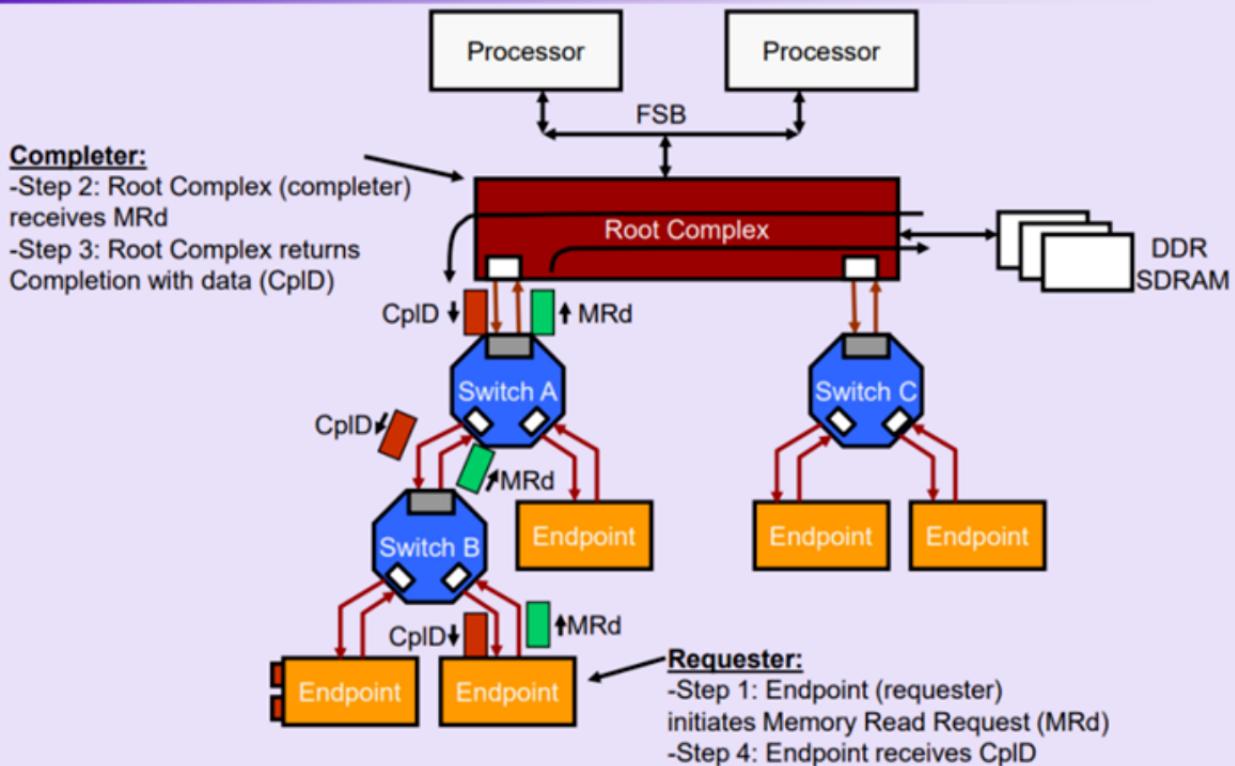
## PCIe

PIO	PCIe	DMA P2P
-----	------	---------

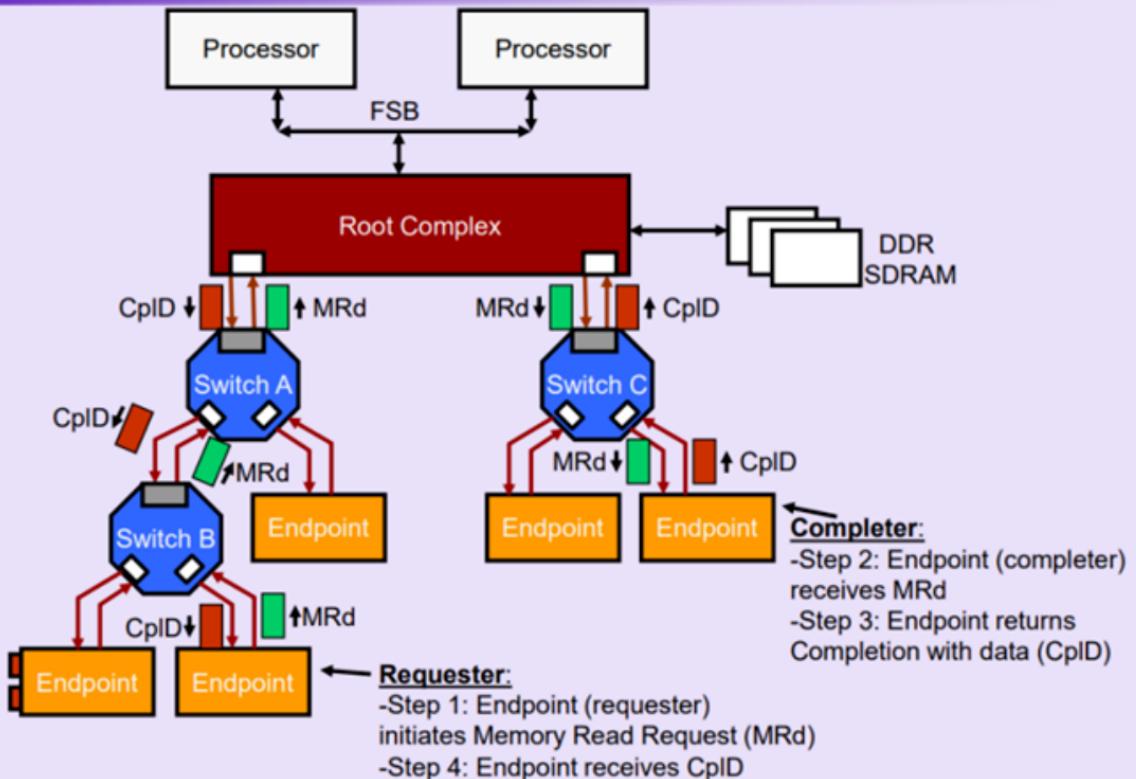
## Programmed I/O Transaction



## DMA Transaction



## Peer-to-Peer Transaction



### RP PIO

Completion with Unsupported Request status (UR Cpl)

Completion with Completer Abort status (CA Cpl)

Completion Timeout (CTO) errors

Configuration Requests I/O Requests Memory Requests RP PIO 9

---

**Cfg UR Cpl** - Configuration Request received UR Completion

---

**Cfg CA Cpl** - Configuration Request received CA Completion

---

**Cfg CTO** - Configuration Request Completion Timeout

---

**I/O UR Cpl** - I/O Request received UR Completion

---

**I/O CA Cpl** - I/O Request received CA Completion

---

**I/O CTO** - I/O Request Completion Timeout

---

**Mem UR Cpl** - Memory Request received UR Completion

---

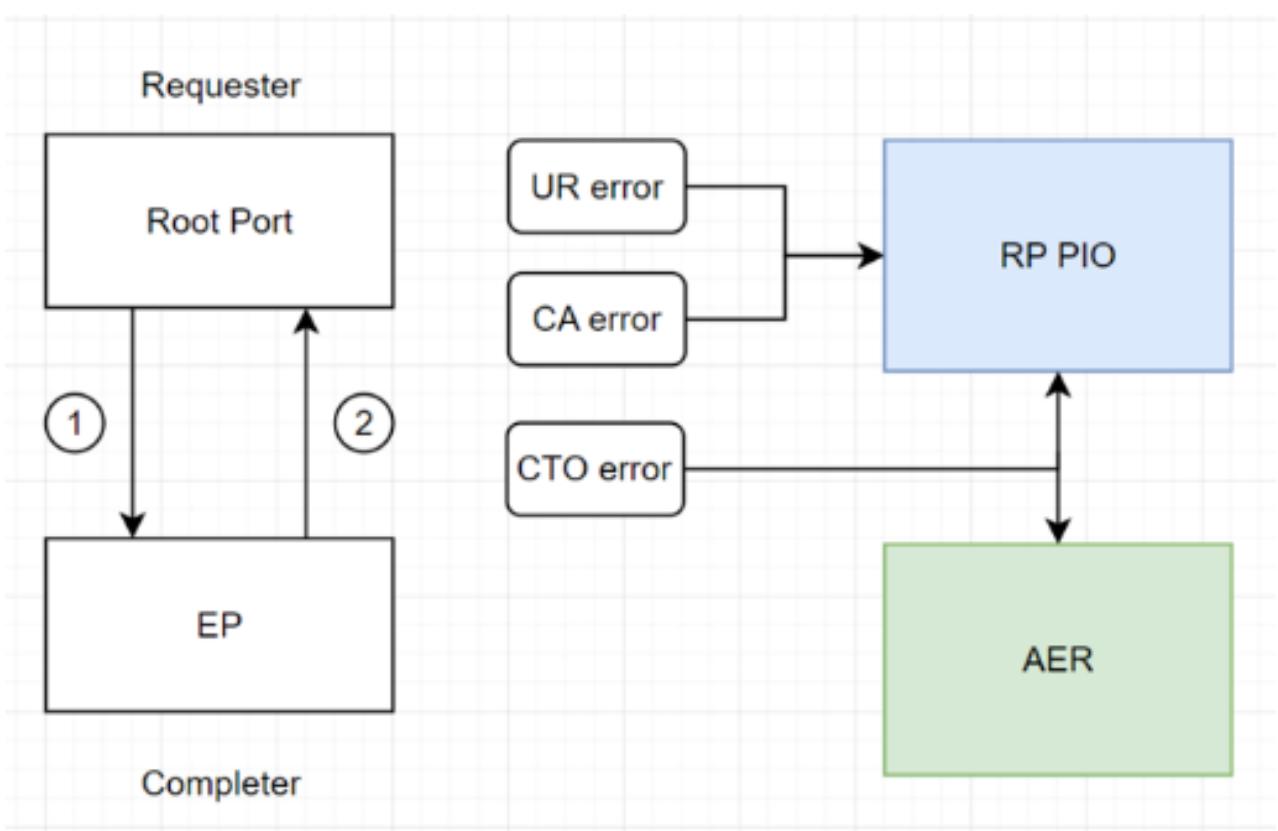
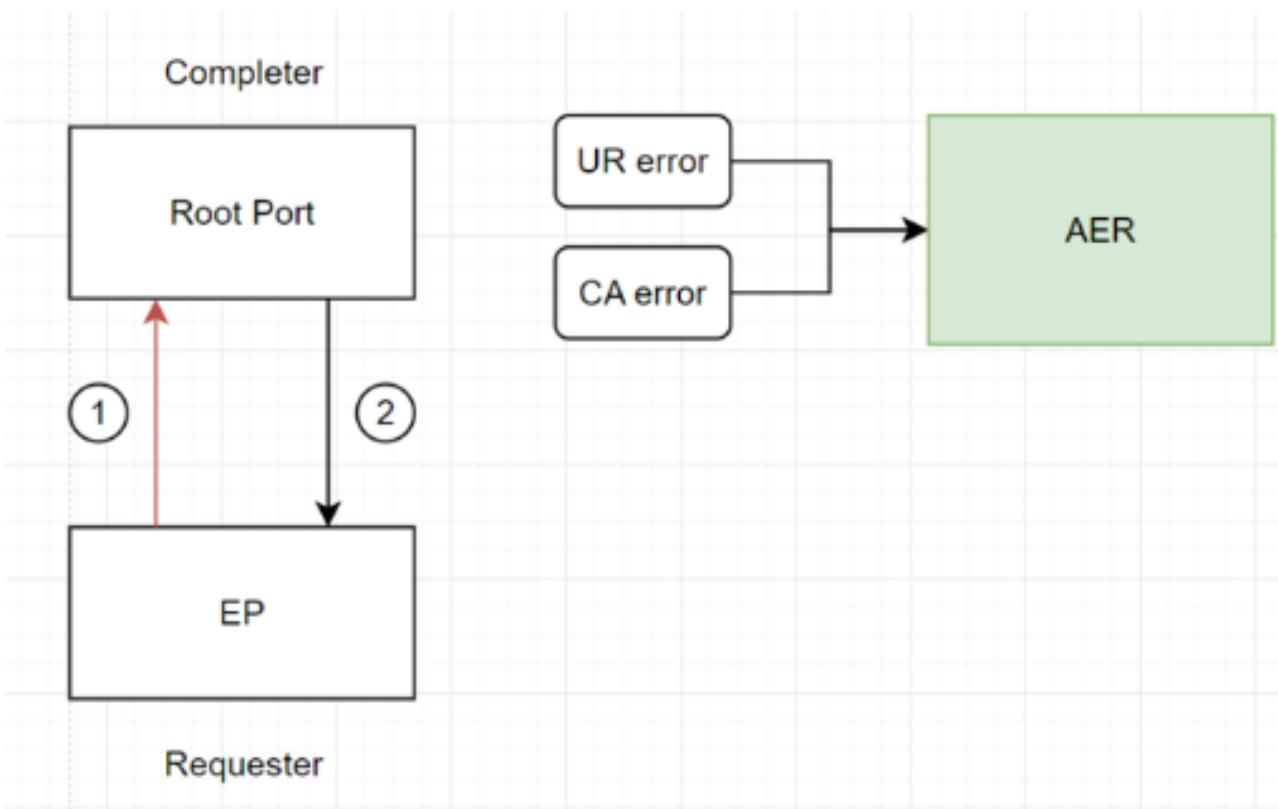
**Mem CA Cpl** - Memory Request received CA Completion

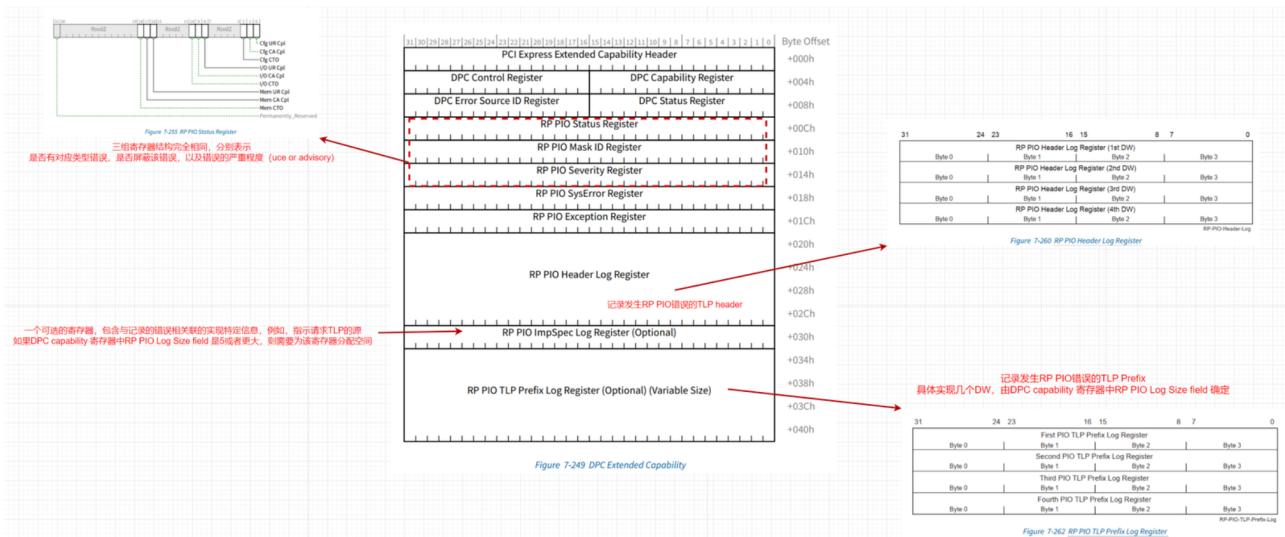
---

**Mem CTO** - Memory Request Completion Timeout

---

AER UR CA CTO                    RP PIO error control  
RP Completer UR CA error     AER  
RP Requester UR CA error     RP PIO CTO error     AER     RP PIO  
                                      RP PIO CTO error     AER     CTO error





RP PIO error control      DPC extended capability

RP PIO      AER capability

The RP PIO Status, Mask, and Severity registers      AER the Uncorrectable Error Status, Mask, and Severity registers

RP PIO      RP PIO Status Register      RP PIO log registers      RP PIO

RP PIO      RP PIO Status Register      DPC      RP PIO log registers

RP PIO      uncorrectable or advisory      RP PIO Severity Register

Severity bit      uncorrectable      DPC( DPC )      DPC / ERR\_COR(if enabled)

Severity bit Clear      ( DPC )      ERR\_COR(if enabled)

The RP PIO Header Log Register, RP PIO ImpSpec Log Register, and RP PIO TLP Prefix Log Registers  
RP PIO log registers

The RP PIO First Error Pointer, RP PIO Header Log, and RP PIO TLP Prefix Log      AER the First Error Pointer, Header Log, and TLP Prefix Log

RP PIO Header Log      RP PIO error TLP header;

RP PIO ImpSpec Log Register      request TLP      DPC capability      RP Log Size field 5

RP PIO TLP Prefix Log Register      RP PIO error TLP Prefix      DW DPC capability      RP Log Size field

Number = RP Log Size - 5 if RP Log Size <=9;

Number = 4 if RP Log Size >9;

## Linux Kernel DPC

### DPC

Linux kernel 5.10.201 DPC RAS      C2000 DPC RAS

```

static struct pcie_port_service_driver dpcdriver = {
    .name          = "dpc",
    .port_type     = PCIE_ANY_PORT,
    .service       = PCIE_PORT_SERVICE_DPC,
    .probe         = dpc_probe,
    .remove        = dpc_remove,
};


```

dpc\_probe DPC

```

349 static int dpc_probe(struct pcie_device *dev)
350 {
351     struct pci_dev *pdev = dev->port;
352     struct device *device = &dev->device;
353     int status;
354     u16 ctl, cap;
355
356     if (!pcie_aer_is_native(pdev) && !pcie_ports_dpc_native)
357         return -ENOTSUPP;
358
359     status = devm_request_threaded_irq(device, dev->irq, dpc_irq,
360                                         dpc_handler, IRQF_SHARED,
361                                         "pcie-dpc", pdev);
362     if (status) {
363         pci_warn(pdev, "request IRQ%d failed: %d\n", dev->irq,
364                  status);
365         return status;
366     }
367
368     pci_read_config_word(pdev, pdev->dpc_cap + PCI_EXP_DPC_CAP, &cap);
369     pci_read_config_word(pdev, pdev->dpc_cap + PCI_EXP_DPC_CTL, &ctl);
370
371     ctl = (ctl & 0xffff4) | PCI_EXP_DPC_CTL_EN_FATAL | PCI_EXP_DPC_CTL_INT_EN;
372     pci_write_config_word(pdev, pdev->dpc_cap + PCI_EXP_DPC_CTL, ctl);
373     pci_info(pdev, "enabled with IRQ %d\n", dev->irq);
374
375     pci_info(pdev, "error containment capabilities: Int Msg #d, RPExt%c PoisonedTLP%c SwTrigger%c RP PIO Log %d, DL_ActiveErr%c\n",
376             cap & PCI_EXP_DPC_IRQ, FLAG(cap, PCI_EXP_DPC_CAP_RP_EXT),
377             FLAG(cap, PCI_EXP_DPC_CAP_POISONED_TLP),
378             FLAG(cap, PCI_EXP_DPC_CAP_SW_TRIGGER), pdev->dpc_rp_log_size,
379             FLAG(cap, PCI_EXP_DPC_CAP_DL_ACTIVE));
380
381     pci_add_ext_cap_save_buffer(pdev, PCI_EXT_CAP_ID_DPC, sizeof(u16));
382     return status;
383 }


```

line356 PCIe AER DPC -ENOTSUPP

line359 devm\_request\_threaded\_irq dev->irq status

devm\_request\_threaded\_irq Linux IRQ dpc\_irq dpc\_handler top  
half bottom half

line362

line368 DPC capability control

line371 control enable Fatal error control

dpc\_irq

```

static irqreturn_t dpc_irq(int irq, void *context)
{
    struct pci_dev *pdev = context;
    u16 cap = pdev->dpc_cap, status;

    pci_read_config_word(pdev, cap + PCI_EXP_DPC_STATUS, &status);

    if (!(status & PCI_EXP_DPC_STATUS_INTERRUPT) || status == (u16)(~0))
        return IRQ_NONE;

    pci_write_config_word(pdev, cap + PCI_EXP_DPC_STATUS,
                         PCI_EXP_DPC_STATUS_INTERRUPT);
    if (status & PCI_EXP_DPC_STATUS_TRIGGER)
        return IRQ_WAKE_THREAD;
    return IRQ_HANDLED;
}

```

DPC status                status     1  
 status 0x8 dpc     1b  
 DPC                return IRQ\_WAKE\_THREAD;  
 IRQ\_WAKE\_THREAD  
 IRQ\_HANDLED

#### dpc\_handler

```

298 static irqreturn_t dpc_handler(int irq, void *context)
299 {
300     struct pci_dev *pdev = context;
301
302     dpc_process_error(pdev);
303
304     /* We configure DPC so it only triggers on ERR_FATAL */
305     pcie_do_recovery(pdev, pci_channel_io_frozen, dpc_reset_link);
306
307     return IRQ_HANDLED;
308 }

```

dpc\_process\_error     pdev                DPC  
 pcie\_do\_recovery     pdev pci\_channel\_io\_frozen dpc\_reset\_link     DPC  
 IRQ\_HANDLED

```

265 void dpc_process_error(struct pci_dev *pdev)
266 {
267     u16 cap = pdev->dpc_cap, status, source, reason, ext_reason;
268     struct aer_err_info info;
269
270     pci_read_config_word(pdev, cap + PCI_EXP_DPC_STATUS, &status);
271     pci_read_config_word(pdev, cap + PCI_EXP_DPC_SOURCE_ID, &source);
272
273     pci_info(pdev, "containment event, status:%#06x source:%#06x\n",
274             status, source);
275
276     reason = (status & PCI_EXP_DPC_STATUS_TRIGGER_RSN) >> 1;
277     ext_reason = (status & PCI_EXP_DPC_STATUS_TRIGGER_RSN_EXT) >> 5;
278     pci_warn(pdev, "%s detected\n",
279             (reason == 0) ? "unmasked uncorrectable error" :
280             (reason == 1) ? "ERR_NONFATAL" :
281             (reason == 2) ? "ERR_FATAL" :
282             (ext_reason == 0) ? "RP PIO error" :
283             (ext_reason == 1) ? "software trigger" :
284             "reserved error");
285
286     /* show RP PIO error detail information */
287     if (pdev->dpc_rp_extensions && reason == 3 && ext_reason == 0)
288         dpc_process_rp_pio_error(pdev);
289     else if (reason == 0 &&
290             dpc_get_aer_uncorrect_severity(pdev, &info) &&
291             aer_get_device_error_info(pdev, &info)) {
292         aer_print_error(pdev, &info);
293         pci_aer_clear_nonfatal_status(pdev);
294         pci_aer_clear_fatal_status(pdev);
295     }

```

status source ID

raw data

line287	DPC RP	reason 3 ext_reason 0	DPC RP PIO	dpc_process_rp_pio_error	
line289	reason 0	unmasked uncorrectable error	AER Advanced Error Reporting	uce	AER

dpc\_process\_rp\_pio\_error

pcie\_do\_recovery(pdev, pci\_channel\_io\_frozen, dpc\_reset\_link);

err.c PCIe

I/O

/ I/O /

“ ”  
**dpc\_reset\_link** dpc.c DPC  
reset\_subordinates  
PCIe pcie\_do\_recovery  
Root Port Downstream Port RCEC  
Endpoint Port  
reset\_subordinates  
dpc\_reset\_link

```

146     pci_ers_result_t dpc_reset_link(struct pci_dev *pdev)
147     {
148         pci_ers_result_t ret;
149         u16 cap;
150
151         set_bit(PCI_DPC_RECOVERING, &pdev->priv_flags);
152
153         /*
154          * DPC disables the Link automatically in hardware, so it has
155          * already been reset by the time we get here.
156          */
157         cap = pdev->dpc_cap;
158
159         /*
160          * Wait until the Link is inactive, then clear DPC Trigger Status
161          * to allow the Port to leave DPC.
162          */
163         if (!pcie_wait_for_link(pdev, false))
164             pci_info(pdev, "Data Link Layer Link Active not cleared in 1000 msec\n");
165
166         if (pdev->dpc_rp_extensions && dpc_wait_rp_inactive(pdev)) {
167             clear_bit(PCI_DPC_RECOVERED, &pdev->priv_flags);
168             ret = PCI_ERS_RESULT_DISCONNECT;
169             goto out;
170         }
171
172         pci_write_config_word(pdev, cap + PCI_EXP_DPC_STATUS,
173                               PCI_EXP_DPC_STATUS_TRIGGER);
174
175         if (pci_bridge_wait_for_secondary_bus(pdev, "DPC",
176                                              PCIE_RESET_READY_POLL_MS)) {
177             clear_bit(PCI_DPC_RECOVERED, &pdev->priv_flags);
178             ret = PCI_ERS_RESULT_DISCONNECT;
179         } else {
180             set_bit(PCI_DPC_RECOVERED, &pdev->priv_flags);
181             ret = PCI_ERS_RESULT_RECOVERED;
182         }
183     out:
184         clear_bit(PCI_DPC_RECOVERING, &pdev->priv_flags);
185         wake_up_all(&dpc_completed_waitqueue);
186         return ret;
187     }

```

line151 DPC

line163	DPC	pcie_wait_for_link	inactive	100	true	false
---------	-----	--------------------	----------	-----	------	-------

line172	status	bit0 1b	DPC	PCIe 5.0 Spec	spec	disabled	DPC	100ms	warnin
---------	--------	---------	-----	---------------	------	----------	-----	-------	--------

Disabled state or at least to bring the Link down under a variety of error conditions, software must leave the Downstream Port in DPC until the Data Link Layer Link Active bit in the Link Status Register reads 0b; otherwise, the result is undefined. See [Section 7.5.3.8](#). See [Section 2.9.3](#) for other important details on Transaction Layer behavior during DPC.

line175            PCI\_DPC\_RECOVERED        PCI\_ERS\_RESULT\_DISCONNECT

line179 PCI\_DPC\_RECOVERED        PCI\_ERS\_RESULT\_RECOVERED

line184    DPC

# Chapter 6

## CXL

### 6.1 CXL

[toc]

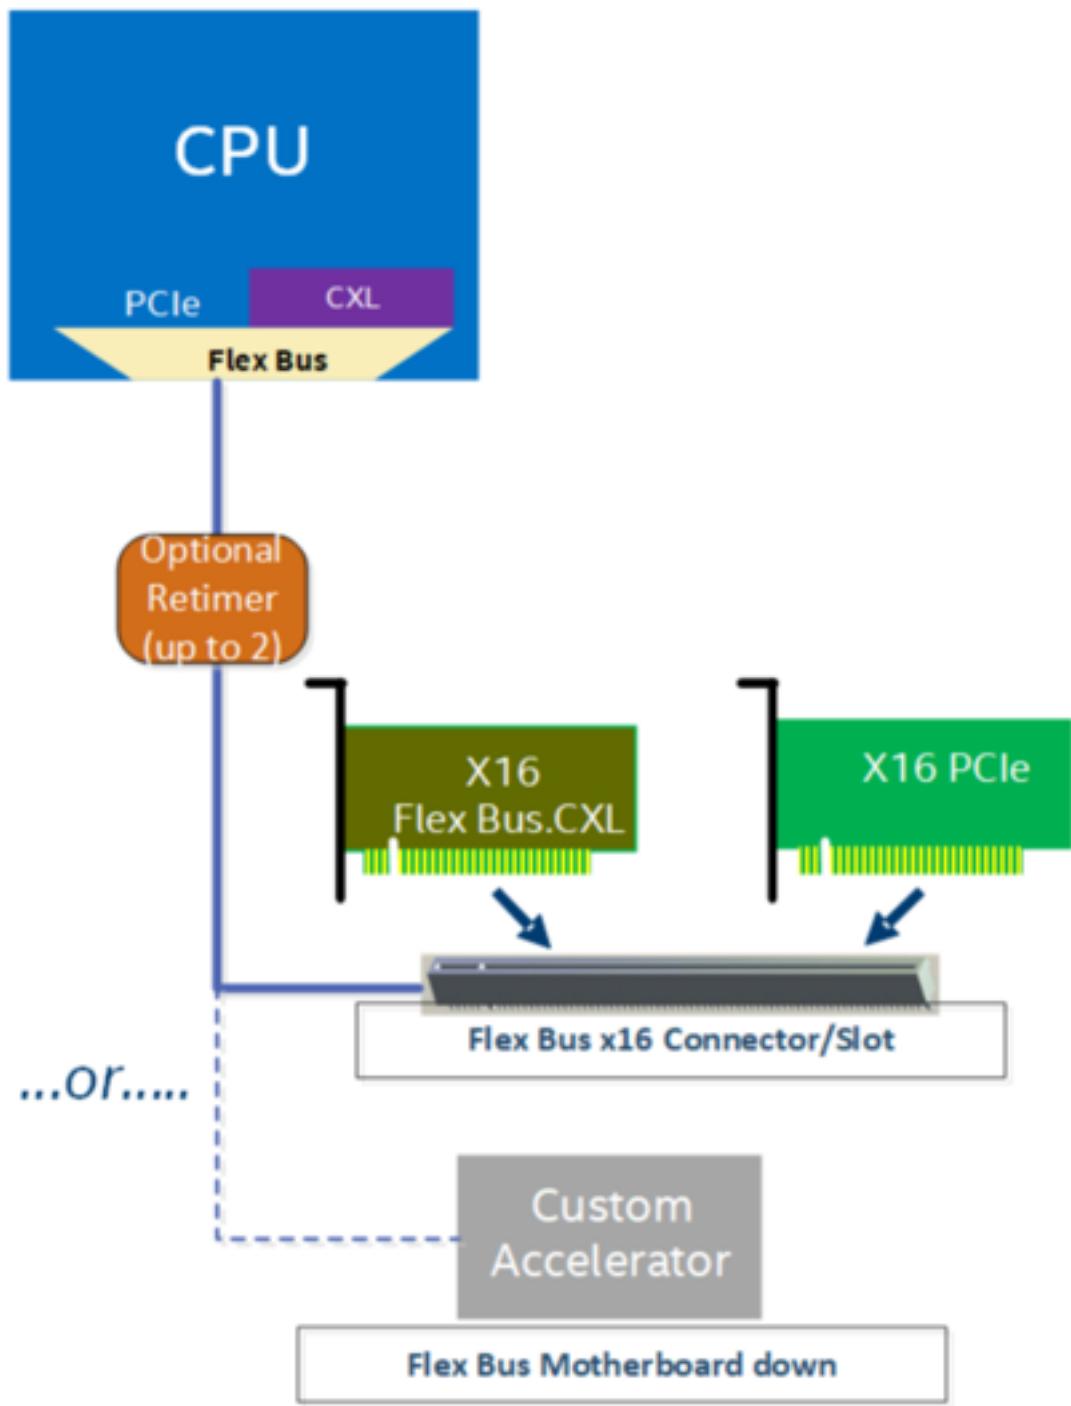
#### Overview

CXL PCI Express 5.0

CXL            I/O(CXL.io PCIe) (CXL.cache) (CXL.memory)

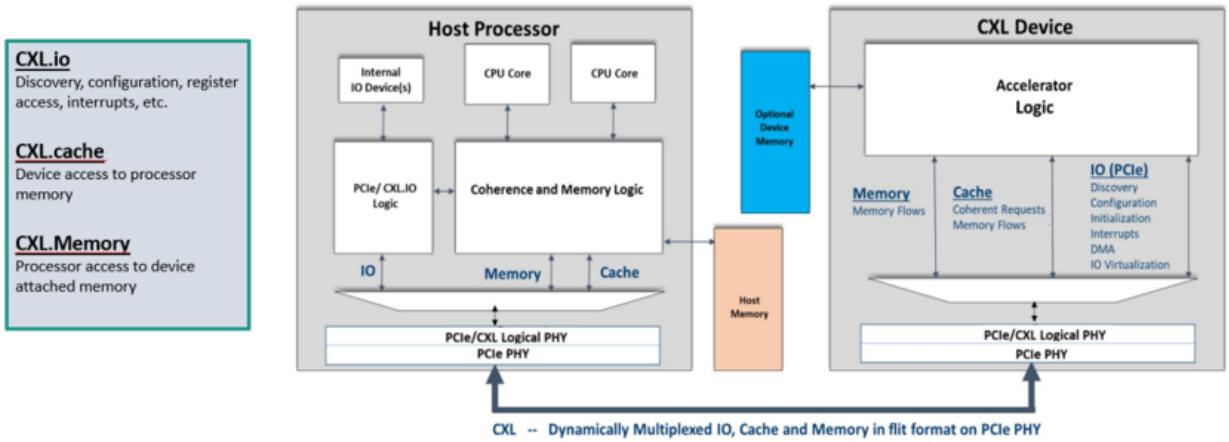
CXL            I/O(CXL.io PCIe) (CXL.cache) (CXL.memory)

CXL PCIe      CXL PCIe    CXL PCIe Gen1     Link Training    PCIe 5.0     Alternate Negotiation  
Protocol CXL        CXL PCIe 5.0                    CXL        PHY



## CXL Device Type

CXL Device	CXL.io	CXL.cache	CXL.mem		CXL.io	CXL.cache
CXL.io	CXL.mem					
CXL.io	PCIe	I/O	(DMA)	-	PCIe	CXL.cache CXL.mem
CXL.cache				Snoop Messages		
CXL.mem		CXL Memroy	Device	CXL.mem	Load Store	

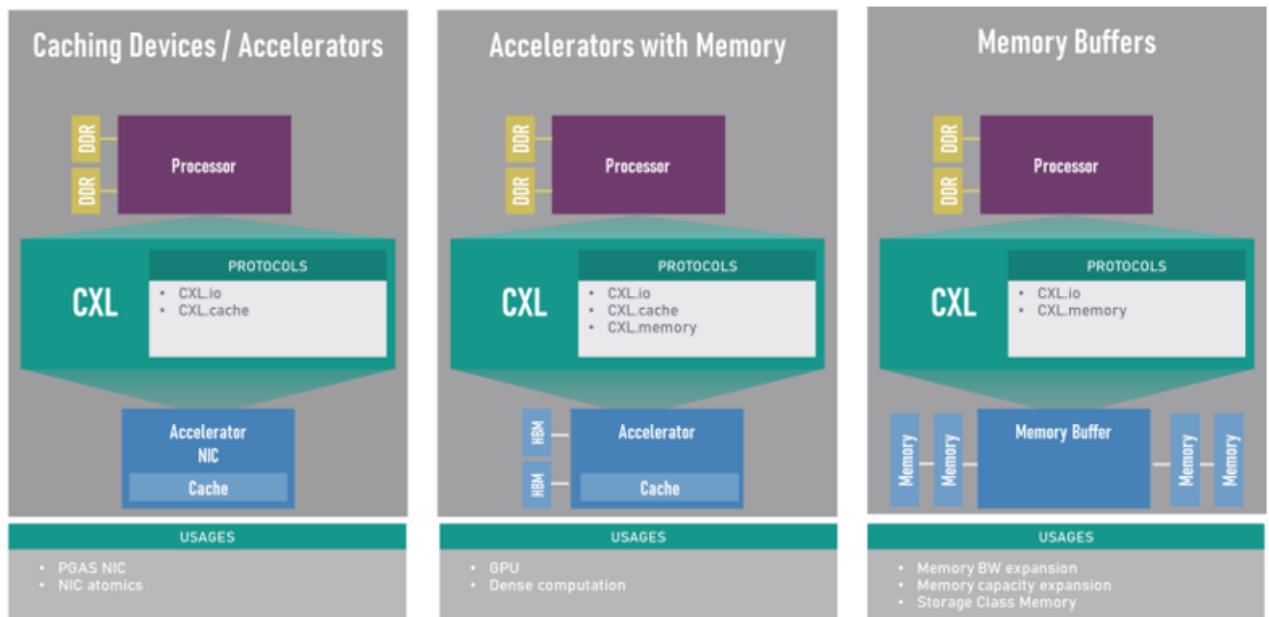


CXL Spec    CXL Device    CXL Device

Type 1 CXL Device:    CXL.io    CXL.cache    CXL Device

Type 2 CXL Device:    CXL Device    HBM    GPU

Type 3 CXL Device:    CXL.io    CXL.mem    CXL Device    CXL Memory Expander

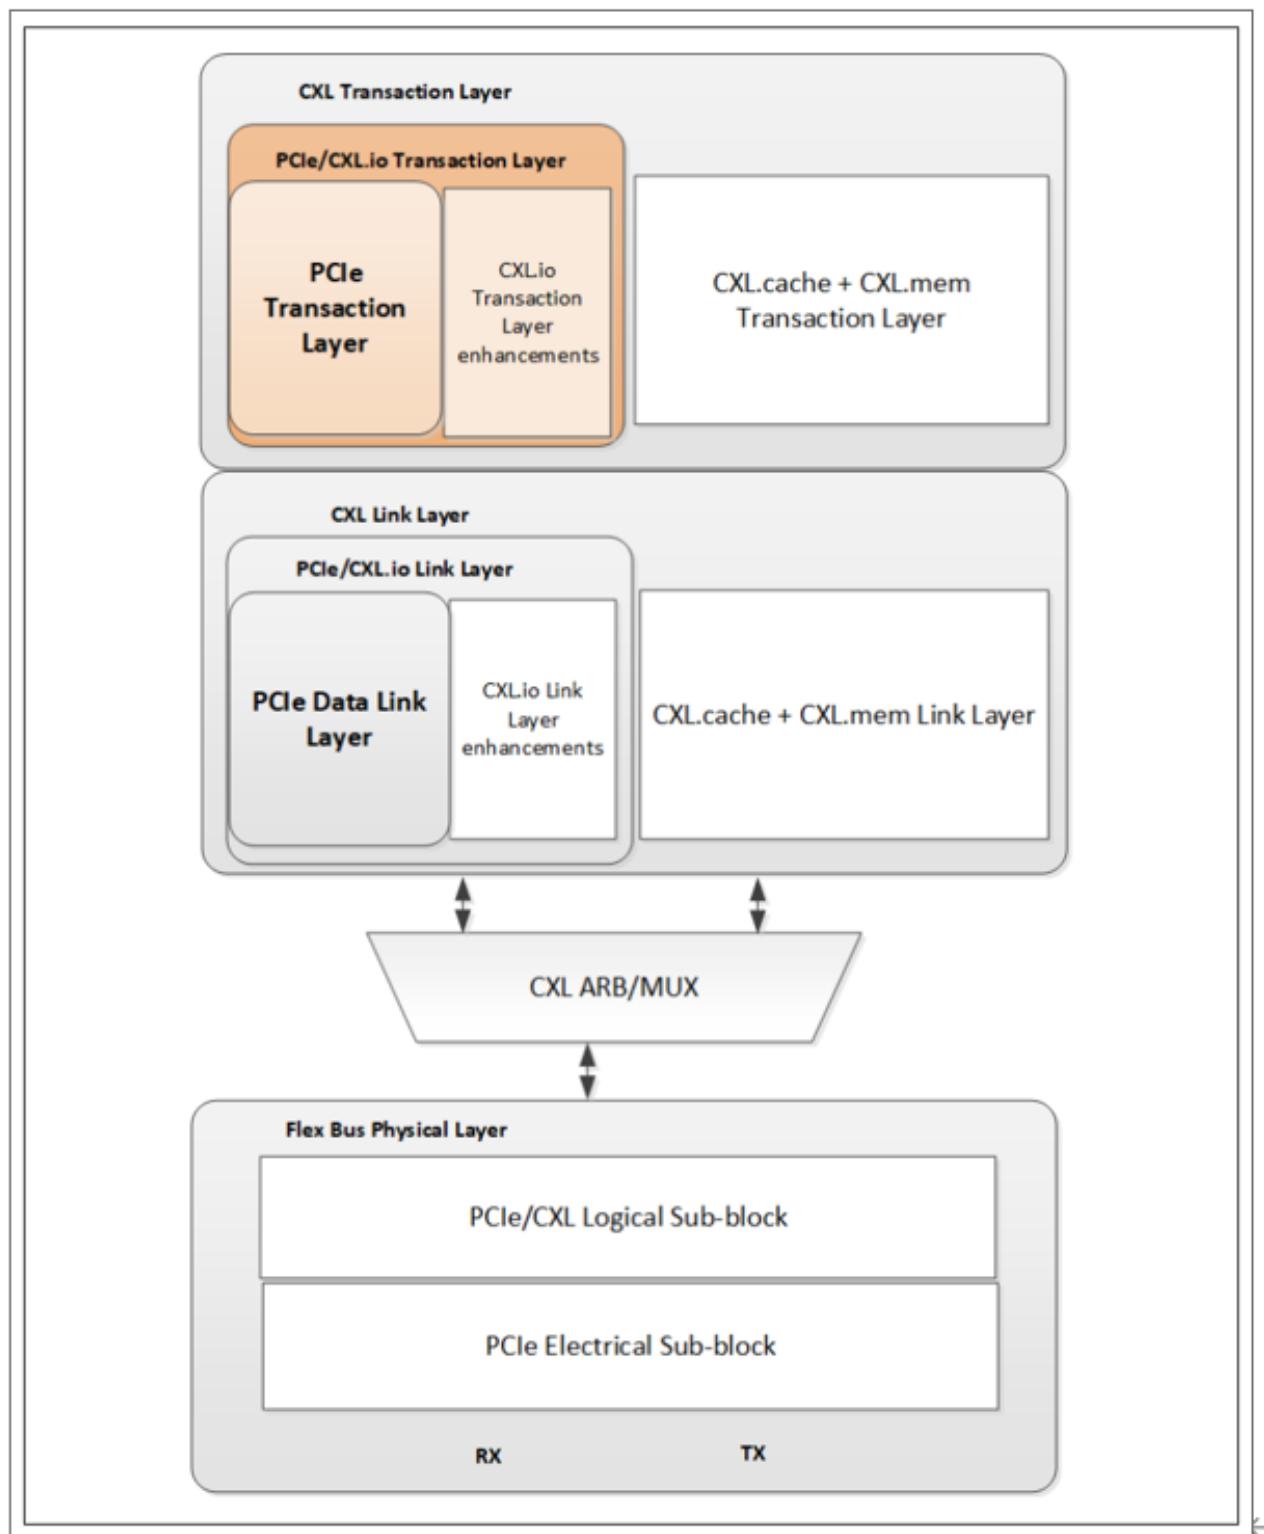


## CXL.io Overview

CXL.io    I/O    non-coherent load/store  
PCI Express 5.0 Base Specification

CXL.io    Flex Bus

PCIe



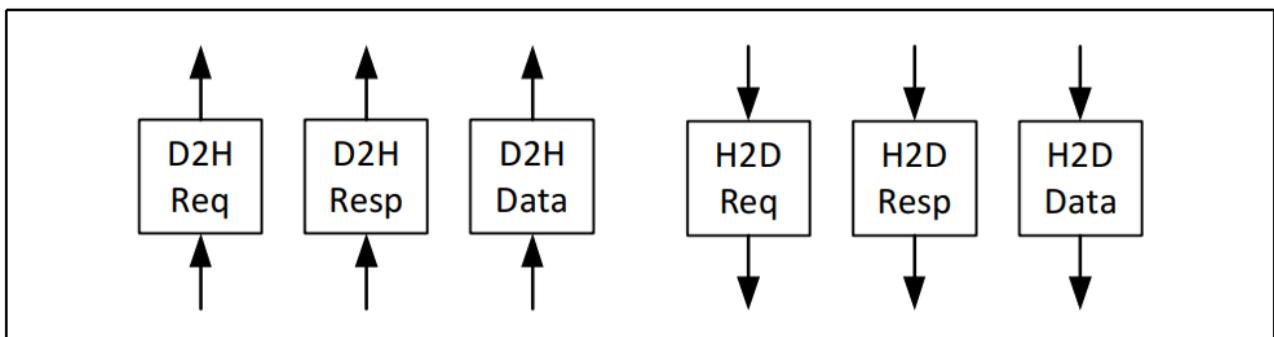
## CXL.cache Overview

CXL.cache  
H2D

CXL.cache Channels D2H

CXL.cache  
H2D

CXL.cache Channels D2H



CXL.mem Overview

CXL Memory Protocol	CXL.mem	CPU Device Memory	Die	Compute Express Link(CXL)	Phy Link
Layer	Memory Controller	Host CPU	Memory Controller	Accelerator	Memory Controller
Persistent	Flat	Hierarchical			

Subordinate Device Accelerator CXL.mem Protocol      Device Coherency Engine(DCOH) DCOH  
CXL.mem

CXL.mem Master Subordinate     "M2S" Subordinate Master     "S2M"

M2S

- -- Req
  - -- RwD
  - S2M
  - -- (NDR)
  - -- DRS

CXL

Features	CXL 1.0 / 1.1	CXL 2.0	CXL 3.0
Release date	2019	2020	2022
Max link rate	32GTs	32GTs	64GTs
Flit 64 byte (up to 32 GTs)	✓	✓	✓
Flit 256 byte (up to 64 GTs)			✓
Type 1, Type 2 and Type 3 Devices	✓	✓	✓
Memory Pooling w/ MLDs		✓	✓
Global Persistent Flush		✓	✓
CXL IDE		✓	✓
Switching (Single-level)		✓	✓
Switching (Multi-level)			✓
Direct memory access for peer-to-peer			✓
Enhanced coherency (256 byte flit)			✓
Memory sharing (256 byte flit)			✓
Multiple Type 1/Type 2 devices per root port			✓
Fabrics (256 byte flit)			✓
			Not supported
			✓ Supported

## Firmware Device

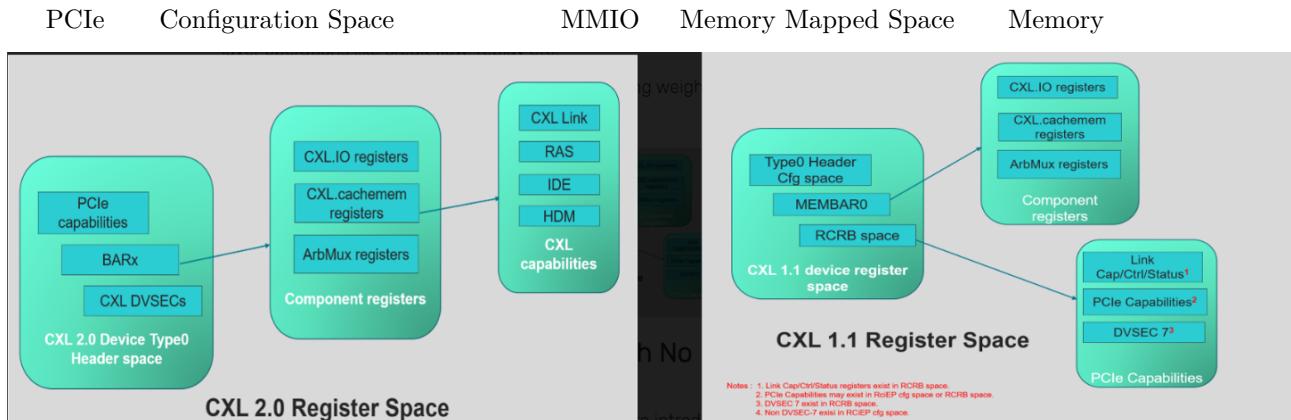
- CXL2.0 Device Component Register HDM Decoder Component Register CXL HDM Decoder Global Control Register (Offset 04h) ->HDM Decoder Enable DVSEC ID 0 Range Size Range Base Memory Address
- CXL 1.1 Device DVSEC ID 0 Range Size Range Base
- CXL Root Port CXL Switch Upstream Port HDM Decoder
- PCIe DVSEC for CXL Device CXL 2.0 Extensions DVSEC for Ports Mem\_Enable

## CXL HDM Decoder

- HDM Decoder Host HDM Decoder
- HDM Decoder
- HDM Decoder DRAM NVRAM

## 6.2 CXL

### CXL



## DVSEC

PCIe DVSEC CXL CXL CXL vendor ID 0x1E98.

CXL 3.0 9 CXL DVSEC

- PCIe DVSEC for CXL Devices CXL PCIe DVSEC RCD LD SLD FMLD CXL Device D0F0 DVSEC Device DVSEC Device PCIe Device CXL Device
- Non-CXL Function Map DVSEC Device Function CXL.cachemem Device Function
- CXL Extensions DVSEC for Ports RP DSP USP PCIe RCH-RCD
- GPF DVSEC for CXL Ports CXL Port GPF GPF Phase1 Phase2 Timeout
- GPF DVSEC for CXL Devices CXL Device GPF GPF Phase2 Timeout Phase2
- PCIe DVSEC for Flex Bus Port Flex Bus Port DVSEC Flex Bus Multi-Device CXL 1.1 CXL Port/Device CXL Modified TS Flex Bus Modified TS Info DVSEC RCRB RCH/RCD DVSEC RCRB
- Register Locator DVSEC CXL Entry

- MLD DVSEC      FM      LD      MLD      Capability      LD      LD-ID
- PCIe DVSEC for Test Capability      CXL Compliance      CXL Compliance      DVSEC      DVSEC

## 9 CXL DVSEC      DVSEC ID

CXL Capability (Name, DVSEC ID)	Devices	PCIe	RCH DP	CXL RP	RCD	RCD UP	LD	SLD	FMLD	USP <sup>1</sup>	DSP <sup>2</sup>
PCIe DVSEC for CXL Devices	0x0000	F	F	F	M	F	M	M	M	F	F
Non-CXL Function Map DVSEC	0x0002	F	F	F	O	F	O	O	O	O	F
CXL Extensions DVSEC for Ports	0x0003	F	F	M	F	F	F	F	F	M	M
GPF DVSEC for CXL Ports	0x0004	F	F	M	F	F	F	F	F	F	M
GPF DVSEC for CXL Devices	0x0005	F	F	F	O	F	M	M	F	F	F
PCIe DVSEC for Flex Bus Port	0x0007	F	M	M	M	M	M	M	M	M	M
Register Locator DVSEC	0x0008	F	O	M	O	O	M	M	M	M	M
MLD DVSEC	0x0009	F	F	F	F	F	F	F	M	F	F
PCIe DVSEC for Test Capability	0x000a	F	F	F	M	F	F	O	F	F	F

<sup>1</sup>USP: Upstream Switch Port  
<sup>2</sup>DSP: Downstream Switch Port

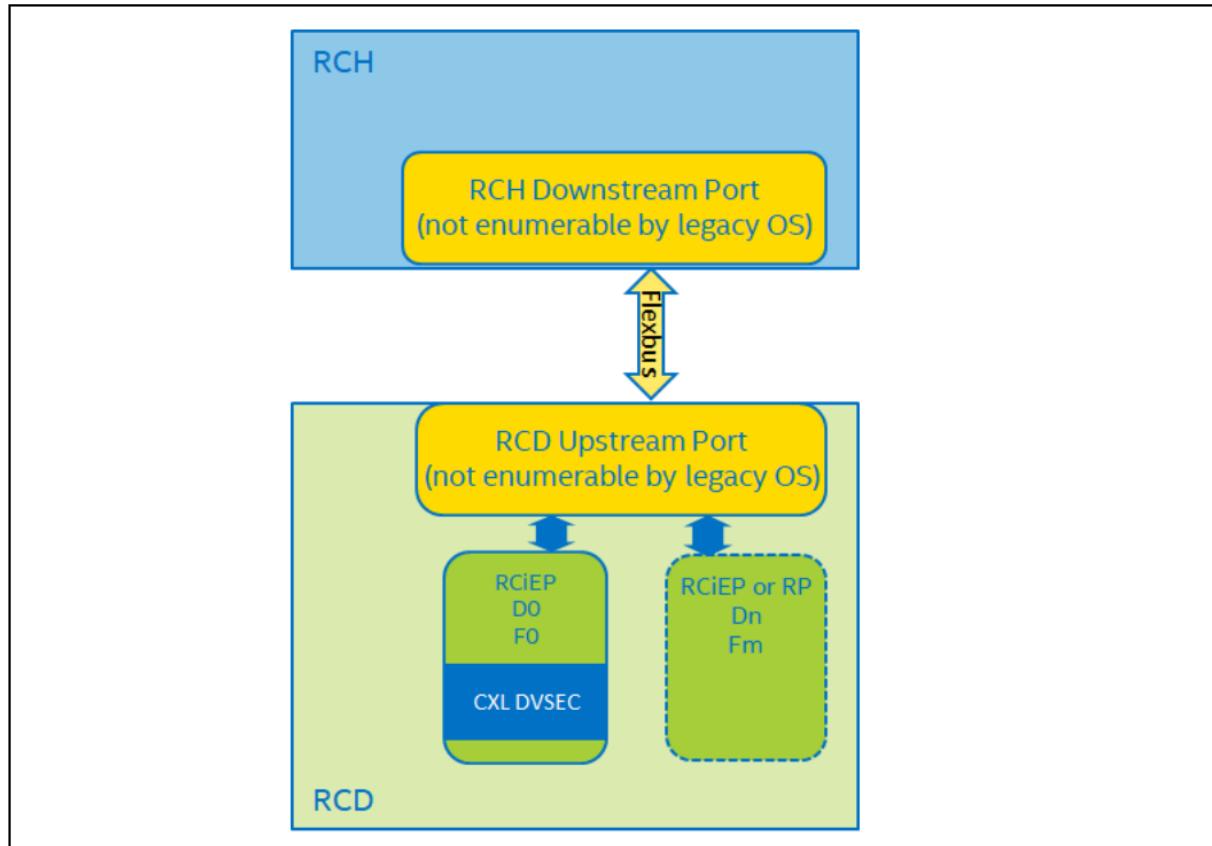
Mandatory      Forbidden      Optional

CXL Capability	DVSECID	HighestDVSECRevisionID	Mandatory	Not Permitted	Optional
PCIe DVSEC for CXL Device ( <a href="#">Section 8.1.3</a> )	0	1	D1, D2, LD, FMLD	P, UP1, DP1, R, USP, DSP	
Non-CXL Function Map DVSEC ( <a href="#">Section 8.1.4</a> )	2	0	P, UP1, DP1, R, DSP	D1, D2, LD, FMLD, USP	
CXL 2.0 Extensions DVSEC for Ports ( <a href="#">Section 8.1.5</a> )	3	0	R, USP, DSP	P, D1, D2, LD, FMLD, UP1, DP1	
GPF DVSEC for CXL Ports ( <a href="#">Section 8.1.6</a> )	4	0	R, DSP	P, D1, D2, LD, FMLD, UP1, DP1, USP	
GPF DVSEC for CXL Devices ( <a href="#">Section 8.1.7</a> )	5	0	D2, LD	P, UP1, DP1, R, USP, DSP, FMLD	D1
PCIe DVSEC for Flex Bus Port ( <a href="#">Section 8.1.8</a> )	7	1	D1, D2, LD, FMLD, UP1, DP1, R, USP, DSP	P	
Register Locator DVSEC ( <a href="#">Section 8.1.9</a> )	8	0	D2, LD, FMLD, R, USP, DSP	P	D1, UP1, DP1
MLD DVSEC ( <a href="#">Section 8.1.10</a> )	9	0	FMLD	P, D1, D2, LD, UP1, DP1, R, USP, DSP	
PCIe DVSEC for Test Capability ( <a href="#">Section 14.16.1</a> )	0Ah	0	D1	P, LD, FMLD, DP1, UP1, R, USP, DSP	D2

DVSEC

<b>PCIe DVSEC for CXL Device ID=0</b>	<b>CXL 2.0 Extensions DVSEC for Ports ID=3</b>
DVSEC CXL Capability (Offset 0Ah)	CXL Port Extension Status (Offset 0Ah)
DVSEC CXL Control (Offset 0Ch)	Port Control Extensions (Offset 0Ch)
DVSEC CXL Status (Offset 0Eh)	Alternate Bus Base (Offset 0Eh)
DVSEC CXL Control2 (Offset 10h)	Alternate Bus Limit (Offset 0Fh)
DVSEC CXL Status2 (Offset 12h)	Alternate Memory Base (Offset 10h)
DVSEC CXL Lock (Offset 14h)	Alternate Memory Limit (Offset 12h)
DVSEC CXL Capability2 (Offset 16h)	Alternate Prefetchable Memory Base (Offset 14h)
DVSEC CXL Range 1 Size High (Offset 18h)	Alternate Prefetchable Memory Limit (Offset 16h)
DVSEC CXL Range1 Size Low (Offset 1Ch)	Alternate Memory Prefetchable Base High (Offset 18h)
DVSEC CXL Range 1 Base High (Offset 20h)	Alternate Prefetchable Memory Limit High (Offset 1Ch)
DVSEC CXL Range 1 Base Low (Offset 24h)	CXL RCRB Base (Offset 20h)
	CXL RCRB Base High (Offset 24h)
<b>GPF DVSEC for CXL Port ID=4</b>	<b>GPF DVSEC for CXL Device ID=5</b>
GPF Phase 1 Control (Offset 0Ch)	GPF Phase 2 Duration (Offset 0Ah)
GPF Phase 2 Control (Offset 0Eh)	GPF Phase 2 Power (Offset 0Ch)
<b>PCIe DVSEC for Flex Bus Port ID=7</b>	<b>Register Locator DVSEC ID=8</b>
DVSEC Flex Bus Port Capability (Offset 0Ah)	Register Offset Low (Block 1)
DVSEC Flex Bus Port Control (Offset 0Ch)	Register Offset High (Block 1)
DVSEC Flex Bus Port Status (Offset 0Eh)	Register Offset Low (Block 2)
DVSEC Flex Bus Port Received Modified TS Data Phase1 (Offset 10h)	Register Offset High (Block 2)
	Register Offset Low (Block 3)
	Register Offset High (Block 3)
<b>MLD DVSEC ID=9</b>	<b>CXL Device Test Capability Advertisement DVSEC ID=A</b>
Number of LD Supported (Offset 0Ah)	DVSEC CXL Test Lock (offset 0Ah)
LD-ID Hot Reset Vector (Offset 0Ch)	DVSEC CXL Test Capability1 (offset 0Ch)
	Device CXL Test Capability2 (Offset 10h)
	DVSEC CXL Test Configuration Base Low (Offset 14h)
	DVSEC CXL Test Configuration Base High (Offset 18h)

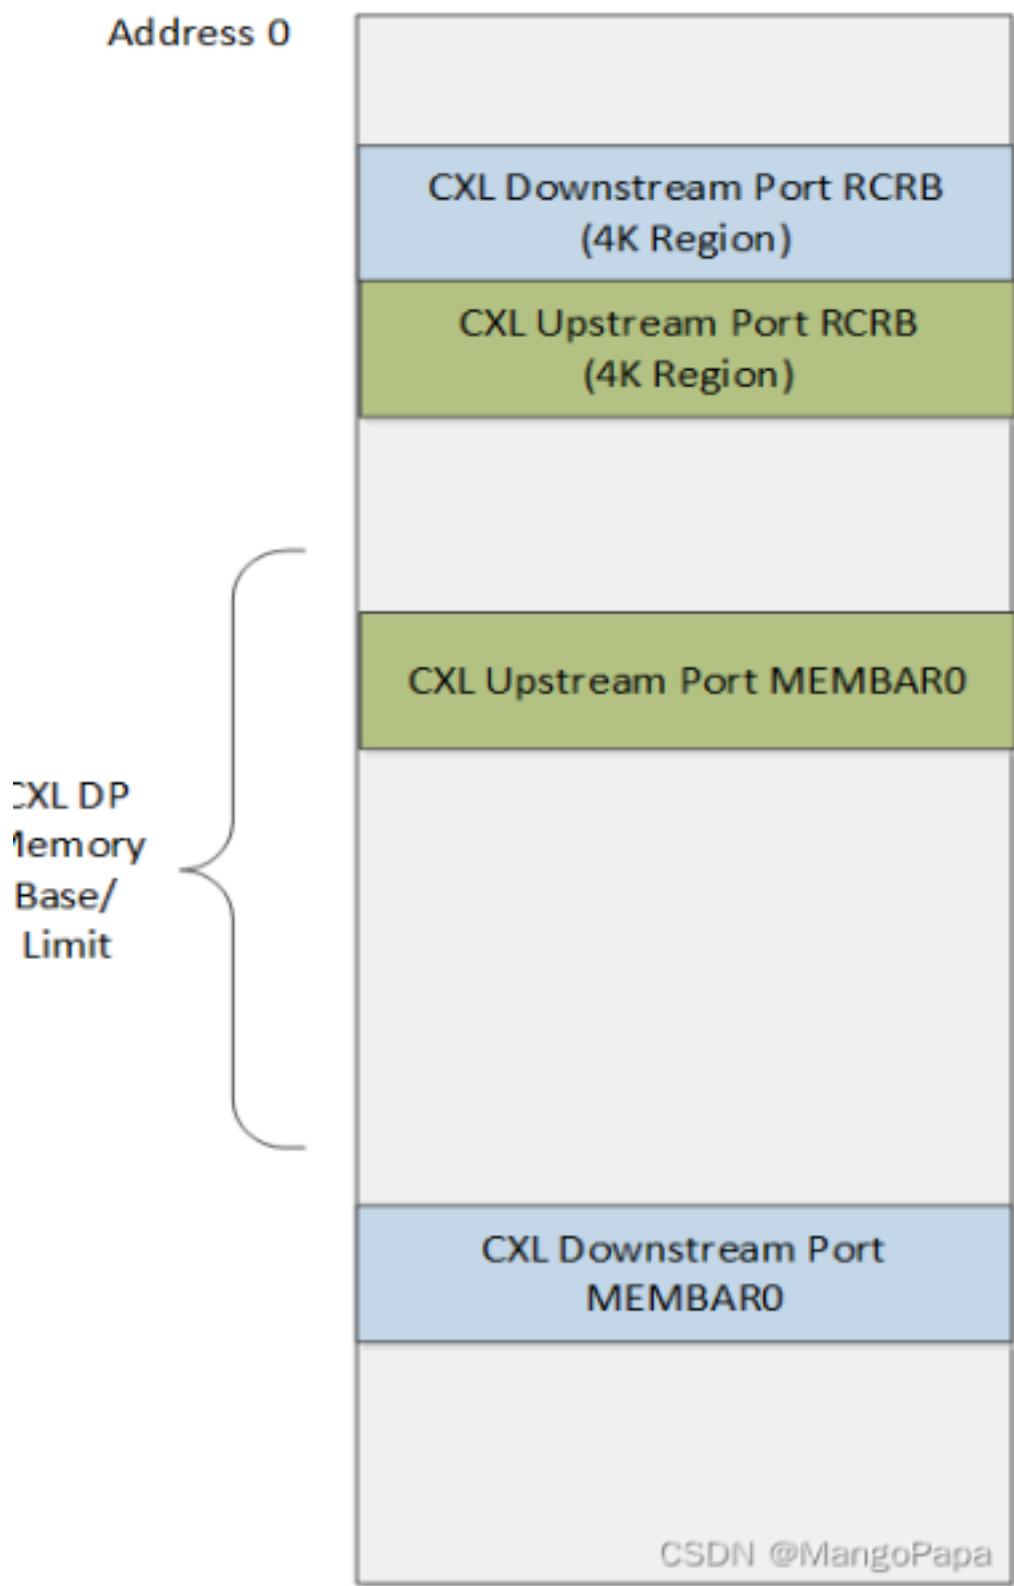
## PCIe Software View of an RCH and RCD



Because the CXL link is not exposed to CXL-unaware OSs, the System Firmware view of the hierarchy is different than that of the CXL-unaware OS.

## Reg in MMIO

	9	PCIe	CXL	DVSEC	CXL	PCIe	PCIe	4KB	CXL Component Register	Size
	CXL	DVSEC	Component Register			Memory	Mapped	MMIO	Memory	
CXL 3.0	CXL	6	MMIO							
• <b>RCH DP RCRB</b>		RCH	4KB		MEMBAR0		ACPI			
• <b>RCH DP Component Register</b>			CXL			ARB/MUX		64KB	RCH DP RCRB	64b
		MEMBAR0	RCH							
• <b>RCD UP RCRB</b>		RCD	4KB		RCH DP RCRB	4KB		RCH DP	Memory	
		MEMBAR0	RCD UP RCRB		RCD	RCH DP	MRd	CXL 1.1	MRd/MWr	
• <b>RCD UP Component Register</b>		CXL			ARB/MUX		64KB	RCRB	RCD	
		RCRD UP RCRB	64b	MEMBAR0	RCD	RCH DP	Memory	MEMBAR0	RCRB	
		RCD			Register Locator DVSEC					
• <b>Component Register for ALL Other CXL Components</b>					CXL Component				Register	
					Locator DVSEC					
• <b>CXL Host Bridge Component Register (CHBCR)</b>					CXL Host Bridge			RP	Memory	
					Interleaving	ACPI CEDT				
RCRB Components Register					CHBCR Register Locator DVSEC			Component Registers for All		
Other CXL Components										



PCIe CXL blog

[https://blog.csdn.net/weixin\\_40357487/article/details/132553156](https://blog.csdn.net/weixin_40357487/article/details/132553156)

## 6.3 CXL in CMN

[toc]

### CXS

ARM RD\_N2 chip to chip CCIX2.0 CMN-700 TRM ARM CML\_SMP

S CXL CXS CCIX2.0 CML\_SMP CCIX2.0

CMN CCG chip CXL

Coherent Multichip Link (CML) device = CCG

A given multi-chip link can be used for: • SMP (CML\_SMP) connection • CXL device attachment

A CML device (CCG) can be configured to be used for **CML\_SMP** connection or **CXL device attachment**.

For SMP systems, CCG block is required to enable multi-chip SMP communications over a **CXS issue B interface**.

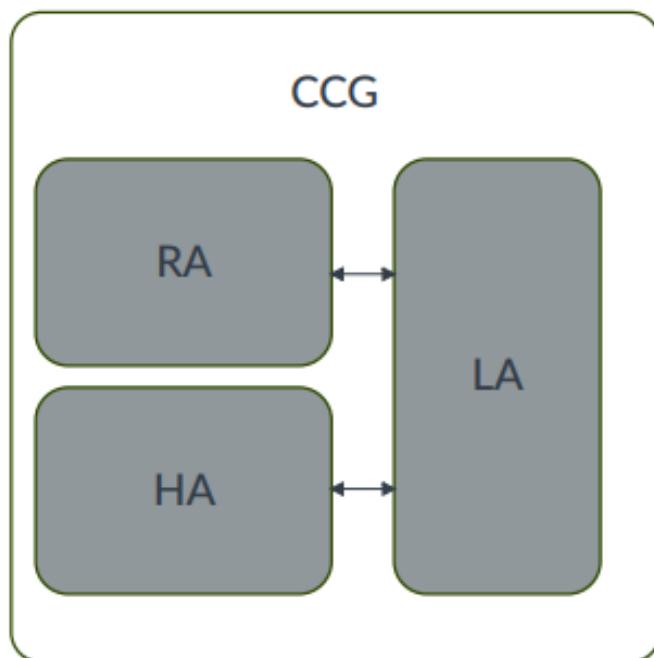
CXL CXS issue A interface

### CCG

#### CML components

The following figure shows a CCG block diagram with RA, HA, and LA.

**Figure 3-17: CCG block diagram with RA, HA, and LA**

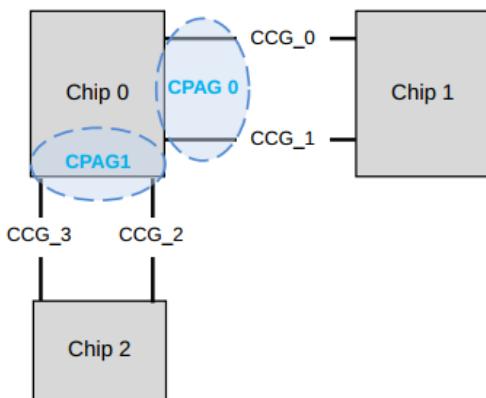


### 3.1.3.2 CML Port Aggregation Groups

The CMN-700 CML configuration supports up to 32 gateway blocks (CCGs). These CML gateways can be grouped in up to 16 CML Port Aggregation Groups (CPAGs).

This feature can be used when connecting two or more chips together with multiple ports between the chips. For example, the following figure shows three chips that are connected by four CCGs that are grouped into two CPAGs.

**Figure 3-19: CML Port Aggregation Groups**



## CCG

The **CML RA** node type has the following MPAM modes:

- SMP mode The CML RA passes the MPAM field on the USER field of the request. When snooped, the CML RA receives the MPAM field and passes it through the CHI SNP MPAM field.
- Non-SMP mode The CML RA drops the MPAM field that is received on the CHI request. The CML RA also does not receive MPAM field on CML snoops in this mode.
- CXSA mode The CML RA passes the MPAM field on the USER field of the request, even though CXSA is in non-SMP mode. You can use a configuration bit to enable passing of MPAM attributes when in CXSA mode.

The **CML HA** node type has the following MPAM modes:

- SMP mode The CML HA receives MPAM fields through the USER field of the request and passes them through the CHI MPAM field. On incoming CHI transactions, the CML HA passes CHI MPAM values through on the USER field.
- Non-SMP mode The CML HA drops the MPAM values that it receives on CHI snoop. The CML HA does not receive the MPAM attributes on CML requests.

CXL      enable **CXSA** mode in RA

LA      enable CXL

## CCG SAM

CMN CCG      CXL      CMN700 type3      CXL type3      memory      CMN SAM      CCG

## 6.4 CXL in SCP and EDK2

SCP      CXL

SCP EDK2

### CCG

- SHA-1: bc14e780c120fc656ce905d6e219078a0006da88

– module/cmn700: configuring CCG for mapping Host address to CXL mem A Memory region is reserved for CXL Memory. CXL.Mem (0x3fe\_0000\_0000) comes under 4TB Chip-0 memory and the whole region is by default configured as SCG. Configured CXL.Mem region in HNF-SAM HTG and CCG SA node IDs for HTGs in following order -

HNF\_SAM\_CCG\_SA\_NODEID\_REG HNF\_SAM\_HTG\_CFG3\_MEMREGION HNF\_SAM\_HTG\_CFG2\_MEMREGION  
HNF\_SAM\_HTG\_CFG1\_MEMREGION

CXL Memory region is accessible as Normal memory with above configuration.

This patch maps Host address space to CXL device mem area through CCG node, based on the CXL device memory size, which is discovered by CXL module. CXL module invokes runtime CMN700 API for mapping the host address space and configuring CCG node.

This patch also adds a flag "cxl\_mem" in CCG\_Config structure for identifying host region and configuration reserved for CXL device memory purpose and thus differentiating from Remote chip memory.

Signed-off-by: Sayanta Pattanayak [sayanta.pattanayak@arm.com](mailto:sayanta.pattanayak@arm.com) Change-Id: I988f471db6a6a55f97320519413daedd8ea

config\_cmn700.c

N2      0-0x400000000000      SCG

HNSAM 0-0x3fe0000000000 CCG-cxl\_mem

```
#if (PLATFORM_VARIANT == 1)
static const struct mod_cmn700_ccg_config ccg_config_table_chip_0[] = {
{
    .remote_rnf_count = 0,
    .remote_mmap_table = {
        .base = UINT64_C(0x3fe000000000),
        .size = UINT64_C(8) * FWK_GIB,
        .type = MOD_CMN700_REGION_TYPE_CCG,
    },
    .ra_mmap_table = {
        {
            .base = UINT64_C(0x3fe00000000),
            .size = UINT64_C(8) * FWK_GIB,
            .remote_haid = 0,
        },
        { 0 }
    },
    .remote_agentid_to_linkid_map = {
        {
            .remote_agentid_start = (RNF_PER_CHIP * CHIP_0),
            .remote_agentid_end = (RNF_PER_CHIP * CHIP_0)
        }
    }
},
```

```

        },
    },
    .smp_mode = false,
    .cxl_mem = true,
},
};

#endif

entry 8G      flag cxl_mem  CCG

static int cmn700_setup_rnsam_ccg_regions(void)
{
    const struct mod_cmn700_config *config;
    const struct mod_cmn700_mem_region_map *region;
    struct cmn700_rnsam_reg *rnsam;
    unsigned int count;
    unsigned int cxra_ldid;
    unsigned int cxra_node_id;
    unsigned int idx;
    uint32_t bit_pos;
    uint32_t group;

    config = ctx->config;
    /* Do configuration for CCG Nodes */
    for (idx = 0; idx < config->ccg_table_count; idx++) {
        region = &config->ccg_config_table[idx].remote_mmap_table;
        if (region->type != MOD_CMN700_REGION_TYPE_CCG) {
            return FWK_E_DATA;
        }
        FWK_LOG_INFO(
            MOD_NAME " [0x%llx - 0x%llx] %s",
            region->base,
            region->base + region->size - 1,
            mmap_type_name[region->type]);
        /* If the region is for extended memory area like CXL.Mem
         * and connected through CCG then the region shouldn't be
         * marked as Non-Hash region. CXL.Mem region should be part
         * of Hashed cache group area.
         */
        if (config->ccg_config_table[idx].cxl_mem == true)
            continue;

        for (count = 0; count < ctx->internal_rnsam_count; count++) {
            rnsam = ctx->internal_rnsam_table[count];
CCG      CXL.Mem      non-hash      hash

static int map_ccg_for_cxl_mem(uint64_t size)
{

```

```

uint32_t idx;
const struct mod_cmn700_config *config = ctx->config;
const struct mod_cmn700_ccg_config *ccg_config;

cmn700_rnsam_stall();

/* Do configuration of CCG Node for mapping remote CXL Mem area. */
for (idx = 0; idx < config->ccg_table_count; idx++) {
    ccg_config = &(config->ccg_config_table[idx]);
    if (ccg_config->cxl_mem == true)
        ccg_setup_for_remote_mem(size, ctx, ccg_config);
}

cmn700_rnsam_uninstall();

return FWK_SUCCESS;
}

```

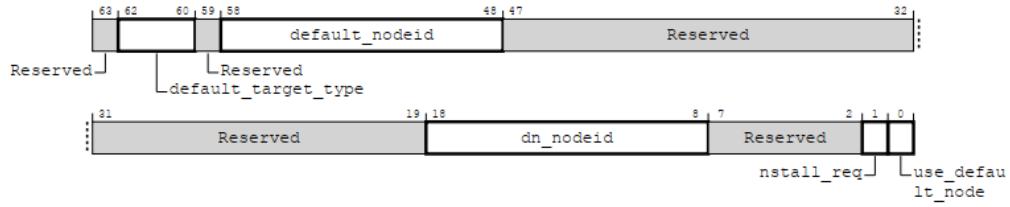
API CCG CXL.Mem

rnsam stall unstall rnsam default node

STATUS 0x2 unstall enable RNSAM

online RNSAM stall

**Figure 4-647: rnsam\_status**



**Table 4-663: rnsam\_status attributes**

Bits	Name	Description	Type	Reset
[63]	Reserved	Reserved	RO	-
[62:60]	default_target_type	Indicates node type 3'b000 HN-F 3'b001 HN-I 3'b010 CXRA 3'b011 HN-P 3'b100 PCI-CXRA 3'b101 HN-S <b>Others</b> Reserved  CONSTRAINT: Only applicable for RN-I	RW	3'b001
[59]	Reserved	Reserved	RO	-
[58:48]	default_nodeid	Default Node ID	RW	Configuration dependent
[47:19]	Reserved	Reserved	RO	-
[18:8]	dn_nodeid	DN Node ID for DN operations	RW	11'b0
[7:2]	Reserved	Reserved	RO	-
[1]	nstall_req	Indicates RN SAM is programmed and ready 1'b0 STALL requests 1'b1 UNSTALL requests	RW	1'b0
[0]	use_default_node	Indicates target ID selection mode 1'b0 Enables RN SAM to hash address bits and generate target ID 1'b1 Uses default target ID	RW	1'b1

```

63
64 static void cmn700_rnsam_stall(void) {
65 {
66     struct cmn700_rnsam_reg *rnsam;
67     uint32_t rnsam_idx;
68
69     for (rnsam_idx = 0; rnsam_idx < ctx->internal_rnsam_count; rnsam_idx++) {
70         rnsam = ctx->internal_rnsam_table[rnsam_idx];
71         rnsam->STATUS |= (rnsam->STATUS & ~CMN700_RNSAM_STATUS_UNSTALL) | CMN700_RNSAM_STATUS_USE_DEFAULT_TARGET_ID;
72     }
73 }
74 __DSB();
75 }
76
77 static void cmn700_rnsam_uninstall(void) {
78 {
79     struct cmn700_rnsam_reg *rnsam;
80     uint32_t rnsam_idx;
81
82     for (rnsam_idx = 0; rnsam_idx < ctx->internal_rnsam_count; rnsam_idx++) {
83         rnsam = ctx->internal_rnsam_table[rnsam_idx];
84         rnsam->STATUS |= (rnsam->STATUS | CMN700_RNSAM_STATUS_UNSTALL) & ~CMN700_RNSAM_STATUS_USE_DEFAULT_TARGET_ID;
85     }
86 }
87 __DSB();
88 }

```

ccg\_setup\_for\_remote\_mem

```

int ccg_setup_for_remote_mem(
    uint64_t size,
    struct cmn700_device_ctx *ctx,
    const struct mod_cmn700_ccg_config *ccg_config)
{
    uint64_t reg_val;
    unsigned int index;
    unsigned int ccg_ldid;
    struct cmn700_hnf_reg *hnf_reg;
    struct cmn700_ccg_ra_reg *ccg_ra_reg;

    cmn700_ccg_ctx.is_prog_for_port_agg = false;
    /* Enable CXSA */
    ccg_ldid = get_ldid(ctx, cmn700_ccg_ctx.is_prog_for_port_agg);
    ccg_ra_reg = ctx->ccg_ra_reg_table[ccg_ldid].ccg_ra_reg;

    for (index = 0; index < ctx->hnf_count; index++) {
        /* Programming sequence to enable CXL.mem regions inside HNSAM */

        hnf_reg = (struct cmn700_hnf_reg *)ctx->hnf_node[index];

        /* Configuring CCG SA node IDs for HTGs in the HNSAM */
        reg_val = hnf_reg->HNF_SAM_CCG_SA_NODEID_REG[0];
        reg_val &= ~(CMN700_HNF_SAM_CCG_SA_NODEID_MASK);
        reg_val |= get_node_id(ccg_ra_reg);

```

```

    hnf_reg->HNF_SAM_CCG_SA_NODEID_REG[0] = reg_val;

    /* CXSA/CXLSA aggregated SA selection function */
    /*      region SN NODE 8SN node, CXSA      */
    hnf_reg->HNF_SAM_HTG_CFG3_MEMREGION[0] |=
        (CMN700_HNF_SAM_HTG_MODE_CXSA <<
         CMN700_HNF_SAM_HTG_SN_MODE_POS);

    /* 64B interleaved */
    /* CXSA      interleaved          0 64B */
    reg_val = hnf_reg->HNF_SAM_HTG_CFG3_MEMREGION[0];
    reg_val &= ~(CMN700_HNF_SAM_HTG_SA_DEVICE_INTERLEAVE_CNTL_MASK <<
                  CMN700_HNF_SAM_HTG_SA_DEVICE_INTERLEAVE_CNTL_POS);
    hnf_reg->HNF_SAM_HTG_CFG3_MEMREGION[0] |= reg_val;

    /* 1 CXSA/CXLSA port used */
    /*      0 */
    reg_val = hnf_reg->HNF_SAM_HTG_CFG3_MEMREGION[0];
    reg_val &= ~(CMN700_HNF_SAM_HTG_SA_PORTS_CNT_MASK <<
                  CMN700_HNF_SAM_HTG_SA_PORTS_CNT_POS);
    hnf_reg->HNF_SAM_HTG_CFG3_MEMREGION[0] |= reg_val;

    /* htg_region_end_addr[51:20] = end address of HTG region */
    reg_val = ccg_config->ra_mmap_table[0].base +
              (size - 1);
    reg_val &= ~(CMN700_HNF_SAM_HTG_REGION_ADDR_RES_MASK);
    hnf_reg->HNF_SAM_HTG_CFG2_MEMREGION[0] |= reg_val;

    /* htg_region_base_addr[51:20] = start address of HTG region */
    reg_val = ccg_config->ra_mmap_table[0].base;
    reg_val &= ~(CMN700_HNF_SAM_HTG_REGION_ADDR_RES_MASK);
    hnf_reg->HNF_SAM_HTG_CFG1_MEMREGION[0] |= reg_val;

    /* Configuring HTG region as Valid */
    hnf_reg->HNF_SAM_HTG_CFG1_MEMREGION[0] |=
        (UINT64_C(0x1) << CMN700_HNF_SAM_HTG_REGION_VALID_POS);
}

/*
 * Program the CXRA SAM with the address range and the corresponding
 * remote HAID.
 */
program_ccg_ra_sam_addr_region(ctx, ccg_config);

update_cxl_mem_region(ccg_config->ra_mmap_table[0].base,
                      ccg_config->ra_mmap_table[0].size);

return FWK_SUCCESS;
}

```

## link

cxl cxsa enable

- RA SAM region:
- RA SAM valid valid

```
ccg_ra_reg->CCG_RA_SAM_ADDR_REGION_REG[0] |=
    fwk_math_log2(size) |
    (cxl_cfg->base & CCG_RA_REMOTE_SAM_BASE_ADDR_MASK) |
    ((uint64_t)0 << CCG_RA_REMOTE_SAM_TGT_HAID_POS);

ccg_ra_reg->CCG_RA_SAM_ADDR_REGION_REG[0] |= CCG_RA_REMOTE_SAM_VALID_MASK;
```

[5]	cxl_cache_en	Enable CXL .cache mode, by default is disabled	RW	1'b0
[4]	cxl_mem_en	Enable CXL .mem mode, by default is enabled	RW	1'b1
[3:2]	cxl_type	Used to program CXL Type for RA	RW	2'b11
		2'b00                      Reserved 2'b01                      Type1 2'b10                      Type2 2'b11                      Type3		
[1]	la_device_mode_en	Enable the Device mode, by default set to Host mode	RW	1'b0
[0]	la_cxl_mode_en	When set enables CXL mode	RW	1'b0

- CCLA flex bus control MEM CACHE
- CCLA config control MEM CACHE

```
reg_val1 = ccla_reg->CCLA_DVSEC_FLEX_BUS_PORT_CONTROL;
reg_val2 = ccla_reg->CCLA_CFG_CTL;

if (cxl_cfg->is_cache_en) {
    reg_val1 |= CCLA_CFG_CTL_CXL_CACHE_EN_MASK;
    reg_val2 |= CCLA_DVSEC_FLEX_BUS_PORT_CTRL_CACHE_EN_MASK;
}
if (cxl_cfg->is_mem_en) {
    reg_val1 |= CCLA_CFG_CTL_CXL_MEM_EN_MASK;
    reg_val2 |= CCLA_DVSEC_FLEX_BUS_PORT_CTRL_MEM_EN_MASK;
}
if (cxl_cfg->is_device_mode) {
    reg_val1 |= CCLA_CFG_CTL_LA_DEVICE_MODE_EN_MASK;
}
```

```
if (cxl_cfg->is_multi_logical_dev) {
    reg_val2 |= CCLA_DVSEC_FLEX_BUS_PORT_CTRL_MULTI_LOGICAL_DEVICE_EN_MASK;
}
```

```
reg_val1 |= (cxl_cfg->cxl_type << CCLA_CFG_CTL_CXL_TYPE_POS);
reg_val1 |= CCLA_CFG_CTL_LA_CXL_MODE_EN_MASK;
```

```
ccg_ra_reg->CCG_RA_CFG_CTRL |= 0x200;
```

[9]	cxsa_mode_en	When set, enables the CCIX Subordinate Agent mode. In this mode RA functions as a CCIX Subordinate Agent  1'b1 CCIX Subordinate Agent 1'b0 CCIX Requesting Agent	RW	1'b0
-----	--------------	---	----	------

link

```
ccg_ra->LINK_REGS[0].CCG_CCPRTCL_LINK_CTRL |= CCPRTCL_LINK_CTRL_LINK_EN_MASK;  
ccg_ha->LINK_REGS[0].CCG_CCPRTCL_LINK_CTRL |= CCPRTCL_LINK_CTRL_LINK_EN_MASK;
```

[0]	lnk0_link_en	agent may still be in link_up state.  Enables CCIX Link 0 when set  1'b0 Link is disabled 1'b1 Link is enabled	RW	1'b0
-----	--------------	---	----	------

```
int exchange_protocol_credit
```

[7:4]	lnk0_num_snpcrds	Controls the number of CCIX snoop credits assigned to Link 0  4'h0 Total credits are equally divided across all links 4'h1 25% of credits assigned 4'h2 50% of credits assigned 4'h3 75% of credits assigned 4'h4 100% of credits assigned 4'hF 0% of credits assigned	RW	4'b0
-------	------------------	---	----	------

## 6.5 CXL

[toc]

### CXL

CXL PCIe

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	Byte Offset
PCI Express Extended Capability Header																												+000h				
32.0 GT/s Capabilities Register																												+004h				
32.0 GT/s Control Register																												+008h				
32.0 GT/s Status Register																												+00Ch				
Received Modified TS Data 1 Register																												+010h				
Received Modified TS Data 2 Register																												+014h				
Transmitted Modified TS Data 1 Register																												+018h				
Transmitted Modified TS Data 2 Register																												+01Ch				
32.0 GT/s Eq Ctl: Lane 3	32.0 GT/s Eq Ctl: Lane 2	32.0 GT/s Eq Ctl: Lane 1	32.0 GT/s Eq Ctl: Lane 0																										+020h			
32.0 GT/s Eq Ctl: Lane 7	32.0 GT/s Eq Ctl: Lane 6	32.0 GT/s Eq Ctl: Lane 5	32.0 GT/s Eq Ctl: Lane 4																										+024h			
32.0 GT/s Eq Ctl: Lane 11	32.0 GT/s Eq Ctl: Lane 10	32.0 GT/s Eq Ctl: Lane 9	32.0 GT/s Eq Ctl: Lane 8																										+028h			
32.0 GT/s Eq Ctl: Lane 15	32.0 GT/s Eq Ctl: Lane 14	32.0 GT/s Eq Ctl: Lane 13	32.0 GT/s Eq Ctl: Lane 12																										+02Ch			
32.0 GT/s Eq Ctl: Lane 19	32.0 GT/s Eq Ctl: Lane 18	32.0 GT/s Eq Ctl: Lane 17	32.0 GT/s Eq Ctl: Lane 16																										+030h			
32.0 GT/s Eq Ctl: Lane 23	32.0 GT/s Eq Ctl: Lane 22	32.0 GT/s Eq Ctl: Lane 21	32.0 GT/s Eq Ctl: Lane 20																										+034h			
32.0 GT/s Eq Ctl: Lane 27	32.0 GT/s Eq Ctl: Lane 26	32.0 GT/s Eq Ctl: Lane 25	32.0 GT/s Eq Ctl: Lane 24																										+038h			
32.0 GT/s Eq Ctl: Lane 31	32.0 GT/s Eq Ctl: Lane 30	32.0 GT/s Eq Ctl: Lane 29	32.0 GT/s Eq Ctl: Lane 28																										+03Ch			

Figure 7-87 Physical Layer 32.0 GT/s Extended Capability

Table 7-77 Transmitted Modified TS Data 2 Register

Bit Location	Description	Attributes
23:0	<p><b>Transmitted Modified TS Information 2</b> - If <u>Modified TS Received</u> is Set, this field contains the <u>Modified TS Information 2</u> field from the <u>Modified TS2 Ordered Set</u> transmitted during the <u>Configuration.Complete LTSSM State</u> (see <u>Section 4.2.6.3.6</u> ).</p> <p>Bits 23:16 contain the value of Symbol 14.</p> <p>Bits 16:8 contain the value of Symbol 13.</p> <p>Bits 7:0 contain the value of Symbol 12.</p> <p>If PCI Express (Usage Mode 0) is the only one supported, this field is permitted to be hardwired to 00 0000h.</p> <p>Default is 00 0000h.</p>	RO
25:24	<p><b>Alternate Protocol Negotiation Status</b> - Indicates the status of the Alternate Protocol Negotiation. Encodings are:</p> <ul style="list-style-type: none"> <li><b>00b</b>      Alternate Protocol Negotiation not supported - <u>Modified TS Usage Mode 2 Supported</u> - <u>Alternate Protocol</u> is Clear.</li> <li><b>01b</b>      Alternate Protocol Negotiation disabled - <u>Modified TS Usage Mode 2 Supported</u> - <u>Alternate Protocol</u> is Set but <u>Modified TS Usage Mode Selected</u> was not 2 during the appropriate LTSSM State.</li> <li><b>10b</b>      Alternate Protocol Negotiation failed - <u>Alternate Protocol Negotiation</u> was attempted and did not locate a protocol that was supported on both ends of the Link.</li> <li><b>11b</b>      Alternate protocol Negotiation succeeded - <u>Alternate Protocol Negotiation</u> located one or more protocols that were supported on both ends of the Link and the Downstream Port selected one of those protocols for use.</li> </ul> <p>If Set, Alternate Protocol Negotiation completed successfully. If Clear, Alternate Protocol Negotiation negotiation has not completed successfully. If <u>Modified TS Usage Mode 1 Supported</u> - <u>Training Set Message</u> and <u>Modified TS Usage Mode 2 Supported</u> - <u>Alternate Protocol</u> are both Clear, this register is permitted to be hardwired to 0000 0000h.</p> <p>If <u>Modified TS Usage Mode 2 Supported</u> - <u>Alternate Protocol</u> is Clear, this bit is hardwired to 0b.</p> <p>If <u>Modified TS Usage Mode Selected</u> does not equal 2, this bit contains 0b.</p> <p>This bit is Cleared on <u>Detect LTSSM State</u>.</p> <p>Default is 0b.</p>	RO

PCIe DVSEC for CXL Devices      CXL PCIe DVSEC      RCD LD SLD FMLD      CXL Device      D0F0  
 DVSEC      Device      DVSEC      Device      PCIe Device      CXL Device

## CXL2.0

PCIe

## CXL1.1

RCiEP

## HDM Decoder

## 6.6 DDR Introduction

[toc]

## Introduction

SDRAM vs. SRAM

- Synchronous Dynamic Random Access Memory

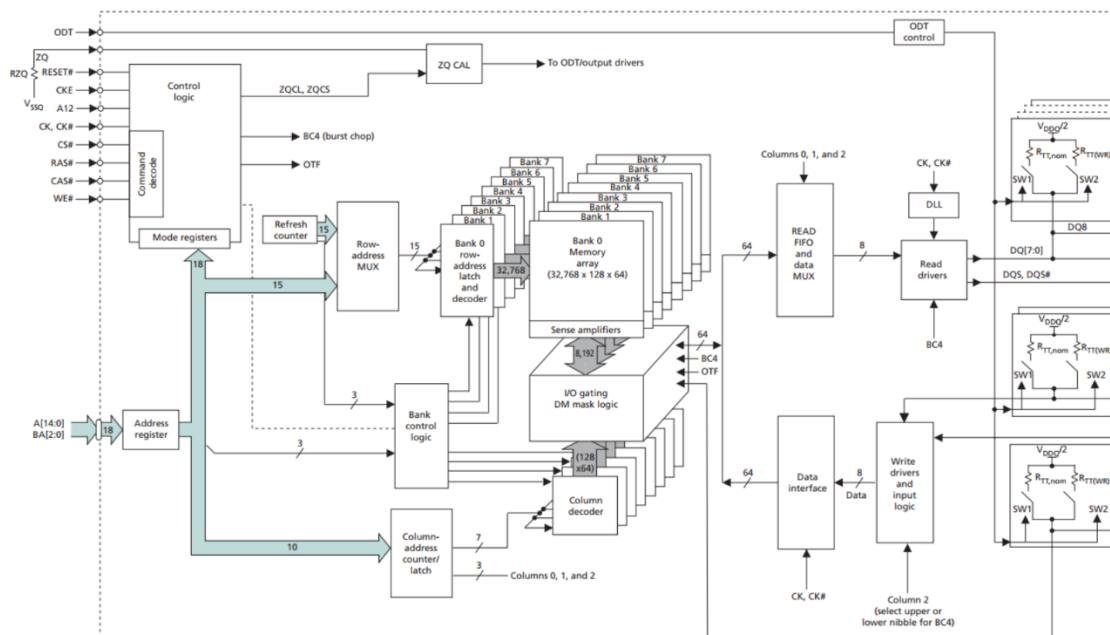
- Static RAM

Type of RAM	Cost per cell	Size of each cell	Power Dissipated	Speed
Dynamic RAM	Low	Very small	High	Slow
Static RAM	High	Large	Low	High

SRAM SDRAM



Figure 4: 256 Meg x 8 Functional Block Diagram



DDR3 SDRAM

- 18b

- 8 bank -> BA[2:0], 3 bit 8
- A 15b
- 10b;
- 8192

**Row Address** 16 A0-A15

**Column Address** 10 A0-A9

**Bank Bank Address** 3 BA0-BA2

4 8

### 512Mb, 1Gb, 2Gb, 4Gb, 8Gb

#### 1Gb

Configuration	256Mb x 4	128Mb x 8	64Mb x 16
# of Banks	8	8	8
Bank Address	BA0 - BA2	BA0 - BA2	BA0 - BA2
Auto precharge	A10/AP	A10/AP	A10/AP
BC switch on the fly	A12/BC#	A12/BC#	A12/BC#
Row Address	A0 - A13	A0 - A13	A0 - A12
Column Address	A0 - A9,A11	A0 - A9	A0 - A9
Page size <sup>1</sup>	1 KB	1 KB	2 KB

#### 2Gb

Configuration	512Mb x 4	256Mb x 8	128Mb x 16
# of Banks	8	8	8
Bank Address	BA0 - BA2	BA0 - BA2	BA0 - BA2
Auto precharge	A10/AP	A10/AP	A10/AP
BC switch on the fly	A12/BC#	A12/BC#	A12/BC#
Row Address	A0 - A14	A0 - A14	A0 - A13
Column Address	A0 - A9,A11	A0 - A9	A0 - A9
Page size <sup>1</sup>	1 KB	1 KB	2 KB

256Mb Configuration

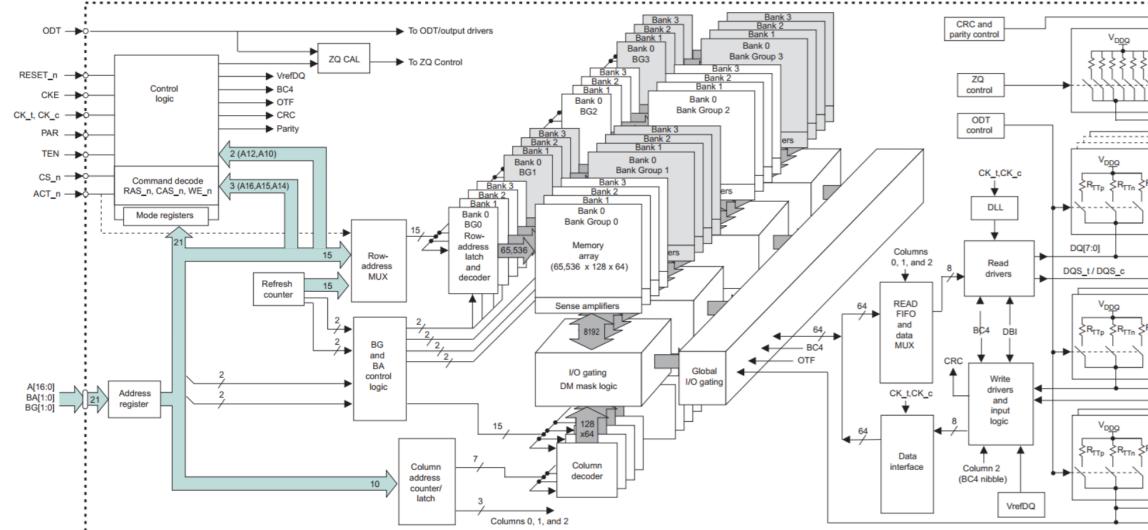
	256Mb x 4	128Mb x 8	64Mb x 16
# of Banks	8	8	8
Bank Address	BA0 - BA2	BA0 - BA2	BA0 - BA2
Auto precharge	A10/AP	A10/AP	A10/AP
BC switch on the fly	A12/BC#	A12/BC#	A12/BC#
Row Address	A0 - A13	A0 - A13	A0 - A12
Column Address	A0 - A9,A11	A0 - A9	A0 - A9
Page size	1 KB	1 KB	2 KB

2Gb Configuration

	512Mb x 4	256Mb x 8	128Mb x 16
# of Banks	8	8	8
Bank Address	BA0 - BA2	BA0 - BA2	BA0 - BA2
Auto precharge	A10/AP	A10/AP	A10/AP

	512Mb x 4	256Mb x 8	128Mb x 16
BC switch on the fly	A12/BC#	A12/BC#	A12/BC#
Row Address	A0 - A14	A0 - A14	A0 - A13
Column Address	A0 - A9,A11	A0 - A9	A0 - A9
Page size	1 KB	1 KB	2 KB

**Figure 3: 512 Meg x 8 Functional Block Diagram**



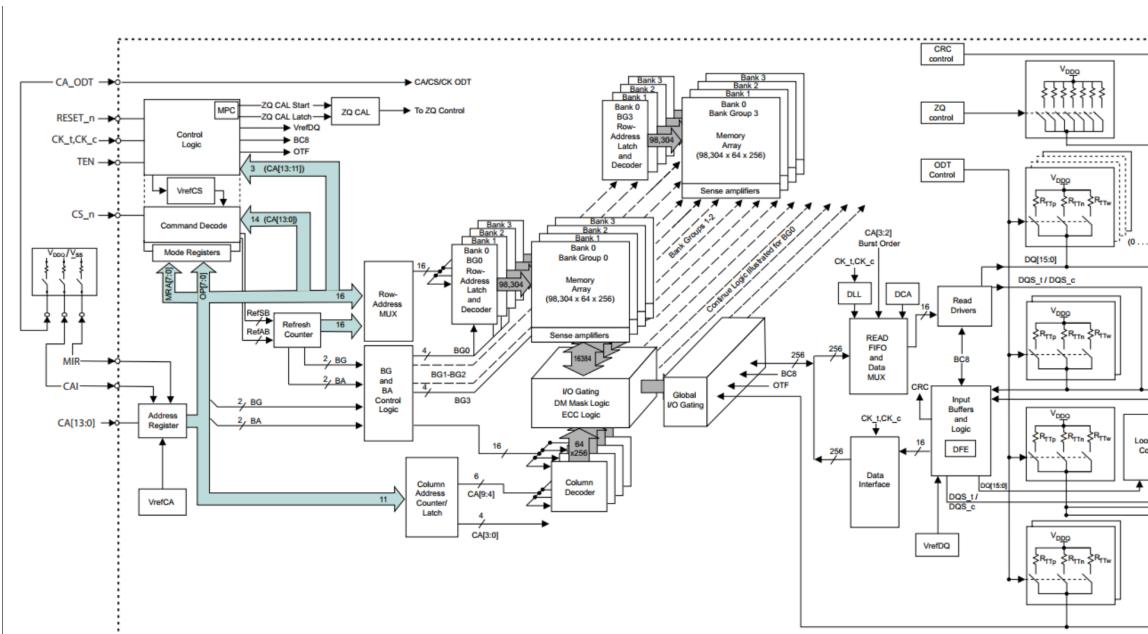
**DDR4 SDRAM**

**Row Address** 16 A0-A15

**Column Address** 10 A0-A9

**Bank** **Bank Address** 2 BA0-BA1

**Bank** **Bank Group Address** 2 BG0-BG1



**DDR5 SDRAM**

**Row Address** 16 A0-A15

**Column Address** 10 A0-A9

**Bank Bank Group Address 2 BA0-BA1**

**Bank Bank Address 3 BG0-BG2**

**Command Address CA**

1. Channel

2. Module              Channel

3. Rank    DDR SDRAM              Module              Rank    Rank              Rank              CS\_n              SDRAM

4. Chip      DDR SDRAM              64    Rank    8    x8    Chip

5. Bank Group DDR4      Bank Group      DDR4      Bank Group              tCCD\_L

6. Bank      Bank      Row      Bank      Activate/Precharge

7. Row Activate/Precharge

8. Column    Column    n    Cell n    SDRAM

9. Cell    Cell    1 bit

## **Page**

- **Row Column DDR**
- **Page**
- 1024                      1024  
2K page      1024      x16      2 B      2KB

## **8N**

32    4 8      8N

32

DQ DQS    DM

Figure: Point-to-Point (DQ/DQS) and Fly-by (CAC) Routing

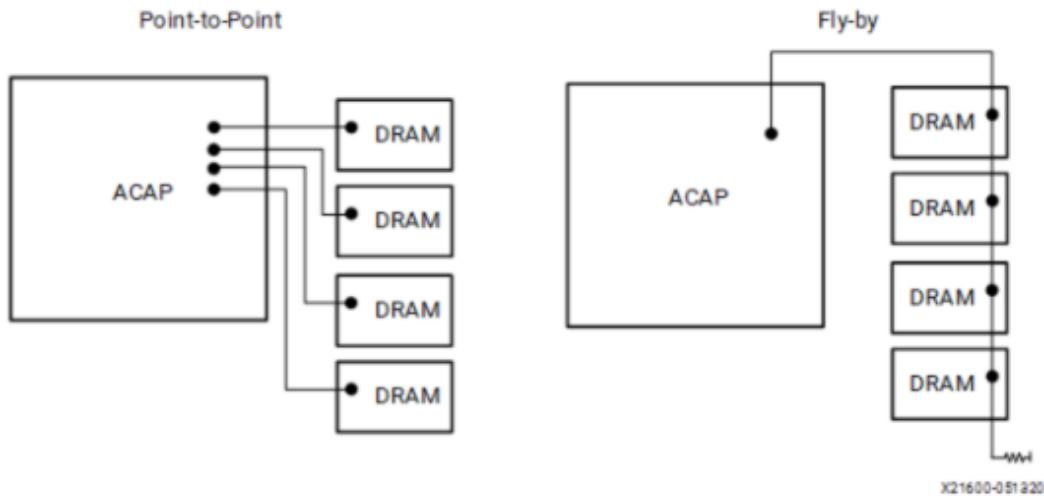


Figure 9: SDRAM 的两种信号连接方式 (图源 Versal ACAP PCB Design User Guide (UG863))

fly-by      DRAM

CLK/CLK# DQS AD CMD

DQ

DQ 8                  kua DQ

#### Write Leveling

Fly-by Topology

SDRAM      DQS      CK              DQ      DQS

DQS Data Strobe    CK Clock

**DQS**

**CK**

**DQ**                  DQ

DQS	DQS-DQS#	CK	DQ	DDR3	01010101	DQS-DQS#	DQ
0 1	Write leveling						

1. SDRAM Write Leveling      SDRAM DQS CK      DQ
2. DQS      DQS      DQ      0-1      001111111111111111110000      0      1      0
3. DQS      DQ      0 1      DQS      CK
4. SDRAM Write Leveling

#### Read Leveling

SDRAM

	SDRAM	Pattern
1.		
2.	0	
3.		
4.		1 0
5.	1	

0-1      Write Leveling    0-1                  0-1      Read Leveling  
Read Centering

Spec

SDRAM JEDEC

- [JESD79F: DDR SDRAM](#)
  - [JESD79-2F: DDR2 SDRAM](#)
  - [JESD79-3F: DDR3 SDRAM](#)
  - [JESD79-4D: DDR4 SDRAM]([https://www.jedec.org/document\\_search?search\\_api\\_views\\_fulltext=jesd79-4 ddr4](https://www.jedec.org/document_search?search_api_views_fulltext=jesd79-4%20ddr4))
  - [JESD79-5B: DDR5 SDRAM](#)

## DDR      LPDDR

- JESD209B: LPDDR SDRAM
  - JESD209-2F: LPDDR2 SDRAM
  - JESD209-3C: LPDDR3 SDRAM
  - JESD209-4D: LPDDR4 SDRAM
  - JESD209-5B: LPDDR5 SDRAM

## HBM SDRAM

- JESD235D: HBM
  - JESD238A: HBM3

GDDR SGRAM

- SDRAM3.11.5.8 R16.01: GDDR4 SGRAM
  - JESD212C.01: GDDR5 SGRAM
  - JESD232A.01: GDDR5X SGRAM
  - JESD250D: GDDR6 SGRAM

<https://www.cnblogs.com/sky-heaven/p/15948268.html>

Blog

<https://zhuanlan.zhihu.com/p/26327347>

## **Part III**

# Chapter 7

## LeetCode

---

(Bubble Sort)	$O(n)$	$O(n^2)$	$O(n^2)$	$O(1)$
(Selection Sort)	$O(n^2)$	$O(n^2)$	$O(n^2)$	$O(1)$
(Insertion Sort)	$O(n)$	$O(n^2)$	$O(n^2)$	$O(1)$
(Quick Sort)	$O(n \log n)$	$O(n \log n)$	$O(n^2)$	$O(\log n)$
(Merge Sort)	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$	$O(n)$
(Heap Sort)	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$	$O(1)$

### 7.1 C basic

[toc]

C

static

extern

C

C

C

C

## 2. 头文件中的全局变量声明

正确的做法是在头文件中声明全局变量，并在一个源文件中定义它。使用`extern`关键字声明全局变量，而不是定义它。

示例：

```
c                                     ⚒ 复制代码
// global.h
extern int globalVar; // 声明全局变量

c                                     ⚒ 复制代码
// file1.c
#include "global.h"

int globalVar = 42; // 定义全局变量

c                                     ⚒ 复制代码
// file2.c
#include "global.h"

void someFunction() {
    globalVar = 100; // 使用全局变量
}
```

这样，`globalVar` 只在 `file1.c` 中定义一次，但可以在其他任何包含 `global.h` 的源文件中使用。

C

```
#include <stdio.h>

int check_endianness() {
    union {
        unsigned int i;
        unsigned char c[4];
    } test_union;

    test_union.i = 0x01020304;

    if (test_union.c[0] == 1) {
        return 1; //
    } else {
        return 0; //
```

```

    }
}

int main() {
    if (check_endianness()) {
        printf("(Big-endian)\n");
    } else {
        printf("(Little-endian)\n");
    }
    return 0;
}

```

test\_union.i 0x01020304 32 4

Big-endian 01 02 03 04

Little-endian 04 03 02 01

**64 0 1**

Brian Kernighan 1

Brian Kernighan

x=x&(x-1)

$$\begin{array}{r}
 x = 15 \quad \begin{array}{r} 1 \\ 1 \\ 1 \\ 1 \end{array} \\
 x-1 = 14 \quad \begin{array}{r} 1 \\ 1 \\ 1 \\ 0 \end{array}
 \end{array}$$

$$\begin{array}{r}
 x \& (x-1) \\
 \hline
 \begin{array}{r} 1 \\ 1 \\ 1 \\ 0 \end{array}
 \end{array}$$

对数字x去除二进制最右侧的1

```

#include <stdio.h>
#include <stdint.h>

void count_bits(uint64_t number, int *count_zeros, int *count_ones) {
    *count_ones = 0;
}

```

```

while (number) {
    number &= (number - 1);
    (*count_ones)++;
}
*count_zeros = 64 - *count_ones;
}

int main() {
    uint64_t number = 0b11010101; //
    int count_zeros, count_ones;
    count_bits(number, &count_zeros, &count_ones);
    printf("Number of 0s: %d, Number of 1s: %d\n", count_zeros, count_ones);
    return 0;
}

```

## MSB LSB

64        Highest Significant Bit, MSB        Least Significant Bit, LSB

```

#include <stdio.h>
#include <stdint.h>

//      LSB
int find_lsb(uint64_t number) {
    if (number == 0) return -1; // 0
    int position = 0;
    while ((number & 1) == 0) {
        number >>= 1;
        position++;
    }
    return position;
}

//      MSB
int find_msb(uint64_t number) {
    if (number == 0) return -1; // 0
    int position = 63;
    while ((number & (1ULL << position)) == 0) {
        position--;
    }
    return position;
}

int main() {
    uint64_t number = 0b11010101; //
    int lsb_position = find_lsb(number);
    int msb_position = find_msb(number);

    if (lsb_position != -1) {
        printf("LSB position: %d\n", lsb_position);
    }
}

```

```

} else {
    printf("The number has no LSB (number is 0).\n");
}

if (msb_position != -1) {
    printf("MSB position: %d\n", msb_position);
} else {
    printf("The number has no MSB (number is 0).\n");
}

return 0;
}

```

## memory copy

```

561 void *memcpy(void *dest, const void *src, size_t count)
562 {
563     char *tmp = dest;
564     const char *s = src;
565
566     while (count--)
567         *tmp++ = *s++;
568     return dest;
569 }

```

## 7.2 Bit

[toc]

338.

```

vector<int> ans;
int countOnes(int nums) {
    int numOnes = 0;
    while(nums){
        nums &= (nums - 1);
        numOnes++;
    }
    return numOnes;
}
vector<int> countBits(int n) {
    for (int index = 0; index <= n; index++) {
        ans.push_back(countOnes(index));
    }
    return ans;
}

```

Brian Kernighan 1

Brian Kernighan

```
x=x&(x-1)
```

## 7.3 LeetCode

[toc]

**35.** while

- mid
- 

```
class Solution {  
public:  
    int searchInsert(vector<int>& nums, int target) {  
        int left = 0;  
        int right = nums.size() - 1;  
        while(left <= right) {  
            //      left + right      int max  
            //  
            int mid = left + ((right - left) >> 1);  
            if (nums[mid] > target) {  
                //  
                right = mid - 1;  
            } else if (nums[mid] < target) {  
                //  
                left = mid + 1;  
            } else {  
                return mid;  
            }  
        }  
        //  
        return left;  
    }  
};
```

**74.** for

mid

```
class Solution {  
public:  
    bool searchMatrix(vector<vector<int>>& matrix, int target) {  
        int left = 0;  
        int right = matrix.size() * matrix[0].size() - 1;  
  
        while (left <= right) {  
            int mid = left + ((right - left) >> 1);  
        }
```

```

        int mid_value = matrix[mid/matrix[0].size()][mid%matrix[0].size()];
        // printf("%d \n", mid_value);

        if (target == mid_value) {
            return true;
        } else if (target > mid_value) {
            //right
            left = mid + 1;
        } else {
            right = mid - 1;
        }
    }

    // for (int x = 0; x < matrix.size(); x++) {
    //     for (int y = 0; y < matrix[0].size(); y++) {
    //         // printf("x-y: %d - %d \n", x, y, matrix[x][y]);
    //         if (target == matrix[x][y]) {
    //             return true;
    //         }
    //     }
    // }

    return false;
}

};


```

### 34.

```

first      last

class Solution {
public:
    vector<int> ans;
    vector<int> searchRange(vector<int>& nums, int target) {
        int left = 0;
        int right = nums.size() - 1;

        int first = -1;
        int last = -1;

        //
        if (nums.size() == 0) {
            ans.push_back(-1);
            ans.push_back(-1);
            return ans;
        }

        while (left<=right)
        {

```

```

/* code */
int mid = left + ((right - left) >> 1);
if (nums[mid] == target){
    first = mid;
    right = mid -1;
} else if (nums[mid] > target) {
    //
    right = mid -1;
} else {
    left = mid + 1;
}
}

left = 0;
right = nums.size() - 1;

while (left<=right)
{
    /* code */
    int mid = left + ((right - left) >> 1);
    if (nums[mid] == target){
        last = mid;
        left = mid +1;
    } else if (nums[mid] > target) {
        //
        right = mid -1;
    } else {
        left = mid + 1;
    }
}

ans.push_back(first);
ans.push_back(last);
return ans;
}

};


```

33.

```

class Solution {
public:
    int search(vector<int>& nums, int target) {
        int left=0,right=nums.size()-1;
        while(left<=right){

```

```

int mid=left+(right-left)/2;
if(target==nums[mid]) return mid;
//
if(nums[left]<=nums[mid]){
    //
    if(target>=nums[left]&&target<nums[mid]){
        right=mid-1;
    }else{
        //
        left=mid+1;
    }
}else{
    //
    if(target>nums[mid]&&target<=nums[right]){
        left=mid+1;
    }else{
        right=mid-1;
    }
}
return -1;
};


```

153.

## 7.4

[toc]

[160.](#)

•

```

class Solution {
public:
    ListNode *getIntersectionNode(ListNode *headA, ListNode *headB) {

        if (headA == NULL || headB == NULL) {
            return NULL;
        }
        ListNode *pa = headA;
        ListNode *pb = headB;
        while (pa != pb) {
            pa = (pa == NULL ? headB : pa->next);
            pb = (pb == NULL ? headA : pb->next);
        }
        return pa;
    }
};


```

```

    }
    return pa;
}
};

```

## 206.

```

ListNode* reverseList(ListNode* head) {
    ListNode *pre = nullptr;
    ListNode *curr = head;

    while (curr = NULL) {
        /*      temp      */
        ListNode *temp = curr->next;
        curr->next = pre;
        pre = curr;
        curr = temp;
    }

    return pre;
}

```

## 92. II

## 234.

```

index

vector<int> nums;

ListNode* p = head;

while (p != nullptr) {
    nums.push_back(p->val);
    p = p->next;
}

for (int i = 0, j = (int)nums.size() - 1; i < nums.size() / 2; i++, j--) {
    if (nums[i] != nums[j]) {
        return false;
    }
}

```

```

        }
    }

    return true;
}

```

## 141.

```

false.

if (head == NULL || head->next == NULL) {
    return false;
}

ListNode * fast = head->next;
ListNode * slow = head;

while(fast != slow) {
    if (fast == NULL || fast->next == NULL) {
        return false;
    }

    fast = fast->next->next;
    slow = slow->next;
}

return true;

```

```

class Solution {
public:
    bool hasCycle(ListNode *head) {
        unordered_set table;
        while(head) {
            if (table.find(head) != table.end()) {
                return true;
            }
            table.insert(head);
            head = head->next;
        }
        return false;
    }
};

```

## 142. II

a b

f=2s

f=s+nb f s n

s=nb

Jb

```
class Solution {
public:
    ListNode *detectCycle(ListNode *head) {
        if (head == NULL || head->next == NULL) {
            return NULL;
        }
        ListNode* fast = head;
        ListNode* slow = head;

        while(true) {
            if (fast == NULL || fast->next == NULL) {
                return NULL;
            }
            fast = fast->next->next;
            slow = slow->next;
            if (fast == slow) {
                break;
            }
        }

        if (fast == slow) {
            ListNode* ptr = head;
            while (ptr != slow) {
                slow = slow->next;
                ptr = ptr->next;
            }
            return ptr;
        }
        return NULL;
    }
};
```

```

class Solution {
public:
    ListNode *detectCycle(ListNode *head) {
        unordered_set<ListNode*> table;
        while(head){
            if (table.find(head) != table.end()) {
                return head;
            }
            table.insert(head);
            head = head->next;
        }
        return NULL;
    }
};

```

## 21.

```

ListNode* mergeTwoLists(ListNode* list1, ListNode* list2) {
    if (list1 == nullptr) {
        return list2;
    } else if (list2 == nullptr) {
        return list1;
    } else if (list1->val < list2->val) {
        list1->next = mergeTwoLists(list1->next, list2);
        return list1;
    } else {
        list2->next = mergeTwoLists(list1, list2->next);
        return list2;
    }
}

```

## 24.

- 1.
- 2.
- 3.

```

class Solution {
public:
    ListNode* swapPairs(ListNode* head) {
        if (head == nullptr || head->next == nullptr) {
            return head;
        }
        // 
        ListNode* newHead = head->next;
        //      newhead
        head->next = swapPairs(newHead->next);
    }
};

```

```

//  

newHead->next = head;  

//  

return newHead;  

}  

};  


```

## 25. K

$K$       head

$K$

```

class Solution {  

public:  

    ListNode* reverseKGroup(ListNode* head, int k) {  

        ListNode *p = head;  

        // *K* head  

        // K  

        for(int i = 0; i < k; i++) {  

            if(!p) return head;  

            p = p->next;  

        }  

        // *K*  

        ListNode *q = head;  

        ListNode *pre = nullptr;  

        while(q != p) {  

            ListNode *tmp = q->next;  

            q->next = pre;  

            pre = q;  

            q = tmp;  

        }  

        //  

        head->next = reverseKGroup(p, k);  

        return pre;  

    }  

};  


```

## 2.

```

l1 = [2,4,3], l2 = [5,6,4]
[7,0,8]
342 + 465 = 807.

```

```

class Solution {  

public:  

    ListNode* addTwoNumbers(ListNode* l1, ListNode* l2) {  


```

```

ListNode *head = nullptr, *tail = nullptr;
int carry = 0;
while (l1 || l2) { //
    //          0
    int n1 = l1 ? l1->val: 0;
    int n2 = l2 ? l2->val: 0;
    //
    int sum = n1 + n2 + carry;
    // head
    // 10      10
    if (!head) {
        head = tail = new ListNode(sum % 10);
    } else {
        tail->next = new ListNode(sum % 10);
        tail = tail->next;
    }
    //
    carry = sum / 10;
    //
    if (l1) {
        l1 = l1->next;
    }
    if (l2) {
        l2 = l2->next;
    }
}
//
if (carry > 0) {
    tail->next = new ListNode(carry);
}
return head;
}
};


```

## 19. N

dummy head

```

class Solution {
public:
    ListNode* removeNthFromEnd(ListNode* head, int n) {
        if (head == nullptr || head->next == nullptr) {
            return nullptr;
        }
        ListNode* dummy = new ListNode(0, head);
        ListNode* p1 = head;
        ListNode* p2 = dummy;

        for (int i = 0; i < n; i++) {
            p1 = p1->next;
        }
        p2->next = p1->next;
        delete p1;
        return dummy->next;
    }
};

```

```

    }

    while(p1){
        p1 = p1->next;
        p2 = p2->next;
    }

    p2->next = p2->next->next;

    ListNode* ans = dummy->next;
    delete dummy;

    return ans;
}

};


```

## 146. LRU

### 138.

hash      value

DNA

```

class Solution {
public:
    unordered_map<Node*,Node*>hmap;
    Node* copyRandomList(Node* head) {
        Node *p=head;
        //      DNA
        //      value
        while(p){
            hmap.insert({p,new Node(p->val)});
            p=p->next;
        }
        // 
        p=head;
        while(p){
            //  next
            hmap[p]->next=hmap[p->next];
            //  random
            hmap[p]->random=hmap[p->random];
            p=p->next;
        }
        return hmap[head];
    }
};


```

```
    }
};
```

## 147.

```
class Solution {
public:
    ListNode* insertionSortList(ListNode* head) {
        //
        if (head == nullptr) {
            return head;
        }

        //
        ListNode* dummyHead = new ListNode(0);
        dummyHead->next = head; //

        ListNode* lastSorted = head; //
        ListNode* curr = head->next; //

        //
        while (curr != nullptr) {
            //           lastSorted
            if (lastSorted->val <= curr->val) {
                lastSorted = lastSorted->next;
            } else {
                //
                ListNode* prev = dummyHead;
                while (prev->next->val <= curr->val) {
                    prev = prev->next;
                }

                //
                lastSorted->next = curr->next;
                curr->next = prev->next;
                prev->next = curr;
            }
        }

        //
        curr = lastSorted->next;
    }

    //
    return dummyHead->next;
}
};
```

PCIe

148.

```
head = [4,2,1,3]
[1,2,3,4]
```

23. K

7.5

C ctype.h

---

```
int isalpha(int c)
int isdigit(int c)
int islower(int c)
int tolower(int c)
int toupper(int c)
```

---

125.

```
: s = "A man, a plan, a canal: Panama"
true
"amanaplanacanalpanama"
```

```
class Solution {
public:
    bool isPalindrome(string s) {
        string temp;
        for (auto c:s) {
            if (isdigit(c)) temp += c;
            if (islower(c)) temp += c;
            if (isupper(c)) temp += tolower(c);
        }
        std::cout<<temp<<std::endl;

        int left = 0;
        int right = temp.length() - 1;
        while (left < right) {
            if (temp[left] == temp[right]) {
                left++;
                right--;
            } else {
                return false;
            }
        }
        return true;
    }
}
```

```

    }
};

392.

        "ace" "abcde"      "aec"
i   j   s   t           i   j   s       j   i   t   s

class Solution {
public:
    bool isSubsequence(string s, string t) {
        int l_s = s.length();
        int l_t = t.length();
        int i = 0, j = 0;

        while(i < l_s && j < l_t) {
            if (s[i] == t[j]) {
                i++;
            }
            j++;
        }
        return i == l_s;
    }
};

```

## 7.6

[toc]

### 1.

```

nums = [2,7,11,15], target = 9
[0,1]
nums[0] + nums[1] == 9      [0, 1]

```

```

class Solution {
public:
    vector<int> twoSum(vector<int>& nums, int target) {
        unordered_map<int ,int> table;
        for (int i = 0; i < nums.size(); i++) {
            auto it = table.find(target - nums[i]);
            if (it != table.end()) {
                return {it->second, i};
            }

            table[nums[i]] = i;
        }
    }
};

```

```

        return {};
    }

};

49.

: strs = ["eat", "tea", "tan", "ate", "nat", "bat"]
: [[["bat"], ["nat", "tan"], ["ate", "eat", "tea"]]

class Solution {
public:
    vector<vector<string>> groupAnagrams(vector<string>& strs) {

        // key value
        unordered_map<string, vector<string>> map;

        //
        for (auto str:strs){
            string key = str;
            sort(key.begin(), key.end());
            map[key].push_back(str);
        }

        // ans
        vector<vector<string>> ans;
        for (auto it = map.begin(); it != map.end(); it++) {
            ans.push_back(it->second);
        }
        return ans;
    }

};

```

## 128.

```

nums = [100,4,200,1,3,2]
4
[1, 2, 3, 4]      4

class Solution {
public:
    int longestConsecutive(vector<int>& nums) {
        //
        unordered_set<int> num_set;
        for (const int& num : nums) {
            num_set.insert(num);
        }

        int longestStreak = 0;

```

```

for (const int& num : num_set) {
    // num - 1 num_set num
    if (!num_set.count(num - 1)) {
        int currentNum = num;
        int currentStreak = 1;

        // num
        // currentStreak
        while (num_set.count(currentNum + 1)) {
            currentNum += 1;
            currentStreak += 1;
        }

        // max longestStreak
        //
        longestStreak = max(longestStreak, currentStreak);
    }
}

return longestStreak;
}
};


```

## 7.7

[toc]

### 283.

```

nums      0
: nums = [0,1,0,3,12]
: [1,3,12,0,0]

0          0          0

class Solution {
public:
    void moveZeroes(vector<int>& nums) {
        int n = nums.size(), left = 0, right = 0;
        while (right < n) {
            if (nums[right]) {
                swap(nums[left], nums[right]);
                left++;
            }
            right++;
        }
    }
};

```

11.

```
class Solution {
public:
    //
    int maxArea(vector<int>& height) {
        int left = 0, right = height.size() - 1; //
        int maxArea = 0; // 0

        //
        while (left < right) {
            //
            int currentArea = (right - left) * min(height[left], height[right]);
            //
            maxArea = max(maxArea, currentArea);

            //
            if (height[left] < height[right]) {
                left++;
            } else {
                right--;
            }
        }

        return maxArea; //
    }
};
```

15.

## 7.8 Stack

[toc]

20.

```
'()''''{''}'''['']'    s
([{}])}
```

```
class Solution {
public:
    bool isValid(string s) {
        int n = s.size();
        if (n % 2 == 1) { //      false
            return false;
        }

        //
        unordered_map<char, char> pairs = {
```

```

        {')', '('},
        {']', '['},
        {'}', '{'}
    };

    stack<char> stk; // 

    //

    for (char ch: s) {
        if (pairs.count(ch)) { // 
            //           false
            if (stk.empty() || stk.top() != pairs[ch]) {
                return false;
            }
            stk.pop(); //
        } else {
            stk.push(ch); //
        }
    }

    return stk.empty(); //
}
};


```

## 155.

push pop top

```

class MinStack {
public:
    stack<int> x_stack;
    stack<int> min_stack;

    MinStack() {
        min_stack.push(INT_MAX);
    }

    void push(int val) {
        x_stack.push(val);
        min_stack.push(min(min_stack.top(), val));
    }

    void pop() {
        x_stack.pop();
        min_stack.pop();
    }
};


```

```

int top() {
    return x_stack.top();
}

int getMin() {
    return min_stack.top();
}
};

/*
 * Your MinStack object will be instantiated and called as such:
 * MinStack* obj = new MinStack();
 * obj->push(val);
 * obj->pop();
 * int param_3 = obj->top();
 * int param_4 = obj->getMin();
 */

```

### 394.

```

s = "3[a]2[bc]"
"aaabcbc"

string decodeString(string s) {
    string ans = "";    //

    stack<int> num_stk;    //
    stack<string> str_stk;  //

    int temp_num = 0;        //
    int len_s = s.size();   //

    for (int i = 0; i < len_s; i++) {
        if (isdigit(s[i])) {
            temp_num = temp_num * 10 + s[i] - '0';  //
        } else if (isalpha(s[i])) {
            ans = ans + s[i];    //
        } else if (s[i] == '[') {
            num_stk.push(temp_num);    //
            temp_num = 0;             //    temp_num
            str_stk.push(ans);        //    ans
            ans = "";                //    ans
        } else if (s[i] == ']') {
            int times = num_stk.top();    //
            num_stk.pop();             //
            string temp = str_stk.top();  //
            str_stk.pop();             //

            //    ans    times    temp
            for (int j = 0; j < times; ++j) {

```

```

        temp += ans;
    }

    ans = temp; // ans
}

}

return ans; //
}

```

### 739.

```

: temperatures = [73,74,75,71,69,72,76,73]
: [1,1,4,2,1,1,0,0]

```

```

vector<int> dailyTemperatures(vector<int>& temperatures) {
    int n = temperatures.size(); //

    vector<int> ans(n, 0); // 0
    stack<int> st; //

    for (int i = 0; i < temperatures.size(); i++) {
        //
        while (!st.empty() && temperatures[i] > temperatures[st.top()]) {
            auto t = st.top(); //
            st.pop(); //
            ans[t] = i - t; //
        }
        st.push(i); //
    }

    return ans; //
}

```

# Chapter 8

## System Design

[system-design-primer](#)

Tool

### 8.1 Git

```
git bash git status
git core.quotepath false quotepath --global
git bash
git config --global core.quotepath false
```

## **Part IV**

# Chapter 9

[toc]

- 1.
- 2.
- 3.
- 4.
- 5.
- 6.
- 7.

## Meeting Minutes

- 1.
2. Jira HSD      unit test
3.                    BIOS      OS      BIOS

## One on One

manager      manager      second manager arch Jun one one

# **Chapter 10**

**2024**

[toc]

**1566**

---

—— GPT-4 ChatGPT

---

**ARM**

---